




ConceptBase.cc User Manual

Version 8.1 ¹

Manfred A. Jeusfeld (ed.)

University of Skövde, 54128 Skövde, Sweden

2019-12-10

¹Copyright 1987-2019 by The ConceptBase Team. All rights reserved. See <http://conceptbase.cc/cbteam.html> for details. This report is published by The ConceptBase Team via <http://conceptbase.cc>. You can use this report for private and academic purposes. Use in a commercial context is only permitted for user's self-study of ConceptBase.cc. If you publish research results or sell products or services based on ConceptBase.cc, then include a reference to this user manual. Unless you are member of the ConceptBase Team, you may not alter this report and you may not republish it, neither electronically nor as hardcopies. Rather point to the original location of this report at <http://conceptbase.cc/CB-Resources.html>. A version of the user manual is included in the ConceptBase.cc installation directory in sub-directory doc/UserManual. This report is published under a Creative Commons License Attribution-NonCommercial-NoDerivatives 4.0 International (<https://creativecommons.org/licenses/by-nc-nd/4.0/>).

Abstract. ConceptBase.cc (in short ConceptBase) is a multi-user deductive object manager intended for conceptual modeling, metamodeling, and coordination in design environments. The system implements O-Telos, a dialect of Telos integrating deductive and object-oriented paradigms. It uniformly represents all information regardless of its abstraction level (data, class, meta class, meta meta class etc.) in a single data structure called P-facts. The deductive query language is seamlessly integrated into the meta class hierarchy. Modeling is supported by meta classing, deduction and integrity checking, active rule specification, functional definition of computation, a module concept, and a rollback mechanism for querying past database states. These principles are combined orthogonally, e.g. deductive rules can be restricted to modules, formulated for meta classes, employed in active rules using functional definitions to compute properties, and be revised without overwriting the earlier definitions. The Java-based user interface CBIva offers a palette of graphical interfaces, such as an Telos editor and a graph editor. The communication between the user clients and the object base is organized in a client-server architecture using TCP/IP. The Java-based command line interface CBSHELL allows for easy scripting of commands directed to a ConceptBase server. The CBSHELL source code can also be used as a programming interface to ConceptBase.

Contributions to this manual were made by: Manfred Jeusfeld, Martin Staudt, Hans Nissen, Christoph Quix, Lutz Bauer, Rainer Gellersdörfer, Michael Gebhardt, Matthias Jarke, Thomas List, René Soiron.

Contributions to the source code of ConceptBase were made by: Masoud Asady, Lutz Bauer (module system), Markus Baumeister, Ulrich Bonn, Stefan Eherer, *Rainer Gellersdörfer* (object store), Michael Gebhardt, Dagmar Genenger, Michael Gocek, Rainer Hermanns, *Manfred Jeusfeld* (CB server, logic foundation, function component), David Kensche, André Klemann, Eva Krüger (integrity component), Rainer Langohr, Farshad Lashgari, Tobias Latzke, Xiang Li, Yong Li, *Thomas List* (object store), Andreas Miethsam, *Hans Nissen* (module system), Martin Pöschmann, *Christoph Quix* (CB server, view component), Christoph Radig (object base interface), Achim Schlosser, Tobias Schöneberg, René Soiron (optimizer), *Martin Staudt* (CB server, query component), *Kai von Thadden* (query component), Hua Wang (answer formatting), Claudia Welter, Thomas Wenig, and others. If you ever contributed to the ConceptBase source code and you are not listed here, then give us a note!

Contents

Chapter 1

Introduction

ConceptBase.cc is a deductive object base management system for meta databases. Its data model is a conceptual modeling language making it particularly well-suited for design and modeling applications. Its underlying data model allows to uniformly represent data, classes, meta classes, meta meta classes etc. yielding a powerful metamodeling environment. The system has been used in projects ranging from development support for data-intensive applications [?], requirements engineering [?, ?, ?], electronic commerce [?], and version&configuration management [?] to co-authoring of technical documents [?]. It has mostly been used in academia for developing specialized modeling languages by means of metamodeling [?, ?, ?].

The key features distinguishing ConceptBase.cc from other extended DBMS and meta-modeling systems are:

- Unlimited meta class hierarchy, allowing to represent information at any abstraction level (data, class, meta class, meta meta class)
- Uniform data structure (called *P-fact*) for concepts, their attributes, their class memberships, and their super- and sub-concepts; all four types of information are full-fledged objects
- Clean formal integration of deductive and object-oriented abstraction by Datalog logical theories
- Complex computations can be user-defined by recursive function definitions, e.g. the length of the shortest path between two nodes
- Queries are defined as classes with user-defined membership constraints; queries can range over any type of object at any abstraction level
- Active rules can be used to define the system's reaction to events; active rules can change the state of the database and can trigger each other
- Client-server architecture with wide-area Internet access

ConceptBase.cc implements the version O-Telos (= Object-Telos) of the knowledge representation language Telos [?]. O-Telos integrates a thoroughly axiomatized structurally object-oriented kernel with a predicative assertion language in the style of deductive databases. A complete formal definition can be found in [?, ?]. O-Telos is purely based on deductive logic but it also supports a frame-like representation of facts.

This user manual is tightly integrated with the ConceptBase.cc Forum. The ConceptBase.cc Forum is an Internet-based workspace where ConceptBase.cc developers and users share knowledge and example models. It contains numerous examples on how to solve certain modeling problems. It is highly recommended to join the workspace. More details are available at <http://conceptbase.sourceforge.net/CB-Forum.html>.

ConceptBase.cc is mainly used for metamodeling and for engineering customized modeling languages. The textbook [?]

Jeusfeld, M.A., M. Jarke, and J. Mylopoulos:
Metamodeling for Method Engineering.
Cambridge, MA, 2009. The MIT Press, ISBN-10: 0-262-10108-4.

introduces into the topic and provides six in-depth case studies ranging from requirements engineering to chemical device modeling. The book and this user manual are complementary to each other.

1.1 Background: Telos and O-Telos

The knowledge representation language *Telos* has been one of the earliest attempts to integrate deduction, object-orientation and metamodeling [?, ?]. The O-Telos [?] dialect supported in ConceptBase.cc has as design goals the semantic simplicity, symmetry of deductive and object-oriented views, metamodeling flexibility, and extensibility at any abstraction level. This emphasis, technically supported by a careful mapping of O-Telos to Datalog with negation (DATALOG⁻), has paid off both in user acceptance and ease of implementation. In essence, ConceptBase.cc is based on deductive database technology with object-oriented abstraction principles like object identity, class membership, specialization and attribution being coded as pre-defined deductive rules and integrity constraints.

Development of ConceptBase.cc started in late 1987 in the context of ESPRIT project DAIDA [?] and was continued within ESPRIT Basic Research Actions Compulog 1 and 2 (1989 – 1995), the ESPRIT LTR project DWQ (Foundations of Data Warehouse Quality, 1996-1999), and the ESPRIT project MEMO (Mediating and Monitoring Electronic Commerce, 1999-2001). Versions have been distributed for research experiments since early 1988. ConceptBase.cc has been installed at more than eight hundred sites worldwide and is seriously used by about a dozen research projects in Europe, Asia, and the Americas.

The direct predecessor of O-Telos is the knowledge representation language *Telos* (specified by John Mylopoulos, Alex Borgida, Martin Stanley, Manolis Koubarakis, Dimitris Plexousakis, and others). *Telos* was designed to represent information about information systems, esp. requirements. *Telos* was based on CML (Conceptual Modeling Language) developed in the mid/late 1980-ties. A variant of CML was created under the label SML (System Modeling Language)¹ and implemented by John Gallagher and Levi Solomon at SCS Hamburg. CML itself was based on RML (Requirements Modeling Language) developed at the University of Toronto by Sol Greenspan and others. Neither RML nor CML were implemented in 1987. They were regarded as theoretic knowledge representation languages with 'possible world semantics'. SML was implemented as a subset of CML using Prolog's SLDNF semantics.

In 1987, we decided to start an implementation of *Telos* and quickly realized that the original semantics was too complex for an efficient implementation. The temporal component of *Telos* included both valid time (when an information is true in the domain) and transaction time (when the information is regarded to part of the knowledge base). The temporal reasoner for the dimensions in ConceptBase.cc V3.0 only to see that there were undesired effects with the query evaluator and the uniform representation of information into objects. Specifically, the specialization of a class into a subclass could have a valid time which could be incomparable to the valid time of an instance of the subclass. Any change in the network of valid time intervals could change the set of instances of a class. Because of that, we dropped the valid time as a built-in feature of objects but we kept the transaction time. A few other features like the declarative TELL, UNTELL and RETELL operations as proposed by Manolis Koubarakis in his master thesis on *Telos* were only implemented in a rather limited way - essentially forbidding direct updates to derived facts. On the other hand, O-Telos extends the universal object representation to any piece of explicit information and reduces the number of essential builtin objects to just five. So, some of the roots of *Telos* in artificial intelligence were abandoned in favor of a clear semantics and of better capabilities for metamodeling. More specifically, there are some important differences of O-Telos to the original *Telos* specification.

- Labels are not used as object identifiers in O-Telos. For example, *Telos* would represent an object like "bill" as P(bill,bill,-,bill). In O-Telos, we represent such objects as P(id123,id123,bill,id123). Each object in O-Telos has a system-generated identifier.

¹ The acronym SML survived in ConceptBase as the filetype .sml for source models.

- Telos promotes the use of level classes such as "Token", "SimpleClass", "MetaClass" etc. to classify user-defined objects into abstraction levels similar to the OMG UML Infrastructure. Besides, Telos defines a so-called omega-level that is parallel to the abstraction levels and defines in particular the object "Proposition". O-Telos does not pre-define the level classes and does not need them. It is up to the user to demand the existence of such levels.
- There are only 5 pre-defined "omega" objects in Telos to define "individual" objects, instantiations, specializations, attributions, and the most-generic object "Proposition". Telos has some more pre-defined such objects like "Class" and "AttributeClass". Such objects are regarded in O-Telos as user-definable objects.
- There are about 50% more builtin axioms in O-Telos (see appendix ??) compared to Telos, in particular to define attribute specialization and to forbid certain pathological databases with groups of entangled objects.

The ConceptBase.cc server adds further capabilities to O-Telos that were not defined/implemented in Telos:

- There are deductive rules and integrity constraints that are mapped to DATALOG⁺ and evaluated by a DATALOG⁺ engine. The engine supports stratified negation, checked at runtime.
- There is a query language that is embedded into the notion of class specialization and realized by the DATALOG⁺ engine. Query classes work like classes whose instances are derived by deductive rules.
- There are active rules that allow to trigger actions when certain events occur.
- Arithmetic is supported as well as recursive function definitions. The implementation utilizes the DATALOG⁺ engine for evaluating recursive function expressions.
- ConceptBase provides a partial evaluator for rules and constraints that range over multiple instantiation levels. Example of such formulas were also proposed for Telos, in particular for cardinality constraints. The partial evaluator makes sure that the DATALOG⁺ engine can always detect stratification violations.
- ConceptBase.cc provides an elaborate module concept for O-Telos databases. Sub-modules see all definitions from super-modules but not of sibling modules (unless imported).
- Stratification in ConceptBase.cc is solely defined in terms of logical stratification (recursive rules in the presence of negation). There is no need on ConceptBase.cc to demand fixed abstraction levels to constrain the instantiation relation between objects. Since ConceptBase.cc is mostly used for metamodeling, the user can define such constraints if they are necessary for the modeling domain.

Since we wanted to be able to manage large knowledge bases (millions of concepts rather than a few hundred), we decided to select a semantics that allowed efficient query evaluation. Telos included already features for deductive rules and integrity constraints. Thus, the natural choice was DATALOG⁺ with perfect model semantics. The deductive rule evaluator and the integrity checker were ready in 1990. A query language ("query classes") followed shortly later. O-Telos exhibits an extreme usage of DATALOG⁺:

- There is only one base relation P(o,x,l,y) called P-facts used for all objects, classes, meta classes, attributes, class membership relationships, and specialization relationships as well as for deductive rules, integrity constraints and queries.
- The semantics of class membership, specialization, and attribution is encoded by around 30 axioms, which are either deductive rules or integrity constraints.
- Deductive rules ranging over more than one classification level (instances, classes, meta classes, etc.) are partially evaluated to a collection of rules (or constraints) ranging over exactly one classification level.

O-Telos should be regarded as the data model of a database for metamodeling. It is capable to represent semantic features of (data) modeling languages like entity-relationship diagrams and data flow diagrams. Once modeled as meta classes in O-Telos, one simply has to tell the meta classes to ConceptBase.cc to get an environment where one can manipulate models in these modeling languages. Since all abstraction levels are supported, the models themselves can be represented in O-Telos (and thus be managed by ConceptBase.cc).

1.2 The architecture of ConceptBase.cc

ConceptBase.cc 8.1 follows a client-server architecture. Clients and servers run as independent processes which interact via inter-process communication (IPC) channels (Fig. 1-1). Although this communication channel was initially meant for use in local area networks, it has been used successfully for nationwide and even transatlantic collaboration of clients on a common server.

The ConceptBase.cc server (CBserver) offers programming interfaces that allow to build clients and to exchange messages in particular for updating and querying object bases using the Telos syntax. We provide support for Java and to a very limited degree for C/C++. Descriptions of the interfaces and the corresponding libraries that are delivered with ConceptBase.cc can be found in the **ConceptBase.cc Programmers Manual**, available via the CB-Forum at <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/885553>. We like to note that the C/C++ interfaces were not maintained since we switched the user interface to Java.

Besides the Java/C API, the CBShell client (see section ??) can be used to interact with a CBserver via the command line or in shell scripts. CBShell is indeed a Java client of the CBserver. The CBShell can also serve as an example client for programming own application specific client tools. There is also a tool that creates an HTTP interface to a CBserver, see section BrokerReproxy in the CB-Forum <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/895647>. Clients would then interact with ConceptBase via HTTP requests.




Figure 1.1: The client-server architecture of ConceptBase.cc

ConceptBase.cc comes with a standard usage environment implemented in Java which supports editing, ad-hoc querying and browsing (CBIVA). The tool CBGraph supports editing diagrams extracted from the database, and CBShell is a command line shell for interacting with the database.

Although ConceptBase.cc provides multi-user support and an arbitrary number of clients may be connected to the same server process, ConceptBase.cc does not yet support concurrency control beyond a forced serialization of messages.

A performance comparison [?] of ConceptBase.cc with Protegé/Racer found that ConceptBase.cc is orders of magnitude faster for queries. It lacks however the reasoning capabilities of Protegé/Racer.

1.3 Hardware and software requirements

The ConceptBase.cc server (CBserver) can be compiled on at least the following platforms² including

- i386 CPUs under Linux Kernel 3.0 or higher,
- x86_64 (AMD64) CPUs under Linux Kernel 3.0 or higher,
- x86_64 (AMD64) CPUs under Windows 10 (Creators Update or later) with Linux sub-system enabled.

Pre-compiled binaries are provided for Linux (and thus also Windows 10). Compilation from the ConceptBase.cc sources on Mac OS-X and other platforms should be possible in principle, though we cannot provide support for them. See instructions distributed with the ConceptBase.cc source files for further details.

Implementation languages for the CBserver are Prolog³ (in particular for logic-based transformation and compilation tasks) and C/C++ (in particular for persistent object storage and retrieval).

The ConceptBase.cc usage environment (CBIva, CBGraph, CBShell) executes on any platform with a compatible Java Virtual Machine. Java 6, Java 7, or Java 8 should all work. We recommend the most recent stable version of Java 7 or Java 8.

The CBserver is dynamically linked with a couple of shared libraries. Under Linux/Unix can check whether all required libraries are installed by

```
export PATH=$CB_HOME/bin:$PATH
ldd $CB_HOME/`CBvariant`/bin/CBserver
```

The installation of ConceptBase.cc requires about 50 MB of free hard disk space. The main memory requirements depend on the size of the object base loaded to the ConceptBase.cc server. The initial main memory footprint is just about 8 MB. We recommend about 20 MB of free main memory for small applications and 200 MB and higher for large applications of ConceptBase.cc. The server can handle relatively large databases consisting of a few million objects. Response times depend on the size of the database and even more on the structure of the query.

Since clients connect to a CBserver via Internet, the server requires the TCP/IP protocol to be available on both the client and the server machine (can be the same computer for single-user scenarios). Note that a firewall installed on the path between the client and the server machine might block remote access to a CBserver. The default port number used for the communication between server and client is 4001. It can be set to another port number by a command line parameter.

The CBserver is by default multi-user capable, i.e. multiple clients can connect to the same CBserver. This feature is by default disabled when you start the CBserver from within the user interface. See section ?? for more details.

The standard ConceptBase.cc client are CBIva and CBGraph (see section ??). The distribution also contains a client CBShell that can be used to interact with a ConceptBase.cc server using a command/shell window. The CBShell client can also be used to run non-interactive scripts, e.g. for loading a sequence of files with Telos source models into the CBserver.

1.3.1 Installation

The download and installation instructions are available from the ConceptBase.cc home page at <http://conceptbase.sourceforge.net/CB-Download.html>. The binaries are installed via a self-extracting Java Archive (CBinstaller.jar).

The sources are made available as a ZIP archive CBPOOL.zip. Compilation from sources on platforms different from Linux requires in-depth expertise due to the manifold of programming languages used for ConceptBase (Prolog, C, C++, Java).

²All trademarks are property of their respective owners.

³ConceptBase.cc now relies on SWI-Prolog [<http://www.swi-prolog.org/>]. Formerly, ProLog by BIM had been used. We only use constructs of SWI-Prolog 5.6 but later versions, in particular SWI-Prolog 6.x, should be compatible.

You can also install a virtual appliance (Linux) that includes the binaries, sources, and the complete development environment.

The default installation directory under Windows is `c:\conceptbase`. Under Linux, ConceptBase is installed by default in the user's home directory `$HOME/conceptbase`.

1.4 Overview of this manual

This manual provides detailed information about using ConceptBase.cc. Information about the installation procedure can be found in the Installation Guide in directory `doc/TechInfo`. New users are advised to follow the installation guide for getting the system started and then to work through the ConceptBase.cc Tutorial. More information about the knowledge representation mechanisms, the applications, and the implementation concepts can be found in the references. Chapter ?? describes the ConceptBase.cc version of the Telos language and gives some examples for its usage. Chapter ?? discusses the parameters that can be set when starting the CBserver. Finally, section ?? describes the ConceptBase.cc Usage Environment.

Appendices contain a formal definition of the Telos syntax and internal data structures (??). Appendix ?? summarizes the mechanism for assigning graphical types to objects and adapting the graphical browsing tool for specific application needs. Appendix ?? contains the full Telos notation of an example model (??) and a case study on the modeling of entity-relationship diagrams with Telos (??). Plenty of further examples for particular application domains and add-ons for metamodeling can be retrieved from the ConceptBase.cc Forum at <http://conceptbase.sourceforge.net/CB-Forum.html>.

1.5 Differences to earlier versions

ConceptBase.cc 8.1 should be largely source-compatible to its direct predecessor. The binary database files and the graph files have a new format and are not compatible with their counterparts created by earlier releases.

The CBserver has now the ability to maintain module sources and query results formatted in external formats in the file system. The module sources allow to co-develop models both via the ConceptBase.cc user interface and by external text editors. Exporting query results in external formats is useful when they are post-processed by external tools. For example, one can generate program source code from a ConceptBase.cc model and have that code processed by a compiler.

The active rule component now supports constructs to enforce a transactional execution of delayed triggers. Triggers can be passed to different queues that are processed with different priorities. This feature allows to delay certain triggers until the consequences of the current trigger are all processed.

The graph editor is now a stand-alone tool and has the ability to store connection details plus a snapshot of the module sources in its graph files. The graph files are then self-contained and can be viewed and updated without having to maintain the module sources elsewhere. It also can display a background image in the graph window. Nodes can now be configured to be resizable. Lines are now drawn with anti-aliasing, yielding much better graph images.

The CBSHELL utility now behaves more like a Linux/Unix shell. It provides easy shortcuts like `'cd'` for changing the module context. It also allows the use of positional parameters, hence making it a better companion to Linux/Unix scripts.

The release notes to ConceptBase.cc 8.1 lists all major changes and issues. You find the release notes in the subdirectory `doc/TechInfo` of your ConceptBase.cc installation directory or via the web site <http://conceptbase.cc>.

The system still has about the same memory footprint as it used to be 10 years ago. You can easily install the complete system for all supported platforms on a 32 MB memory stick.

1.6 License terms

ConceptBase.cc is distributed under a FreeBSD-style copyright license since June 2009. Both binary and source code are available via <http://sourceforge.net/projects/conceptbase/files>.

The FreeBSD-style copyright license of ConceptBase.cc reads like follows:

The ConceptBase.cc Copyright

Copyright 1987-2019 The ConceptBase Team. All rights reserved.

Redistribution and use in source and binary forms, with or without modification, are permitted provided that the following conditions are met:

1. Redistributions of source code must retain the above copyright notice, this list of conditions and the following disclaimer.
2. Redistributions in binary form must reproduce the above copyright notice, this list of conditions and the following disclaimer in the documentation and/or other materials provided with the distribution.

THIS SOFTWARE IS PROVIDED BY THE CONCEPTBASE TEAM ``AS IS'' AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE CONCEPTBASE TEAM OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

The views and conclusions contained in the software and documentation are those of the authors and should not be interpreted as representing official policies, either expressed or implied, of the ConceptBase Team.

This license makes ConceptBase.cc free software as promoted by the Free Software Foundation. The license is upwards-compatible to the GNU Public License (GPL), i.e. developers can combine GPL-ed software with ConceptBase.cc software as long as they include the above FreeBSD-style license for the ConceptBase.cc components. It should also be compatible with many other free license models. The source code is copyrighted by The ConceptBase Team, consisting of all contributors to the source code at the central code repository of the system.

Binary and source distributions of ConceptBase.cc may contain third-party software. Their licenses are listed by default in the directory `doc/ExternalLicenses` of the installation directory.

You are welcome to contribute the ConceptBase.cc project! Join the ConceptBase.cc Forum at <http://conceptbase.sourceforge.net/CB-Forum.html> to do so. We are also happy to learn about your research/application, in which ConceptBase.cc plays a role. If you publish results of your work, then please include a reference pointing to this user manual and/or to the standard ConceptBase.cc reference [?].

The ConceptBase.cc system is published under the liberal FreeBSD-style license, which allows commercial use and modification of the source code. This user manual is however published under a less liberal copyright license. Use is permitted for private and academic purposes. Moreover, commercial users may use the manual to self-study the ConceptBase.cc system. Only members of The ConceptBase Team, see <http://conceptbase.cc/cbteam.html>, are allowed to modify the user manual. Re-publication of the user manual in print or online is not permitted. If you make changes to your own copy of the source code of ConceptBase.cc, then you may not document them in this user manual (or its sources). You rather should write a companion report that lists the differences of your version of ConceptBase.cc to the user manual distributed via the home page of ConceptBase.cc (<http://conceptbase.cc>).

The ConceptBase.cc Forum (<http://conceptbase.sourceforge.net/CB-Forum.html>) contains material submitted by different authors. The default license for the material on the CB-Forum is "Creative Commons BY-NC 4.0", which does not permit commercial use of its content.

The ConceptBase.cc logo at <http://conceptbase.sourceforge.net/conceptbase-cc-logo.gif> is created and copyrighted by Manfred Jeusfeld. It may be used on official ConceptBase.cc web sites, the ConceptBase.cc User Manual and Tutorials, and the ConceptBase.cc Forum. It may also be included in binary distributions created by members of the ConceptBase Team. If you are not a member of the ConceptBase.cc Team and like to use the logo for your own binary distribution of ConceptBase.cc, then you need to ask the copyright holder for a permission.

1.7 Disclaimer

All trademarks are owned by their respective owners. This report may contain flaws based on human errors. We disclaim liability for any such flaws. It may also be that the ConceptBase.cc does not provide all functionality described in this report, or that the functionality is provided by other mechanisms as described in this report. Links to external websites are provided for informational (academic) purposes. We disclaim responsibility for the views expressed on these websites.

We sometimes use the short form ConceptBase. We then always refer to ConceptBase.cc.

Chapter 2

O-Telos by ConceptBase.cc

ConceptBase.cc is an implementation of the O-Telos data model. O-Telos is derived from the knowledge representation language Telos as designed by Borgida, Mylopoulos and others [?]. While Telos was geared more to its roots in artificial intelligence, O-Telos is more geared to database theory, in particular to deductive databases. Nevertheless, O-Telos is to a large degree compatible to the original Telos specification. In some respects, it generalizes Telos, for example by removing the requirement to classify objects into the levels for tokens, simple classes, and meta classes. In O-Telos, we have just five predefined objects (see appendix on the axioms of O-Telos).

Telos (and O-Telos) as well also have strong links to the semantic web, in particular to the triple predicates used for defining RDF(S) statements. The main difference is that O-Telos is based on quadruples where the additional components identifies the statement. While RDF(S) has to use special link types to reify triple statements, i.e. to make statements about statements, O-Telos statements are simply referred to by their identifier.

Telos' structurally object-oriented framework generalizes earlier data models and knowledge representation formalisms, such as entity-relationship diagrams or semantic networks, and integrates them with predicative assertions, temporal information, and in particular metamodeling. This combination of features seems to be particularly useful in software information applications such as requirements modeling and software process control. A formal description of O-Telos can be found in [?, ?]. The following example is used throughout this section to illustrate the language:

A company has employees, some of them being managers. Employees have a name and a salary which may change from time to time. They are assigned to departments which are headed by managers. The boss of an employee can be derived from his department and the manager of that department. No employee is allowed to earn more money than his boss.

This section is organized as follows: first, the "logical" and "frame" representations of O-Telos are explained. Then, the predicative sublanguage for deductive rules and integrity constraints are presented. Subsection ?? presents a declarative query language which introduces queries as classes with optional predicative membership specification.

2.1 Propositions and frames

As a hybrid language O-Telos supports two different representation formats: a logical ("propositions") and a frame representation. The latter format is based on the logical one. As explained in the next subsections the logical representation also forms the base for integrating a predicative assertion language for deductive rules, queries, and integrity constraints into the frame representation. We start with the so-called P-fact representation of a O-Telos database.

A *historical* O-Telos database is a finite set of propositions (=P-facts=objects):

$$OB = \{P(oid, x, n, y, tt) | oid, x, y, tt \in ID, n \in LABEL\}$$

where `oid` has the key property within the set, `ID` is an infinite but countable set of identifiers. The set `LABEL` is a set of names over some alphabet. The components `oid`, `x`, `n`, `y`, `tt` are called identifier, source, label (or name), destination and transaction time of the proposition¹. We read them as follows:

The object `x` has a relationship called `n` to the object `y`. This relationship is believed by the system for the time interval `tt`.

The transaction time `tt` is represented by two time points $tt(a, b)$, where a is the time when the *tell* time and b the *untell* time of the P-fact. The historical O-Telos database is the basis for the *rollback* O-Telos database that is visible at a given point of time.

$$OB_{rbt} = \{P(oid, x, n, y) | P(oid, x, n, y, tt) \in OB, rbt \ll tt\}$$

The clause $t_1 \ll t_2$ expresses that the time interval t_1 is contained in time interval t_2 , i.e. t_1 is during t_2 . The value of the rollback time depends on the kind of formula to be processed: integrity constraints are always evaluated against the *current* database OB_{now} ² (`now`=the smallest time interval that contains the current time). The rollback time of queries is usually provided together with the query when it is submitted from a user interface to a ConceptBase server. By default, it is *now* as well. Subsequently, we shall use OB_{rbt} rather than OB . Note that OB_{rbt} strips the transaction time `tt` from the P-facts. We shall still call both a P-fact. Any P-fact from OB_{rbt} has a unique counterpart in OB , hence the transaction time can always be looked up in OB . We demand that any OB_{rbt} is consistent but do not demand it for the database OB , which stores the whole history of updates.

O-Telos imposes some structural axioms on databases, e.g. referential integrity, correct instantiation and inheritance ([?]). The complete list of axioms is contained in appendix ???. The axioms are linked to predefined objects that are part of each O-Telos database. There are five predefined O-Telos objects for five patterns of propositions:

- **Proposition** contains all propositions as instances. A proposition is any P-fact in an O-Telos database OB_{rbt} that has the form $P(oid, x, n, y)$. Any proposition must fall into exactly one of the subsequent cases.
- **Individual** is a the class of all P-Facts that have the form $P(oid, oid, n, oid)$. Such P-facts are denoted as nodes in the graphical representation of an O-Telos database.
- **InstanceOf** contains all explicit instantiation objects as instances. This is exactly the set of P-facts matching the pattern $P(oid, x, *instanceof, c)$. We say that x is an (explicit) instance of c . In the graphical representation, an instantiation object is a link between some object x and its class c . Note that any such explicit instantiation is also an object/proposition with identifier `oid`.
- **IsA** contains all explicit specialization objects as instances. Specialization P-facts match the pattern $P(oid, c, *isa, d)$. A specialization object is graphically displayed as a link between a subclass c and its superclass d .
- **Attribute** contains all explicit attribution/relation objects as instances. An attribution/relation object matches the pattern $P(oid, x, m, y)$ where m must be different from `*instanceof` and `*isa`. It is displayed by a link between the source object x and the destination object y . The label of the attribution object is the name of the attribute link starting from x . The object y is also called the value or destination of the attribute. In O-Telos, the class **Attribute** subsumes relations between objects since values are just objects in O-Telos. In ConceptBase, **Attribute** is a shortcut for **Proposition!attribute**, i.e. the link **attribute** of **Proposition**.

¹We will see in section ?? that the predicative language operates on a snapshot of the database, i.e. on those propositions that are believed at a specified reference time called rollback time. This time is an interval. The of the interval is the time when the object has been told/created. The end of the interval is either the time when the object was untold/deleted, or it is a special symbol *infinity*, indicating that the object is currently believed, i.e. it is not deleted.

²Only objects that have a right-open belief time shall be visible OB_{now} . This is due to the fact that the end time of an object can only be changed once, namely when the object is untold. The UNTELL operation can only have happened in the past when OB_{now} is built.

ConceptBase supports deductive rules for deriving the instantiation of an object to a class and attributes/relations between objects. This derived information has no object property, i.e. it is not identified and it is not represented as a proposition. Specifically, the instantiation of propositions to the above five pre-defined objects is derived by deductive rules, specifically by axioms 18-22 in appendix ??.

Additional to the above predefined classes³, there are the builtin classes `Class`, `Integer`, `Real` and `String`. `Class` contains all so-called classes (including itself) as instances. The only special property of `Class` is the definition of two attribute categories `rule` and `constraint`. Hence, instances of classes can have deductive rules and integrity constraints. Integer and real numbers are written in the usual way, strings are character sequences, e.g. "this is a string". These three classes are supported by comparison predicates like $(x < y)$ discussed in section ??, and by functions like `PLUS`, `MINUS` discussed in section ??.

As legacy support, ConceptBase provides the pre-defined classes `Token`, `SimpleClass`, `MetaClass`, and `MetametaClass` to structure the database into objects that have no instances (tokens), objects that have only tokens as instances (simple classes), objects that have only simple classes as instances (meta classes), and finally objects that have only meta classes as instances (meta meta classes). These classes are provided only for compatibility with older Telos specifications. In fact, an absolute hierarchy from tokens to simple classes to meta classes etc. is not an essential ingredient of O-Telos and in many situations too restrictive.

Instead, meta class levels are implicitly expressed via instantiation. If an object x is an instance of object c and object c is an instance of object m_c , then m_c is also called a *meta class* of x , and c a *class* of x . Being a class or a meta class is **relative** to the object x that we consider. For example, m_c is the class of c . This implicit definition of the meta class concept is far more flexible than a fixed structure:

1. There is virtually no limit in the meta class hierarchy: there can be meta classes, meta meta classes, meta meta meta classes etc.
2. A class can have object from different meta class levels as instances. This is in particular important for extending the capabilities of the O-Telos language. An example of a class that has objects from different levels as instances is `Proposition`: it has *all* objects as instances.
3. A user does not need to decide to which meta class level an object belongs.

Strict conformance to the membership to meta class levels can still be enforced by user-definable integrity constraints.

As a user, you don't work directly with propositions but with textual (frame) and graphical (semantic networks) views on them. Both are not based on the oid's of objects but on their label components. To guarantee a unique mapping we need the following naming axiom:

Naming axiom (see also axioms 2,3,4 in appendix ??)

1. The label ("name") of an individual object must be unique, i.e. if two objects have the same label then they are the same.
2. The label of an attribute must be unique within all attributes with a common source object, i.e. no two explicit attributes of the same object can have the same label. However, two different objects can well have attributes sharing the same label.
3. The source and destination of an instantiation object are unique, i.e. between two objects x and y may be at most one explicit instantiation link.
4. The source and destination of a specialization object are unique.

The **frame syntax** of O-Telos groups the labels of propositions with common source o around the label of o . The exact syntax is given in appendix ??. In this section we introduce it by modeling the employee example:

³Strictly speaking, we should better use the term predefined object or predefined proposition.

```

Employee in Class with
  attribute
    name: String;
    salary: Integer;
    dept: Department;
    boss: Manager
end

Manager in Class isA Employee end

Department in Class with
  attribute
    head: Manager
end

```

The label of the “common source” in the first frame is `Employee`. It is declared as instance of the class `Class` and has four attributes. The class `Manager` is a subclass of `Employee`.

Oid’s (preceded by ‘#’ in our examples) are generated by the system. This leads to the following set of propositions corresponding to the frames above. The transaction time inserted by the system is denoted by omission marks.

```

P (#E, #E, Employee, #E)
P (#1, #E, *instanceof, #Class)
P (#3, #E, name, #String)
P (#4, #E, salary, #Integer)
P (#5, #E, dept, #D)
P (#6, #E, boss, #M)
P (#M, #M, Manager, #M)
P (#7, #M, *instanceof, #Class)
P (#8, #M, *isa, #E)
P (#D, #D, Department, #D)
P (#9, #D, *instanceof, #Class)
P (#10, #D, head, #M)

```

Instantiation to the pre-defined class `Individual` is implicitly given by the structure of the three individual propositions named `Employee`, `Manager`, and `Department`. Analogously, the attributes #3, #4, #5, #6 and #10 are automatically regarded as instances of the class `Attribute`. The instances of `Attribute` are also called *attribution objects* or *explicit attributes*. Propositions #1, #2, #7 and #9 are instances of the class `InstanceOf` (holding explicit instantiation objects), and #8 is an instance of the class `IsA` (explicit specialization objects). Note that all relationships are declared by using the identifiers (not the names) of objects. Thus, #Class, denotes the identifier of the object `Class` etc.

The identifiers are maintained internally by `ConceptBase`’s object store. Externally, the user refers to objects by their name. A standard way to describe objects together with their classes, subclasses, and attributes is the frame syntax. Frames are uniformly based on object names.

The next frames establish two departments labelled `PR` and `RD` and state that the individual object `mary` is an instance of the class `Manager`. `Mary` has four attributes labelled `hername`, `earns`, `advises` and `currentdept` which are instances of the respective attribute classes of `Employee` with labels `name`, `salary` and `dept`.

```

mary in Manager with
  name
    hername: "Mary Smith"
  salary
    earns: 15000
  dept

```

```

    advises:PR;
    currentdept:RD
end

PR in Department end

RD in Department end

```

The corresponding propositions for the frame describing mary are:

```

P(#mary,#mary,mary,#mary)
P(#E1,#mary,*instanceof,#M)
P(#E3,#mary,hername,"Mary Smith")
P(#E4,#E3,*instanceof,#3)
P(#E5,#mary,earns,15000)
P(#E6,#E5,*instanceof,#4)
P(#E7,#mary,advises,#PR)
P(#E8,#E7,*instanceof,#5)
P(#E10,#mary,currentdept,#RD)
P(#E11,#E10,*instanceof,#5)

```

The attribute categories `name`, `salary` and `dept` must be defined in one of the classes of `mary`. In this case `mary` is also instance of `Employee` due to the following axiom which defines the inheritance of class membership in O-Telos, and hence can instantiate these attributes:

Specialization axiom (axiom 13 in appendix ??)

The destination (“superclass”) of a specialization inherits all instances of its source (“sub-class”).

An example is the specialization #8: all instances of `Manager` (including `mary` are also instances of `Employee`. O-Telos enforces **typing** of the attribute values by the following general axiom:

Instantiation axiom (axiom 14 in appendix ??)

If `p` is a proposition that is instance of a proposition `P` then the source of `p` must be an instance of the source of `P`, and the destination of `p` must be an instance of the destination of `P`.

For example, “`Mary Smith`” must be an instance of `String`. The individual `mary` also shows another feature: attribute classes specified at the class level do not need to be instantiated at the instance level. This is the case for the `boss` attribute of `Employee`. On the other hand, they may be instantiated more than once as e.g. `dept`.

In some cases for attribute categories occurring in a frame the corresponding objects which are instantiated by the concrete attributes, can not uniquely be determined⁴. This multiple generalization/instantiation problem is solved⁵ by the following condition which must hold for O-Telos databases:

Multiple generalization/instantiation axiom (axiom 17 in appendix ??)

If `p1` and `p2` are attributes of two classes `c1` and `c2` which have the same label component `l`, and `i` is a common instance of `c1` and `c2` which has an attribute with category `l`, then there must exist a common specialization `c3` of `c1` and `c2` with an `l` labelled attribute `p3` which specializes `p1` and `p2`, and `i` is instance of `c3`. Particularly if `c1` is specialization of `c2` and `p1` is specialization of `p2`, `c1` and `p2` already fulfill the conditions for `c3` and `p3`.

⁴Subsection ?? contains an example for this problem in the context of linking logical formulas to O-Telos objects.

⁵For specialization relationships between two objects we need an axiom similar to the instantiation axiom which requires specialization relationships between their sources and destination components. [?] contains the complete axiomatization.

O-Telos treats all three kinds of relationships (attribute, isa, in) as objects. Thus each attribute, instantiation or generalization link of `Employee` may have its own attributes and instances. For example, each of the four `Employee` attributes is an instance⁶ of an attribute class denoted by the label `attribute` but can also have instances of its own. The attribute with label `earns` of `mary` is an instance of attribute `salary` of class `Employee`. Syntactically, attribute objects are denoted by appending the attribute label with an exclamation mark to the name of some individual. The relationship between `salary` and `earns` could be expressed as

```
mary!earns in Employee!salary
end
```

Instantiation links are denoted by the operator "`->`" and specialization links by "`=>`". They should always be enclosed in parentheses:

```
(mary->Manager)
end

(Manager=>Employee)
end
```

The operators can be combined to form complex expressions. The next example shows how to reference the instantiation link between the attribute `mary!earns` and its attribute class `Employee!salary`. The second frame shows that arbitrarily complex expressions are possible. The parentheses have to be used to make the operator expressions unique. The attribution operator "`!`" has a stronger binding than the instantiation and specialization operators. According to our own experience, complex expressions for denoting objects are rare in modeling. It is good to know that any object in O-Telos can be uniquely referenced in the frame syntax.

```
(mary!earns->Employee!salary) with
  comment
    com1: "This is a comment to an instantiation between attributes"
end

(mary!earns->Employee!salary)!com1 with
  comment
    com2: "This is a comment to the the previous comment attribute"
end
```

The labels `InstanceOf`, `IsA` and `Attribute` for the three Telos system classes are indeed alias names for the following object expressions:

```
Attribute    <--->  Proposition!attribute
InstanceOf   <--->  Proposition->Proposition
IsA          <--->  Proposition=>Proposition
```

Hence, `Attribute` as the alias name for the attribute with label `attribute` of `Proposition`. `InstanceOf` is the alias name for the instantiation link between `Proposition` and `Proposition`. The object `Proposition` is indeed an instance of itself because it has the shape of an individual object, which is a special case of the shape of a proposition. Finally, `IsA` is an alias name for the specialization link between `Proposition` and `Proposition`. This representation deviates slightly from the axioms in appendix ?? because it was originally implemented in `ConceptBase` in this way. The reflexive definition of `InstanceOf` (`IsA`) as instantiation (specialization) links is redundant since O-Telos axioms 18-22 derive these instantiations anyway.




Figure 2.1: Graph representing the example database

Figure ?? shows the graphical representation of mary and her relationships to the other example objects. Labelled links are attributes/relations. The thicker link from Manager to Employee is a specialization. The other links are instantiations. If a link is dotted, then it is derived. Individual objects are displayed as nodes.

O-Telos propositions have a temporal component: the transaction time⁷. The transaction time of a proposition is not assigned by the user but by the system at the time of an update (TELL, UNTELL, RETELL). ConceptBase uses right-open and closed predefined time intervals. Right-open time intervals are represented like in the subsequent example:

```
P(#mary, #mary, mary, #mary, tt(milliseconds(1992, 1, 11, 17, 5, 42, 102),
infinity))
```

The object `mary` is believed since 17:05:42 on January 11, 1992. The label 'infinity' denotes that the end time of the object lies in the future and is not yet known. In any case, the current time 'now' is regarded to be smaller than 'infinity'. Right-open transaction times indicate objects that are part of the "current" knowledge base.

Closed intervals (denoted by binary `tt`-terms) indicate "historical" objects, i.e. objects that have been untold. Example:

```
P(#E1, #mary, *instanceof, #M, tt(milliseconds(1992, 1, 11, 17, 5, 42, 0),
milliseconds(1995, 12, 31, 23, 59, 59, 999))
```

The object `#E1`, i.e. the instantiation of `mary` to the class `Manager` is believed from 17:05:42 on January 11, 1992, until the end of the last millisecond of the year 1995. We call the first component of the transaction time also the start time object and the second component the end time. Start and end time of an object can be retrieved by the predicates `Known`, and `Terminated` (see section ??). Transactions in ConceptBase usually add or terminate several propositions. At the begin of the transaction, ConceptBase reads the current time and uses it to set the transaction time of all affected propositions. Consequently, all

⁶These instantiations were left out in the set of propositions for the employee example above.

⁷The original specification of Telos has two temporal components. The *valid time* and the *transaction time*. The valid time is defined as the time interval when the statement made by a Telos proposition is true in the world. The transaction time is the time when this statement is part of the knowledge base. O-Telos skipped the valid time because it is virtually impossible to have a tractable implementation of Allen's interval calculus [?], when in interplays with deductive rules. Earlier versions of ConceptBase until version 3.1 did however implement both time components.

inserted propositions get the same start time of their transaction time and all terminated propositions get the same end time of their transaction time. The database does not change between two transactions, hence the finite sequence of transaction times can be used to enumerate all updates to the database. Examples on how to inspect the database using the transaction time are in the CB-Forum at <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/1921789>.

2.2 Rules and constraints

The ConceptBase predicative language CBL [JK90] is used to express integrity constraints, deductive rules and queries. The variables inside the formulas have to be quantified and assigned to a “type” that limits the range of possible instantiations to the set of instances of a class. ConceptBase offers a set of predicates for the predicative language defined on top of an O-Telos database as visible for a given rollback time, i.e. OB_{rbt} for some rbt . Any rule, constraint or query is run against OB_{rbt} rather than the full database OB .

2.2.1 Basic predicates

The following predicates provide the basic access to an O-Telos database. Some have both an infix and a prefix notation. As usual we employ the object identifier to refer to an object.

1. $(x \text{ in } c) \text{ or } \text{In}(x, c)$
The object x is an instance of class c .
2. $(c \text{ isA } d) \text{ or } \text{Isa}(c, d)$
The object c is a specialization (subclass) of d
3. $(x \text{ m } y) \text{ or } \text{A}(x, m, y)$
The object x has an attribution link to the object y and this link has the attribute category m . Structural integrity demands that the label m belongs to an attribute of a class of x .
4. $\text{Ai}(x, m, o)$
The object x has an explicit attribute o . This attribute is instance of an attribute category with label m .
5. $(x \text{ m/n } y) \text{ or } \text{AL}(x, m, n, y)$
The object x has an attribution link labelled n to the object y . The attribution has the category m .
6. $\text{From}(p, x)$
The object p has source x .
7. $\text{To}(p, y)$
The object p has destination y .
8. $\text{Label}(p, l)$
The object p has label l . If l is used as a variable, it must be quantified over the class Label .
9. $\text{P}(p, x, n, y)$
There is an object $P(p, x, n, y)$ in the database OB_{rbt} .
10. $\text{Known}(p, t)$ The object p is known in OB_{rbt} since t , i.e. an object $P(p, x, n, y, tt)$ is part of the database OB and t is the start time of tt . The argument t is a string of the format " $tt(\text{millisecond}(\text{yr}, \text{mo}, \text{d}, \text{h}, \text{m}, \text{s}))$ ". It is regarded as an instance of the class TransactionTime .
11. $\text{Terminated}(p, t)$ The object p is unknown in OB_{rbt} after t , i.e. an object $P(p, x, n, y, tt)$ is part of the database OB and t is the end time of tt . The argument t is represented like with Known . An object that has not yet been untold has the end time " $tt(\text{infinity})$ ".

12. $(x \text{ [in] } mc) \text{ or } In2(x, mc)$
The object x is an instance of class c and c is an instance of class mc . In other words, $(x \text{ [in] } mc)$ is equivalent to $\text{exists } c/\text{VAR } (x \text{ in } c) \text{ and } (c \text{ in } mc)$
13. $(x \text{ [m] } y) \text{ or } A2(x, m, y)$
The object x and y are linked by an attribute $a1$. The attribute $a1$ is an instance of an attribute $a2$ which itself is an instance of an attribute $a3$ with label m . The predicate is equivalent to the formula $\text{exists } c, d, n/\text{VAR } (x \text{ in } c) \text{ and } (y \text{ in } d) \text{ and } (c \text{ m/n } d) \text{ and } (x \text{ n } y)$.

The predicates $In2$ and $A2$ are also called *macro predicates* since they are standing for sub-formulas. They are fully supported in constraints of query classes. The predicate $A2$ is not yet supported for deductive rules and integrity constraints due to limitations of the formula compiler. You can use the AL predicate instead. Examples on using macro predicates are available from the CB-Forum (<http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/877047>). They are also discussed in more detail in section ??.

The relation of the above predicates and the P-facts of the database is defined by the O-Telos axioms (appendix ??). For example, axiom 7 states

$$\forall o, x, n, y, p, c, m, d \ P(o, x, n, y) \wedge P(p, c, m, d) \wedge In(o, p) \Rightarrow AL(x, m, n, y)$$

So, if an attribute object o of an object x is an instance of an attribute object p of the object c , then $AL(x, m, n, y)$ (also written as $(x \text{ m/n } y)$) can be derived. This axiom provides those solutions to the AL predicate that are directly based on P-facts. Further solution can be derived via user-defined deductive rules. The other predicates are based on P-facts as well. The Ai predicate is for historical reasons not included in the list of axioms. It is defined as

$$\forall o, x, n, y, p, c, m, d \ P(o, x, n, y) \wedge P(p, c, m, d) \wedge In(o, p) \Rightarrow Ai(x, m, o)$$

There are a few variants for the predicates for instantiation, specialization and attribution to check whether a fact is actually stored or deduced:

1. $In_s(x, c)$
The object x is an explicit instance of class c .
2. $In_e(x, c) \text{ or } : (x \text{ in } c) :$
The object x is an explicit instance of class c , or of one of the sub-classes of c , or of the system class of x . The system class of individual objects is `Individual`, attribution objects have the system class `Attribute`, instantiation objects the system class `InstanceOf`, and specialization objects have the system class `IsA`.
3. $A_e(x, m, y) \text{ or } : (x \text{ m } y) :$
The objects x and y are linked by an explicit attribute with attribute category m . The attribute category is either explicitly assigned to the attribute or derived by a rule (see subsection ??).
4. $Isa_e(c, d) \text{ or } : (c \text{ isA } d) :$
The class c is a direct subclass of class d .
5. $AD(p, x, y)$
The objects x and y are linked by an explicit or derived relation/attribute of category m , where $P(p, c, m, d)$ defines the attribute category. This predicate is mostly used for internal purposes of ConceptBase. The predicate is computationally expensive when the first argument is a variable.

The above predicates can be used, for example, to define default values (see <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/2396075>). Since deduction should be transparent to the user, one should avoid using the above predicates if the proper predicates $In(x, c)$ and $A(x, m, y)$ can do the job.

The predicate $A(x, m, y)$ is translated by the ConceptBase to the $AD(p, x, y)$ predicate, where the first argument is a constant. For example,

$$\forall e/Employee \exists s/Integer A(e, salary, s)$$

would be translated to

$$\forall e/Employee \exists s/Integer AD(Employee!salary, e, s)$$

2.2.2 Notes on attribution

The attribution of objects in O-Telos (axioms 7 and 8 in appendix ??) is more generic than in object-oriented approaches, in particular UML. In O-Telos, an attribution relates two arbitrary objects. In languages such as UML, attributes are defined at classes to declare which states an object (instance of the class) may have. This is well possible in O-Telos as well, e.g. by declaring the integer-valued `salary` attribute of a class `Employee` and using it for instances of the class. However, O-Telos does not restrict attributes to just values. The target of an attribute can be any object. Hence, the concept of an attribute in O-Telos is the generalization of an UML association and an UML attribute. A second difference is that an O-Telos attribute has essentially several labels tagged to it: its own label (object label) and the labels of its attribute categories (class labels). The latter are the labels of the attributes declared at the classes of an object, the first is the label of the attribution at the level of the object that has the attribute. We illustrate this subsequently.

Attributes at the instance level are instances of attributes at the class level (=attribute categories). An attribute category at the class level can be instantiated several times at the instance level. For example, consider the frame for `Mary`:

```
mary in Manager with
  name, aliasname
    hername: "Mary Smith"
  salary
    earns: 15000
  dept
    advises: PR;
    currentdept: RD
end
```

The object `mary` has four attributes with object labels `hername`, `earns`, `advises`, and `currentdept`. The attribute categories are `name`, `aliasname`, `salary`, and `dept`. The last category is instantiated twice. ConceptBase uses the following predicate facts (infix variant of the AL predicate) to express the content of the frame:

```
(mary in Manager)
(mary name/hername "Mary Smith")
(mary aliasname/hername "Mary Smith")
(mary salary/earns 15000)
(mary dept/advises PR)
(mary dept/currentdept RD)
```

So, there are four attributes using four attribute categories. Like an object can have multiple classes, an attribute can have multiple categories. In fact, explicit attributes in O-Telos are just objects and their attribute categories are their classes. At the lowest abstraction level (tokens), the object labels of the attributions frequently do not carry a specific meaning and can then be neglected when formulating logical expressions. The attribution predicate $(x \ m \ y)$ performs just this projection. In the example, the following attributions facts are true:

```

(mary name "Mary Smith")
(mary aliasname "Mary Smith")
(mary salary 15000)
(mary dept PR)
(mary dept RD)

```

The class labels `name`, `aliasname` etc. are defined at an abstraction level where the meaning of some application domain is captured. The class label (attribute category) of an attribute is defined as an object label of an attribute at the class level. For example, the `name` and `aliasname` attributes could be defined for the class `Employee` as follows:

```

Employee in Class with
  attribute, single
    name: String
  attribute
    aliasname: String
end

```

Here, the following predicate facts would be true:

```

(Employee in Class)
(Employee attribute/name String)
(Employee single/name String)
(Employee attribute/aliasname String)
(Employee attribute String)
(Employee single String)

```

The mechanism for attribution is exactly the same as for instances like `mary`. Note that the 3-argument attribution predicate expressing `(mary name "Mary Smith")` represents a meaningful statement for some reality to be modeled. On the other hand, the predicate fact `(Employee attribute String)` is less significant because the label `attribute` does not transport a specific domain meaning. Here, the 4-argument attribution predicate such as used for the fact `(Employee attribute/name String)` is required. Still, from a formal point of view, there is no different treatment of predicates at the class and instance level. This uniformity is the basis for meta-modeling, i.e. the definition of modeling languages by means of meta classes. The class labels `attribute` and `single` need to be defined at the classes of `Employee`. Those are `Class` and the pre-defined class `Proposition`, to which any object including `Employee` and `mary` is instantiated. In this case, both `attribute` and `single` are defined for `Proposition`:

```

(Proposition attribute/attribute Proposition)
(Proposition attribute/single Proposition)

```

Note that `attribute` has itself as category. This is the most generic attribute category and applies to any (explicit) attribution. For this reason, the category `attribute` can also be omitted in frame definitions of objects. The above definition of `Employee` is equivalent to

```

Employee in Class with
  single
    name: String
  attribute
    aliasname: String
end

```

Both the attribution predicate $(x \text{ } m \text{ } y)$ and its long form $(x \text{ } m/n \text{ } y)$ can be derived, i.e. occur as conclusion of a deductive rule. In such cases, there are no explicit attribute objects between x and y . ConceptBase demands, that in such cases one of the classes of x has an attribute with label m . Deductive rules for $(x \text{ } m/n \text{ } y)$ are introduced with ConceptBase V7.1. They allow to simulate multi-sets, i.e. derived attributes where the same value can occur multiple times. Examples are available in the CB-Forum (<http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/2330042>).

2.2.3 Assigning attribute categories to explicit attributes

The instantiation of an explicit attribute to an attribute category can be explicit (see above), or via inheritance, or via a user-defined rule. Explicit instantiation is typically established when telling a frame like the `Employee` example to the database. *Instantiation by inheritance* is more rarely used but is in fact just the application of the specialization principle to attribution objects:

```
Employee with
  attribute
    salary: Integer
end
Manager isA Employee with
  attribute
    bonus: Integer
end
Manager!bonus isA Employee!salary end
```

Here, the `bonus` attribute is declared as specialization of the `salary` attribute. Any instance of the `bonus` attribute will then be an instance of the `salary` attribute via the usual class membership inheritance of O-Telos. For example,

```
mary in Manager with
  bonus
    bon1: 10000
end
```

shall make the following attribution facts true:

```
(mary bonus/bon1 10000)
(mary bonus 10000), A_e(mary,bonus,10000)
(mary salary/bon1 10000)
(mary salary 10000), A_e(mary,salary,10000)
(mary attribute/bon1 10000)
(mary attribute 10000), A_e(mary,attribute,10000)
```

The third method to instantiate an explicit attribute to an attribute category is via a *user-defined rule*. We use the `employee` example again:

```
Employee in Class with
  attribute
    salary: Integer;
    premium: Integer;
    country: String
  rule
    premrule: $ forall e/Employee prem/Employee!premium
              (e country "NL") and Ai(e,premium,prem)
              ==> (prem in Employee!salary) $
end
```

Now, consider the following instances:

```
marijke in Employee with
  salary sal: 50000
  premium pr: 3000
  country ctr: "NL"
end
```

This makes the following attribution facts true:

```
(marijke salary/sal 50000)
(marijke salary 50000), A_e(marijke,salary,50000)
(marijke premium/pr 3000)
(marijke premium 3000), A_e(marijke,premium,3000)
(marijke salary/pr 3000)
(marijke salary 3000), A_e(marijke,salary,3000)
(marijke country/ctr "NL")
(marijke country "NL"), A_e(marijke,country,"NL")
```

Hence, any explicit premium attribute of an employee of the Netherlands is regarded as an explicit salary as well.

Note that the three cases discussed here are for *explicit* attribution objects. You may also define rules that derive $(x \text{ m } y)$ or $(x \text{ m/n } y)$ directly. In such cases, there is no need for an explicit attribute between x and y . The attribution is completely derived.

2.2.4 Reserved words

In order to avoid ambiguity, neither `in` and `isa` nor the logical connectives `and` and `or` are allowed as attribute labels⁸. Likewise, names of predicates such as `A`, `Ai`, `In` should not be used as object names or variable names. The same holds for the keywords `with` and `end`, which are used in the frame syntax.

2.2.5 Comparison predicates

The next predicates are second class citizens in formulas. In contrast to the above predicates they cannot be assigned to classes of the O-Telos database base. Consequently, they may only be used for testing, i.e. in a legal formula their parameters must be bound by one of the predicates 1 - 8.

1. $(x < y)$, $(x > y)$, $(x \leq y)$, $(x \geq y)$
 x and y may be instances of any class. If they are instance of `Integer` or `Real`, they are ordered numerically. If they are instance of `TransactionTime` they are ordered according to the time they are representing (newer times are greater than older times). Otherwise, they are ordered alphabetically.
2. $(x = y)$
The objects x and y are equal.
3. $(x \neq y)$ or $(x \neq y)$
The objects x and y are not the same.

All comparison predicates may use functional expressions as operands. They are evaluated before the comparison predicates is evaluated. See section ?? for examples. The predicates $(x == y)$, `UNIFIES`(x, y) and `IDENTICAL`(x, y) defined in earlier releases of ConceptBase are deprecated. It is recommended to use $(x = y)$ instead.

2.2.6 Typed variables

The exact syntax of CBL is given in appendix ??. The types of variables (i.e. quantified identifiers) are interpreted as instantiations:

• forall $x/C \text{ F}$ $\rightarrow$ forall $x (x \text{ in } C) \Rightarrow F$

⁸For the example of subsection ?? among others the ground predicates `(mary in Manager)`, `(Manager isA Employee)` and `(mary earns 15000)` are valid facts describing the contents of the database. We suggest to choose verbs (e.g. `earns` in our example) for attribute labels to get more natural and readable predicates.

- $\text{exists } x/C \ F \rightarrow \text{exists } x \ (x \text{ in } C) \text{ and } F$

The class C attached to variable x is called the *variable range*. The anonymous variable range VAR is treated as follows.

- $\text{forall } x/VAR \ F \rightarrow \text{forall } x \ F$
- $\text{exists } x/VAR \ F \rightarrow \text{exists } x \ F$

Anonymous variable ranges are only permitted in meta formulas, see section ??.

2.2.7 Semantic restrictions on formulas

We demand that each variable is quantified exactly once inside a formula. This is no real restriction: in case of double quantification rename one of the variables. More important is a restriction similar to static type checking in programming languages that demands a strong relationship between formulas and the knowledge base:

Predicate typing condition

- (1) Each constant (= arguments that are not variables) in a formula F must be the name of an existing object in the O-Telos database, or it is a constant of the builtin classes Integer, Real, or String.
- (2) For each attribution predicate $(x \ m \ y)$ (or $Ai(x, m, o)$, resp.) occurring in a formula there must be a unique attribute labelled m of some class c of x in the knowledge base, the so-called *concerned class*.
- (3) For each instantiation predicate $(x \text{ in } c)$, the argument c must be a constant.

All instantiation and attribution predicates need to be "typed" according to the predicate typing condition. Formally, we don't assign types to such predicates but *concerned classes*. Any instantiation predicate and any attribution predicate in a formula must have a unique concerned class. It is determined as follows:

- The concerned class of an instantiation predicate $(x \text{ in } c)$ is the class c . The argument c may not be a variable.
- The concerned class of attribution predicates $(x \ m \ y)$ and $Ai(x, m, o)$ is principally the most special attribute with label m of all classes of x ⁹. The O-Telos axioms listed in appendix ??, in particular axiom 17, make sure that there may not be more than one candidate attribute if x is the name of an existing object. If x is a variable, we demand that there is at most one candidate in the variable range of x and its superclasses. If no class of x (i.e. also no superclass of any class of x) defines such an attribute and the CBserver has been started with the predicate typing mode 'extended', then the concerned class is determined from the subclasses of the classes of x . Theoretically, one can choose the common superclass of all such attributes of subclasses of the classes of x (if existent). However, ConceptBase currently demands that there must be a single such attribute in the subclass hierarchy.

Example: The concerned class of $(e \text{ boss } b)$ in the `SalaryBound` constraint in subsection ?? is the `Employee!boss`. The class of variable e is `Employee`. This is the most special superclass of itself and indeed defines the attribute `Employee!boss`.

The purpose of the predicate typing condition is to allow ConceptBase to compile attribution predicates $(x \ m \ y)$ to an internal form `Adot(cc, x, y)` that replaces the attribute label m by the object identifier `cc` of the concerned class. This enormously speeds up the computation of predicate extensions. A similar effect is applicable to instantiation predicates. Here, the concerned class of $(x \text{ in } c)$ is c . Another effect of the predicate typing condition is that certain semantically meaningless predicate occurrences are

⁹Since any object is an instance of `Proposition`, ConceptBase will include this class when searching the concerned class of an attribution predicate.

detected at compile time. For example, $(x \text{ m } y)$ can only have a non-empty extension, if some class of x defines an attribute with label m .

If the argument x in a predicate $(x \text{ m } y)$ is a variable, then the initial class of x is determined by the the variable range in the formula. The variable `this` of query class constraints can have multiple initial classes, being the set of superclasses of the corresponding query class. All superclasses of c are also regarded as classes of x . If x is a constant, then the classes of x are determined by a query to the database. A formula violating the first clause of the predicate typing condition would make a statement about something that is not part of the database. As an example, consider the following formula:

```
forall x/Emplye not (x boss Mary)
```

With the example database of section ??, we find two errors: There are no objects with names `Emplye` and `Mary`.

There are two possible cases to violate the second part of the restriction. The first case is illustrated by an example:

```
forall x/Proposition y/Integer (x salary y) ==> (y < 10000)
```

In this case the classes of x , `Proposition` and any of its superclasses, have no attribute labelled `salary`. Therefore, the predicate $(x \text{ salary } y)$ cannot be assigned to an attribute of the database. Instead, one has to specify

```
forall x/Employee y/Integer (x salary y) ==> (y < 10000)
```

or

```
forall x/Manager y/Integer (x salary y) ==> (y < 10000)
```

depending on whether the formula applies to managers or to all employees.

The second clause of the predicate typing condition is closely related to multiple generalization/instantiation. Suppose, we add new classes `Shop`, `Guest` and `GuestEmployee` to the given class `Employee`:

```
Shop in Class
end
```

```
Guest in Class with
    attribute
        dept: Shop
end
```

```
GuestEmployee in Class isA Guest,Employee
end
```

The following formula refers to objects of class `GuestEmployee` and their `dept` attribute. The problem is that two different attributes, `Employee!dept` and `Guest!dept`, apply as candidates for the predicate $(x \text{ dept } PR)$:

```
forall x/GuestEmployee (x dept PR) ==> not (x in Manager)
```

In order to solve this ambiguity, we demand that in such cases a common subclass exists that defines an attribute `dept` which conforms to both definitions, e.g.

```
Shop in Class
end
```

```
GuestEmployee with
    attribute
        dept: ShopDepartment
end
```

```
ShopDepartment in Class isA Shop,Department
end
```

The third clause of the predicate typing condition is forbidding instantiation predicates with a variable in the class position. The restriction is a pre-condition for an efficient implementation of the incremental

formula evaluator of ConceptBase. Without a constant in the class position of $(x \text{ in } c)$ any update of the instances of any class matches the predicate. Hence, ConceptBase would need to re-evaluate the formula that contains the predicate. Since any update (TELL,UNTELL,RETELL) is containing instantiation facts, any formula with an unrestricted predicate $(x \text{ in } c)$ has to be re-evaluated for any update. This inefficiency can be avoided by demanding that the class position is a constant. A relaxation to this clause (and clause 2) is discussed in sub-section ??.

When compiling the frames, ConceptBase will make sure that the attribute `GuestEmployee!dept` is specializing the two `dept` attributes of `Shop` and `Department`. As a consequence, the attribution predicate $(x \text{ dept } PR)$ can be uniquely attached to its so-called *concerned class* `GuestEmployee!dept`.

The predicate typing condition holds for all formulas, regardless whether they occur as constraints or rules of classes or within query classes¹⁰.

2.2.8 Rule and constraint syntax

A legal *integrity constraint* is a CBL formula that fulfills predicate typing condition. A legal *deductive rule* is a CBL formula fulfilling the same condition and having the format:

```
forall x1/c1 ... xn/cn R ==> lit(a1,...,am)
where
```

- `lit` is a predicate of type 1 or 3, and
- the variables in `a1, ..., am` are contained in `x1, ..., xn`

In O-Telos, rules and constraints are defined as attributes of classes. Use the category `constraint` for integrity constraints, and the category `rule` for deductive rules. The text of the formula has to be enclosed by the character '\$'. The choice of the class for a rule or constraint is arbitrary (except for query classes which use the special variable 'this').

Continuing our running example, the following formula is a deductive rule that defines the `boss` of an `Employee`. Note that the variables `e, m` are both `forall`-quantified.

```
Employee with
  rule
    BossRule : $ forall e/Employee m/Manager d/Department
               (e dept d) and (d head m)
               ==> (e boss m) $
  constraint
    SalaryBound : $ forall e/Employee b/Manager x,y/Integer
                  (e boss b) and (e salary x) and (b salary y)
                  ==> (x <= y) $
end
```

The second formula is an integrity constraint that uses the `boss` attribute defined by the above rule. The constraint demands a `salary` of an `Employee` does not exceed the `salary` of his `boss`. Note that you can define multiple salaries for a given instance of `Employee`. The constraint is on each individual `salary`, not on the sum¹¹! Also note that the arguments of the `<=` predicate are bound by the two predicates with attribute label `salary`.

¹⁰The enforcement of the restriction has been extended to query classes as of ConceptBase release 6.1. To support applications that were written for earlier releases, a CBL option `-cc` (predicate typing) has been introduced to disable the check for query classes. Details are in section ???. With ConceptBase 7.2 (March 2010), the predicate typing has been further extended and now will scan subclasses of the classes of `x` in attribution predicates $(x \text{ m } y)$ in case that superclasses do not provide a matching attribute class. You need to set the CBL option `-cc` to 'extended' to activate this behavior. The extended mode creates more cases that the predicate type (=concerned class) is found. It should be noted that objects like `x` that are not instance of a class that defines and attribute with label `m` will lead to a failure of the predicate $(x \text{ m } y)$, i.e. its negation is then true.

¹¹Use multi-sets as discussed in <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/2330042> if you want to constrain the sum of salaries.

2.2.9 Meta formulas

Some formulas violating the predicate typing condition can be re-written to a set of formulas that do not violate the condition. The so-called *meta formulas* are a prominent category of such formulas. They have occurrences of predicates with so-called *meta variables*. There are two cases. First, an instantiation predicate $(x \text{ in } c), : (x \text{ in } c) :$ or $\text{In_s}(x, c)$ where the class argument c is a variable. Second, an attribution predicate $(x \text{ m } y)$ or $: (x \text{ m } y) :$ where the label argument m is a variable. In such cases, the concerned class cannot be determined directly even though the formula as such is meaningful. ConceptBase relies on predicate typing for the sake of efficiency and static stratification. The concerned class is internally used as predicate name. This increases the selectivity and reduces the chance on non-stratified deduction rules. Fortunately, all meta formulas can be re-written to formulas fulfilling the predicate typing condition. The re-writing replaces the meta variables by all possible value. Since all variables are bound to finite classes, the re-writing yields a finite set of formulas. However, if a meta variable is bound to a class with a large extension, the re-writing will also yield a large set of generated formulas.

Meta formulas allow to specify assertions involving objects from different levels and hence significantly improve flexibility of O-Telos models. An example for the usage of meta formulas can be found in the appendix ?? where the enforcement of constraints in ER diagrams is solved in an elegant way.

As instructional example, assume we want to define that a certain attribute category M is transitive, i.e. if $(x \text{ M } y)$ and $(y \text{ M } z)$, then $(x \text{ M } z)$ shall hold. Many attribute categories are supposed to be transitive, for example the `ancestor` relation of persons, or the `connection` relation between cities in a railway network.

The following meta formula defines transitivity once and forever:

```

Proposition in Class with
  attribute
    transitive: Proposition
  rule
    trans_R:
    $ forall x,y,z,M/VAR
      AC/Proposition!transitive
      C/Proposition
      P(AC,C,M,C) and (x in C) and
        (y in C) and (z in C) and
        (x M y) and (y M z) ==> (x M z) $
  end
end

```

The rule is a meta formula because C and M are meta variables. In this case, one can re-write the formula by replacing all possible fillers for AC , i.e. by the instances of `Proposition!transitive`. A filler for AC will determine fillers for C and M since the first argument of a proposition $P(AC, C, M, C)$ is identifying the proposition.

As a consequence, one can define the ancestor relation to be transitive by simply telling

```

Person in Proposition with
  transitive
    ancestor: Person
end

```

ConceptBase will match the attribute `Person!ancestor` with the variable AC in the above meta formula. This yields $P(\text{Person!ancestor}, \text{Person}, \text{ancestor}, \text{Person})$, which binds the meta variable C to `Person` and M to `ancestor`. The resulting generated formula is:

```

forall x,y,z/VAR (x in Person) and (y in Person) and (z in Person)
  and (x ancestor y) and (y ancestor z)
  ==> (x ancestor z)

```

which can be shortened to

```
forall x,y,z/Person (x ancestor y) and (y ancestor z)
==> (x ancestor z)
```

The technique to generate such 'shortened' formulas is called *partial evaluation*. Its input are facts like (Person!ancestor in Proposition!transitive) and the output are formulas that specialize the original meta formulas for the case of the input facts.

The above formula is fulfilling the O-Telos predicate typing condition. Likewise, the connection relation of cities gets transitive via:

```
City in Proposition with
  transitive
  connection: City
end
```

The advantage of meta formulas is that they save coding effort by re-using them in different modelling contexts. If a meta formula is linked to an attribute category (like `transitive` in the example above, then the semantic of several such attribute category can be combined by just specifying that a certain attribute has multiple categories. Assume for example that we have defined acyclicity with a similar meta formula:

```
Proposition in Class with
  attribute
    acyclic: Proposition
  constraint
    acyclic_IC:
      $ forall x,y,M/VAR
        AC/Proposition!acyclic
        C/Proposition
        P (AC,C,M,C) and (x in C) and
          (y in C) and
          (x M y) ==> not (y M x) $
end
```

Then, the `ancestor` attribute can be specified to be both transitive and acyclic by

```
Person in Proposition with
  transitive,acyclic
  ancestor: Person
end
```

The more categories like `transitive` and `acyclic` are defined with meta formulas, the greater is the productivity gain for the modeler. Not only does it save coding effort. It also reduces coding errors since formula specification is a difficult task. Meta formulas are a natural extension to classical metamodeling. They allow to specify the meaning of modeling constructs at the meta class level. The mapping to simple formulas allows an efficient evaluation. It also allows to retrieve the specialized semantics definition of a model (instance of a metamodel) since the generated simple formulas are attached to the constructs of the model (in the example above they are attached to classes `Person` and `City`). The meta formula compiler is fully incremental, i.e. if the database is updated, then the set of generated simple formulas is also updated if necessary. For example, if one removes the category `transitive` from the `connection` attribute of `City`, then the generated simple formula will also be removed.

Meta formulas that contain meta variables under existential quantification cannot be compiled directly, but there is an elegant trick to circumvent this restriction. Consider for example the formula:

```
$ forall x/VAR SC/CLASS spec/ISA_complete
  (spec super SC) and (x in SC) ==>
  exists SUBC/CLASS (spec sub SUBC) and (x in SUBC) $
```

The meta variable SUBC is under an existential quantifier. To circumvent the problem, we write an intermediary rule replacing the predicate (x in SUBC):

```
$ forall x/Proposition spec/ISA SUBC/CLASS
    (spec sub SUBC) and (x in SUBC) ==> (x inSubRel SUBC) $
```

and then re-write the original constraint to

```
$ forall x/VAR SC/CLASS spec/ISA_complete
    (spec super SC) and (x in SC) ==>
    exists SUBC/CLASS (spec sub SUBC) and (x inSubRel SUBC) $
```

So essentially, we pass the meta variable to the condition of the intermediary rule. The attribute inSubRel is just used to be able to specify a dedicated conclusion predicate for the intermediary deductive rule. It is defined as attribute of Proposition. The complete example is at <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/d3070600/mp-ISA-complete.sml.txt>.

Many more re-usable examples for meta formulas are in the ConceptBase-Forum at <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/1042523>.

Tell order for meta formulas

Meta formulas are compiled by ConceptBase using the partial evaluation strategy. Facts matching certain predicates of the meta formula trigger the partial evaluation, which then leads to new formulas that need to be compiled by ConceptBase. In certain cases, the generated formulas are themselves meta formulas, that need to be further partially evaluated when other matching facts are inserted to the system.

The implementation of the incremental compilation suggests that the meta formula and the input facts should not be defined in the same TELL transactions. When creating your models that involve meta formulas, you should put the definition of the meta formula in a different file than the class/object definitions that utilize the meta formulas. It may also be wise to store the definitions in different modules (see section ??). Specifically, the meta formulas should be defined in a super-module of the modules that use these definitions.

2.2.10 Further object references

In addition to the so called *select expressions* !, =>, -> already introduced above for directly referring to attributes, specializations and instantiations as objects, three other basic constructors may be used within frames and assertions.

- $\wedge$ is the counterpart of ! and denotes the target of an attribute instead of the attribute object itself, e.g. `mary $\wedge$ advises` is the same as `PR`.
- The set valued . operator has the commonly used meaning as in paths in object-oriented models and relates an object with the set of all attribute values of a certain category, i.e. `mary.dept` contains both `PR` and `R&D`.
- As . can be understood as the set variant of $\wedge$ but employing the attribute *category* instead of the concrete attribute label, the same holds for | (with respect to !). Thus `mary|dept` is the set of all attributes (as objects) that belong to category `dept` and have source `mary`.

Note, that . and | are only allowed to occur within assertions wherever classes may be interpreted as range restrictions, e.g. in quantifications or at the right hand side of in predicates. The full syntax which allows combinations of all basic constructors can be found in the appendix. For illustration we just give two examples here. The first is an alternative representation for the rule above, the second could be a constraint stating that all bosses of Mary earn exactly 50.000.

1. `forall e/Employee m/Manager (m in e.dept.head) ==> (e boss m).`
2. `forall b/Mary.dept.head (b salary 50000)`

2.2.11 User-definable error messages for integrity constraints

ConceptBase provides a couple of errors messages in case of an integrity violation. These errors messages refer to the logical definition of the constraint and are sometimes hard to read. To provide more readable error messages, one can attach so-called *hints* to constraint definitions. These hints are attached as comments with label *hint* to the attribute that defines the constraint.

Consider the salary bound constraint above. A hint could look like:

```
Employee!SalaryBound with
  comment
    hint: "An employee may not earn more than her/his manager!"
end
```

It is also possible to attach hints to meta-level constraints. In this case, the hint text can refer to the meta-level variables occurring in the meta-level constraint. These variables will be replaced by the correct fillers when the meta-level constraint is utilized in some modeling context.

Assume, for example, we want to have a symmetry category and attach a readable hint to it:

```
Proposition with
  attribute
    symmetric: Proposition
end

RelationSemantics in Class with
  constraint
    symm_IC: $ forall AC/Proposition!symmetric C/Proposition x,y/VAR M/VAR
              P(AC,C,M,C) and (x in C) and (y in C) and
              (x M y) ==> (y M x) $
end

RelationSemantics!symm_IC with
  comment
    hint: "The relation {M} of {C} must be symmetric,
          i.e. (x {M} y) implies (y {M} x)."
```

Note that the references to the *meta variables*¹² M and C are surrounded by curly braces, and that these meta variables are also occurring in the meta-level constraint. Now, use the *symmetric* concept in some modeling context, e.g. to define that the *marriedTo* attribute of *Person* should be symmetric:

```
Person with
  symmetric
    marriedTo: Person
end
```

At this point of time, ConceptBase will find the hint text for the *symmetric* constraint and will adapt it to the context of C=Person and M=marriedTo. When an integrity violation occurs, the *substituted* hint

```
"The relation marriedTo of Person must be symmetric,
i.e. (x marriedTo y) implies (y marriedTo x)."
```

will be presented to the user. An example violation is:

```
bill in Person with
  marriedTo m1: eve
end

eve in Person end
```

¹²Meta variables are those variables that occur in the class position of (x in c), or in the label position of (x m y) or Label(x,m). In the running example, C and M are meta variables.

One can also define a hint for the meta-level constraint that refers only to a (non-empty) subset of the meta variables. If a hint for a meta formula cannot be substituted as shown above, ConceptBase will not issue the hint but rather the text of the generated formula.

Examples of user-defined error messages can be found in the ConceptBase-Forum at <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/1543277>.

2.2.12 Immutable properties

Immutable attributes cannot be changed (retold) once they are defined for their respective source object. For example, the two spouses in a marriage contract cannot be changed once the marriage contract object has been created. Key attributes in the entity relationship model are another example. Once an entity gets its key, the key may never be changed. Of course, the object as a whole can be removed.

ConceptBase provides an attribute category for such objects:

```
Proposition with
  attribute
    immutable: Proposition
end
```

The semantics cannot be expressed by a static integrity constraint but by an active rule that guards the deletion of immutable attributes. See also <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/d3452001/Immutable.sml.txt> in the CB-Forum. The immutable attribute category is pre-defined in ConceptBase, but the active rule implementing its semantics is not. Include it from the CB-Forum and add it to your source models when needed. The definition below shows the use of immutable attributes. The spouses are created when the source object `marriage1` is created. Afterwards, they shall not be updated for the whole time span of `marriage1`.

```
Marriage in Class with
  immutable
    spouse1 : Person;
    spouse2 : Person
end

marriage1 in Marriage with
  spouse1 s1 : mary
  spouse2 s2 : bill
end
```

The immutable attribute category instructs the integrity constraint compiler to prune unnecessary integrity checks on updates of these attributes. This feature can be disabled by setting the CBserver parameter `-o` to a value smaller than 4, see also section ??.

2.3 Query classes

ConceptBase realizes queries as so-called query classes, whose instances fulfill the membership constraint of the query [?]. This section first defines the structural properties of the query language CBQL and then introduces the predicative component. Queries are instances of a system class `QueryClass` which is defined as follows:

```
QueryClass in Class isA Class with
  attribute
    retrieved_attribute: Proposition;
    computed_attribute: Proposition
  single
    constraint: MSFOLquery
end
```

A super classes of query class imposes a range condition of the set of possible instances of the query class: any instance of the query class must be an instance of the superclass. Example: “socially interested” are those managers that are member of a union.

```

Union in Class
end

UnionMember in Class with
    attribute
        union: Union
end

SI_Manager_0 in QueryClass isA Manager, UnionMember
end

QueryClass SI_Manager in QueryClass isA Manager, UnionMember with
    retrieved_attribute
        union: Union;
        salary: Integer
end

```

Super classes themselves may be query classes, which is the first kind of query recombination. The second frame shows the feature of **retrieved attributes** which is similar to projection in relational algebra. Example: one wants to see the name of the union and the salary of socially interested managers. The attributes must be present in one of the super-classes of the query class. In this example, the union attribute is obviously inherited from the class UnionMember and salary is inherited from Manager. CBQL demands that retrieved attributes are necessary: each answer must have at least one value for them. If an object does not have such an attribute then it will not be part of the solution. As usual with attribute inheritance, one may specialize the attribute value class, e.g.

```

Well_off_SI_Manager in QueryClass isA SI_Manager with
    retrieved_attribute
        salary: HighSalary
end

HighSalary in Class isA Integer with
    rule
        highsalaryrule: $ forall m/Integer
            (m >= 60000)
            ==> (m in HighSalary) $
end

```

The new attribute value class HighSalary is a subclass of Integer so that each solution of the restricted query class is also a solution of the more general one. It should also be noted that HighSalary also could have been another query class. This is the second way of query recombination.

Retrieved attributes can well be derived by a deductive rule. In such cases, ConceptBase generates a label for the derived attribute. For example, consider a boss rule discussed earlier for instances of Employee

```

Employee in Class with
    attribute
        salary: Integer;
        boss: Manager
    rule
        BossRule : $ forall e/Employee m/Manager d/Department
            (e dept d) and (d head m)
            ==> (e boss m) $
end

```

The following query shall then return employees together with their salaries and bosses:

```
EmpSalBoss in QueryClass isA Employee with
    retrieved_attribute
        salary: Integer;
        boss: Manager
end
```

The derived boss attributes get a system-generated label in the answer produced by ConceptBase. Note that retrieved attributes are necessary but there may be more than one value per attribute category, e.g. more than one boss.

Retrieved attributes and super-classes already offer a simple way of querying a knowledge base: projection and set intersection. For more expressive queries there is an predicative extension, the so-called **query constraint**. We use the same many-sorted predicative language as in section ?? for deductive rules and integrity constraints and introduce a useful abbreviation:

Let Q be a query class with a constraint F that contains the predefined variable `this`. Then, the query class is essentially an abbreviation for the two deduction rules

```
forall this  $F' \implies Q(\text{this})$ 
forall this  $Q(\text{this}) \implies (\text{this in } Q)$ 
```

The deduction rules are generated by the query compiler and only listed here for discussing the meaning of a query class. The variable `this` stands for any answer object of Q . We call `this` also the answer variable. The sub-formula F' is combined from the query constraint F and the structural properties of the query, in particular the super-classes and the retrieved attributes. Each super-class C of Q contributes a condition $(\text{this in } C)$ to the sub-formula F' . Each retrieved attribute like $a:D$ contributes a condition $((\text{this } a \ v) \text{ and } (v \text{ in } D))$ to F' . Moreover, each retrieved attribute add the new argument v to the predicate Q . The following example shows the translation.

```
QueryClass Well_off_SI_Manager1 isA SI_Manager with
    retrieved_attribute
        union: Union
    constraint
        well_off_rule: $ exists s/HighSalary
                        (this salary s) $
end
```

The generated deduction rules for this query class are:

```
forall this,v
    (this in SI_Manager) and
    (this union v) and (v in Union) and
    (exists s (s in HighSalary) and (this salary s))
==> Well_off_SI_Manager1(this,v)

forall this,v Well_off_SI_Manager1(this,v)
==> (this in Well_off_SI_Manager1)
```

Classes occurring in a query constraint may be query classes themselves, e.g. `HighSalary`. This is the third way of query recombination.

The next feature introduces so-called **computed attributes**, i.e. attributes that are defined for the query class itself but not for its super-classes. The assignment of values for the solution is defined within the query constraint. Like retrieved attributes, computed attributes are included in the answer predicate so that the proper answer can be generated from it.

The following example defines a computed attribute `head_of` that stands for the department a manager is leading. The attribute `head_of` is supposed to be computed by the query. It is not an attribute of `SI_Manager` or its super-classes. We expect that an answer to the query includes the computed attribute. Note that a reference `~head_of` to the computed attribute occurs inside the query constraint.

```

QueryClass Well_off_SI_Manager2 isA SI_Manager with
    retrieved_attribute
        union: Union
    computed_attribute
        head_of: Department
    constraint
        well_off_rule: $ exists s/HighSalary
                        (this salary s) and
                        (~head_of head this) $
end

```

The variable `~head_of` in the constraint is prefixed with `~` to indicate that it is a placeholder for the computed attribute with the same label `head_of`. We recommend to use the prefix to avoid confusion of the placeholder variable in query constraints and corresponding attribute label in the query definitions. Analogously, you can use the prefixed answer variable `~this` instead of the plain version `this`. Concept-Base will accept both the prefixed and the non-prefixed version for the answer variable and the placeholder variable of computed attributes. Non-prefixed placeholders in constraints are replaced internally by the prefixed counterparts.

The generated deduction rules for above query would be:

```

forall this,v1,v2
    (this in SI_Manager) and
    (this union v1) and (v1 in Union) and
    (v2 in Department) and
    (exists s (s in HighSalary) and (this salary s)
             and (v2 head this))
==> Well_off_SI_Manager2(this,v1,v2)

forall this,v1,v2 Well_off_SI_Manager2(this,v1,v2)
==> (this in Well_off_SI_Manager2)

```

Computed attributes are treated differently from retrieved attributes. The retrieved attribute `union` causes the inclusion of the condition `(this union v1) and (v1 in Union)`. The corresponding variable `v1` does not occur in the sub-formula generated for the query class constraint. The computed attribute causes the inclusion of the condition `(v2 in Department)` but typically also occurs in the query constraint. Like retrieved attributes computed attributes are necessary, i.e. any solution of a query with a computed attribute must assign a value for this attribute. There is no limit in the number of retrieved and computed attributes. The more of them are defined for a query class, the more arguments shall the answer predicate have.

Recursion can be introduced to queries by using recursive deductive rules or by referring recursively to query classes. The example asks for all direct or indirect bosses of `bill`:

```

QueryClass BillsMetaBoss isA Manager with
    constraint
        billsBosses:
            $ (bill boss this) or
            exists m/Manager
              (m in BillsMetaBoss) and
              (m boss this)$
end

```

Further examples can be found in the directory

`$CB_HOME/examples/QUERIES.`

Queries are represented as O-Telos classes and consequently they can be stored in the knowledge base for future use. It is a common case that one knows at design time **generic queries** that are executed at run-time with certain parameters. CBQL supports such parameterizable queries:

```

GenericQueryClass isA QueryClass with
    attribute
        parameter: Proposition
end

```

Generic queries are queries of their own right: they can be evaluated. Their speciality is that one can easily derive specialized forms of them by substituting or specializing the parameters. An important property is that each solution of a substituted or specialized form is also a solution of the generic query. This is a consequence of the inheritance scheme. The example shows that parameters can also be retrieved and computed attributes. Note, that variable for the parameter in the constraint is prefixed here with `~`; you may also omit the prefix in the constraint as explained above).

```

What_SI_Manager in GenericQueryClass isA Manager,UnionMember with
    retrieved_attribute,parameter
        salary: HighSalary;
        union: Union
    computed_attribute,parameter
        head_of: Department
    constraint
        well_off_rule: $ (~head_of head this) $
end

```

There are two kinds of specializing generic query classes:

1. Specialization of a parameter `[a:C']`

Example: `What_SI_Manager[salary:TopSalary]`

In this case `TopSalary` must be a subclass of `HighSalary`. The solutions are those managers in `What_SI_Manager` that not only have a high but a top salary.

2. Instantiation of a parameter `[v/a]`

Example: `What_SI_Manager[Research/head_of]`

The variable `head_of` is the replaced by the constant `Research` (which must be an instance of `Department`).

One may also combine several specializations, e.g.

```
What_SI_Manager[salary:TopSalary,Research/head_of].
```

The specialized queries can occur in other queries in any place where ordinary classes can occur, e.g.

```

110000 in Integer end
QueryClass FavoriteDepartment isA Department with
    retrieved_attribute
        head: What_SI_Manager[110000/salary]
end

```

Parameters that don't occur as computed or retrieved attributes are interpreted as existential quantifications if they are not instantiated. Note that parameters need to be known as objects before using them in query calls.

2.3.1 Query definitions versus query calls

Telling a frame that declares an instance of `QueryClass` (as well as its sub-classes `GenericQueryClass`, and `Function`) constitutes the definition of a query. It shall be compiled internally into Datalog code not visible to the user. Once defined, a query can be called simply by referring to its name. Hence, if `Q` is the name of a defined query class, then `Q` is also an admissible query call. It results in the set of all

objects that fulfill the membership constraint of Q . ConceptBase regards these objects as derived instances of the query class Q .

If a query class has parameters, then any of its specialized forms is also an admissible query call. For example, if Q has two parameters p_1, p_2 in its defining frame, then $Q[v_1/p_1, v_2/p_2]$ is the name of a class whose instances is the subset of instances of Q where the parameters p_1 and p_2 are substituted by the values v_1 and v_2 . The substitution yields a simplified membership constraint that precisely defines the extension of $Q[v_1/p_1, v_2/p_2]$.

If a generic query class is called with all parameters substituted by fillers, then one can omit the parameter labels. Assume that the query Q has just the parameters p_1 and p_2 . Then the expression $Q[v_1, v_2]$ is equivalent to $Q[v_1/p_1, v_2/p_2]$. ConceptBase uses the alphabetic order of parameter labels to convert the shortcut form to the full form.

2.3.2 Query classes and deductive integrity checking

ConceptBase regards query classes as ordinary classes with the only exception that class membership cannot be postulated (via a TELL) but is derived via the class membership constraint formulated for the query class. A consequence of this equal treatment is that a constraint formulated for an ordinary class can refer directly or indirectly to a query class, e.g.:

```
Unit in Class with
  attribute
    sub: Unit
end
BaseUnit in QueryClass isA Unit with
  constraint
    c1: $ not exists s/Unit!sub From(s,~this) $
end
SimpleUnit in Class isA Unit with
  constraint
    c: $ forall s/SimpleUnit (s in BaseUnit) $
end
```

Here, the constraint in the class `SimpleUnit` refers to the query class `BaseUnit`.

ConceptBase supports references to query classes without parameters¹³ in ordinary class constraints and rules. A prerequisite is that the query class is an instance of the builtin class `MSFOLrule`. Membership to this builtin class is necessary to store the generated code for an integrity constraint (or a rule that an integrity constraint might depend upon) and to enable the creation of a dependency network between the query class and the integrity constraints. There are two simple methods to achieve membership to `MSFOLrule`.

Method 1: Make sure that any query class is an instance of `MSFOLrule`. This can simply be achieved by telling the following frame prior to your model:

```
QueryClass isA MSFOLrule end
```

Method 2: Decide for each query class individually. You tag only those query classes that are used in rules or constraints. This individual treatment saves some code generation at the expense of being less uniform. Such an individual tagging would look like

```
BaseUnit in QueryClass,MSFOLrule isA Unit with
  constraint
    c1: $ not exists s/Unit!sub From(s,~this) $
end
```

¹³If the CBserver option `-cc` is set to off, we also allow calls to generic query classes in rules and constraints. In such cases, incremental integrity checking will be incomplete and thus potentially wrong. Only experienced users should employ them.

ConceptBase will reject an integrity constraint or rule if it refers to a query class that is not an instance of MSFOLrule.

If a query class is defined as instance of MSFOLrule, then it should not have a meta formula as constraint! This is a technical restriction that can easily be circumvented by using normal deductive rule.

For example, instead of the query class

```
UnitInstance in QueryClass,MSFOLrule isA Proposition with
  constraint
    c1: $ (~this [in] Unit) $
end
```

you should define

```
UnitInstance in Class with
  rule
    r1: $ forall x/VAR (x [in] Unit) ==> (x in UnitInstance) $
end
```

The example uses the macro predicate `(x [in] Unit)` explained earlier in this section. It is equivalent to the sub-formula `exists c (x in c) and (c in Unit)`.

2.3.3 Nested query calls and shortcuts

ConceptBase has capabilities to form nested expressions from generic query classes. The idea is to combine them like nested functional expressions, e.g. $f(g(x), h(y))$. The problem is however that queries stand for predicates and nested query calls are thus formally higher-order logic (predicates occur as arguments of other predicates), and consequently outside Datalog. Still, the feature is so useful that we provide it. A nested query call is like an ordinary parameterized query call except that the parameters can themselves be query calls. For example, `COUNT[What_Sl_Manager[10000/salary]/class]` counts the instances of the parameterized query call `What_Sl_Manager[10000/salary]`. Syntactically, query calls can be arbitrarily deep, e.g.

```
Union[Intersec[EmpMinSal[800/minsal]/X,
              EmpMaxSal[1400/maxsal]/Y]/X,
      Manager/Y]
```

ConceptBase does perform the usual type check on the parameters by analyzing the instantiation of the *core class* of a query call. For example, the core class of `EmpMinSal[800/minsal]` is `EmpMinSal`. Thus, ConceptBase will check whether `EmpMinSal` is an instance of the class expected for the parameter `X`.

Nested query calls are mostly used in combination with functional expressions, i.e. nested query calls where queries are functions (see section ??). A function in ConceptBase is a query class that has at most one answer object for any combination of input parameters. Of particular interest for nested query calls are functions that do not operate on values (such as integers) but rather on classes such as the class of all employees with more than two co-workers. ConceptBase provides a collection of aggregate functions that operate on classes. For example, the `COUNT` function returns the number of instances of a class. The input of such a function can be any nested query call.

```
QueryClass EmployeeWith2RichCoworkers isA Employee with
  constraint
    c2: $ (COUNT[RichCoworker[this/worker]/class] = 2) $
end
```

The outer predicate (here: `COUNT`) is an instance of `Function`, i.e. delivers at most one value for the given argument. It is also possible that both operands of a comparison predicate are nested expressions:

```

QueryClass EmployeeWithMoreRichCoworkersThanWilli isA Employee with
  constraint
    c2: $ (COUNT[RichCoworker[this/worker]/class] >
          COUNT[RichCoworker[Willi/worker]/class]) $
end

```

ConceptBase supports shortcuts for query calls and function calls (see section ??) in case that all parameters of a query (or function) have fillers in the query call. In such cases, one can write $Q[x_1, x_2, \dots]$ instead of $Q[x_1/p_1, x_2/p_2, \dots]$. ConceptBase shall replace the actual parameters x_1, x_2 etc. for the parameter labels p_1, p_2 in the *alphabetic order* of the parameter labels. For example, the expression `RichCoworker[this]` is equivalent to `RichCoworker[this/worker]` since `worker` is the only parameter label of the query. Likewise, `COUNT[c]` is a shortcut for `COUNT[c/class]`. Since `COUNT` is also a function, we support `COUNT(c)` as well to match the usual notation for function expressions. The last query class is thus equivalent to:

```

QueryClass EmployeeWithMoreRichCoworkersThanWilli isA Employee with
  constraint
    c2: $ (COUNT(RichCoworker[this]) > COUNT(RichCoworker[Willi])) $
end

```

Since the `COUNT` function is frequently used, ConceptBase provides the shortcut `#c` for `COUNT(c)`. Consequently, the shortest form of the above query would be:

```

QueryClass EmployeeWithMoreRichCoworkersThanWilli isA Employee with
  constraint
    c2: $ (#RichCoworker[this] > #RichCoworker[Willi]) $
end

```

The shortcut is also applicable to the `Union` example above. The expression below computes the numbers of instances of the set expression.

```
#Union[Intersec[EmpMinSal[800], EmpMaxSal[1400]], Manager]
```

The definitions for `Union` and `Intersec` can be found in the ConceptBase-Forum at <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/896920>.

Besides `COUNT`, ConceptBase supports aggregation function for finding the minimum, maximum and average of a set. Aggregation functions are not limited to numerical domains. For example, one can define a function that returns an arbitrary instance of a class:

```

selectrnd in Function isA Proposition with
  parameter
    class: Proposition
end

```

The membership constraint has to be provided as so-called CBserver plug-in, see chapter ??.

```
selectrnd(RichCoworker[Willi/worker])
```

would then return an arbitrary instance of `RichCoworker[Willi/worker]`. Random functions can be useful in the context of active rules (section ??), e.g. to initiate the firing of a rule with an arbitrary candidate out of the set of candidates. The code for `selectrnd` is accessible via the ConceptBase-Forum at <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/1694234>.

2.3.4 Reified query calls

You might want to memorize certain query calls that you want to call over and over again. ConceptBase provides a built-in class `QueryCall`, which you can instantiate by such query calls as ordinary objects, i.e. *reified query calls*. The following example defines the class `count` as a query call object:

```
COUNT[Class/class] in QueryCall end
```

Of course, you can ask the query call `COUNT[Class/class]` without having told it as an object. Reifying `COUNT[Class/class]` additionally allows you to use it as an attribute of another object, or to browse it with the graph editor. Examples for query calls, in particular for using integer intervals as class attributes, are available in the CB-Forum at <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/2571997>.

2.4 View definitions

The view language of ConceptBase is an extension of the ConceptBase Query Language CBQL. Besides some extensions that allow an easier definition of queries, views can also be nested to express n-ary relationships between objects.

The system class `View` is defined as follows:

```
Class View isA GenericQueryClass with
attribute
    inherited_attribute : Proposition;
    partof : SubView
end
```

Attributes of the category `inherited_attribute` are similar to retrieved attributes of query classes, but they are not necessary for answer objects of the views, i.e. an object is not required to have a filler for the inherited attribute for being in the answer set of the view.

The `partof` attribute allows the definition of complex nested views, i.e. attribute values are not only simple object names, they can also represent complex objects with attributes. The following view retrieves all employees with their departments, and attaches the head attribute to the departments.

```
View EmpDept isA Employee with
retrieved_attribute, partof
    dept : Department with
        retrieved_attribute
            head : Manager
    end
end
```

As the example shows, the definition of a complex view is straightforward: for the “inner” frame the same syntax is used as for the outer frames. The answers of this view are represented in the same way, e.g.

```
John in EmpDept with
dept
    JohnsDept : Production with
        head
            ProdHead : Phil
    end
end
```

```
Max in EmpDept with
dept
    MaxsDept : Research with
        head
            ResHead : Mary
    end
end
```

To make the definition of views easier, we allow some shortcuts in the view definition for the classes of attributes.

For example, if you want all employees who work in the same departments as John, you can use the term `John.dept` instead of `Department`. In general, the term `object.attrcat` refers to the set of attribute values of `object` in the attribute category `attrcat`. This path expressions may be extended to any length, e.g. `John.dept.head` refers to all managers of departments in which John is working.

A second shortcut is the explicit enumeration of allowed attribute values. The following view retrieves all employees, who work in the same department as John, and earn 10000, 15000 or 20000 Euro.

```
View EmpDept2 isA Employee with
retrieved_attribute
  dept : John.dept;
  salary : [10000,15000,20000]
end
```

As mentioned before, “inner” frames use the same syntax as normal frames. You can also specify constraints in inner frames which refer to the object of an outer frame.

```
View EmpDept_likes_head isA Employee with
retrieved_attribute,partof
  dept : Department with
    retrieved_attribute, partof
      head : Manager with
        constraint c : $ A(this,likes,this::dept::head) $
      end
    end
  end
end
```

The rule for using the variable “this” in nested views is, that it always refers the object of the main view, in this case an employee. Objects of the nested views can be referred by `this::label` where `label` is the corresponding attribute name of the nested view. In the example, we want to express that the employees must like their bosses. Because the inner view for managers is already part of the nested view for departments we must use the double colon twice: `this::dept` refers to the departments and `this::dept::head` refers to the managers.

If you reload the definition of a view into the Telos Editor, the complex structure of it is lost. During compilation of the view, the view is translated into several classes and some additional constraints are generated, so the resulting objects might look quite strange if you reload them.

2.5 Functions

Functions are special queries which have mandatory input parameters and return at most one result for a given input. Functions can either be user-defined by a membership constraint like for regular query classes, or they may be implemented by a PROLOG code, which is defined either in the `OB.builtin` file (this file is part of every ConceptBase database) or in a LPI-file (see also section ??).

A couple of aggregation functions are predefined for counting, summing up, and computing the minimum/maximum/average. Furthermore, there are functions for arithmetic and string manipulation. See section ?? for the complete list. Since functions are defined as regular Telos objects, you can load their definition with the Telos editor of the ConceptBase User Interface.

Unlike as for user-defined generic query classes, you have to provide fillers for all parameters of a function. We will refer to any query expression whose outer-most query is a function as a *functional expression*.

The intrinsic property of a function is that it returns *at most one* answer object¹⁴ for a given combination of input parameters. This property allows to form complex functional expressions including arithmetic expressions. Functions are also special query classes, hence you use them wherever a query class is expected. Subsequently, we introduce first how to define and use functions like queries. Then, we define the syntax of functional expressions and the definition of recursive functions such as the computation of the length of the shortest path between two nodes.

¹⁴A function returns the empty set *nil* if it is undefined for the provided input values.

2.5.1 Functions as special queries

Assume that an attribute (either explicit or derived) has at most one filler. For example, a class `Project` may have attributes `budget` and `managedBy` that both are single-valued. A third attribute `projMember` is multi-valued (default for category 'attribute').

```
Project with
  single
    budget: Integer;
    managedBy: Employee
  attribute
    projMember: Employee
end
```

The two functions `getBudget` and `getManager` return the corresponding objects:

```
getBudget in Function isA Integer with
  parameter
    proj: Project
  constraint
    c1: $ (proj budget this) $
end

getManager in Function isA Employee with
  parameter
    proj: Project
  constraint
    c1: $ (proj managedBy this) $
end
```

The two functions share all capabilities of query classes, except that the parameters are required (one cannot call a function without providing fillers for all parameters) and that there is at most one return object per input value.

Function can be called just like queries, for example `getBudget [P1/proj]` shall return the project budget of project P1 (if existent). You can also use the shortcut `getBudget [P1]` like for any other query, and the functional form `getBudget (P1)`. The latter is the preferred form for function calls. Note that ConceptBase adopts the mathematical style for function calls rather than the object-oriented one¹⁵.

If your function has several arguments, then use alphabetically sorted parameter names in the function definition:

```
f in Function isA D with
  parameter
    a1: R1;
    a2: R2
  constraint
    c: $ ... $
end
```

This corresponds to the mathematical function signature

$$f : R_1 \times R_2 \rightarrow D$$

Functions without parameters are also possible, e.g. a function `magicnumber` that returns a fixed number:

¹⁵An object-oriented style for the first function would be `P1.getBudget ()` rather than `getBudget (P1)`.

```

magicnumber in Function isA Integer with
  constraint
    cmagic: $ (this = 42) $
end

```

You can call it with an empty argument list: `magicnumber()`.

2.5.2 Shortcuts for function calls and functional expressions

Since functions require fillers for all parameters, ConceptBase offers also the functional syntax `f(x)` to refer to function calls. The expression `f(x)` is a shortcut for `f[x/param1]`, where `param1` is the only parameter of function `f`. If a function has more than one parameter, then they are replaced according their alphabetic order:

```

g in Function isA T with
  parameter
    x: T1;
    y: T2
  constraint
    ...
end

```

A call like `g(bill,1000)` is a shortcut for `g[bill/x,1000/y]` because `x` is occurs before `y` in the ASCII alphabet. The shortcut can be used to form complex functional expressions such as `f(g(bill,getBudget(P1)))`. There is no limitation in nesting function calls. Function calls are only allowed as left or right side of a comparison operator. They are always evaluated before the comparison operator is evaluated. For example, the equality operator

```
$ ... (f(g(bill,getBudget(P1))) = f(g(mary,1000))) $
```

will be evaluated by evaluating first the inner functions and then the outer functions. Note that the parameters must be compliant with the parameter definitions.

As a special case of functional expressions, ConceptBase supports arithmetic expressions in infix syntax. The operator symbols `+`, `-`, `*`, and `/` are defined both for integer and real values. ConceptBase shall determine the type of a sub-expression to deduce whether to use the real-valued or integer-valued variant of the operation. Examples of admissible arithmetic expressions are

```

a+2*(b-15)
n+f(m)/3

```

Provided that `a` and `b` are variables holding integers, the first expression is equivalent to the function shortcut

```
IPLUS(a, IMULT(2, IMINUS(b, 15)))
```

and to the query call

```
IPLUS[a/i1, IMULT[2/i1, IMINUS[b/i1, 15/i2]/i2]/i2]
```

The second arithmetic expression includes a division which in general results in a real number. Hence, the function shortcut for this expression is

```
PLUS(n, DIV(f(m), 3))
```

The whole expression returns a real number.

2.5.3 Example function calls and definitions

1. The following three variants all count the number of instances of Class:

```
COUNT[Class/class]
COUNT(Class)
#Class
```

The result is an integer number, e.g.

```
119
```

The operator # is a special shortcut for COUNT.

2. The subsequent query sums up the salaries of an Employee:

```
SUM_Attribute[bill/objname, Employee!salary/attrcat]
SUM_Attribute(Employee!salary, bill)
```

Note that the parameter label attrcat is sorted before objname for the function shortcut. The result is returned as a real number, even if the input numbers were integers.

```
2.5001000000000e+04
```

You can also use functions also in query class to assign a value to a “computed_attribute”:

```
QueryClass EmployeesWithSumSalaries isA Employee with
computed_attribute
    sumsalary : Real
constraint
    c: $ (sumsalary = SUM_Attribute(Employee!salary, this)) $
end
```

3. Complex computations can be made by using multiple functions in a row. This query returns the percentage of query classes wrt. the total number of classes.

```
Function PercentageOfQueryClasses isA Real with
constraint
    c: $ exists i1,i2/Integer r/Real
        (i1 = COUNT[QueryClass/class]) and (i2 = COUNT[Class/class]) and
        (r = DIV[i1/r1,i2/r2]) and (this = MULT[100/r1,r/r2]) $
end
```

The query can be simplified with the use of function shortcuts to

```
Function PercentageOfQueryClasses isA Real with
constraint
    c: $ (this = MULT(100, DIV(COUNT(QueryClass),
                                COUNT(Class)))) $
end
```

and with arithmetic expressions to

```
Function PercentageOfQueryClasses isA Real with
constraint
    c: $ (this = 100 * #QueryClass / #Class) $
end
```

The function `PercentageOfQueryClasses` has zero parameters. You can use it as follows in logical expressions

```
$ ... (PercentageOfQueryClasses() > 25.5) ... $
```

So, a function call `F()` calls a function `F` that has no parameter.

Functions that yield a single numerical value can directly be incorporated in comparison predicates. For example, the following query will return all individual objects that have more than two attributes:

```
ObjectWithMoreThanTwoAttributes in QueryClass isA Individual with
  constraint
    c1 : $ (COUNT_Attribute(Proposition!attribute,this) > 2) $
end
```

The functional expression used in the comparison can be nested. See section ?? for details. You can also re-use the above query to form further functional expressions, e.g. for counting the number of objects that have more than two attribute. You find below all three representations for the expression.

```
COUNT[ObjectWithMoreThanTwoAttributes/class]
COUNT(ObjectWithMoreThanTwoAttributes)
#ObjectWithMoreThanTwoAttributes
```

2.5.4 Programmed functions

If your application demands functional expressions beyond the set of predefined-functions, you can extend the capabilities of your ConceptBase installation by adding more functions. There are two ways: first, you can extend the capabilities of a certain ConceptBase database, or secondly, you can add the new functionality to your ConceptBase system files. We will discuss the first option in more details using the function `sin` as an example, and then give some hints on how to achieve the second option.

A function (like any builtin query class) has two aspects. First, the ConceptBase server requires a regular Telos definition of the function declaring its name and parameters. This can look like:

```
sin in Function isA Real with
  parameter
    x : Real
  comment
    c : "computes the trigonometric function sin(x)"
end
```

The super-class `Real` is the range of the function. i.e. any result is declared to be an instance of `Real`. The parameters are listed like for any regular generic query class. The comment is optional. We recommend to use short names to simplify the constructions of functional expressions. The above Telos frame must be permanently stored in any ConceptBase database that is supposed to use the new function.

The second aspect of a function is its implementation. The implementation can be in principal in any programming language but we only support PROLOG because it can be incrementally added to a ConceptBase database. An implementation in another programming language would require a re-compilation of the ConceptBase server source code. The syntax of the PROLOG code must be conformant to the Prolog compiler with which ConceptBase was compiled. This is in all cases SWI-Prolog (www.swi-prolog.org). For our `sin` example, the PROLOG code would look like:

```
compute_sin(_res,_x,_C) :-
    cbserver:arg2val(_x,_v),
    number(_v),
    _vres is sin(_v),
    cbserver:val2arg(_vres,_res).

tell: 'sin in Function isA Real with parameter x : Real end'.
```

The first argument `_res` is reserved for the result. then, for each parameter of the function there are two arguments. The first is for the input parameter (`_x`), the second holds the identifier of the class of the parameter (here: `_C`). It has to be included for technical reasons. The clause 'tell:' instructs ConceptBase to tell the Telos definition when the LPI file is loaded. Instead of this clause you may also tell the frame manually via the ConceptBase user interface.

There are a few ConceptBase procedures in the body of the `compute.sin` that are of importance here. The procedure `cbserver:arg2val` converts the input parameter to a Prolog value. ConceptBase internally always uses object identifiers. They have to be converted to the Prolog representation in order to enter them into some computation. The reverse procedure is `cbserver:val2arg`. It converts a Prolog value (e.g. a number) into an object identifier that represents the value. If necessary, a new object is created for holding the new value.

The above code should be stored in a file like `sin.swi.lpi`. This file has to be copied into the ConceptBase database which holds the Telos definition of `sin`. You will have to restart the ConceptBase server after you have copied the LPI file into the directory of the ConceptBase database.

If you want the new function to be available for all databases you construct, then you have to copy the file `sin.swi.lpi` to the subdirectory `lib/SystemDB` of your ConceptBase installation. Note that your code might be incompatible with future ConceptBase releases. If you think that your code is of general interest, you can share it with other ConceptBase users in the Software section of the CB-Forum (<http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/2768063>).

2.5.5 Recursive function definitions

Some functions like the Fibonacci numbers are defined recursively. ConceptBase supports such recursive definitions. If the function is defined in terms of itself, then express the recursive definition in the membership constraint of the function:

```
fib in Function isA Integer with
  parameter
    n: Integer
  constraint
    cfib: $ (n=0) and (this=0) or
           (n=1) and (this=1) or
           (n>1) and (this=fib(n-1)+fib(n-2))
    $
end
```

The variable `this` stands for an answer object that fulfills the constraint `cfib`. Note that ConceptBase regards integers also as objects. ConceptBase shall internally compile the disjunction into three formulas:

```
forall n,this/Integer (n=0) and (this=0) ==> fib(this)
forall n,this/Integer (n=1) and (this=1) ==> fib(this)
forall n,this/Integer (n>1) and (this = fib(n-1)+fib(n-2)) ==> fib(this)
```

ConceptBase employs a bottom-up query evaluator to evaluate the recursive function. Thus, the result of a function call `fib(n)` shall only be computed once and then re-used in subsequent calls.

If the recursion is not inside a single function definition but rather a property of a set of function/query definitions, then you must use so-called *forward declarations*. They declare the signature of a function/query before it is actually defined. A good example is the computation of the length of the shortest path between two nodes.

```
spSet in GenericQueryClass isA Integer with
  parameter
    x: Node;
    y: Node
end
sp in Function isA Integer with
```

```

parameter
  x: Node;
  y: Node
constraint
  csp: $ (x=y) and (this=0) or
        (x nexttrans y) and (this = MIN(spSet[x,y])+1)
        $
end
spSet in GenericQueryClass isA Integer with
parameter
  x: Node;
  y: Node
constraint
  csps: $ exists x1/Node (x next x1) and (this=sp(x1,y)) $
end

```

Here, the query class `spSet` computes the set of length of shortest path between the successors of a node `x` and a node `y`. The length of the shortest path is then simply 0 if `x=y` or the minimum of the `spSet[x,y]` plus 1 if there is a path from `x` to `y`, and undefined else. The signature of `spSet` must be known for compiling `sp` and vice versa. `ConceptBase` has a single-pass compiler. Hence, it requires the forward declaration. The query `spSet` is not a function because it returns in general several numbers. The complete example for computing the length of the shortest path is in the CB-Forum, see <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/1694234>.

Recursive function definitions require much care. Deductive rules shall always return a result after a finite computation. This does not hold in general for recursive function definitions when they use arithmetic subexpressions. These subexpressions can create new objects (numbers) on the fly and could thus force `ConceptBase` into an infinite computation. On the other hand, they are more expressive than pure deductive rules and thus useful to analyze large models in a quantitative way.

2.6 Query evaluation strategies

`ConceptBase` employs an SLDNF-style query evaluation method, i.e. query predicates are evaluated top-down much like in standard Prolog. This is known to cause infinite loops for certain recursive rule sets. To overcome this, the SLDNF evaluator is augmented by a *tabling sub-system* [?], which detects recursive predicate calls and answers them from the cached results of a query (the so-called table) rather than entering an infinite loop. This tabled evaluation computes the fixpoint (=answer) of a query provided that the overall rule set is stratified. Even more, dynamically stratified rule sets are supported as well. Other than with the static stratification test, a violation is detected at run time of a query rather than at compile time.

For a precise definition of stratification, we refer you to the literature on deductive databases. For the purposes of this manual, consider the following rule:

```

forall p/Position (exists p1/Position (p moveTo p1) and not (p1 in Win))
==> (p in Win)

```

`ConceptBase` internally compiles such rules into a representation where `Position`, `moveTo`, and `Win` are predicate symbols:

```

forall p
  (exists p1 Position(p) and Position(p1) and moveTo(p,p1) and not Win(p1))
==> Win(p)

```

Static stratification requires that one can consistently assign stratification levels (=numbers) to the set of predicate symbols such that

1. If there is a rule with conclusion predicate `A` and positive condition predicate `B` (=not negated), then the level of `A` must be greater or equal the level of `B`.

2. If there is a rule with conclusion predicate A and negated condition predicate B , then the level of A must be strictly greater the level of B .

In the example above, the conclusion predicate `Win` depends on the condition predicate `not Win`. Since we only can assign one level to `Win`, we cannot find a static stratification for the above rule. The same argument also works in case of multiple inter-dependent rules. Static stratification can be tested at compile-time of a rule.

Dynamic stratification is an extension of static stratification, i.e. any statically stratified rule set is also dynamically stratified. It is not only considering predicate symbols but also the arguments with which a predicate is called at run-time. Obviously, this depends on the database state at a certain point of time. The global rule of dynamic stratification is that the answer to a predicate call $A(x)$ may not depend on its negation $\text{not } A(x)$. Such a clash can be detected by maintaining a stack of active predicate calls.

ConceptBase reports a violation of dynamic stratification in the log window of the CB client with a message indicating the predicate that participated in the stratification violation. There dynamic stratification test of ConceptBase catches different cases that result in slightly different error messages. Essentially, they all are reduced to the pattern that P and $\neg P$ cannot be true at the same time for a predicate that is part of a recursive chain of calls.

In practice, most rule sets are already statically stratified, i.e. no violation can occur regardless of the data in the database. Counter examples are in the CB-Forum (see <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/888832>) in the models `Russel.sml` and `Win.sml`. These examples are neither statically nor dynamically stratified. Note also the example `WinNim.sml` which uses the same query as `Win.sml` but is dynamically stratified. Even in the case of stratification violations, ConceptBase will display an answer to a query. The user can then decide which parts of the answer are usable. The stratification test can be enabled or disabled for the ConceptBase server via the parameter `-st` (section ??).

2.7 Multi-level modeling with ConceptBase

Multi-level modeling is about the use of multiple abstraction levels (classes, metaclasses, metametaclasses, etc.) to define the types of a database. Traditionally, there were only two abstraction levels (database schema= $M1$ and database instance= $M0$) but even very early proposals like Abrial's binary data model had already three abstraction levels (adding metaclass schema= $M2$).

In multi-level modeling, the metaclasses are regarded as first-class citizens of the database, i.e. they are objects too. This coincides with the Telos data model, which regards any explicit information as object. A feature of multi-level modeling that is not part of the Telos data model is the so-called "potency" of attributes and relation. The potency roughly specifies how many times the source (or target) of an attribute/relation must be instantiated in order to form a factual instance of the attribute/relation.

Assume that there is a metaclass `Product`, which defines an attribute `serialnumber` with potency 2. Assume further that the class `Car` is an instance of `Product` and `myCar` is an instance of `Car`. Then, `myCar` can form a factual attribute for `serialnumber`, since it is 2 instantiation levels below `Product`. While this form of attribution is not part of Telos, one can well axiomatize potencies using rules and constraints in ConceptBase [?]. The example specifications are also available at <http://conceptbase.cc/ddi/>.

As stated in [?], ConceptBase was one of the earliest systems supporting a simple form of multi-level modeling. Before the dual deep instantiation formalization [?], the materialization construct, a precursor of multi-level modeling, was also formalized with ConceptBase by Dahchour et al. [?].

2.7.1 Expressing semantics at the metamodel level

ConceptBase does not have builtin support for explicit potencies of attributes (and relations). Instead, all attributes have source and target potency 1: when instantiating the attribute, one has to instantiate both the source and the target of the attribute, yielding a new attribute with its own label. However, one can use the "macro" predicates $(x \text{ [in] } c)$ and $(x \text{ [m] } y)$ to specify the semantics of modeling language constructs such as the key property for attributes of entity types.

Consider the following definitions about the ERD modeling language:

```
Entity in Class end
Domain in Class end

EntityType in Class with
  attribute field: Domain
  single key: Domain
  rule
    r1: $ forall e/VAR (e [in] EntityType) ==> (e in Entity) $
  end

Value in Class end
Domain in Class with
  rule
    r1: $ forall v/VAR (v [in] Domain) ==> (v in Value) $
  end
```

Figure ?? visualizes the classification of entities and values. The class `Entity` subsume all instances of instances of `EntityType` and `Value` subsumes all instances of instances of `Domain`. The directed links with broken lines are instantiations.




Figure 2.2: Classifying entities and values

The rules for classifying entities and values exploit a fundamental principle of ConceptBase/Telos: all explicit facts in the database are uniformly represented as propositions $P(id, x, n, y)$, and consequently they have a system-generated object identifier, serving as a persistent memory address. This principle allows deriving predicates like $(e \text{ in } Entity)$. The variable 'e' is standing for the object identifier of the respective object. This stands in contrast to the relational data model, in which tuples are identified by the key attributes defined for their respective relation. The case of `Domain` is even more interesting. In ConceptBase, all values like numbers and strings are also objects with system-generated identifiers. Hence, they can be classified just like entities are classified. In relational databases (and most other data models), there is a strict separation of objects and values. In ConceptBase, all stored information has object identifiers, i.e. an address where it can be looked up and linked. The dichotomy of values and identifiers in object-oriented (programming) languages is not present in ConceptBase¹⁶.

In the OMG terminology [?], the constructs `EntityType` and `Domain` are at the M2 level (meta classes). The predicate $(e \text{ [in] } EntityType)$ is equivalent to

```
exists ET/VAR (e in ET) and (ET in EntityType)
```

Thus, rule r1 of `Entity` designates it as a simple class, or a class technically at the OMG M1 level. However, it is indeed a construct of the ERD modeling notation since certain properties of ERDs such as the key property require to refer to objects at the data level (OMG level M0):

¹⁶There is however a different dichotomy in ConceptBase that resembles the object-vs-value separation: Each proposition $P(id, x, n, y)$ has an identifier 'id' and a label 'n'. A number like '120' is for example represented by a proposition $P(id123, id123, 120, id123)$. The label is always the representation of the object as used in the universe of discourse. The identifier is used to look it up and lets it participate in formulas.

```

EntityWithSharedKey in QueryClass isA Entity with
  computed_attribute
    entity2: Entity;
    keyvalue: Value
  constraint
    cshare: $ (~this [key] ~keyvalue) and
              (~entity2 [key] ~keyvalue) and (~entity2 \= ~this) $
end

```

The above query shall return all those entities that share the key with another entity. The result objects are all at the data level, however, we do not refer to any specific entity types here. This is an example of defining a construct like 'key' at the meta class level and using it to query objects at the data level. The 'key' attribute is available to all ERD diagrams. It shows that three OMG levels (M2, M1, M0) are considered at the same time to define and evaluate the semantics of a modeling language. From an ontological point of view, the concept `Entity` is interesting. It subsumes all entities that are instances of some entity type (M1 level). Thus, it is the super class of all possible entity types. Telos represents all factual information as propositions with object identity. This allows us to store facts at any abstraction level (objects, classes, meta classes, and even objects not belonging to any of the classical OMG levels) in the same uniform data model.

The class `Entity` is semantically an M1-level object, because its instances (a.k.a. its extension) are M0-level objects representing some objects of a reality. It is the most general entity type. Figure ?? contrasts the deductive definition of `Entity` (left side) with the definition using sub-class relations (right side). The deductive rule defining the instances of `Entity` is equivalent to placing `Entity` as superclass of all entity types, i.e. all instances of `EntityType`. While `Entity` is semantically an M1-level object, it is not part of the domain model (here the domain of employees and projects). It is rather part of the definition of the ERD language, enriching it with the semantics of its constructs.




Figure 2.3: Entity classification vs. sub-classing

The pattern of entity classification can also be applied to attributes and relations. For example, the metaclass `EntityType` has an attribute `key` for qualifying attributes that identify entities. On the M1 level, this becomes an attribute `identifier` of `Entity`. The classifying rule is

```
forall kv/VAR (kv [in] EntityType!key) ==> (kv in Entity!identifier)
```

Figure ?? shows an example that leads to the classification of key attributes. Note that this allows to retrieve the identifiers of any entity, regardless of the schema definition at the M1 level. The complete example is available from the CB-Form at <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/d2230805/MacroFormulas.sml.txt>. The complete example also introduces the class `Relationship` as M1 counterpart of the M2 metaclass `RelationshipType` in the same style as `Entity` and `Value` as counterparts of `EntityType` and `Domain`. All three classes are connected (role and value links) and thus allow to query the M0 data level without referring to any schema class like `Employee`. It thus supports schema-less querying.




Figure 2.4: Classifying key attributes

2.7.2 DeepTelos

DeepTelos is a set of three axioms to realize multi-level modeling in harmony with the existing O-Telos axioms on instantiation, specialization, attribution (see appendix ??). The idea is use a new predicate

(m IN c)

to declare m as most general instance [?] of the class c. Any ordinary instance of c then becomes a sub-class of m. The three rules are defined as follows:

```

Proposition with
  attribute
    ISA: Proposition; IN: Proposition
end
DeepTelosRules in Class with
  rule
    mrule1: $ forall m,x,c/Proposition (x in c) and (m IN c) ==> (x ISA m) $;
    mrule2: $ forall x,c,d/Proposition (c ISA d) and (x in c) ==> (x in d) $;
    mrule3: $ forall c,d,m,n/Proposition
              (m IN c) and (n IN d) and (c ISA d) ==> (m ISA n) $
end

```

Further discussion and examples of DeepTelos are available at <http://conceptbase.cc/deeptelos>.

2.7.3 Crossing abstraction levels

OMG demands strict separation between models at different abstraction levels. The only allowed relation is the instantiation between objects at two neighboring levels. For example, the object Employee at the M1 (model) level can be declared as instance of the object EntityType at the M2 (model) level. Other relations are forbidden. ConceptBase/O-Telos does not have such restrictions. It has so-called omega classes (Proposition is the most important such class), which have objects from any abstraction level as instance. Even more, there is no builtin notion of abstraction level in O-Telos. Abstraction levels introduce a form of rigor into metamodeling that is beneficial to avoid semantic confusions, e.g. to avoid instantiating real-world objects into meta-classes. One can enforce such a rigor in ConceptBase by defining level objects like Token, SimpleClass, MetaClass, MetametaClass and so forth and then enforcing constraints that only allow instantiation between neighbor levels. This was in fact discussed with the original Telos specification but abandoned with O-Telos.

The following example shows that there are applications where level-crossing relations are useful. The principle idea is that there is a modeling level that describes the reality) and a parallel level of the creation process of the models. Each construct is man-made, regardless of the abstraction level. Hence, it makes perfect sense to specify who has created a given construct. Such information is very common in software engineering, where the updates to the code base are associated to members of the development team. Consider the Figure ?? about an excerpt of the ERD language and its history (see also <http://conceptbase.sourceforge.net/meta-modeling.html>).

Concept with

```

    attribute created: "1-Jan-2004: 12:03"    { * from M3 to M0 *}
end

EntityType in Concept with
    attribute attr : Domain
end

Person in EntityType with
    attr name: String
end

PC in Person with
    name pname: "Peter Chen"
    attribute proposed: EntityType { * from M0 to M2 *}
end

```

There are two relations that cross abstraction levels. First, the relation `created` of `Concept` (M2 level) points to a time object at M0 level. Second, the relation `proposed` links the object `PC` to the object `EntityType`. Both relations use the builtin attribute category attribute of the omega class `Proposition`.




Figure 2.5: Links crossing abstraction levels

To state that Peter Chen proposed the ERD language is a different type of statement than the statement that Peter Chen is a person, which is a type of entity. Still, both types of statements can co-exist and allow for more expressive metamodels. For example, one can define that certain constructs of a modeling language are only available to experienced modelers, or that development projects from a certain domain should use a subset of the available constructs.

2.8 Datalog queries and rules

The definition of queries in `ConceptBase` is often complicated by the limitations of the expressiveness of the query language or by the limitations of the query optimizer to find the best solution. The concept of datalog queries and rules was introduced to overcome these limitations. Datalog queries and rules give the

experienced user the possibility to define the executable code of query (or rule) directly, including the use of standard PROLOG predicates such as `ground/1` to improve the performance of a query or rule.

Although datalog queries and rules can be used as any other query (or rule), they can cause an inconsistent database. This is due to the fact, that the datalog queries and rules will usually not be evaluated while the semantic integrity of the database is checked.

2.8.1 Extended query model

Datalog queries are defined in a similar way as standard query classes. They must be declared as instance of the class `DatalogQueryClass`.

```
Class DatalogQueryClass isA GenericQueryClass with
attribute
    code : String
end
```

The attribute `code` defines the executable code of the query as string.

Datalog rules have to be defined as an instance of `DatalogInRule` or `DatalogAttrRule`, depending on whether their conclusion should be an In-predicate or an A-predicate.

```
Class DatalogRule with
attribute
    concernedClass : Proposition;
    code : String
end

Class DatalogInRule isA DatalogRule
end

Class DatalogAttrRule isA DatalogRule
end
```

The attribute `concernedClass` specifies the class for the In-predicate or the attribute class for the A-predicate.

2.8.2 Datalog code

The datalog code is a list of predicates, separated by commas (,). As in Datalog or Prolog, this will be interpreted as a conjunction of the predicates. To use disjunction, the `code` attribute has to be specified multiple times.

All predicates that may be used in standard rules and queries may also be used in datalog queries (see section ?? for a list). An argument of a predicate may be one the following:

- an object name: If the object name starts with an upper case letter or includes special characters such as `!` or `"`, it must be written in single quotes (`'`). This also holds for string object, e.g. `"a string"` must be written as `'a string'`.
- a predefined variable, defined by the context of the query or rule: If the query has a super class, the variable `this` refers to the instances of this class. If the query has parameters or computed attributes, variables with the names of the parameters or computed attributes will be predefined. In a `DatalogInRule`, the variable `this` refers to the instances of the concerned class. In a `DatalogAttrRule`, the variables `src` and `dst` refer to the source and destination object of the attribute. Note, that **all predefined variables have to be prefixed with `~` and must be encoded in single quotes (`'`)**, e.g. `'~this'`, `'~src'`, `'~param'`.

- existential variables: These variables must be declared in a special predicate `vars(list)`, which has to be the first predicate of the code. For example, the predicate `vars([x,y])` defines the variables `x` and `y`.

A query expression of the form `query(q)` may be also used as predicate, or as second argument of an In-predicate. `q` may be any valid query expression, e.g. just the name of a query class, or a derive expression including the specification of parameters (for example, `find_instances[Class/class]`).

In addition, PROLOG predicates can be used as predicates. You can define your own PROLOG predicates in a LPI-file (see section ?? for an example).

2.8.3 Examples

This section defines a few datalog queries and rules for the standard example model of Employees, Departments and Managers (see `$CB_HOME/examples/QUERIES`).

The first example defines a more efficient version of the recursive `MetabossQuery`.

```
DatalogQueryClass MetabossDatalogQuery isA Manager with
  parameter
    e : Employee
  code
    r1 : "In('~e', 'Employee'), A('~e', boss, '~this') ";
    r2 : "vars([m]),
        In('~e', 'Employee'),
        In(m, query('MetabossDatalogQuery[~e/e]')),
        A(m, boss, '~this') "
end
```

Note that the disjunction of the original query is represented by two code-attributes. The example shows also the use of query expressions and existential variables.

The second example is the datalog version of the rule for the `HighSalary` class. The infix-predicate `>=` is represented by the predicate `GE`.

```
Class HighSalary2 isA Integer
end

DatalogInRule HighSalaryRule2 with
  concernedClass
    cc: HighSalary2
  code
    c: "In('~this', 'Integer'), GE('~this', 60000) "
end
```

The last example shows the definition of a rule for an attribute. It also shows, how the performance of a rule can be improved by specifying different variants for different binding patterns. The example defines two rules, depending on the binding of the variable `src`. The rule defines the transitive closure of the `boss` attribute. Rule `r1` is applied, if both arguments `src` and `dst` are unbound. The second rule is used, if at least `src` is bound, and the last rule will be applied, if we have a binding for `dst` but not for `src`.

```
DatalogAttrRule MetabossRule with
  concernedClass
    cc : Employee!boss
  code
    r1 : "vars([m]), var('~src'), var('~dst'), In(m, 'Manager'),
        A(m, boss, '~dst'), A('~src', boss, m) ";
    r2 : "vars([m]), ground('~src'),
```

```

        A('~src',boss,m),A(m,boss,'~dst')";
r3 : "vars([m]),ground('~dst'),var('~src'),
      A(m,boss,'~dst'),A('~src',boss,m)"
end

```

Note, that the predicates `var` and `ground` are builtin predicates of PROLOG. Thus, this is also an example for calling PROLOG predicates in a query or rule.

Chapter 3

Answer Formats for Queries

The ConceptBase server provides an ASK command in its interface, which allows to specify in which text-based format the answer should be returned. There are two pre-defined formats: one for returning a list of object names, and one for returning a list of object frames. These two formats can be extended by user-defined answer formats.

Examples for answer formats are available from the CB-Forum, see <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/861803>.

To understand ConceptBase answer formats, we first look at the syntax of the ASK command at the programming interface. We use here the syntax of CBshell (section ??) since it can be tested directly in a command terminal:

```
ask "<querydefinition>" FRAMES <answer-format> <roll-backtime>
ask <querycall> OBJNAMES <answer-format> <roll-backtime>
```

The first variant is for queries where the query is provided as a Telos frame, the second one is for parameterized query calls like `Q[abc/param1]`. If the query call contains special characters or is computes of several query calls then surround them by double quotes. The rollback time is typically `Now`, i.e. the query refers to the current database state. There are four options for the answer format:

LABEL: the answer shall be returned as a comma-separated list of object names, e.g. `bill, mary, john`

FRAME: the answer shall be returned as a list of Telos frames.

JSONIC: the answer shall be returned as a list of JSON-like frames (experimental).

answer-format-name: Finally, the name of a user-defined answer format can be provided, like `MyFormat`. This will override any answer format that is assigned via the `forQuery` attribute for the given query. Hence, one can maintain several answer formats for the same query.

default: This value is leaving the choice to the ConceptBase server. For function calls, the answer format `LABEL` is selected, for other query calls, ConceptBase first checks if an answer format was defined for the query (attribute `forQuery`). If that exists, it is selected, otherwise the format `FRAME` is selected. If there are more than one, the first one is selected.

The answer format `JSONIC` is an alternative to `FRAME`. Here is an example contrasting the two formats. The first two entries are in `FRAME` format, the latter two in `JSONIC`.

```
Employee in Class isA Person,Agent with
  attribute
    salary: Integer;
    name: String
```

```

    rule
      r1: $ forall e/Employee exists s/Integer (e salary s) $
    end

bill in Employee with
  salary
    bsal: 1000
  name
    firstname: "William";
    lastname: "Smith"
  attribute
    creator: mjeu
end

{ "id" : "Employee",
  "type" : ["Class"],
  "super" : ["Person", "Agent"],
  "salary" : "Integer",
  "rule/r1" : "$ forall e/Employee exists s/Integer (e salary s) $"
}

{ "id" : "bill",
  "type" : ["Employee"],
  "salary/bsal" : "1000",
  "name/firstname" : "\"William\"",
  "name/lastname" : "\"Smith\"",
  "attribute/creator": "mjeu"
}

```

The command `ask <querycall>` is a shortcut for

```
ask <querycall> OBJNAMES LABEL Now
```

Assume, we want to execute a query call `Q[abc/param1]` on the current database and have the answer formatted according as frames, we would issue:

```
ask Q[abc/param1] OBJNAMES FRAME Now
```

We assume here that there is no user-defined answer format for query `Q`. If there is a user-defined answer-format like `MyFormatA` and `MyFormatB`, but the answer formats are not assigned via `forQuery` to `Q`, we can call:

```
ask Q[abc/param1] OBJNAMES MyFormatA Now
```

and also

```
ask Q[abc/param1] OBJNAMES MyFormatB Now
```

If we assign `MyFormatB` to `Q` via the `forQuery` attribute, the last call is equivalent to the first call using default as answer format name.

Subsequently, the specification of such user-defined answer formats is presented.

3.1 Basic definitions

By default, ConceptBase displays answers to queries in the FRAME format (see 'A' and 'B' below). For many applications, other answer representations are more useful. For example, relational data is more

readable in a table structure. Another important example are XML data. If ConceptBase is integrated into a Web-based information system, then answers in HTML format are quite useful. For this reason, answer format definitions are provided.

Answer formats in ConceptBase are based on term substitutions where terms are evaluated against substitution rules defined by the answers to a query. A substitution rule has the form $L \rightarrow R$ with the intended meaning that a substring L in a string is replaced by the substring R . The object of a term substitution is a string in which terms may occur, for example:

this is a string called {a} with a term {b}

Assume the substitution rules:

- {a} $\rightarrow$ string no. {x}
- {b} $\rightarrow$ that was subject to substitution
- {x} $\rightarrow$ 123

The derivation of a string with terms proceeds from left to right. First, the term occurrence {a} is dealt with. The next term in the string is then {x} which is evaluated to 123. Finally, {b} is substituted and the result string is this is a string called string no. 123 with a term that was subject to substitution.

We denote a single derivation step of a string S_1 to a string S_2 by $S_1 \Rightarrow S_2$. It is defined when there occurs a substring L in S_1 , i.e. $S_1 = V + L + W$ and a substitution rule $L \rightarrow R$ and $S_2 = V + R + W$. The substrings V and W may be empty. A string S is called ground when no substitution rule can be applied. A sequence $S \Rightarrow S_1 \Rightarrow \dots \Rightarrow S_n$ is called a derivation of S . A complete derivation of S ends with a ground string. In our example, the complete derivation is:

this is a string called {a} with a term {b}.
 $\Rightarrow$ this is a string called string no. {x} with a term {b}.
 $\Rightarrow$ this is a string called string no. 123 with a term that was
 subject to substitution.

An exception to the left-to-right rule are complex terms like $\{do(\{y\})\}$. Here, the inner term $\{y\}$ is first evaluated (e.g. to 20) and then the result $\{do(20)\}$ is evaluated.

In general, term substitution can result in infinite loops. This looping can be prevented either by restricting the structure of the substitution rule or by terminating the substitution process after a finite number of steps. The end result of a substitution process of a string is called its derivation. In ConceptBase, the substitution rules are guaranteeing termination except for the case of external procedures. The problem with the exception is solved by prohibiting cyclic calls of the same external procedure during the substitution of a call. A cyclic call is a call that has the same function name (e.w. query class) and the same arguments (expressed as parameter substitutions).

In ConceptBase, an answer format is an instance of the new pre-defined class 'AnswerFormat'.

```
Individual AnswerFormat in Class with
attribute
  forQuery : QueryClass;
  order : Order;
  orderBy : String;
  head : String;
  pattern : String;
  tail : String;
  fileType: String
end
```

The first attribute assigns an answer format to a query class (a query may have at most one answer format). The second and third attribute specify the sorting order of the answer, i.e. one can specify by which field an answer is sorted, and whether the answer objects are sorted 'ascending' or 'descending'

(much like in SQL). The 'orderBy' attribute specifies the property by which the answer shall be sorted. The most common value is the expression "this", i.e. sort the answer by the name of the objects in the answer. You can also specify an attribute expression such as "this.name" referring to an answer variable. If you specify "none" for 'orderBy', then the answer is not sorted. If the number of objects in an answer exceeds 5000, then no sorting is applied due to memory limitations.

The 'head', 'pattern', and 'tail' arguments are strings that define the substring substitution when the answer is formatted. They contain substrings of type `expr` that are replaced. The head and tail strings are evaluated once, independent from the answer to the query. Usually, they do not contain expressions but only text. The response to a query is a set of answer objects A1,A2,... The pattern string is evaluated against each answer object. For each answer object, the derivation of the pattern is put into the answer text. Hence, the complete answer consists of

- derivation of head string
- + derivation of pattern string for answer object A1
- + derivation of pattern string for answer object A2
- ...
- + derivation of tail string

The `fileType` attribute is explained in section ??.

In the next sections, we will explain more details about answer formats using the following example: An answer object A to a query class QC has by default a 'frame' structure

```
A in QC with
  cat1
    label11: v11;
    label12: v12
    [...]
  cat2
    label21: v21
    [...]
end
```

In case of a complex view definition VC, the values v_{ij} can be answer objects themselves, e.g.

```
B in VC with
  cat1
    label11: v11;
    label12: v12
    [...]
  cat2
    label21: v21 with
      cat21
        label211: v211
        [...]
      end
    [...]
end
```

3.2 Constructs in answer formats

3.2.1 Simple expressions in patterns

We first concentrate on the pattern attribute, i.e. they are not applicable in the head or tail attribute of an answer format. The pattern of an answer format is applied to *each* answer object of a given query call, effectively transforming it according to the pattern. The following expressions are allowed. Capital letters in the list below indicate that the term is a placeholder for some label occurring in the query definition or the answer objects.

{this} denotes the object name of an answer object (e.g. 'A'). The syntax for object names is defined in section ???. In particular, attribution objects have names like `mary!earns`.

{this.ATTRIBUTE.CAT} denotes the value(s) of the attribute 'ATTRIBUTE.CAT' of the current answer object 'this'. The attribute must be defined in the query class (`retrieved_attribute`, `computed_attribute`) or view (`inherited_attribute`). Note that some attributes are multi-valued. For the answer object 'A', `{this.cat1}` evaluates to `v11, v12` and `{this.cat2}` evaluates to `v21`. For the complex object 'B', a path expression like `{this.cat2.cat21}` is allowed and yields `v21`. Note that all such expressions are set-valued.

{this^ATTRIBUTE.LABEL} denotes the value of the attribute 'ATTRIBUTE.LABEL' of the current answer object 'this'. The attribute_label is at the level of the answer object, i.e. not at the class level but at the instance level. Therefore, this expression is rarely used because the instance level attribute labels are usually unknown at query definition time. Example A and B: `{this^label112}` evaluates to `v12`.

{this|ATTRIBUTE.CAT} denotes all attribute labels of the answer object 'this' that are grouped under the category 'ATTRIBUTE.CAT' (defined in the query class). Example: `{this|cat1}` evaluates to `label111, label112`.

The derivation of pattern is performed for each answer object that is in the result set of a query. The answer object 'A' induces the following substitution rules:

```

{this}      → A
{this.cat1} → v11, v12
{this.cat2} → v21
{this^label11} → v11
{this^label12} → v12
{this^label21} → v21
{this|cat1}  → label11, label12
{this|cat2}  → label21

```

Extended examples for these simple expressions are given in `simple_answerformats1.sml` and `simple_answerformats2.sml` in the directory `$CB_HOME/examples/AnswerFormat/`.

Note that only objects that match the query constraint are in the answer. Particularly, the categories `computed_attribute` and `retrieved_attribute` require that at least one filler for the respective attribute is present in an answer object! Use ConceptBase views ('View') and `inherited_attribute` in case that zero fillers are also allowed for answers.

There are few other simple expressions that may be useful. They just list attributes without having to refer to specific attributes.

{this.attrCategory} denotes all attribute categories that are present in an answer object. Example: For answer object 'A' `{this.attrCategory}` evaluates to `cat1, cat2`.

{this|attribute} lists all attribute labels occurring in an answer object. Example: for answer object 'A', `{this|attribute}` evaluates to `label111, label112, label21`.

{this.attribute} lists all attribute values occurring in an answer object. Example: for answer object 'A', `{this.attribute}` evaluates to `v11, v12, v21`.

{this.oid} displays the internal object identifier of the answer object {this}.

For the answer object A, these expressions induce the following additional substitution rules:

```

{this.attrCategory} → cat1, cat2
{this|attribute}    → label111, label112, label21
{this.attribute}    → v11, v12, v21

```

3.2.2 Pre-defined variables

The following variables can be used in the head, tail, and pattern of an answer format. They do not refer to the variable *this*.

- {user}** outputs the user name of the current transaction; typically has the structure name@address
- {transactiontime}** the time when the current transaction was started; has format YYYY-MM-DD hh:mm:ss and is based on Coordinated Universal Time (UTC), formerly known as Greenwich Mean Time
- {cb.version}** the version number of ConceptBase
- {cb.date.of.release}** the version number of ConceptBase
- {currentmodule}** is expanded to the name of the current module.
- {currentpath}** is expanded to the complete module path (starting with root module `System`) that was active when starting the transaction
- {database}** is expanded to relative or absolute path of the database that was specified with the `-d` option of the CBserver (section ??); if no database was specified, the variable is expanded to `<none>`

ConceptBase also adds those command line parameters that deviate from their defaults to the set of pre-defined variables. The most common ones are (see also section ??):

- {loadDir}** : directory from which the CBserver loads Telos source files at start-up; command-line parameter `-load`
- {saveDir}** : directory into which the CBserver saves Telos sources of modules at shut-down or client logout; command-line parameter `-save`
- {viewDir}** : directory into which the CBserver materializes results of certain queries command-line parameter `-views`

Moreover, if the current transaction was a call of a query like `MyQuery[v1/param1, ...]` then `{param1}` will be evaluated to `v1`. This makes all parameter substitutions of a query call available to answer formatting.

3.2.3 Iterations over expressions

In case of expressions with multiple values, the user may want to generate a complex text that uses one value after the other as a parameter. This is in particular useful to transform multiple attribute values like `this.cat1`. The 'Foreach' construct has the format:

```
{Foreach( (expr1,expr2,...), (x1,x2,...), expr )}
```

The expression `expr1` is evaluated yielding each a list of solutions `s11, s12, ...`. The same is applied to `expr2` yielding a list `s21, s22, ...`. Then, the variables `x1, x2, ...` are matched against the first entries of all lists, i.e. `x1=s11, x2=s21, ...`. This binding is then applied to the expression `expr` which should contain occurrences of `{x1}`, `{x2}`, ... This replacement is continued with the second entries in all lists yielding bindings `x1=s12, x2=s22, ...`. This is continued until all elements of all lists are iterated. If some lists are smaller than others, the missing entries are replaced by `NULL`.

During each iteration, the new bindings induce substitution rules for the binding

Iteration 1:

`{x1}` $\longrightarrow$ `s11`

`{x2}` $\longrightarrow$ `s21`

...

Iteration 2:

`{x1}` $\longrightarrow$ `s12`

`{x2}` $\longrightarrow$ `s22`

...

Note that the third argument `expr` may contain other subexpressions, even a nested ‘Foreach’. An example for iterations is given in `$CB_HOME/examples/AnswerFormat/iterations.sml`.

The Foreach construct contains three arguments separated by commas. These two commas are used by ConceptBase to parse the arguments of the Foreach-construct and similar answer formatting expressions. They are not printed to the answer stream.

3.2.4 Special characters

If one wants a comma inside an expression that shall be visible in the answer, then one has to escape it ‘\,’. The same holds for the other special characters like ‘(’, ‘)’ etc. Here is a short list of supported special characters. Some require a double backslash.

```

\\n : new line (ASCII character 10)
\\t : tab character
\\b : backspace character
\\0 : empty string (no character)
\\( : left parenthesis
\\) : right parenthesis
\\, : comma

```

In principal, any non-alphanumerical character like ‘(’, ‘)’’, ‘{’, ‘}’, ‘[’, ‘]’ can be referred to by the backslash operator. Note that the vanilla versions of these characters are used to denote expressions in answer formats. Hence, we need to ‘escape’ them by the backslash if they shall appear in the answer.

3.2.5 Function patterns

The substitution mechanism for answer formats recognizes patterns such as

```
{F(expr1,expr2,...)}
```

as function calls. An example is the `ASKquery` construct from section ?? . The mechanism is however very general and can be used to realize almost arbitrary substitutions. The parentheses and the commas separating the arguments of `F` are parsed by ConceptBase and not placed on the output. The following simple function patterns are pre-defined:

- `{QT(expr)}` puts the `expr` into double quotes if not already quoted.
- `{UQ(expr)}` removes double quotes from `expr` if present.
- `{ALPHANUM(expr)}` outputs an alphanumeric transcription of `expr`. This is useful when an object names contains special characters but the output format requires an alphanumeric label. If `expr` already evaluates to an alphanumeric label, then it is inserted unchanged.
- `{From(expr)}` inserts the source object of `expr`. The source object is computed by the predicate `From(x, o)` of section ?? . This pattern is useful for printing attributes. ConceptBase will first apply pattern substitution to `expr` and then to the whole term `{From(expr)}`.
- `{To(expr)}` inserts the destination object of `expr`. The destination object is computed by the predicate `To(x, o)` of section ?? .
- `{Label(expr)}` inserts the label of the object referenced by `expr`. The label is computed by the predicate `Label(x, l)` of section ?? .
- `{Oid(expr)}` inserts the identifier of the object referenced by `expr`.

Note that the argument `expr` can be another pattern such as `{this}`. Specifically, the expression `{Oid({this})}` is equivalent to `{this.oid}`. However, the `Oid` pattern is also applicable to patterns not including `{this}`.

Examples are available from the CB-Forum, see <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/2502545>.

The above set of patterns can be extended by user-defined functions via LPI plugins. In principle, any routine that can be called from the ConceptBase server, can also be called in an answer format. The programming interface is not documented here since this requires extensive knowledge of the ConceptBase server source code. For the experienced user, we provide an example in the subdirectory `examples/AnswerFormat` of the ConceptBase installation directory, see files `externalcall.sml` and `externalcall.swi.lpi` (`externalcall.bim.lpi` for BIM-Prolog variant). Note that one has to create a persistent database, load the model `externalcall.sml` into it, then terminate the ConceptBase server, and then copy the file `externalcall.swi.lpi` or `externalcall.bim.lpi` into the database directory (see also appendix ??). Thereafter, restart ConceptBase and call the query `EmpDept`.

3.2.6 Calling queries in answer formats

A query call within an answer format is an example of a so-called external procedure. The pattern as well as head and tail of an answer format may contain the call to a query (possibly the same for which the answer format was defined for). This allows to generate arbitrarily complex answers.

```
{ASKquery(Q[subst1,subst2,...],formatname)}
```

The argument `Q` is the name of a query. The arguments `subst1`, `subst2` are parameter substitutions (see section ??). The argument `formatname` specifies the answer format for the query call `Q[subst1,subst2,...]`. The answer format may also have parameters (see below). If you use `default` as answer format name, then the ConceptBase server will pick the default format. This is `LABEL` (list of object names) for function calls. For other queries, it is the first answer format that list the query `Q` in its `forQuery` attribute. If no such answer format exists, the format `FRAME` (Telos frames) is chosen.

The effect of `ASKquery` in an answer format is that the above query call is evaluated and the `ASKquery` expression is replaced by the complete answer to the query. In terms of the substitution, the following rule is applied:

```
{ASKquery(Q[subst1,subst2,...],default)} → X
```

where `X` is the result of the query call `Q[subst1,subst2,...]` after derivation of the arguments `subst1`, `subst2` etc. This sequencing is important since an `ASKquery` call can contain terms that are subject to substitution, e.g.

```
{ASKquery(MyQuery[{this.cat1}/param1,{this.{x1}}/{x2}],default)}.
```

ConceptBase will always start to evaluate left to right and the innermost terms before evaluating the terms that contain inner terms. Hence, the derivation sequence is

```
{ASKquery(MyQuery[{this.cat1}/param1,{this.{x1}}/{x2}],default)}
⇒ {ASKquery(MyQuery[alpha/param1,{this.{x1}}/{x2}],default)}      (r2)
⇒ {ASKquery(MyQuery[alpha/param1,{this.name}/{x2}],default)}      (r1)
⇒ {ASKquery(MyQuery[alpha/param1,"smith"/{x2}],default)}          (r3)
⇒ {ASKquery(MyQuery[alpha/param1,"smith"/param2],default)}        (r4)
⇒ The answer is ...                                                (r5)
```

where we assume the following example substitution rules

```
r1: {x1} → name
r2: {this.cat1} → alpha
r3: {this.name} → "smith"
r4: {x2} → param2
r5: {ASKquery(MyQuery[...],default)} → The answer is ...
```

This guarantees that the query call is ‘ground’, i.e. does not contains terms which are subject to substitution.

The `ASKquery` construct allows to introduce recursive calls during the derivation of a query since there can (and should) be an answer format for `Q` which may contain expressions `ASKquery` itself. In principle, this allows infinite looping. However, the answer format evaluator prevents such loops by halting the expansion when a recursive call with same parameters has occurred. The answer then contains an ⁴⁷

character at the position where the loop was detected. Additionally, an error message is written on the console window of the ConceptBase server (tracemode must be at least low).

A simple example for use of `ASKquery` is given in `recursive-answers.sml`. The example uses a view instead of a query class in order to include also answers into the solution which not have a filler for the requested attribute, i.e. `hasChild` is `inherited_attribute`, not `retrieved_attribute`.

It is common practice to combine the `ASKquery` construct with 'Foreach' in order to display an iteration of objects in the same way. The user should define an answer format for the iterated query `Q` as well.

Do not mix the use of `ASKquery` with the view definitions in ConceptBase! The nesting depth of a view is determined by the view definition. The nesting depth of an answer generated by expansion of `ASKquery` is only limited by the complexity of the database. For example, one can set up an ancestor database and display all descendants of a person and recursively their descendants in a single answer string for that person. The nested `ASKquery` inside an answer format usually results in the unfolding also using an answer format (possibly the same as used for the original query). This feature allows the user to specify very complex structured answers that might even contain the complete database. In particular, complex XML representations can be constructed in this way.

3.2.7 Expressions in head and tail

The features 'head' and 'tail' are similar to pattern. The difference is that any expression using 'this' (the running variable for answer objects) is disallowed. This only leaves function patterns such as `ASKquery` expressions and pre-defined patterns such as `{user}`. Of course, the head and tail strings can contain multiple occurrences of `ASKquery` or other function patterns.

3.2.8 Conditional expressions

Conditional expressions allow to expand a substring based on the evaluation of a condition. The syntax is:

```
{ IFTHENELSE (predicate, thenstring, elsestring) }
```

The 'predicate' can be one of

- `{ GREATER (expr1, expr2) }`
- `{ LOWER (expr1, expr2) }`
- `{ EQUAL (expr1, expr2) }`
- `{ AND (expr1, expr2) }`
- `{ OR (expr1, expr2) }`
- `{ ISFIRSTFRAME () }`
- `{ ISLASTFRAME () }`

Example: `{ GREATER ({this.salary}, 10000) }`. Note that the arguments may also contain expressions. The predicate `{ ISFIRSTFRAME () }` is true, when ConceptBase starts with processing answer frames. It is false, when the first frame has been processed. The predicate `{ ISLASTFRAME () }` is true when ConceptBase starts with processing the last answer frame for a given query. Otherwise, it is false. An example for conditional expressions is provided in the CB-Forum, see file "csv.sml" in <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/861803>. In most cases, the `IFTHENELSE` construct can be avoided by a more elegant query class formulation.

3.2.9 Views and path expressions

If the answer format is defined for a complex view, then path expressions like `this.cat2.cat21...` for the parts of the complex answer can be defined. An example for use of answer formats for views is given in `views.sml`.

The reader should note that complex path expressions can only refer to components that were defined as retrieved, computed, or inherited attributes in the view definition. For example, one cannot refer to `this.dept.budget` in the example view `EmpDept` in `views.sml` since it is not a retrieved attribute of the `dept` component of the view definition `EmpDept`. The second expression of the answer format `EmpDeptFormat` uses the builtin procedure `UQ`. It removes the quotes `'` from a string. Analogously, a procedure `QT` can be used to put quotes around a term.

3.3 Parameterized answer formats

The general way to use an answer format for a query is to define the attribute `forQuery`. Another possibility is to specify the answer format for a query is to use the answer representation field of the `ASK` method in the `IPC` interface.

The following code is an example for specifying a user-defined answer in the `ASK` method. This example is written in Java and uses the standard Java API of `ConceptBase` (see the Programmer's Manual for details).

```
import i5.cb.api.*;

public class CBAnswerFormat {

    public static void main(String[] argv) throws Exception {

        CBclient cb=new CBclient("localhost",4001,null,null);

        CBAnswer ans=cb.ask("find_specializations[Class/class,TRUE/ded]",
                           "OBJNAMES","AFParameter[bla/somevar]","Now");
        System.out.println(ans.getResult());
        cb.cancelMe();
    }
}
```

In the example, a connection is made to a `ConceptBase` server on `localhost` listening on port 4001. The `ask`-method of the `CBclient` class sends a query to the server. The first argument is the query, the second argument is the format of the query (in this example, it is just one object name), the third argument is the answer representation, and the last argument is the rollback time.

The third argument, is the the answer representation. There are four predefined answer representations. `FRAME` returns the answers as Telos frames, including retrieved and computed attributes. `LABEL` returns only the names of the answer objects as a comma-separated list. Thirdly, the format `JSONIC` returns the answer in JSON-like frames. Finally, `default` lets the `CBserver` choose between `LABEL` (for function calls), the explicit answer format assigned for a query (attribute `forQuery`), and `FRAME` (otherwise). Besides these pre-defined answer representations, one can specify user-defined answer formats. This is also the preferred way. In our case, it is a parameterized answer format: `AFParameter[somevalue/somevar]`. This means that the result of the query will be formatted according to the answer format `AFParameter` and the variable `somevar` will be replaced with `somevalue`. The variable can be used like any other expression, i.e. it must be enclosed in `{ }`.

The following definition of `AFParameter` is an example, how the parameter can be used in the pattern. If the parameter is not specified, the string `{somevar}` will not be replaced.

```
Individual AFParameter in AnswerFormat with
    head hd : "<result>"
    tail tl : "</result>"
```

```

pattern
  p : "
<object>
  <type>{somevar}</type>
  <name>{this}</name>
</object>"
end

```

Note, that you can use any answer format (with or without parameters) as answer representation in the ASK method.

3.4 File type of answer formats

The optional `fileType` attribute of answer formats is used by the server-side materialization of query results (section ??). ConceptBase will use the specified file type when storing the query results in the file system. The default value is "txt". The attribute is single-valued though single-valuedness is not enforced.

3.5 Bulk query calls

It is sometimes useful to call the same query class with multiple arguments in a single call rather than in a sequence of calls. Each individual call from a ConceptBase client to the server comes with a certain latency time. Thus, if one would have to call the same query for dozens of arguments in a sequence, most of the answer time would actually be the latency time.

To address this problem, the CBserver offers a query call pattern for bulk queries:

```
bulk[q, x1, x2, x3, ...]
```

The query `q` stands for a query class with a single parameter. The bulk query is converted by the CBserver into the following sequence of query calls:

```
q[x1], q[x2], q[x3], ...
```

The answers to the query call are collected into a single answer set and then transformed with the answer format of the query call. Hence, the answer format is applied to the whole answer and not to the part answers.

Example:

```
ask bulk[Q, abc, def] OBJNAMES MyFormatA Now
```

Sorting of answers is disabled for bulk queries in order to return the answers in the sequence indicated by the arguments of the bulk query call. Here, the answer to the argument `abc` shall precede the answer to argument `def`. Arguments that do not reference an existing object are removed from the argument list by the CBserver before answering the query. Bulk queries are only supported for generic query classes with a single parameter. The query class may not be a builtin query class.

The main purpose of bulk queries is to speed up the interaction between the ConceptBase clients, such as CBGraph, and the CBserver. You can however use them with the CBShell.

Chapter 4

Active Rules

Active rules specify certain actions that have to be executed if an event occurs and a condition is fulfilled at this time. Because active rules consists of an **Event**, a **Condition** and an **Action**, they are abbreviated as *ECARule*.

Events (ON-part) of ECARules are insertions and deletions of objects (Tell/Untell) and queries (Ask). The events are detected during the processing of the input frames for a Tell or Untell or Ask operation. For example, if you tell 2 frames at a time, and the first frame matches an event for an ECARule, then the ECARule is executed before the second frame is processed. You can also control the sequencing of the firing of ECA rules by the so-called ECAMode and by priority orderings in the set of defined ECARules.

The condition of the ECARule (IF-part) is a logical expression over the database. It will bind free variables occurring in the condition (if any) and these bindings together with the bindings of the event are passed to the action part of the ECARule.

The action (DO-part) of an ECARule is evaluated for each evaluation of the IF-part that has is true in the database. The elements in a DO-part can be Tell, Untell, Retell, Ask and Call actions. Call actions can call any Prolog predicate, for example a Prolog predicate defined as a CBserver plug-in (see appendix ??). Retell actions are combining an Untell and a Tell, in particular for assigning a new attribute value, e.g. `Retell((e salary newsalary))`. Optionally, one can specify an ELSE-part which consists of actions that are executed when the IF-part is not satisfiable for any binding of the free variables. The Ask action is only useful for ECARules that have an ELSE-part. In this case, the Ask will retrieve information from the database that cannot be retrieved within the IF-part. The special action 'reject' will abort the complete transaction that directly or indirectly triggered the ECARule.

The effect of ECARules is subject to the regular integrity checking of ConceptBase. If an integrity violation is detected, then the whole transaction including all updates by ECARules is rolled back. The integrity test is started after all enabled ECARules have fired.

4.1 Definition of ECARules

In ConceptBase, active rules are defined similar to query classes. The user has to create an instance of the builtin class `ECARule`. The following frame shows the Telos definition of the class `ECARule`.

```
Class ECARule with
attribute
    ecarule : ECAAssertion;
    priority_after: ECARule;
    priority_before : ECARule;
    mode : ECAMode;
    active : Boolean;
    depth : Integer;
    rejectMsg : String
constraint
```

```

    { ... }
end

```

A correct ECARule must specify at least the attribute `ecarule`, the other attributes are optional. The language for ECAassertions is a extension of the assertion language, it is specified as text between `$` signs in the same way as rules and constraints.

4.1.1 ECAassertion

An ECAassertion has the following structure (the syntax is described in section ??).

```

$ x1,x2/C1 y1/C2 ...
ON [TRANSACTIONAL] event [FOR x]
[IF|IFNEW] condition
DO action1, action2 ...
[ELSE action3]
$

```

The TRANSACTIONAL modifier cannot be used in combination with the FOR clause. The ELSE-part is optional. The first line contains the declaration of all variables used in the ECAassertion. The specified classes of the variables (here: C1 and C2) are only used for compilation of the rule, during the evaluation of the rule it is *not* tested if the variables are instances of the specified classes. If necessary, include predicates like `In(z, Class)` in the IF-part of the ECARule. The variables can occur in the ON-, IF-, DO-, and ELSE-part of the ECARule. When the IF-part begins with `IFNEW`, then the whole condition shall be evaluated against the newest database state. See also 'new' tag for conditions.

There is also a variant of ECARules without an IF-part. It is equivalent to the longer form on the right side. The shortcut can also be used in combination with the TRANSACTIONAL modifier or the FOR clause.

<pre> \$ x1,x2/C1 y1/C2 ... ON event DO action1, action2 ... \$ </pre>	<pre> \$ x1,x2/C1 y1/C2 ... ON event IF TRUE DO action1, action2 ... \$ </pre>
--	--

4.1.2 Events

Possible events are the insertion (Tell) or deletion (Untell) of attributes (A), instantiation links (In), or specialization links (Isa). For example, if the rule should be executed if an object is inserted as instance of Class, then the event statement is: `Tell In(x, Class)`. Furthermore, an event may be a query, e.g. if you specify the event `Ask find_instances[Class/class]` the ECA rule is executed before the evaluation of the query `find_instances` with the parameter `Class`. Potential updates to the database caused by the ECA rules will be persistent, i.e. in such cases an Ask can well update the database. It is possible to use variables as placeholders for parameters in the Ask event clause.

The event detection algorithm takes only extensional events into account. Events that can be deduced by a rule or a query are not detected. However, the algorithm is aware of the predefined Telos axioms, e.g. if an object is declared as an instance of a class, the object is also an instance of the super classes of this class.

4.1.3 Conditions

The condition (IF-part) of an ECARule consists of predicates combined by the logical operators 'and', 'or', and 'not'. Quantified sub-expressions (`forall`, `exists`) are not allowed. You can however use query classes to encode such sub-expressions¹. The arguments of the predicates are either bound by the ON-part of the ECA rule or they are free variables. When an event occurs that fires an ECARule, then the condition is

¹An example for a quantified sub-expressions is `not exists y/D (x m y)`.

evaluated against the database yielding bindings for the free variables. Each such binding will be passed to the action part of the ECArule. Note that ECArules without any free variable are also possible. By default, predicates are evaluated against the old state of the object base (i.e., before the transaction started). If a predicate has to be evaluated on the new database state, i.e. the intermediate state representing the updates processed so far during the transaction, then it has to be quoted by the backward apostrophe, for example `'(x in Class)` instead of `(x in Class)`. The syntax `new((x in Class))` is supported as well and equivalent to the use of the backward apostrophe. Note that only conditions of ECA rules can see intermediate database states. If the whole condition shall be evaluated against the new database state, the use the clause `IFNEW` instead `IF`.

4.1.4 Actions

Actions are specified in a comma-separated list. The syntax is similar to that one of events, except that you can also ask queries (`Ask`) and call Prolog predicates (`Call`). The standards actions are as follows:

Tell *predicate* : The predicate fact is told to the system. The predicate must be either an attribution, instantiation, or specialization predicate. All arguments of the predicate must be bound at execution time of the action.

Untell *predicate* : The predicate fact is untold from the system.

Retell *predicate* : Old facts are first untold and then the new predicate fact is told. See below for restrictions.

Ask *predicate* : The predicate is evaluated, possibly binding free variables. Note that this action can fail if the predicate is not true in the database.

Call *proc* : The Prolog predicate matching *procedure* is called. This can bind free variables. You can define your own Prolog code as CBserver plugin. You may also use `CALL` as keyword (deprecated).

Raise *query* : The query call *query* is raised as event. The query is not evaluated. This action can be used to trigger other ECArules that have a matching Ask event. All parameters should be bound. The *query* may not be a builtin query class or function.

noop : This action stands for "no operation", i.e. nothing is done. It can be useful for certain constructions with empty DO parts.

reject : The current transaction is aborted and the database state is rolled back.

Instead of the capitalized action names `Tell` you can also use the small caps variant `tell`. The same rule is applicable for `Untell`, `Retell`, `Ask`, `Raise`, and `Call`. Analogously, the event statements in the ON-part of an ECArule can also use the small caps variants.

All variables in `Tell`, `Untell` and `Retell`² actions must be bound. The insertion of an attribute `A(x, ml, y)` is only done, if there is no attribute of category `ml` with value `y` for object `x`. Then, a new attribute with a system-generated label is created. If an attribute `A(x, ml, y)` should be deleted, then all attributes of category `ml` with value `y` for object `x` are deleted³. If the argument of `Retell` is a fact `AL(x, ml, n, y)`, then there will be at most one stored astribute with category `ml` and label `n` that has to be deleted before the new fact is told. It is well possible that an attribute is updated (untell+tell) several times by action parts of ECArules during a single ConceptBase transaction. Only the state of the attribute after all ECArule firings will be visible after a successful commit.

There are a few special procedures, which may be called within the DO- or ELSE-part of an ECArule:

²Retell is currently restricted to attribution predicates `A(x, ml, y)` and `AL(x, ml, n, y)`, respectively their infix versions `(x ml y)` and `(x ml/n y)`. It will replace the old value of attribute `ml` of `x` by `y`. ConceptBase realizes the `Retell(A(x, ml, y))` by a combination of `Untell(A(x, ml, z))` removing the first stored attribution fact `A(x, ml, z)` and a subsequent `Tell(A(x, ml, y))`. If there are several facts such as `A(x, ml, z1)` and `A(x, ml, z2)`, then only the first one is removed by the `Untell`.

³Note, that an object can have more than one attribute value in one attribute category, but the attributes must have different labels.

CreateIndividual(Prefix,ID): A new individual object with the given prefix and a system generated suffix is created. The object identifier of the created object is returned in the second argument, which must be therefore a free variable. The prefix must be an existing object name (e.g. the class name) otherwise the ECArule compiler will report an error.

CreateNew(ClassName,ID): A new instance of class `ClassName` is created. The new instance carries the label `ClassName` as prefix with a system-generated suffix appended to it. The identifier of the created object is returned in the argument `ID`.

newLabel(Prefix,L): A new label that is not yet used as object name is created. The prefix must be an existing object name. The 2nd argument should be a variable. The type of the variable can be `Label` or a more specific class name.

CreateAttribute(AttrCat,x,y,ID): This predicate creates a new attribute for object `x` with value `y` in the given attribute category (e.g. `Employee!salary`). The attribute will get a system-generated label. In contrast to the action `Tell(A(x,ml,y))` the attribute is also created, if another attribute with the same attribute category already exists.

Other predicates to be invoked via `Call` can be defined in a LPI-file (see Counter example below).

Events and actions may also be specified in a prefix syntax, e.g. `Tell (x in Class)` instead of the longer form `Tell ((x in Class))`. Furthermore, there are two simple builtin actions: `noop`⁴ is not doing anything (except that it succeeds), and `reject` aborts the current transaction.

If the execution of an ECArule leads to an update to the database, i.e. via `Tell` or `Untell`, then the updated database is subject to integrity checking. If a violation is detected, then the whole transaction is rolled back including the updates done by ECArules.

4.1.5 Priorities

The attributes `priority_after` and `priority_before` ensure, that this ECArule is executed *after* or *before* some other ECArules, if several rules can be fired at the same time. There can be multiple values for each of these attributes. As you may have noted one of the two attributes is redundant since `(r1 priority_after r2)` is equivalent to `r2 priority_before r1`). Still, ConceptBase provides both. Furthermore, ConceptBase does not automatically check the consistency of the priority declarations (if `r1` is before `r2` then `r2` cannot be before `r1`). ConceptBase also does not provide for the transitivity of the priority. You can however define this yourself via appropriate deductive rules.

4.1.6 Coupling mode of an ECA rule

The coupling mode of an ECArule determines the point of time when the condition and the action of the ECArule are evaluated and executed. Possible values are:

Immediate: The condition is evaluated immediately after the event has been detected. If it evaluates to a non-empty answer, the DO-action is executed immediately, too. If the answer is empty, then the ELSE-actions are executed (provided that the ECArule has an ELSE part).

ImmediateDeferred: The condition is evaluated immediately after the event has been detected. If it evaluates to a non-empty answer, the actions of the DO-part are executed towards the end of the current transaction.

Deferred: The condition is evaluated towards the end of the current transaction. If it evaluates to a non-empty answer, the actions of the DO-part of the ECArule are executed immediately after the evaluation of the condition. Otherwise, the ELSE-actions are executed (provided that the ECArule has an ELSE part)

⁴Earlier ConceptBase releases used the keyword `commit` instead of `noop`. The semantics was the same. We continue to support the use of `commit` in legacy ECArules.

The answer to the evaluation of the condition of an ECARule is the set of all combinations of variable fillers that make the condition true.

The modes `Immediate` and `ImmediateDeferred` differ considerably from the mode `Deferred`: the condition of the ECARule is evaluated while not all frames of the current transaction are told (resp. untold). So, a quoted predicate like ``(x in Class)` will be evaluated against a database state in which only those frames of the current transaction are visible (resp. invisible) that were told (resp. untold) before the ECARule was triggered!

The default is `Immediate`. `ConceptBase` shall enforce a first-in-first-out sequencing of ECARules with modes `ImmediateDeferred` and `Deferred`. This sequence will enforce the complete execution of a triggered ECARule before the next rule triggering is handled. The strict sequencing avoids intertwining of action executions of multiple ECARule threads. So, if the answer to a condition of an ECARule has multiple entries, then the actions belonging to the respective answers are executed in a sequence in which no action of another ECARule is called.

4.1.7 Execution Semantics

The coupling mode of an ECARule influences its execution semantics. There are three basic steps. First, a Tell/Untell/Ask transaction is translated in a sequence of atomic events. In case of a Tell/Untell, each frame inside the Tell induces a delimiter in the event list that is used later for the event processing. A typical event list might look like

e_1	e_2	e_3	◇	e_4	e_5	◇	...
-------	-------	-------	---	-------	-------	---	-----

Here, the events $e_1 \dots e_3$ were generated for the first frame of a Tell/Untell, the events e_4, e_5 were generated for the second frame. The diamond separates the events generated for subsequent frames of the same Tell/Untell operation. The event list is right-open since the execution of actions from ECARules can lead to further events. Each event has one of the forms *Tell(lit)*, *Untell(lit)*, or *Ask(q)*, where *lit* is an attribution, instantiation, or specialization predicate, and *q* is a query call.

`ConceptBase` will scan the event list and process events as soon as it detects a delimiter (denoted by the diamond above). For example, when `ConceptBase` "sees" the delimiter between e_3 and e_4 , it will start to process e_1 to e_3 , one after the other. Processed events are removed from the event list. The first step is to determine the matching ECARules for a given event e_i . This yields a *working set* of rules for each processed event e_i :

$$ws(e_i) = \text{Set of all ECARules whose ON part matches } e_i$$

The matching of the ON part of the ECARule and the event typically leads to a binding of variables in the rule's IF and DO parts. This binding is stored in the rule's representation within the work set. The rules in the working set are sorted to reflect the priority settings of ECARules. If two rules have no priority defined between them, then the definition order is used to sort them (older rules before newer rules). Each rule r in the working set will then be processed as follows, one after the other.

The rule processing depends on the coupling mode of the ECARule in the working set. If the mode is `Immediate`, then the IF part is evaluated (yielding all possible combinations of variables that make the IF part true). If the answer set is empty, then `ConceptBase` will call the ELSE actions of the ECARule. Otherwise, `ConceptBase` will call the DO part for each variable combination determined in the previous step.

If the mode is `ImmediateDeferred`, then `ConceptBase` will evaluate the IF part of the current rule like before. Instead of calling the DO (or ELSE part), it will however put a trigger *do(r, actions)* on a trigger queue q . The parameter r identifies the rule. The parameter *actions* contains all instantiations of the DO-part of r by the variable substitutions computed by the evaluation of the IF part of r . If there are no such substitutions, then *actions* is set to the ELSE part of r (if existent). Otherwise, no actions would be appended to the waiting queue.

If the mode is `Deferred`, then `ConceptBase` will not immediately evaluate the IF part but will just append a trigger *def(e_i, r)* to the trigger queue q .

When all events in the event list are converted to triggers, i.e. when all frames in a Tell/Untell transactions are transformed and stored, or when the Ask call has been processed, then `ConceptBase` will start to process the trigger queue q . It is processed in a first-in-first-out (FIFO) manner. Each `do` and `def` trigger

is processed according to its type. If the entry has the form $\text{do}(r, \text{actions})$, then ConceptBase will execute the actions in the second parameter, possibly leading to new events and triggers.




Figure 4.1: State of the trigger queue

Figure ?? shows a snapshot of a trigger queue. The triggers t_k have the form $\text{do}(r, \text{actions})$ or $\text{def}(e_i, r)$. There are two pointers "last fired" (initially 0) and "next free" (initially 1) to manage the queue. New triggers are added at the right end (incrementing the next free pointer). The processing of the queue starts from left to right. The last processed item is pointed to by the "last fired" pointer. The queue is empty (resp. completely processed) iff "last fired" plus 1 equals to "next free".

If the entry has the form $\text{def}(e, r)$, then ConceptBase will first determine all combinations of variables in the IF part of r . Then it will call the actions of the DO (or ELSE part) for the computed answers. Note that the event e typically binds some variables in the rule r .

Actions of an ECArule can update the database. They will then lead to new entries in the event list that are processed just like described above. The events are generated per action (Tell/Untell/Retell) and then followed by a delimiter, i.e. ConceptBase will start processing the events after each Tell/Untell/Retell action.

4.1.8 Switching Queues

You can control the execution order of ECA triggers via coupling modes, precedence (priority), and via a feature called *queue switch*. The queue switch utilizes separate trigger queues for ECA triggers:

Main queue q0: The main queue is the default queue. It is processed under a first-in-first-out regime. When all triggers of the main queue are processed, ConceptBase will stop with the ECA execution.

Sub-queue q1: The sub-queue q1 is initially empty and is filled by triggers from actions following a queue switch. As soon as q1 is empty, the system will resume processing of q0. Hence, the end of the sub-transaction is detected by running out of triggers in the sub-queue.

User-defined queues: The FOR-clause allow to generate dedicated trigger queues on the fly (see below).

There are two ways to specify a queue switch. The first is by including the keyword `TRANSACTIONAL` in the ON-part of an ECArule, e.g.

```
ON TRANSACTIONAL Tell (x in A)
```

As soon as an event e is matched against the event clause of a "transactional" ECArule, the trigger queue is switched to q1. Subsequent triggers will then be put on q1 instead of q0. Note that several ECArules can match a given event. The switch occurs when the first transactional rule is encountered.

The *second* method creates user-defined trigger queues via the FOR-clause:

```
ON Tell (x in A) FOR x
```

In the above example, each event that matches the event clause will also fill the variable x in the FOR-clause. This will instruct ConceptBase to switch to the (new) trigger queue labelled x . If you use "FOR q1", then the method yields the same effect as with the `TRANSACTIONAL` clause. As soon the the trigger queue labelled x is empty, ConceptBase will switch back to the newest non-empty trigger queue, or back to q0.

The events on `q1` are processed before the remaining events of `q0`. The queue switch leads essentially to a prioritization of events on `q1` over `q0`.

Examples highlighting the differences between the execution modes and the transaction model are presented in the CB-Forum at <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/2924361>. You can use the tracemode `high` (see chapter ??) to debug ECArules. The ConceptBase server will then write trace messages about the execution of ECArules on the console terminal.

4.1.9 Activate and Deactivate ECA rules

The attribute `active` allows the user to deactivate the rule without untelling it. Possible values are `TRUE` and `FALSE`. The default value is `TRUE`, i.e. by default are ECArules active.

4.1.10 Depth

The attribute `depth` specifies the maximum nesting depth of ECArules. ECArules may be fired by events which are produced by actions of the same or other ECArules. Because this often results in an endless loop, the execution of the ECArule is aborted if the current nesting depth is higher than the specified value. The default of this attribute is 0 (= no limitation).

4.1.11 User-definable Error Messages

If the ECArule rejects the transaction in some cases, it is useful to specify an error message. The value of the attribute `rejectMsg` is returned to the user.

4.1.12 Constraints

The constraints of the class `ECArule` ensure, that the attributes `mode`, `active` and `depth` have only single values and that the attribute `ecarule` has exactly one value.

4.2 Examples

The Telos source files of the following examples can also be found in your ConceptBase installation directory at `$CB_HOME/examples/ECArules`. Further examples are in the CB-Forum at <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/1747992>.

4.2.1 Materialization of views by active rules

Materialization of views means that deduced information is stored in the object base. We provide here an example, how to materialize and maintain simple views.

```
Class Employee with
attribute
    salary : Integer
end
View EmployeeWithHighSalary isA Employee with
constraint
    c : $ exists i/Integer (this salary i) and (i > 100000) $
end
Class EmployeeWithHighSalary_Materialized end
```

The view `EmployeeWithHighSalary` contains all employees who earn more than 100.000. The class `EmployeeWithHighSalary_Materialized` will contain the same employees. This implemented by the following ECArules:

```

ECArule EmployeeWithHighSalary_Materialized_Ins with
ecarule
    er : $ x/Employee
        ON Tell (x in Employee)
        IF `(x in EmployeeWithHighSalary)
        DO Tell (x in EmployeeWithHighSalary_Materialized) $
end
ECArule EmployeeWithHighSalary_Materialized_Del with
ecarule
    er : $ x/Employee
        ON Untell (x in Employee)
        IF (x in EmployeeWithHighSalary)
        DO Untell (x in EmployeeWithHighSalary_Materialized) $
end
ECArule EmployeeWithHighSalary_Materialized_Ins_salary with
ecarule
    er : $ x/Employee y/Integer
        ON Tell (x salary y)
        IF `(x in EmployeeWithHighSalary)
        DO Tell (x in EmployeeWithHighSalary_Materialized) $
end
ECArule EmployeeWithHighSalary_Materialized_Del_salary with
ecarule
    er : $ x/Employee y/Integer
        ON Untell (x salary y)
        IF (x in EmployeeWithHighSalary)
        DO Untell (x in EmployeeWithHighSalary_Materialized) $
end

```

The first rule checks, if the employee belongs to the view, when the employee was inserted. Note, that we don't use the constraint of the view in the ECArules, we just reuse the view definition here. The second rule does the same for deletion of employees. The first rule checks the IF-part on the new database state since the employee's salary is usually told together with the employee. In contrast, the IF-part of the second rule is checked against the old database state, i.e. where the employee was still defined.

The third rule checks, if the employee is an instance of the view class, when the attribute salary was inserted. Again, the fourth rule does the same for deletion of the attribute.

If the number of employees is large it is more efficient to ask for the instances of materialized than to evaluate the view. However, if updates occur quite often, materialization is not good, because materialized view must be maintained for every update transaction.

4.2.2 Counter

This example shows how to call prolog predicates with an ECArule. It implements a counter for a class Employee. The counter is stored as an instance of the object EmployeeCounter. Whenever an employee is inserted or deleted from the object base, the counter is incremented or decremented.

```

Class Employee end
EmployeeCounter end

ECArule EmployeeCounterRule with
ecarule
    er : $ x/Employee i,i1/Integer
        ON Tell(In(x,Employee))
        IF (i in EmployeeCounter)
        DO Untell(In(i,EmployeeCounter)),
            Call(increment(i,i1)),
            Tell(In(i1,EmployeeCounter))
        end
end

```

```

        ELSE Tell(In(1,EmployeeCounter))
    $
end

ECArule EmployeeCounterRule_del with
ecarule
    er : $ x/Employee i,i1/Integer
    ON Untell(In(x,Employee))
    IF (i in EmployeeCounter)
    DO Untell(In(i,EmployeeCounter)),
        Call(decrement(i,i1)),
        Tell(In(i1,EmployeeCounter))
    $
end

```

The files `$CB_HOME/examples/ECArules/counter.*.lpi`⁵ contain the code for the Prolog predicates `increment` and `decrement`. You must copy LPI files to the database directory before you start the ConceptBase server (see also appendix ??). Note, that all free variables of the PROLOG predicate must be bound in its call. Furthermore, the variables must be bound to object identifiers, if you want to use them in a `Tell`, `Untell` or `Ask` action.

The effect of the `increment` and `decrement` procedures can also be achieved using the arithmetic expressions like $i+1$. The simple solution with arithmetic expressions is available from <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/d2773786/EmployeeCount.sml.txt>. The purpose of the example above is only to show that user-defined PROLOG predicates can be called in the DO-part of an ECArule. You can define more interesting PROLOG predicates like sending an email to a user with content derived from the object base. This requires however some knowledge of PROLOG and of the internal features of the ConceptBase server.

You should also note that the count of a class `c` can always (and more correctly) be computed by the function `COUNT(c)`. That does even count inherited and deduced instances.

4.2.3 Timestamps

An often asked requirement in metamodeling applications is the recording of creation and modification dates. ConceptBase stores the creation time of an object in its object base, primary for the use of *Rollback queries*. With the predicate $Known(x,t)$ the time of the creation of x can be made visible in rules or queries. The following frames shows how to use it:

```

Class Employee with
attribute
    salary : Integer;
    createdOn : TransactionTime;
    lastModified : String
rule
    createdOnRule : $ forall t/TransactionTime
        Known(this,t) ==> (this createdOn t) $
end

EmpWithoutLastModified in QueryClass isA Employee with
constraint
    noLM: $ not exists t/String (this lastModified t) $
end

```

⁵Since ConceptBase 6.2, LPI plug-ins can be defined in two different formats. A file with the suffix `.bim.lpi` is intended to be used by a CBserver based on MasterProlog (formerly BIM-Prolog). If the CBserver is based on SWI-Prolog, the server reads files with the suffix `.swi.lpi`. Both Prolog Environments used a slightly different syntax which requires different implementations. ConceptBase 7.0 and later only supports SWI-Prolog, hence providing the swi variant is sufficient.

The limitation of this approach is, that it just records the creation date of an object and not the time when it was modified. i.e. the value of an attribute was changed.

To overcome this restriction, one can use ECARules to update the attribute `lastModified` of the above example, whenever an attribute of the category `salary` is inserted.

```
ECARule LastModified_init with
  mode m: ImmediateDeferred
  ecarule
    er : $ t/TransactionTime y/Employee
        es/Employee!salary
    ON Tell (es in Employee!salary)
    IFNEW
      From(es,y) and (y in EmpWithoutLastModified) and Known(es,t)
    DO Tell (y lastModified t)
    $
end

ECARule LastModified_change with
  mode m: ImmediateDeferred
  ecarule
    er : $ t1,t2/TransactionTime y/Employee i/Integer
        es/Employee!salary lab/Label
    ON Tell (es in Employee!salary)
    IFNEW
      From(es,y) and (y lastModified t1) and Known(es,t2)
    DO Untell (y lastModified t1),
      Tell (y lastModified t2)
    $
end
```

Both ECARules are evaluated against the new database state. The first ECARule is for the case where the object `y` has not yet a `lastModified` attribute. Then, it has to be initialized. The second rule takes care for the updating case. Note that the transaction time is represented as string of the form "`tt (year,month,day,hour,minute,sec,millisecond)`".

A deprecated solution is available from <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/d2936786/ECA-lastmodified.sml.txt>. It uses the `Ask` action in the '`DO`' and '`ELSE`' parts to query perform different actions depending on whether an object already has a filler for the attribute `lastModified`.

4.2.4 Simulation of Petri Nets

ECARules are a powerful tool to express semantics of concepts that are not expressible by deductive rules. For example, the semantics of petri nets, in particular the firing of an enabled transition, can be expressed by a single ECARule.

```
ECARule UpdateConnectedPlaces with
  mode m: Deferred
  rejectMsg rm: "The last firing of a transition failed.
                Check whether the transition was enabled!"
  ecarule
    er : $ fire/FireTransition t/Transition p/Place m/Integer
    ON Tell (fire transition t)
    IF (t in Enabled) and
      (p in ConnectedPlace[t]) and
      (m = M(p)+IM(p,t))
    DO Retell (p marks m)
    ELSE reject $
end
```

The example shows that the condition can also be a complex logical expression. Note that the event (ON-part) binds the two variables `fire` and `t`. The condition (IF-part) additionally binds the free variables `p` and `m`. Each binding of the free variables is passed to the DO-part leading to some update of the `tokenFill` attribute. In case that the IF-part cannot be evaluated to true, the ELSE-part is executed. That will lead to an abortion of the current transaction rolling back all updates and issuing the error message listed under `rejectMsg`. A complete specification of modeling petri nets is in the CB-Forum (<http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/1419080>). Further examples of ECA rules can be found at <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/1747992>.

It should be noted that one can also ask queries in the DO-part (and ELSE-part) of an ECA rule. Other than for the IF-part, such queries are only evaluated once.

4.3 Optimization of ECA rules

ECA rules can be configured to follow various execution semantics. A particular issue is the evaluation of the ECA condition (IF-part). Originally, the predicates in the IF-part were not re-ordered by ConceptBase to gain a better performance. Instead, they were executed exactly in the sequence in which they were defined. This is still the case when the CBserver parameter `-eo` is set to `off`.

Consider the following example as a variant of `UpdateConnectedPlaces` discussed in the previous subsection:

```
ECARule UpdateConnectedPlacesV1 with
  mode m: Deferred
  ecarule
    er : $ fire/FireTransition t/Transition p/Place m/Integer
        ON Tell (fire transition t)
        IF (p in Place) and
           (m = M(p)+IM(p,t)) and
           (t in Enabled) and
           (p in ConnectedPlace[t])
        DO Retell (p marks m) $
  end
```

`UpdateConnectedPlacesV1` is equivalent to `UpdateConnectedPlaces` but it results in significantly longer execution times. The reason is that the condition of `UpdateConnectedPlacesV1` starts with a predicate with an unbound variable `p`. In contrast, the condition of `UpdateConnectedPlaces` starts with a predicate whose variable `t` is bound by the ON-part of the ECA rule.

In cases where the IF-part does not contain a mixture of quoted (new database state) and unquoted (old database state), one can outsource the condition to a query class. The advantage is that the constraint of the query class is automatically optimized by ConceptBase.

The extended example is in the CB-Forum at <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/3242537>.

When the CBserver parameter `-eo` is set to `on` (=default), then ConceptBase will apply various heuristics to re-order the predicates in the condition. This can result in several orders of magnitude better performance. Currently, the optimization only applies to conditions that are conjunctions of predicates and whose predicates are all referring to the same database state.

4.4 Limitations of the current implementation

The current implementation of active rules in ConceptBase has several limitations.

- The dependency graph of ECA rules may contain cycles, i.e. rule `R1` executes an action, which fires the rule `R2` and `R2` has an action which fires `R1`. The current algorithm for detecting cycles tests the dependency graph during compile time. The cycle checker prints a warning message on the console if it has detected a possible cycle, but the compilation of the rule is not aborted, because it is still

possible that the cycle does not occur during the execution of the rule. The user must take care, that endless loops are avoided.

- The ON-part of an ECArule consists of a single atomic event. ConceptBase does not support complex events such as "event 1 occurs after event 2".
- Instantiations and specializations of system classes (Proposition, Attribute, InstanceOf, ...) are not detected in the event manager. Adding those events would affect the performance.
- Like QueryClasses, ECArules can not be modified after they have been told (e.g. change the event, condition or actions). The only updateable attribute of ECArule is `active`. If you set this attribute to `FALSE`, you can deactivate the rule for a certain time. Deleting or setting the attribute to `TRUE` re-activates the ECArule. If you don't want to use an ECArule anymore, you can untell it as a whole.
- The ConceptBase ECArule evaluator is not very efficient, roughly 1000 rule triggerings per second. Response times may be long in case that an ECArule leads to many updates to the object base.
- ECArules that update the database will temporarily disable the cache-based predicate evaluator. By this, certain recursive deductive rules called in the IF-part of ECArules may not terminate if you activate the CBserver parameter `unsafe` (see section ??). The 'unsafe' mode disables checks on the presence of recursion. This leads to faster execution time since the 'safe' mode would refresh the predicate cache after each update to the database.

The expressive power of ECArules exceed the one of deductive rules, which are limited to Datalog. So, one could be tempted to prefer ECArules over deductive rules. The contrary should however be your choice:

- Deductive rules (that are not using arithmetic or functions) will always terminate, even under recursion. Active rules can well run into infinite loops if not programmed carefully.
- Deductive rule evaluate several orders of magnitude faster than ECArules.
- Deductive rules never change the database (=set of the stored objects).
- ECArules are practically always triggered by updates or lead to a database update. Thus, they will lead to emptying the cache of derived facts maintained by the Datalog engine. Calling a deductive rule shall not empty the cache but just lead to extending it.

There are some scenarios, where you do need the power of ECArules. For example, you may want to trigger the call of an external program, when a certain condition becomes true in the database. Or you need to change the database state of certain objects when a certain update occurs. For example, triggering a transition in a petri net shall change the state of the 'places' connected to the triggered transition. Such semantics is beyond Datalog and requires more expressive power, such as provided by ECArules.

Chapter 5

The Module System

Modules divide a *ConceptBase* database into a hierarchy of namespaces that determine which objects are visible inside the scope of a given module. Hence, each module forms a database that is part of the whole database. The scope of visibility also applies to deductive rules, integrity constraints, queries, active rules, and functions. They are also objects of the database and thus subject to the visibility rules. Any object in the database belongs to exactly one module, but it can also be visible to other modules, in particular to the sub-modules. Hence, whenever an object is created, updated or deleted (TELL, UNTELL), then this has to be done in the context of the module that defines the object. You can only query (ASK) an object, if the object is visible in your current module context. Analogously, an object x can only reference other objects that are visible in the module context where x is defined.

Modules that are not visible to each other also do not interfere with each other. Hence, two users can use the same *ConceptBase* database and any change (TELL, UNTELL) done in the one module context does not influence the way how the other module reacts to requests. Modules are organized in a hierarchy. Changes to a supermodule shall impact all its sub-modules. This applies in particular to the integrity of the sub-modules. For example, if an object in a sub-module refers to an object in a super-module, then an update to the object in the super-module could render the sub-module inconsistent. This would lead to a rejection of the update. Analogously, the result of queries in a sub-module also depends on the objects in the super-modules. A module in *ConceptBase* is quite similar to a DATABASE in SQL in terms of isolating work spaces. On the other hand, the sub-module construct allows for controlled sharing of objects.

One useful application of modules in *ConceptBase* are modelling situations where different objects are labeled with identical names. In earlier versions of *ConceptBase* this was prohibited by the *Naming Axiom* which demands that different objects have different names. Now objects with identical names can be stored in different modules without interferences. Another application is to store alternative conceptual models of the same domain in different modules. The alternative versions can also share a common core, e.g. by storing it in a super-module of the version modules.

Another application of modules is to separate a metamodel (defining some constructs) from the models represented in terms of the metamodel. The metamodel shall be defined in a super-module and the models would be stored in sub-modules of this super-module. By this organization, the models do not interfere with each other unless they are explicitly linked via export/import clauses. In multi-perspective modeling, models need to be rather tightly linked to each other. Then, one should better store all such models in a single module.

The set of visible objects in a module context is not limited to the set of objects defined for the module. The *ConceptBase* module concept permits to reuse existing objects from different modules via **import** and **export** interfaces. Furthermore, modules can be **nested**. A nested module object (also called sub-module) is defined in the context of another module, called its super module. A nested module object is only visible within the context of its super module. Objects that are visible in the super module are visible (and can be reused) to all its sub-modules.

ConceptBase users can be assigned to individual **home modules**, i.e. the module they start working with when registering with the *ConceptBase* server. So-called **auto home modules** force every user in her own module to further reduce potential unwanted interferences. A flexible **access control** mechanism

allows to define access rights of users simply by query classes defined either globally or locally to a module.

5.1 Definition of modules

The class `Module` defines an attribute `contains` and thus is the construct to create new modules. Each module is created by creating an instance of the class `Module`.

```
Module in Class with
  attribute
    contains: Proposition
end
```

Thus, a module is a container of objects. Modules create a name space: object names must be unique within one module, but different modules can contain different objects with identical names.

Tell, **Untell** and **Ask** operations work relatively to a specified module. The normal way of specifying the context in which a transaction takes place is using the *Set Module* function of *CBIVA*. See subsection ?? to learn how to change the module context of a transaction. The command line interface *CBShell* has similar operations to switch the module context.

The basic set of predefined objects of *ConceptBase* (such as `Class`, `Proposition`, `QueryClass`, etc.) belongs to the predefined module `System`. The default module of clients logging into a *CBserver* is `oHome`, a direct submodule of `System`. You can set the module context to your other modules in order to manipulate them.

The `contains` attribute of the `Module` object is a derived attribute and is a link to all objects defined inside the context of a certain module. This attribute is not stored explicitly but can be used anywhere in the O-Telos assertion language.

Now we show how to create modules. We introduce a small running example in order to demonstrate all module-related facilities of *ConceptBase*.

Let's assume you've started *ConceptBase* with a fresh database. After telling the following two frames, the server contains two new modules, which are nested inside the pre-defined `oHome` module:

```
Master in Module end
Work in Module end
```

The super-module `oHome` includes the two objects `Master` and `Work` but not the objects contained in them (compare also Figure ??). We call `oHome` the super-module of the sub-modules `Master` and `Work`. One you switch to a sub-module, say `Master`, then all sub-subsequent operations apply to this module context. Telling new objects makes them members of this module context, in other words `Master` contains them.

5.2 Switching between module contexts

The standard way for changing the module context is to choose the *select module* menu entry from the *Options* menu in the user interface. A dialog with a listbox containing all known modules in the current context are shown. Double-clicking a module entry in the listbox sets the current module context to the selected module and lists all modules that are visible in this module (i.e., its submodules and supermodules).

In our example, you should get a window displaying the modules `Master`, `System`, `Work`, and `oHome` in alphabetic order (see Figure ??). Now double-click the entry `Work`. As a result the module context of the *CBserver* is set to the module `Work`. As the `Work` module has got no nested modules, there are no additional modules displayed in the listbox (i.e., all modules that are visible in the `oHome` module are also visible in the `Work` module. Alternatively, you can specify the module context using the *load model* operation and placing a

```
{ $set module=MODNAME }
```




Figure 5.1: The module context selection window

inline command placed before the first Telos frame of a source model file. `MODNAME` stands for the name of the module in which you wish to define the content of the source model file. Module paths like `System/oHome/Work` are allowed as well. Please note that only one inline-command is allowed within one source model file. Specifying module contexts using inline commands is a facility for automatically loading Telos frames of large applications which are spread over different modules – without requiring the user to employ the *select module* function from *CBIVA*.

In CBSHELL, you can use the command `cd` (alias: `setModule`) to switch the module context. Let's set the module context first to `oHome` and then TELL the following frames for a very simple ER notation (using CBSHELL syntax, see section ??):

```
cd oHome

tell "EntityType end
RelationshipType with
  attribute
    role: EntityType
end"
```

Then set the module context first to `Master` and then TELL the `Employee` frame:

```
cd Master

tell "Employee in EntityType,Class with
  attribute
    name : String
  constraint
    nec: $ forall e/Employee exists n/String (e name n) $
end"
```

The object `Employee` is now only visible in the module `Master`. When you set the module context to `Work` (or `System`) and try to load the `Employee` object, you should get an error message from the server stating that the object `Employee` is not visible in that module context. The object `Employee` correctly references `EntityType`, since it is visible via the super-module `oHome`.

5.3 Using nested modules

A nested module object is a module defined in the context of a super module and therefore is contained in the super module. The objects contained in the nested module are *not* contained in the super module.

After the definition of the `Master` and `Work` modules as nested modules to the `oHome` module, let's define a nested module to the `Work` module. We assume that the current module context is set to the `Work` module. Now tell the following module object:

```
cd Work
tell "Test in Module end"
```

As a result we have defined the nesting hierarchy depicted in Figure ??.




Figure 5.2: Nested module hierarchy

A nested module can see the content of all modules on the path to the root module `System`. Therefore, when you set the module context to `Test`, you can reference all objects contained in the modules `Test`, `Work`, `oHome`, and `System`. When you set the module context to `Work`, you can reference all objects contained in `Work`, `oHome`, and `System`.

Any `ConceptBase` database has the pre-defined modules `System` and `oHome`. The `System` module contains the pre-defined objects of `ConceptBase`, i.e. the O-Telos classes `Proposition`, `Individual`, `Attribute`, `InstanceOf`, and `IsA`, but also a large number of other pre-defined objects that are required for a functioning system, e.g. `Integer`, `Class`, `QueryClass`, and many more. The `oHome` module is initially empty. It is the default home module for users and clients that add/update/delete and query objects. A user could delete objects in the `System` module, which could then disable core functions of `ConceptBase`. One can however also prevent such updates by limiting the access to the `System` module.

5.4 Exporting and importing objects

In order to use objects which are not visible in a module, *ConceptBase* offers the export/import facility. The class `Module` defines two further attributes, namely to specify objects exported from a module and to modules imported by a module.

```
Module with
  attribute
    contains      : Proposition;
    exports       : Proposition;
    imports       : Module
end
```

In order to allow other modules to import objects from a module, we need to define an `export` attribute from the module object to those objects. We call the set of objects exported by a module the *export interface*

of a module. In order to include the export interface of another module to a module, we need to define an `imports`-attribute from the module object to the module to be imported.

In our running example we would like to define a specialization of the Class `Employee` within the `Work` module. It is desirable to reuse `Employee` from the `Master` module instead of redefining it in the `Work` module.

In order to import the class `Employee` to the `Work` module, we have to define all objects belonging to the `Employee` class as exported objects. First change the module context to `Master`. Now TELL the following frame within module `Master`:

```
Master with
  exports
    e1 : Employee;
    e2 : Employee!name
end
```

Now you have defined the objects `Employee` and `Employee!name` as exported objects of the `Master` module. Any module in your database, which defines an `imports`-attribute to the `Master` module, can now reference these objects. Now change the module context to `Work` and TELL the following frame within module `Work` (see also Figure ??):

```
Work with
  imports
    i1 : Master
end
```

The objects mentioned above are visible in the context of the `Work` module. Check this by loading the `Employee` object with the *Edit/Load Object* function of *CBIva*. The import declaration for `Work` is done within the module context `Work`. The attribute `Work!imports` is thus part of the `Work` module while the `Work` object is contained in the `oHome` module. The object `Master` is visible in the `Work` module as well. The visibility rules restrict the possible import declarations. You can now define the specialization `Manager` of `Employee` in the context of the `Work` module:

```
Manager isA Employee end
```

When untelling `exports` declarations from a module, *ConceptBase* checks for integrity violation in all concerned modules. Try untelling the `exports` attributes from the `Master` module and you should get an error message saying that referential integrity is violated in the `Work` module. The reason for this violation is simple: since the class `Employee` is no longer exported from `Master`, it is no longer visible inside the `Work` module and therefore the referential integrity (a builtin O-Telos axiom) is violated for the `Manager` specialization of `Employee`.

If you define integrity constraints in the exporting module `Master`, then these constraints are not checked for objects in the importing module `Work`. For example, the constraint `nec` of `Employee` is only visible in `Master` and its sub-modules, not in `Work` and its sub-modules. Hence, you can still declare an object like `bill` in module `Test` that does not need to fulfill the constraint for `Employee`:

```
cd Test
tell "bill in Employee end"
```

5.5 Modules and metamodeling

The module hierarchy in Figure ?? assigns objects belonging to different abstraction levels (meta classes, classes, objects, ...) to different modules. This assignment is not prescribed, but it is recommended. For example, the module `oHome` stores meta classes such as `EntityType`. This metaclass is then used in the sub-module `Work` to define a class like `Manager`. Finally, the module `Test` instantiates the class `Manager` by `bill`. There is a natural reason for such a structure. The persons defining meta classes

are engineering modeling languages. The persons defining classes are conceptual modelers or application programmers, and the persons defining objects at the lowest abstraction level are application users. These activities depend on each other but one should separate the workspaces to shield an update to the modeling language from an update to a conceptual model. Another reason is that a modeling language can be used for many conceptual models. Each conceptual model needs the definitions of the modeling language but they typically should be separated from each other. Hence, each conceptual model can be stored in its own module, being a sub-module of the module defining the modeling language.

The following CShell script shows how a module hierarchy is created. The command "cd" is used to switch to a module, the command "mkdir" creates a new submodule in the current module. The module `oHome` contains the submodule `ERnotation` for the definition of the ER modeling language. The two submodules `UModel` and `LibModel`. The submodule `UData` of `UModel` stores a sample database for the university model.

```
cd oHome
mkdir ERnotation
cd ERnotation
tellModel ERD-Language.sml.txt
mkdir UModel
mkdir LibModel
cd ERnotation/UModel
tellModel UniversityModel.sml.txt
mkdir UData
cd ERnotation/UModel/UData
tellModel UniversityData.sml.txt
```

5.6 Setting user home modules

When ConceptBase is used in a multi-user setting, it makes sense to automatically assign clients of ConceptBase users to a dedicated module context, their so-called *home module*. To use this feature, the database of the ConceptBase server has to contain instances of the pre-defined class `CB_User`. This class is defined as follows:

```
CB_User with
  attribute
    homeModule : Module
end
```

Assume that we have two users `mary` and `john` who need to be assigned to different modules when they log into the CBserver by their favorite user interface. The system administrator should then include the following definitions to the database of the CBserver:

```
Project1 in Module end
Project2 in Module end

mary in CB_User with
  homeModule m1 : Project1
end

john in CB_User with
  homeModule m1 : Project2
end
```

As a consequence, the start module of the two users will be set accordingly when they log into the CBserver. The home module feature is especially useful in a teaching environment. The teacher can put

some Telos models into the shared `oHome` module. Students' home modules would be assigned to sub-modules, e.g. based on group membership. Each student group can then work on an assignment by working on their sub-module without interfering with other student groups.

There is a subclass `AutoHomeModule` of `Module`, which supports applications of `ConceptBase` where by default any user should work in her own module context. Rather than having to define separate modules for each user explicitly, you can just define a certain module to be an instance of `AutoHomeModule`.

```
LectureModule in AutoHomeModule with
  exception e1: mary
end

mary in CB_User with
  homeModule m1 : LectureModule
end

john in CB_User with
  homeModule m1: LectureModule
end
```

You can also define a rule to assign all or a subset of users to this module:

```
CB_User in Class with
  rule
    homeRule : $ forall u/CB_User (u homeModule LectureModule) $
  end
```

Here, user `john` (a student) will be automatically be assigned to a new module `M_john` that is created as sub module of `LectureModule`. User `mary`, presumably a teacher, is defined to be an *exception* to this rule and she will get the home module `LectureModule`. By this, one can reduce the chances of unwanted interferences between users of the module `LectureModule`. Still, all users can read the definitions in the module `LectureModule` and its submodules unless access restrictions are defined (see section ??).

The simplest way to separate the workspaces of *any* user is to tell

```
oHome in AutoHomeModule end
```

In this case, no user needs to be defined explicitly¹ as instance of `CB_User` and still will get assigned her own sub module to work in. The auto-home module becomes active as soon as the module is declared as instance of `AutoHomeModule`. If you want to enable `oHome` as auto-home module from the very beginning when a database is created, then you can activate it by a parameter of the `CBserver`, e.g.

```
cbserver -g public -new MYNEWDB
```

This will instruct the `CBserver` to tell "`oHome in AutoHomeModule end`" when it sets up the new database. See also section ??.

The home module definitions need to be made within module `oHome` because they will be evaluated upon client registration (server method `ENROLL_ME`) in this module context. Please note that the module context is only dependent on the user name, not on the client and not on the network location of the user. It could well be that a user `mary` is defined on multiple computers on the network and that different natural persons are identified by `mary`. `ConceptBase` currently cannot detect such cases.

¹A side effect of the server method `ENROLL_ME` is that the user of the registering client will automatically be defined as an instance of class `CB_User`. The definition is be made in the context of module `oHome`.

5.7 Limiting access to modules

When multiple users work on the same server, their workspaces not only need to be separated in a controlled way by means of the module feature. Users are also interested in controlling who has which rights on their workspace (=module). ConceptBase includes basic support for rights definition and enforcement via a user-definable query class `CB_Permitted`. The signature of this query class has to conform the following format:

```
CB_Permitted in GenericQueryClass isA CB_User with
  parameter
    user: CB_User;
    res: Resource;
    op: CB_Operation
  ...
end
```

A user is allowed to perform the operation `op` on the resource `res` iff the constraint of the query is satisfied. Then, `user` is returned as answer of the query. If not, the answer is `nil` (equals empty set). Some fundamental definitions are pre-defined objects of ConceptBase:

```
Resource with end
Module isA Resource end
CB_Operation end
CB_ReadOperation isA CB_Operation end
CB_WriteOperation isA CB_Operation end
TELL in CB_WriteOperation end
ASK in CB_ReadOperation end
```

Hence, at least two operations `TELL` and `ASK` are pre-defined symbolizing write and read accesses to a resource. Modules are the prime examples of resources to be access-protected. Currently, only access to them is monitored by the CBserver.

When a user wants to switch to a new module, then he must at least have the permission to execute the operation `ASK` on it, i.e. read permission. Otherwise, the module switch is rejected. This check is the main protection scheme offered to module owners. Define the query class `CB_Permitted` in the module that needs protection.

Assume that there is a user `jonny` who wants to protect his module `Mjonny`. To do so, he would first define the module and set the module context to `Mjonny`.

```
Mjonny in Module end
```

Then, he would set the module context to `Mjonny` define his rights management policy, for example:

```
CB_Group with
  attribute
    groupMember: CB_User;
    permitted_read: Resource;
    permitted_write: Resource;
    owner_resource: Resource
  end

CB_User isA CB_Group end

CB_Group in Class with
  rule
    rr1: $ forall p/Resource u/CB_User
```

```

        (u owner_resource p) ==> (u permitted_write p) $;
rr2: $ forall p/Resource u/CB_User
    (u permitted_write p) ==> (u permitted_read p) $;
rr3: $ forall u/CB_User (u groupMember u) $;
rr4: $ forall u/CB_User p/Resource
    ( exists g/CB_Group (g groupMember u) and
      (g owner_resource p) ) ==> (u owner_resource p) $;
rr5: $ forall u/CB_User p/Resource
    ( exists g/CB_Group (g groupMember u) and
      (g permitted_write p) ) ==> (u permitted_write p) $;
rr6: $ forall u/CB_User p/Resource
    ( exists g/CB_Group (g groupMember u) and
      (g permitted_read p) ) ==> (u permitted_read p) $
end

```

CB_Permitted in GenericQueryClass isA CB_User with

```

parameter
  user: CB_User;
  res: Resource;
  op: CB_Operation
constraint
  cperm: $ (
    ( not exists u/CB_User
      (u owner_resource ~res) and
      not (u == ~user) )
    or
    ( (~op in CB_ReadOperation) and
      (~user permitted_read ~res) )
    or
    ( (~op in CB_WriteOperation) and
      (~user permitted_write ~res) )
  )
  and UNIFIES(~user,~this) $
end

```

In the above example, access rights are granted to groups of ConceptBase users. The *owner* of a resource will always have full access via rules rr1 and rr2.

Then, the user would claim ownership to the module via

```

jonny in CB_User with
  owner_resource r1: Mjonny
end

```

Then, only jonny can switch to the module Mjonny. If a second user like mary was to be granted read permission, jonny would define within module Mjonny:

```

mary in CB_User with
  permitted_read r1: Mjonny
end

```

In the above example, rights can also be granted to groups of users and then inherited to its members via rules rr4 to rr6. It should be noted that the definitions of *owner_resource*, *permitted_read* and *permitted_write* are just for illustrating what is possible. ConceptBase only requires the query

class `CB_Permitted` in the module where the access rights need to be enforced. If such a query class (or a local version as explained below) is not defined, then any access is permitted for any user.

When a user attempts to switch to new module context, ConceptBase checks whether the user has at least read permission, i.e. permission for executing the operation `ASK`, on the module. If permission is not granted, the user cannot switch the module context and an error message is presented.

The definition of the query `CB_Permitted` is visible in the module where it is defined and in all sub-modules of this module. One can also define a local version of the query by appending the module name to its name, e.g. `CB_PermittedMjonny`. This version is only tested for access to the module `Mjonny`. The local overrides the general version `CB_Permitted` and its function is not inherited to sub-modules. The subsequent definition prevents updates to the `Test` module, if it is visible in the `Test` module and access control is enabled by the `CBserver`:

```
GenericQueryClass CB_PermittedTest isA CB_User with
  parameter
    user: CB_User;
    res: Resource;
    op: CB_Operation
  constraint
    cperm: $ (~op in CB_ReadOperation) and UNIFIES(~user,~this) $
end
```

There are plenty of ways to combine general and local versions of `CB_Permitted` yielding different access policies. When using access control, at least the module `System` should be protected. Otherwise, users could change essential definitions affecting all other users. Examples for access control policies are in the HOW-TO section of the ConceptBase Forum (<http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/2281940>).

It is very well possible to make access to a ConceptBase database completely impossible by errors in the definition of `CB_Permitted`. For example, one could deny access to any operation by the following simple rule:

```
CB_Permitted in GenericQueryClass isA CB_User with
  parameter
    user: CB_User;
    res: Resource;
    op: CB_Operation
  constraint
    cperm: $ FALSE $
end
```

In such cases, one has to start the `CBserver` with disabled access control and repair the definition of the query `CB_Permitted`. Access control is by default set to level 1, which just constrains the scope of an `UNTELL` to the local module. You need to set the `CBserver` option `-s` to 2 to fully enable access control (see option `-s` in section ??).

5.7.1 Access to System module

The access definition via the query `CB_Permitted` or its localized variant is available in the module `oHome` and its submodules. It is not available for checking permission to access the `System` module. The reason is that user details are stored in `oHome` and signature of `CB_Permitted` requires that the user details are known. The protection of the `System` module is instead configured via the security level of the `CBserver` (option `-s`):

- Level=0: No access control. Read and write operations are allowed to the `System` module.
- Level=1: Read and writes are allowed but only if the current module is set to `System`.

- Level=2: Read is permitted but write operations to `System` are disallowed.

Caution: The access control feature of ConceptBase avoids some *unwanted* interferences in a setting where multiple users work on the same server. The system is *not* save against malicious attacks! Neither does it prevent all unwanted interferences.

5.8 Listing the module content

The `contains` attribute allows to check which objects belong to a module. It can be used by a simple query that lists the module content of all modules that are currently visible:

```
ShowModule in QueryClass isA Module with
  computed_attribute
    cont : Proposition
  constraint
    ccont : $ (~this contains ~cont) $
end
```

A more sophisticated method is to use the builtin query `listModule`. A call of `listModule` without parameters will list the current module as Telos frames. You can also provide the module to be listed as a parameter, e.g.

```
listModule[System]
```

will list the content of the module `System`. ConceptBase will check read permission before a module content is listed. You can also use a module path as parameter, e.g.

```
listModule[System-oHome-Work]
```

A module path is formed much like a directory path in a file system. The root module is `System` and modules names are separated by the character `'-'`. You can also use `'/'` as module separator:

```
listModule[System/oHome/Work]
```

The implementation of the `listModule` query preserves the order in which objects have been created. Note that the `System` module is defining the essential objects that ConceptBase requires to run correctly. You can list the `System` module but you should not change it.

If the content of a module was created by separate transactions, then `listModule` shall indicate them by a separator line `"{---}"`. This separator is disabled when you specify `-mg whole` as CBserver parameter (see section ??). The separator is technically a comment. However, CBGraph and CBIva shall use such separator line to split a sequence of frames into separate TELL transactions.

If a module path does not exist or the current user has no read permission on modules in the path, then the answer `"{* no *}"` is returned.

5.8.1 Restrictions of `listModule`

The query `listModule` extracts all objects of a module into a single Telos source. Since a module is typically created by a sequence of TELL/UNTELL transactions, this Telos source can in rare cases fail to be told by a single transaction. An example is the specification of the ERD model in <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/188651>. The university model (and then university data) depend on rules defined in the model ERD-Semantics, specifically the semantics of the `ISA` type (instances of subclasses are also instances of superclasses). If you tell the university model and the ERD semantics in the same transaction, then the rules for the `ISA` type are not yet usable for checking the consistency of the university model. A way out is to store the university model as sub-module of the module that contains the ERD semantics, and university data as sub-module of the university module.

A way out of this dilemma is to use the module feature of ConceptBase in a thoughtful way. In this case, define a module like `ERnotation` to which the ERD notation is told. Within this module, define a submodule like `UModel`, to which the example ER model is told. Finally, define a submodule `UData` to hold the sample data. By nesting the submodules in this way, the sample data objects can "see" the definitions of the example ER model. And the example ER model can "see" the objects of the ER notation.

5.9 Purging a module

The builtin query `purgeModule` attempts to delete all propositions of the current module or the module that is specified as a parameter. The operation fails if the module contains a non-empty submodule. Further, the operation may non be applied to the `System` module.

The operation requires that the user has appropriate permissions on the module to be purged.

The query `purgeModule` is declared as a hidden object. Hence, it shall not show up in the 'Display queries' dialogue of CBIva. The current implementation is experimental. An alternative is to list the contents of the current module and then applying the `UNTELL` operation to the whole content.

5.10 Saving and loading module sources

The ConceptBase server (see section ??) has two command line options `-save` and `-load` to synchronize with the content of database modules as Telos sources files in the file system of the CBserver. The purpose of this feature is to allow an easy way to save/reload the complete content of a database using a readable source format. These sources can be modified with a regular text editor.

The source files generated by the save function include a `set module` directive that instructs the CBserver to load the file to the original module path when the load function is invoked (or when the file is loaded manually via the `load model` function of CBIva (see section ??). There are two ways to represent saved/loaded module sources in the file system:

flat directory structure: If the module separator is set to `'-'` (see option `-ms` in section ??), the files names consist of a module path starting with the root module name `System` followed by the file type `sml`. For example, the file `System-oHome-AB.sml` will hold the contents of the module `AB` that is a submodule to `oHome` that is a submodule of `System`.

deep directory structure: If the module separator is set to `'/'`, then the module sources are placed in sub-directories named like the modules. For example, the file `System/oHome/AB/AB.sml` will hold the contents of the module `AB` that is a submodule to `oHome` that is a submodule of `System`.

The save function is activated when the CBserver option `-save savedir` is specified. The source files are saved in the directory specified by the `savedir` parameter. This parameter must be the path to an existing directory. If activated, the save function is invoked upon the following events:

1. The CBserver is shutdown. In this case, the complete module tree starting from root module `System` is saved.
2. A client disconnects from the CBserver. Here, the module tree starting at the home module of the user associated to the client tool is saved.
3. A user changes the module context. In this case, the old module is saved.

In all cases, the save function is executed with the rights of the user who started the CBserver. The save function requires at least read permission for the module to be saved. The above rules are also applicable to the server-side materialization of query results, see section ??.

The load function gets activated when the option `-load loaddir` is specified. The directory `loaddir` should contain files with file/directory names being formed as explained for the save function. It loads the files in *alphabetic order* to control the sequence in which the files are loaded. If you manually add Telos

source files to a directory that is about to be loaded, then make sure that its file name is alphabetically sorted after the module file name, e.g. `AB_01extension.sml` is loaded after `AB.sml`. The import is executed once at CBserver startup with the rights of the user who started the CBserver. If a file contains an error, the loading of this module source fails. Error messages shall be displayed in the trace log of the CBserver. Note that the CBserver can be started with a non-empty database. The import of source files will be added to the already existing content of the database.

Examples:

```
cbserver -d DB1 -save /home/meee/DB1SRC
```

This command starts up a CBserver that will eventually save the module sources in the specified directory. Note that the saving takes place either at CBserver shutdown or when a client tool disconnects (partial save).

```
cbserver -u nonpersistent -save /home/meee/SRC
```

This variant will start a CBserver with a non-persistent database but the contents will nevertheless stored as Telos sources file in `/home/meee/SRC`.

```
cbserver -u nonpersistent -load SRC1 -save SRC2
```

This command starts a the CBserver (i.e. only system objects are defined) and then loads module sources from the directory `SRC1`. Then, client tools may modify the contents of the database. Finally, the module contents are saved in directory `SRC2`.

```
cbserver -d DB1 -load /home/meee/DB1SRC -save /home/meee/DB1SRC
```

The load and save directories may also be the same. Note that the save function will eventually overwrite the files that have been loaded at CBserver startup.

```
cbserver -u nonpersistent -load DB1SRC
```

This command will start a non-persistent CBserver and loads the module sources of directory `DB1SRC`.

Module sources do not contain the historic states of objects that are maintained with rollback times in the CBserver database. Hence, a persistent database is containing more information than the saved module tree and is also faster to start up compared to loading module sources. Still, the load/save function offers a simple way to keep a textual representation synchronized with the evolving database state, or to back-up/re-load a database state. The CBserver parameter `-db` combines the function of `-d`, `-load`, `-save`, and `-views`. Hence, all files will be accessed/updated in the database directory.

5.11 Server-side materialization of query results

Similar to the saving of module sources, the CBserver parameter `-views` enables the materialization of certain query results in the file system of the CBserver. To do so, one has to specify the queries to be materialized. Only queries with a single parameter or with no parameter can be materialized. The queries need to be listed with the module that contains the objects that match the query. Example:

```
MyModule with
  saveView
    v1: Q1;
    v2: Q2
end
```

You can also use deductive rules deriving the values for the `saveView` attribute.

The queries `Q1` and `Q2` need to be visible in the module `MyModule`. The queries need to have an answer format (section ??) defined for them (attribute `forQuery`). Assume that the query `Q1` has the single parameter `param:C1`. `ConceptBase` will then call the query `Q1[x/param]` for each instance `x` of class `C1`.

The result is stored in a file with name `x` in the directory specified with the `-views` parameter. If the view is extracted from a module different to `oHome`, then the filename includes the module name as a prefix. The file type of the file is taken from the optional `fileType` attribute of the answer format of `Q1`. The default file type is `".txt"`. The result of queries with no parameter is stored in files carrying the name of the query. If the module separator is set to `'/'`, then the files are stored in sub-directories that reflect the module path, from which the view was extracted, see also section ??.

The materialization of query results is executed at the same events when the saving of module sources takes place (section ??). To enable the materialization, you need to specify the target directory with the `-views` option:

```
cbserver -d MYDB -views /home/meee/MyViews
```

You can also use the `CBSHELL` utility (section ??) to extract the query results and materialize them on the *client side*. This method is more flexible but you need to program the `CBSHELL` scripts for to extract all required views. For example, the `CBSHELL` script

```
connect alpha 4001
ask Q1[x/param] OBJNAMES default Now
showAnswer
exit
```

connects to the `CBserver` running on a host named `alpha` with port number 4001 and will extract just the answer to `Q1[x/param]`. If there are more answers to be extracted, one has to employ separate scripts for each of them and execute them one after the other to save the results in separate files. The *server-side* method using the `-views` option will determine all possible fillers `x` for the parameter `param` and automatically save the results of `Q1[x/param]` in a separate file with filename `x`.

You can use the `-db` option for activating the saving/loading of module sources and the materialization of query results within the database directory:

```
cbserver -db MYDB
```

This creates a single directory with both the binary database files, the module sources, and the materialized query results. Further examples are available from the `CB-Forum` at <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/3097259>.

5.11.1 Post-export command

`ConceptBase` can only export textual query results. Some output formats such as program source code can be further processed by calling appropriate tools. This can be initiated manually, or one can configure the directory specified in the `-views` option with a command file `postExport.sh`. This command file shall then be executed by the `CBserver` whenever a saving operation on the views directory has taken place.

You have to regard that the command is executed in the context of the `CBserver`. Hence, it is executed in the directory in which the `CBserver` was started. You should therefore change to the views directory inside the post-export command file like shown in this example of `postExport.sh`:

```
#!/bin/sh
cd `dirname $0`
cp *.xml /home/meee/public
```

The second line changes to the views directory. The third line does the specific processing of the materialized query results. Here, we just copy the file. More interesting are calls to transformation routines.

You should remove the write permission for the post-export command to prevent that it can be overwritten² by the materialization function. Under Unix/Linux, this can be achieved via the command

```
chmod u-w postExport.sh
```

The above command only has to be executed once within the views directory. If you use the module separator '/' (i.e. the deep structure explained in section ??), then the view files are stored in the directory of the module they belong to. This also may effect the way how you program the postExport script file. In Unix/Linux, you can use the 'find' command to fetch all files that are subject for post-processing:

```
#!/bin/sh
cd `dirname $0`
xmlfiles=`find . -name "*.xml"`
for file in $xmlfiles; do
  cp $file /home/meee/public
done
```

This script also works fine in the case of a flat view directory structure.

²Overwriting the post-export command file may be a desired feature. You can then generate it from within the CBserver like any other file in the views directory. This is however a major security hole and we strongly discourage to use this feature.

Chapter 6

The ConceptBase.cc Server

The ConceptBase.cc server (CBserver) offers its services via a TCP/IP port to client programs. The main services are to TELL or UNTELL O-Telos objects and to ASK queries to the database. The operations are called by clients (for example, the user interfaces described in section ??). An arbitrary number of clients can connect to a CBserver.

The CBserver is started¹ by a command line

```
cbserver <params>
```

assuming that the installation directory of ConceptBase.cc is added to the search path of executable programs. If it is not on the search path, then simply change the current directory to the installation directory of ConceptBase.cc or use its absolute path of the cbserver script.

6.1 CBserver parameters

The following parameters are available for the 'cbserver' command:

- d dbdir** Set database `dbdir` to be loaded. If the database does not exist, it is created and initialized with the O-Telos pre-defined objects. The database is maintained as a directory. Setting the database is mandatory except when the update mode is set to `nonpersistent` (see below). You cannot start two concurrent servers which use the same database directory. To avoid this case, a file `OB.lock` is created in the database directory when the first server is started. If the server crashes during its execution, the file `OB.lock` will still exist in the directory. Before you restart the server, you might have to remove this file manually.
- db dbdir** Like `-d` but also sets the load/save/views directories to `dbdir`, i.e. the CBserver will automatically maintain the module sources in `dbdir` and also materialize the selected queries in the same directory. See section ?? for details.
- new dbdir** Like `-d` but deletes any existing database at location `dbdir` before it creates and initializes it.
- u updatemode** controls update persistency. The allowed values are `persistent` and `nonpersistent`. If no database is provided by parameter `"-d"`, then the default update mode is set to `nonpersistent`. Otherwise, the default is `persistent`. In `nonpersistent` mode, all updates are lost after the ConceptBase server is stopped. In `persistent` mode, updates are stored in the files of the database and will be available for future sessions.

¹You can also start the CBserver from within the ConceptBase.cc user interface CBIva. Details are in the installation guide distributed with ConceptBase.cc and in section ??.

- U untellmode** controls the way how UNTELL is executed by the server. The allowed values are `verbatim` and `cleanup` (default). In `verbatim` mode, the UNTELL operation will only untell the facts directly described by the O-Telos frame being submitted as argument. In `cleanup` mode, UNTELL will also try to remove the instantiation to the O-Telos system classes `Individual`, `Attribute`, `InstanceOf` and `IsA`. By doing so, UNTELL behaves inverse to the TELL operation. More details are explained in subsection ??.
- port portnr** sets the TCP/IP socket portnumber for client connections to the CBserver. The value `portnr` must be between 2000 and 65535. If there is already a process using the portnumber, the CBserver will abort. The default value for the portnumber is 4001.
- p portnr** is the same as "-port portnr". Deprecated since it conflicts with a predefined command line parameter of SWI-Prolog.
- version** display version info and exit.
- help** display list of CBserver options and exit.
- license** display license and exit.
- t tracemode** sets the tracemode of the CBserver. It is one of `silent`, `no`, `minimal`, `low`, `high`, `veryhigh`. The tracemode determines the amount of text displayed by the server during its execution. The tracemode does not influence the function but is used for debugging. The default tracemode is set to `no` (only display CBserver interface). The tracemode `low` will configure the CBserver to trace the CBserver interface calls plus answers, and the tracemode `no` virtually disables tracing. The tracemode `silent` is even suppressing the message 'CBserver ready' when starting up the CBserver. The tracemode `high` and `veryhigh` are useful for debugging the system. In these two modes, an unlikely fatal signal like division by zero will not directly abort the CBserver process but start a debug dialog. Enter "h" for options to diagnose the problem in collaboration with the ConceptBase developers.
- c cachemode** turns on the query cache to allow recursive query evaluation. The value `cachemode` is one of `off`, `transient`, and `keep` (default). In `transient` mode the cache is emptied before each transaction. In `keep` mode, the cache is emptied when the maximum number of entries in the cache is exceeded or an update has invalidated the cache. Further details are explained in section ??.
- o optmode** controls the optimizer for rules, constraints and queries. The value `optmode` is one of 0 (no optimization), 1 (structural optimization by exploiting builtin O-Telos axioms), 2 (optimizing join order), 3 (combines 1 and 2), or 4 (combines 1 and 2 with trigger pruning). Default and recommended is 4.
- r secs** automatically restarts the CBserver after a crash, or when it was started with option `-sm slave` and the last client exits. The value `secs` specifies how many seconds to wait before restart. You may want to use this option in a multi-user setting, where the CBserver runs on a different machine than the user clients. The `-r` option must be handled with great care since it can easily lead to an infinite loop of restarts, e.g. when a database file is corrupted. In such cases you might have to reboot the whole computer!
- s securitylevel** configures the access control mechanism of ConceptBase. The value 0 means that no access control is employed. Any user can ask, tell, untell, retell any object in any module. You can also untell objects defined in a super-module. Level 1 (default) provides a very basic protection: one can only untell objects if they are defined in the current module. This prevents in particular undesired deletions of objects defined in the `System` module. Level 2 fully enables access control. First, untelling of objects must happen in the module where the object has been defined. Second, any transaction submitted by a user to the CBserver is checked against the permission rules as defined in section ??.
- Level 3 enables at most read access to a module. In addition, the permission rules must allow read access. This level makes sense if you want to freeze a database state. Enable

access control when you use ConceptBase in a multi-user setting and you want to avoid erroneous interferences between different users.

- e maxerrors** sets the maximum number of errors to be displayed to a ConceptBase.cc client within one transaction. The value -1 means that no restriction is applied. Set to 0 to suppress any errors message and to a positive number to limit the displayed errors messages to that number. A low positive number can speed up the communication between ConceptBase client and server if a lot of error messages are generated. The default is 20.
- cc ccmode** (predicate typing) controls to which extent the CBserver applies strict typing of attribution predicates ($x\ m\ y$) occurring in the membership constraints of query classes. If the mode is set to `strict` (=default), attribution predicates without a unique concerned class² shall not be accepted. If the mode is set to `extended`, the search for concerned classes shall include subclasses (see section ??). If the mode is set to `off`, ConceptBase.cc also accepts queries with unstrictly-typed attribution predicates. The strict mode is preferable since it avoids certain semantic errors. Deduction rules and integrity constraints may never violate the predicate typing condition, even if the mode is set to `off`. An example for a query using non-strict predicate typing is available from the CB-Forum, see <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/1270138>.
- mu mumode** (multi-user mode) specifies whether the CBserver should run in multi-user mode (value `enabled`) or in single-user model (value `disabled`). By default, the multi-user mode is enabled, allowing multiple users with different user names to connect to the CBserver. In single-user mode, only clients started by the same user (identified by her name) can connect to the CBserver. The single-user mode is recommended if you want to block other users from logging into your CBserver. Since the test is done only on the user name, a malicious attacker could use your user name for an account on another computer and then successfully log into your CBserver. Use Internet firewalls to protect against such attacks. If you specify an administrator user (option -a), then this user can always connect to the CBserver.
- v vmode** controls whether view maintenance rules are generated (`vmode=on`) or not (`vmode=off`). View maintenance rules are used to keep a ConceptBase.cc view up-to-date upon changes to the object base. Default value for `vmode` is `off`.
- mc maxcost** this parameter defines the maximum cost level for a predicate in a binding path that is used to compile a meta formula (see section ??). The evaluation of a binding path yields fillers for the meta variables. Set `maxcost` to 10 if such a predicate should have about one free variable. Set it to 100 if it may have two free variables. Default is 100. The higher the number, the more candidate paths are generated, increasing the likelihood that a binding path is found. On the downside, a high value increases the compile time of meta formulas.
- pl pathlen** sets a maximum length for binding path candidates. In principle, the number of candidates can explode with the path length. Like the previous parameter, the path length influences the ability of ConceptBase.cc to compile meta formulas. The default value is 5. If you set the value to 0, then no meta formula can be compiled.
- im imax** sets the maximum number of iterations used to re-order attribution predicates with one free variable. Evaluating such predicates first can lead to faster elimination of free variables and thus lead to better query and ECA performance. The default value is 3.
- eca emode** controls the ECA sub-system. Possible values are `unsafe` (ECArules are evaluated without safeguarding recursive deductive rules), `off` (ECArules are not evaluated, even if some are defined), and `safe` (ECA rules are evaluated with safeguarding recursive rules; this is the default). Use the mode `unsafe` if none of your ECArules calls recursive predicates on the newest database state. This may lead to a limited speed-up.

²The concerned class is a consequence of the predicate typing condition of section ?? . You can roughly compare it to typing of variables in programming languages.

- eo eomode** controls the optimization of conditions of ECArules. Possible values are `on` (default) and `off`. The optimization is done by a re-ordering of predicates in the condition. Hence, you only want to turn off optimization to gain full control over the order of evaluation in ECA conditions.
- load dir** specifies the directory from which the CBserver will load module sources at start-up time. The module sources must have file names starting with `System` and file type `sml`. Typically, they are generated via the `-save` flag in the preceding session of the CBserver. The default is `none`, i.e. no module sources are loaded at CBserver start-up.
- save dir** specifies the directory into which to save certain textual excerpts of the database, in particular module listings. The parameter has the default value `none`, which disables the saving function. Currently, the CBserver only saves module listings. Each module is stored in one file with file type `sml`. The directory `dir` must exist. The module listing is performed when the CBserver is shut down (complete module tree is listed), or when a client tool logs out (home directory tree of the client tool is listed). The module listings uses the `set module` directive to enable the import of the file to the right module location. See also section ?? for details.
- views dir** specifies the directory into which the results of certain queries are materialized. See section ?? for details.
- ms sep** specifies the module separator to be used for saving module listings and views. If the separator is set to `'-'` (default), then all module sources and views are stored at the top level directory. If the module separator is set to `'/'`, then the files are stored in a deep sub-directory structure that mirrors the module structure.
- mg mgmode** specifies whether module listings are generated with separators `"{---}"` for each transaction occurring in the module (option `split`) or without such separators (option `whole`). The default is `split`. The `split` option better supports cases where metaformulas are defined and used in the same module.
- rl rlmode** controls the way how the CBserver creates labels for generated formulas. The default value is `on`, instructing the CBserver to find a readable label for the generated formula. It typically consists of the labels of the participating metaclass attributes occurring in the metaformula. If set to `off`, the CBserver will just take a system-generated label that contains a unique identifier. This is slightly less readable (if you want to inspect the generated formulas) but safe against certain possibilities of assigning the same label twice.
- ia tmax** sets the maximum number of hours during which a client should interact with the CBserver to be regarded as active. Negative values are interpreted as 'infinity'. This parameter is only used when a CBserver uses the `-sm slave` and `-r` options, or the `-g public` option. The default value for `tmax` is 2.0 hours.
- sm servermode** sets the server mode. Possible values are `master` (default) and `slave`. In slave mode, the last client that leaves the CBserver will also shutdown the CBserver, provided that the CBserver and the client were started by the same user. This option is useful when a CBserver is only needed while still clients are registered. A master CBserver must always be stopped explicitly.
- st stratmode** enables or disables the rule stratification test. Possible values are `on` (default) and `off`. If enabled, then the query evaluator shall dynamically test whether stratification violations occur. They shall then be reported as an error. Disable the test, if you are sure that the answers are correct even though a stratification violation occurs.
- g cmd** provides a special command to the CBserver. There are currently three such commands. The command `nolpi` instructs the CBserver to ignore any plug-in file (see section ??). The command `public` configures the CBserver as a public CBserver (see ??). The command `exit` instructs the CBserver to exit immediately after start-up. This can be useful to combination with the option `-views`, `-db` and `-save` to materialize some excerpts from a stored database.

-a user designates `user` as 'administrator' of this CBserver. For the time being this just gives the right to shutdown the CBserver. By default, the user who started the CBserver is its administrator. This user shall also keep the right to shutdown the server, even when another user is the designated administrator. If you specify the user name with host, e.g. `billy@myhost`, then only the user `billy` on host `myhost` is recognized as additional administrator.

If a CBserver is started without any parameter, then the update mode shall be set to `nonpersistent`, the trace mode to `no`, multi-user mode is disabled, and the server mode to `slave`. The other parameters are set to their defaults. Such a CBserver is useful as companion of tools that need it only while they are running.

```
cbserver
```

A ConceptBase client running on the same computer will then connect to this CBserver, when it uses 'localhost' as host and 4001 as port number. Since the CBserver runs in slave mode, it will shut down when its client disconnects.

6.2 ConceptBase under Windows 10

The CBserver is only compiled for Linux architectures. This means that user of other platforms must rely on a Linux system to utilize ConceptBase. The traditional way is to start the CBserver on such a Linux system and then connect to it, possibly using a so-called public CBserver (see section ??). Since April 2017, this detour is no longer required for users of Windows 10 (64bit, Creators Update). This version of Windows is capable to let Linux programs run under the 'bash' utility, which is basically a whole Linux system under Windows that realizes calls to the Linux API by hooks to the Windows API. Hence, it is not a virtual machine, it lets you run the Linux (64bit) variant of the CBserver natively on Windows 10.

To enable the Linux capability on Windows 10, follow the instructions at <http://conceptbase.sourceforge.net/CB-WinLinux.html>. Note that you must have installed Java (64bit) on your Windows machine, not Java (32bit) to take full advantage of this feature. You can check whether your Java is 64bit by calling the following command in a Windows command window.

```
java -d64 -version
```

Users of older Windows versions and of other operating systems can continue to use the public CBserver (see section ??) to take advantage of ConceptBase.

6.3 Database format

A ConceptBase database is a directory that contains at least the following files:

- **OB.symbol**: a binary file that associates object names (like 'MyClass') with object identifiers.
- **OB.telos**: binary file storing all propositions
- **OB.rule**: text file containing the generated Prolog code for rules, constraints, and queries
- **OB.ruleinfo**: text file containing argument information about queries and some information for the cost-based formula optimizer
- **OB.ecarule**: text file containing the generated Prolog code for active rules

The database files may only be updated via the ConceptBase server. Their initial state is bootstrapped from textual Telos frames that define the pre-defined objects of O-Telos. Since the pre-defined objects can change from version to version, we cannot guarantee binary compatibility of ConceptBase databases. You can easily export the textual definitions from a database via the `-save` option. Those definitions can then be imported to the new ConceptBase version. The database directory may contain further text files with filetype 'lpi'. These are Prolog plugins loaded at startup time, see also section ??.

6.4 Modifying the system database

A new database is created from the database `lib/SystemDB` in your ConceptBase installation directory. The System database contains exactly the objects of the root module `System`. They include for example the definitions of the objects `Proposition`, `Individual`, `Attribute`, `InstanceOf`, and `IsA`. Further the objects `Class`, `QueryClass`, `Function`, `ECARule`, `Module` and many more are defined that are needed to formulate queries and to use the capabilities of the system.

Whenever a new database is created, the files from this System database are copied into the new database directory. This allows experienced ConceptBase users to adapt the System database to their needs. Just start a ConceptBase server with

```
cbserver -d $CB_HOME/lib/SystemDB -s 0
```

Replace `$CB_HOME` by the path to your ConceptBase installation directory. Then start a CBIva user interface, connect to the CBserver and switch to the module `System`. Assume you want to predefine the class `Container` and declare a `Model` as subclass of `Container`:

```
Container with
  attribute
    contains: Proposition
end
Model isA Container end
```

You can also add rules and constraints about containers, e.g. that containers may not contain themselves. A more significant extension would be to add active rules to the system database. For example, the active rules in CB-Forum at <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/3260276> changes the semantics of the UNTELL operation. Such definitions are subsequently included in any new database that you create. Be careful with deleting system objects. The code of the CBserver relies on the existence of certain system objects.

6.5 Tracing and restarting

The trace of the CBserver can be saved by redirecting its output, e.g.

```
cbserver -r 10 -port 4444 -t high -d MYDB >> mylogfile.log
```

The CBserver can also be started directly from the ConceptBase.cc User Interface (see section ??) and most parameters can be specified interactively. The command line version is recommended when one CBserver serves multiple users or when user interface and server shall run on different machines. The parameter `-r 10` instructs the CBserver to restart after 10 seconds if a crash has occurred.

A special error message during the startup of the CBserver is the following:

```
### FATAL ERROR:
Application is locked by hostname, PID 1234
### CBserver aborted
```

This message is printed if there is still a file with the name `OB.lock` in the database directory (option `-d`). The `OB.lock` file should avoid that two servers are using the same database directory. The file may be left over of a previous CBserver if the server was not stopped correctly (e.g. aborted by Ctrl-C or it crashed). If you get this error message, make sure that there is no other server running that uses this directory and then delete the file `OB.lock`. Then, the CBserver should start correctly.

6.6 Public CBservers

A public CBserver is a ConceptBase server that is accessible from the whole network. As such this is a property that any CBserver has, except when the multi-user capability is disabled, or when your firewall prevents external access.

If you work with an existing database, you may want to specify some access rights rules like suggested in section ???. We neglect them in this simple example. Now, start the public CBserver with suitable parameters under Linux/Unix:

```
cbserver -r 2 -a jonny -g public -ia 0.5 -u nonpersistent -d MDB &> log.txt &
```

The option `-g public` enables the slave mode implicitly and tells `oHome` as an instance of `Auto-HomeModule`. The combination of the slave mode and the option `-r` instructs CBserver to stop and restart when the last client exits. The update mode is nonpersistent. As a consequence, the restarted CBserver will use the unchanged state of `MDB`. This is useful, if you want to provide the service of CBserver to a larger group of (anonymous) users. They can log in with a client, operate on the database in nonpersistent mode and eventually leave the CBserver. When the last active client³ leaves the CBserver, then CBserver will freshly start up after 2 seconds. Due to the option `-u nonpersistent` user-defined objects are only stored at the public CBserver while there are still active clients enrolled to the public CBserver.

The parameter `-a` sets the administrator user of the public CBserver. This user is allowed to shutdown the CBserver from a client. The auto home feature will assign different users to individual workspaces. Unless you introduce access rights (and enable them via the CBserver option `-s 2`) the users can also manipulate the modules of other users. However, the CBserver is restarted whenever the last client logs off. Hence, the definitions of different users are not permanently stored on the CBserver. The parameter `-ia` used here instructs the CBserver to regard a client as active if it had its last interaction within 0.5 hours. Clients that were inactive for a longer time will not prevent the CBserver from restarting when another (active) client logs off the CBserver.

Consider the following alternative command to start a public CBserver:

```
cbserver -t no -r 2 -a jonny -g public -ia -1 -d ALLDB &> /dev/null &
```

Here, the updates are stored persistently in the database `ALLDB`. Changes to modules are not lost when the last active client leaves the CBserver. There is no log file created as well. The option `'-ia -1'` used here instructs the CBserver to regard any client as active, regardless of how long ago its last interaction occurred.

Configure the `CBIva` interface to use the public CBserver via the variable `PublicCBserver`. It can be set by the menu item "Options/Edit Options Manually" of `CBIva`, see section ???. A value different from `none` enables the use of the public CBserver by `CBGraph`. You can optionally append a port number like `"cbserver.acme.com/4002"`. The default value for the port number is 4001. Installations of ConceptBase on platforms, for which no binaries of the CBserver exist, may use a default public CBserver. This is the case OS-X and older versions of Windows⁴.

Don't forget to save the options after changing the variable.

The best way to interact with a public CBserver is to use graph files, see section ???. Assume that a public CBserver is running on host `cbserver.acme.com`. and that `cbserver.acme.com` was configured as the public CBserver to be used. Then, calling

```
cbgraph graph1.gel
```

will connect to the public CBserver instead of 'localhost'. The port number for the connection is taken from the graph file. Do not forget to save the graph file before exiting `CBGraph` if the public CBserver was started in non-persistent mode. If you subsequently open the graph file, it will attempt to connect to the same host. You can force it to attempt the connection to localhost instead by

³An active client is a client whose last interaction with the CBserver was less than a certain number of hours ago, specified with the CBserver parameter `-ia`.

⁴Windows 10 (64bit) Creators Update has the ability to activate a Linux sub-system which can then be used to start CBserver transparently. If you had previously used a public CBserver under Windows 10 and now want to utilize a local CBserver, then reset the variable `PublicCBserver` to `none`.

```
cbgraph -host localhost graph1.gel
```

If you are using ConceptBase under Linux, then the CBserver is running by default on your local computer ("localhost"). The same is true for Windows 10 with an enabled Linux subsystem. Users on Mac computers or older Windows versions are by default using a public CBserver at the university where ConceptBase is developed. This public server is meant for testing the system. **DO NOT USE it for managing confidential information!** If you plan to use ConceptBase in a serious way, we recommend that you set up a protected Linux server on your own network that runs the CBserver. Users can set up the variable "PublicCBserver" in CBIva to the address of the protected Linux server to automatically connect to that server, see section ???. You can also connect to that server manually via the "Connect" function of the ConceptBase clients CBIva, CBGraph, and CBSHELL.

6.7 The tabling subsystem

Since version V5.2.4 ConceptBase.cc features a new query evaluation method, which uses a so-called tabling cache to store intermediate results of predicates that are called during the top-down (SLDNF) query evaluation. Assume, for example, that an employee 'bill' has two projects 'p1' and 'p2'. Then, the result of a predicate '(bill hasProject x)' with variable x would be the set {(bill hasProject p1), (bill hasProject p2)} consisting of facts. We call this fact set also the *extension* of the predicate.

After a completed predicate evaluation, the tabling cache of the predicate holds its extension. Tabling speeds up query evaluation and prevents infinite loops when ConceptBase.cc evaluates recursive queries and deductive rules. Essentially, the tabled evaluation allows to compute dynamically stratified semantics of the Datalog database underlying ConceptBase.cc. Plenty of examples for recursive rules and queries are provided in the online ConceptBase Forum.

The CBserver provides three tabling cache modes to control the behavior during query evaluation:

- c off** In this mode, the cache is completely disabled. Use this mode when your models do not include recursive rules. The mode is only provided for backward compatibility and has no advantages.
- c keep** The cache is only emptied when necessary, in particular when the cache has been invalidated by an update to the database, or when the maximum number of facts in the cache is exceeded. The maximum number is currently set to 20000. Exceeding the maximum is not an error. It only indicates that the cache is marked for being emptied. If necessary, the cache emptying takes place before a transaction. The keep mode is on average consuming more main memory than the transient mode but speeds up response time enormously in case of re-use of query results. The 'keep' mode is the default mode for tabling.
- c transient** The tabling cache is emptied before each transaction (ask, tell, untell, retell). A subsequent query is always evaluated starting with an empty cache. This mode is somewhat 'safer' than the 'keep' mode since it starts each query with an empty cache state. While the answer to a query is in principal independent from the cache mode, the cache mode has a certain influence on the persistence of objects created within a transaction. Specifically, results of arithmetic expressions computed during one transaction shall be removed after the transaction when the cache mode is 'transient'. In cache mode 'keep', these objects continue to exist and are visible to future transactions.

ConceptBase will only call tabled evaluation for deductive predicates. Other predicates are evaluated by the regular SLDNF engine of the underlying Prolog engine. By default, the tabling cache mode is activated in mode *keep*. Some statistics on cache usage are written to the terminal window of the ConceptBase server when the tracemode has been set to *veryhigh*.

Acknowledgements: The techniques for the tabled query evaluator of ConceptBase.cc are inspired by the 'tabled evaluation' [?, ?]. We do however not delay the evaluation of negated predicates but rather re-order them at compile time to guarantee that all variables are bound at call time. Tabled evaluation is also implemented in XSB [<http://xsb.sourceforge.net/>] and DES [<http://www.fdi.ucm.es/profesor/fernandes/>].

6.8 Database persistency

The default update mode is 'persistent'. In persistent mode, all changes to the database are written to the file system at the directory specified in the parameter '-d'. Persistent mode is suitable when a CBserver runs for a longer period of time and is directly updated by application programs. If ConceptBase.cc is used for testing and modeling purposes, the update mode 'nonpersistent' is an interesting alternative. We discuss two scenarios for the nonpersistent mode and one for the persistent mode.

Scenario 1: Single-user modeling. When a user needs to model a certain application domain with classes and meta classes, he usually works with external Telos files (aka source models, file extension '.sml'). These files can include comments like usual with program source code. The recommended mode here is '-u nonpersistent' without specifying a database. The user can load the source models into such a non-persistent server and make corrections to the source files in case of errors or design changes. Here, ConceptBase is mostly used to check and analyze the models. Recommended options:

```
cbserver -u nonpersistent -mu disabled
```

Scenario 2: Lab assignments. Assume that a teacher wants students to exercise a certain modelling task using ConceptBase.cc. Then, he would prepare some Telos files with necessary definitions (e.g. some meta classes) and load them into a persistent ConceptBase server. After that, he can restart the ConceptBase server in non-persistent mode on the same database created before. Student can then work on their extensions while the state of the database can easily be set back to the original state defined by the teacher. The module system of ConceptBase.cc can be used to support multiple students to work on the same server without interfering with each other, see section ???. Recommended options:

```
cbserver -d MYDB -s 2 -mu enabled -u nonpersistent
```

The second scenario might also be useful in modeling. If there are some parts that are regarded as stable, the modeller can decide to make them persistent and only add/modify those Telos models that are still subject to change. In particular for large Telos models, this strategy saves time. Note that the updates by the users are lost when the non-persistent CBserver is stopped. This might be useful, if you want to re-use the same initial database MYDB several times, e.g. for different user groups.

Scenario 3: Project work. Here the students work for several days on a given task. Changes shall not be lost. The use of the -db option will not only store the database in binary form but also store the source code of all modules as text files in directory MYDB. Recommended options:

```
cbserver -db MYDB -s 2 -mu enabled
```

If ConceptBase.cc is used in a multi-user setting, then one can combine the update mode with the module feature (see section ??). In this scenario, multiple users access the same CBserver. A common super module (e.g. the module oHome) carries the common objects of the users. Each user can be assigned to her own hown module (a sub module of the common super module) and create and update objects in this workspace without interfering with other users. If several groups of users shall share their definitions, then they would be assigned to the same home module. The home module may have sub modules for testing and releasing definitions. By employing access rights to modules, one can also design which user has which read/write permissions. The builtin query listModule allows to save the contents of a module to a Telos source file (see section ??).

6.9 The UNTELL operation

ConceptBase.cc realizes the concept of a historical database. The TELL operation submits O-Telos frames to the CBserver. The CBserver extracts the 'novelty' of the submitted frames and translates it into a set of P-facts to be stored in the object store. Any P-fact has a so-called *belief time* associated to it (see section ??). The belief time is an interval (t_1, t_2) whose left boundary t_1 is the time point when the P-fact was

inserted to the object store, i.e. the time when the transaction was executed that led to the insertion of the P-fact. The right boundary t_2 specifies the time point after which the CBserver assumes the P-fact to be not true anymore. It is initialized with 'Now' when the P-fact is created. This symbolic value is interpreted as the current time. You may also interpret such a time interval to be right open.

The UNTELL operation terminates the belief time of P-facts specified in an O-Telos frame. The value 'Now' is replaced by the time⁵ when the UNTELL operation is executed.

From the user's perspective, a TELL operation is about creating some objects and the UNTELL operation is about deleting them⁶. Many users expect the UNTELL operation to be symmetric to the TELL operation, i.e. untelling a frame that has been told before should remove the frame completely. This is however not the case for the following reasons:

- An O-Telos frame being argument of a TELL or UNTELL operation is not necessarily all the information about an object but just some.
- Other objects might refer to an object told previously. An UNTELL operation would then be rejected to preserve referential integrity.
- ConceptBase.cc adds instantiation to the builtin classes `Individual`, `Attribute`, `Instance-Of`, and `IsA` depending on the type of the object.

The last reason is most significant in preventing symmetry. As an example consider the O-Telos frame (referred to as frame 1).

```
bill in Employee with
  name
  bname: "William"
end
```

ConceptBase.cc will recognize that `bill` is an individual and that `bill!bname` is an attribute. This information is attached internally to the P-facts, more precisely, it is derivable from the structure of the corresponding P-facts. Assume that you started a CBserver with `untell model verbatim` (see below). When you ASK the CBserver for the frame of `bill` after the TELL operation on frame 1, you will get frame 2:

```
Individual bill in Employee with
  attribute, name
  bname: "William"
end
```

Hence, the instantiation to the builtin classes `Individual` and `Attribute` is added to the frame. If we submit the original frame 1 to an UNTELL operation, ConceptBase.cc assumes by default that only two facts should be untold:

1. The instantiation of `bill` to `Employee`.
2. The instantiation of the attribute `bill!bname` to its attribute category `Employee!name`.

As a consequence, the object `bill` and its attribute continue to exist after the UNTELL on frame 1. It would look like (frame 3):

⁵ ConceptBase takes the system time using timezone Coordinated Universal Time (UTC) of the computer on which it is running and rounds it to milliseconds. The time is captured when the transaction is initiated, i.e. all P-facts told or untold in the transaction will use that transaction time.

⁶ 'Deleted' objects can however be recovered by setting the so-called roll-back time before an ASK transaction is issued. Only ASK transaction are allowed on historical database state. It makes little sense to update a historical state. That would be 'falsifying' the history of stored P-facts.

```

Individual bill with
  attribute
    bname: "William"
end

```

Only by untelling this frame 3 as a second operation or by untelling the completed frame 2, the object `bill` and its attribute `bill!bname` shall be made historical.

This asymmetry of TELL and UNTELL is regarded by some ConceptBase.cc users as unnatural behavior. They expect that an UNTELL is also supposed to affect the objects themselves, not just their instantiation to classes. To support those users, we provide a so-called untell mode (see parameter `-U` in the list of CBserver command line parameters). The untell mode `verbatim` will cause ConceptBase.cc to behave as explained above. The mode `cleanup` (default) will take care to remove the objects themselves provided that they have no other instantiations to classes except instantiation to the four builtin classes `Individual`, `Attribute`, `InstanceOf`, and `IsA`. Furthermore, no other object may be linked to the object subject to be untold.

6.9.1 Cascading UNTELL

You can use ECARules to realize a cascading UNTELL, i.e. if an object is deleted then all its links to other objects are deleted as well, leading to potential follow-up deletions. The required ECARules are provided in <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/3260276>. You only need to add the ECARules to your database to enact the cascading UNTELL. You should be careful with using the cascading UNTELL. For example, untelling a system class will also untell all instantiations to it. Use the CBserver option `-s 1` to prevent such undesired deletions.

6.10 Memory consumption and performance

ConceptBase.cc stores objects in a dedicated object store maintained in main memory. A P-fact $P(o, x, n, y)$ consumes about 800 bytes of main memory. That means that one can store roughly 1 million P-facts in 1 GB of main memory. A typical Telos frame is stored with roughly 10 P-facts. Hence, 1 GB of main memory allows you to store around 100.000 Telos frames. On 32 bit CPUs, this results in a maximum of roughly 400.000 frames that can fit into 4 GB of addressable main memory. This restriction virtually vanishes with 64 bit CPUs.

A single TELL/UNTELL operation submitted to the CBserver should not contain more than about 2000 frames (at about 5 attributes per frame). Otherwise, the compiler can run out of stack memory.

The raw performance of the object store, i.e. the time needed to reconstruct a frame for a given object identifier, is virtually independent from the number of P-facts that it stores. However, if you have defined many rules or integrity constraints, the performance may well degrade significantly with the number of stored P-facts. The same holds for queries. We tested the response times for standard queries such as computing the transitive closure in relation to varying database sizes. Results indicate that ConceptBase apparently approximates in many cases the theoretic optimum.

The performance of the active rule evaluator (section ??) is currently rather limited. We measured around 100 rule firings per second. This can be a performance bottleneck when many active rules are being processed.

6.11 The Java API to the CBserver

The application programming interface (API) to the ConceptBase server (CBserver) is realized by the Java class `LocalCBclient`. Most (Java) application programmers presumably only need the simple String-based part of `LocalCBclient` to interact with a CBserver. `LocalCBclient` uses socket-based data streams to realize the communication. We define here the methods of this String-based API. If the argument starts with an "s", then the data type is `String`. If it starts with an "i", then the data type is `int`.

```
LocalCBclient cbClient = new LocalCBclient();
```

- This constructor provides the API object `cbClient` to be used for the subsequent method calls.

```
sAnswer = cbClient.cbserver();
```

- This method attempts to start a single-user CBserver in "slave mode" on localhost with port number 4001. The method only works on platforms for which the CBserver was compiled. This is currently limited to Linux and Windows 10 (see also section ??). In other cases, you need to start a CBserver on a host system that supports the CBserver and then use the 'connect' method. If the CBserver was successfully started, the method returns "yes", otherwise "no". If the `cbClient` is already connected to a CBserver, then the method also returns "yes". The CBserver shall automatically shut-down when the last client disconnects from it.

```
sAnswer = cbClient.connect(sHost, iPort, sTool, sUser);
```

- This method connects your Java program to the CBserver specified by `sHost` (the domain name of the computer on which your CBserver runs on) and `iPort` (the port number is an integer; usually 4001). The string `sTool` shall be a self-selected name of your tool, e.g. "ModelerXY" and `sUser` is a string containing a user name (typically your user name). If you use `null` as username, then the `cbClient` will use the login name of the computer user that started the Java program. The return value is "yes" if successful and "no" else. Use the boolean-valued function `cbClient.isConnected()` to check whether `cbClient` is currently connected to a CBserver. If the socket connection breaks down, then the `cbClient` will set the connection status to unconnected.

```
sAnswer = cbClient.connect();
```

- Connects to the public CBserver (see section ??) if configured for your installation. Otherwise, attempts to start a single-user "slave" local CBserver on port 4001 if not started. Connects to the local CBserver with the default tool name "LocalCBClient" and `null` as username. It behaves similar to the "connect" command of CBShell.

```
sAnswer = cbClient.disconnect();
```

- This method disconnects connects your Java program from the CBserver. The return value is "yes" if successful and "no" else. If the CBserver was started in "slave mode" and the `cbClient` is the last remaining client of the CBserver, it will shut down.

```
sAnswer = cbClient.pwd();
```

- This method outputs your current working module path, e.g. "System-oHome".

```
sAnswer = cbClient.mkdir(sModule);
```

- This method creates a new submodule (e.g. "MyMod") in the current module.

```
sAnswer = cbClient.cd(sModule);
```

- This method changes the current module to `sModule`, e.g. "MyMod".

```
sAnswer = cbClient.tells(sFrames);
```

- This method tells (=defines) the objects given by the string `sFrames` to the current module of the CBserver. The return value is "yes" if successful and a string containing user-readable error messages else.

```
sAnswer = cbClient.untells(sFrames);
```

- This method untells (=removes) the objects given by the string `sFrames` from the current module of the CBserver. The return value is "yes" if successful and a string containing user-readable error messages else.

```
sAnswer = cbClient.asks(sQuery,sFormat);
```

- This method asks the query call `sQuery` (given by a String) to the CBserver. The return value is a string containing the answer to the query. The return value is "no" in case of errors, e.g. when the query is not defined. The query call can have arguments, e.g. `"get_object [Class/objname]"`. The CBserver shall use the answer format `sFormat` for the result. Thus, if you want to define your own answer format, then use the facilities described in section ???. The query shall be answered in the context of the current module and the current time ("Now"). The answer "nil" stands for an empty answer.

```
sAnswer = cbClient.asks(sQuery);
```

- Same as `cbClient.asks(sQuery, "default")`, i.e. the CBserver determines the applicable answer format (in most cases: "FRAME").

```
sAnswer = cbClient.clientid();
```

- Return the identifier by which this client is registered to a CBserver.

```
sAnswer = cbClient.clearall();
```

- Attempt to delete all objects from the current module. Returns "yes" if successful, otherwise "no". Calls the 'purgeModule' query described in section ???. Use this method with great care since it wipes out the whole module content.

Below is the listing of the Java program `TinyClient.java` that uses the String-valued API:

```
import i5.cb.api.*;
public class TinyClient {
    private static LocalCBclient cbClient = null;
    public static void main(String argv[]) {
        String sAnswer;
        cbClient=new LocalCBclient();
        sAnswer = cbClient.connect("cbserver.acme.org",4001,"TinyClient",null);
        sAnswer = cbClient.tells("Employee in Class end");
        sAnswer = cbClient.asks("get_object [Employee/objname]");
        System.out.println(sAnswer);
        sAnswer = cbClient.disconnect();
    }
}
```

You need to compile the Java program with the `cb.jar` library. This is available from the ConceptBase installation directory (referred here as `CB_HOME`). To compile the Java program call

```
javac -classpath $CB_HOME/lib/classes/cb.jar TinyClient.java
```

Before running the client, make sure that the CBserver runs on the specified Linux computer (here "cbserver.acme.org") under the specified port number (here 4001), and that this port number is accessible from your client computer. You may want to configure that CBserver as a public CBserver (see section ???) to have a save way to connect your Java program to it.

The Java program can then be started under Linux as follows:

```
java -classpath $CB_HOME/lib/classes/cb.jar:. TinyClient
```

Under Windows, you should use:

```
java -classpath c:\conceptbase\lib\classes\cb.jar;. TinyClient
```

The Java source code for `TinyClient.java`, `TinyClient2.java` (uses the `cbserver` method), `TinyClient3.java` (uses the `connect` method), and a more elaborate example `SimpleClient.java` is included in the directory `examples/Clients/JavaClient` of your `ConceptBase` installation directory.

Chapter 7

The CBShell Utility

The ConceptBase.cc Shell (CBShell) is a command line client for ConceptBase.cc. It allows to interact with a CBserver via a text-based command shell. Moreover, it can process commands from a script file without further user interaction. The utility can be employed to automate certain activities such as batch-loading a large number of Telos models into a CBserver, or to extract certain answers from a CBserver.

7.1 Syntax

There are two ways to use the CBShell. The first one processes the commands from a script file (batch mode). The second one prompts the user for commands (interactive mode).

```
cbshell [options] scriptFile [params]
cbshell [options]
```

7.2 Options

- l** This options instructs CBShell to write errors and some statistic information to the files error.log and stat.log.
- f scriptFile** Execute the commands specified in scriptFile rather than requesting commands from the command line interface. If the -f option is used and a scriptFile is specified, the commands of the file will be executed without user interaction, and CBShell will exit at the end. The prefix -f can also be omitted, i.e. "cbshell scriptFile" is equivalent to "cbshell -f scriptFile".
- t** This option can only be used in combination with the -f option. It shall instruct CBShell to confirm each command in the script file before it is executed.
- i** This option modifies the 'cbserver' command by invoking a CBserver compiled directly from its sources. For developers only.
- a** This options instructs CBShell not to directly show the answer to each command in interactive mode; instead you have to manually call showAnswer
- v** This options enables the verbose mode. In this mode, the command and the answer are always displayed on standard output
- p** Disables the display of the command prompt in interactive mode. This may be useful when CBShell is used in a Unix pipe where the preceding program generates the commands and feeds them into CBShell.

-q Instructs CShell to convert single quotes in positional parameters into escaped double quotes. This option is useful when a parameter contains special characters and still shall be regarded as a valid object label by ConceptBase. Useful when calling CShell scripts within regular scripts (e.g. bash) that pass parameters with special characters to the CShell scripts. See also section ??.

params At most nine user-defined positional parameters can be supplied. They are bound to the CShell variables \$1 to \$9. The CShell variable \$0 is bound to the name of the scriptfile.

7.3 Commands

cbscver [*serveroptions*] : starts a CServer with the specified options and connects to it

connect *host port* : connects to an already running CServer; if a public CServer is configured then it is the default host; otherwise localhost is default host; 4001 is default port number; if the connections fails and the local operating system can start a CServer, then a local CServer is started with default settings

disconnect : disconnects from a CServer

stop : stops the CServer which is currently connected

tell frames : tells frames to the CServer; enclose the frames in double quotes

untell frames : untells frames from the CServer; enclose the frames in double quotes

retell *untellFrames tellFrames* : untells and tells frames to a CServer in one transaction; enclose both arguments in double quotes¹

tellModel *file1 file2 ...* : tells files to the CServer; the files can have file types ".sml" and ".txt"; if the first file from the list exists on the computer of the CShell client, then the files will be loaded from the local file system, otherwise the CServer is requested to load the files from its own file system

ask *Query [QueryFormat [AnswerRep [RollbackTime]]]* : asks a query; possible query formats are OBJNAMES and FRAMES; the answer representation can be LABEL, FRAME, JSONIC, default, or a user-defined answer format; the rollback time should be Now. If the query format is OBJNAMES, then *Query* is a string in double quotes containing a query call (or a comma-separated list of query calls). If the query format is FRAMES, then *Query* is a string of Telos frames including a query definition. The answer representations are discussed in section ??.

hypoAsk *frames Query [QueryFormat [AnswerRep [RollbackTime]]]* : tells frames temporarily and asks a query

lpicall *lpicall* : executes the LPI call; only for debugging purposes

prolog *prologstatement* : executes the Prolog statement; only for debugging purposes

why : gets error messages for the last transaction and prints them on stdout

result *completion result* : compares the given result with the last result which has been received; this command hence can be used to check whether the CServer produces the expected completion (ok, error) and result; use this command in combination with the option -l

cd *mod* : changes the module context of this shell; if the parameter *mod* is omitted, the module context will set the module to the user's home module, by default oHome

pwd : display the current absolute module path, e.g. "System-oHome-MyModule"

¹There is a small syntactic restriction for the retell command. You need to avoid a line consisting just of a double quote to terminate the untellFrames. Instead, start the tellFrames in the same line that terminates the untellFrames.

lm *mod* : list all frames defined in a module; shortcut for `ask listModule[mod/module]`; uses `currentmodule` if called without parameter

ls *class* : display the instances of *class*; uses `Individual` as default if called without parameter

mkdir *module* : creates a new module with the given name within the current module; implemented by a tell operation `"mod in Module end"`; so we mimick the navigation within modules by commands known from Unix to manage directories

showAnswer : print the last result on standard output; this can be useful if you employ the Cbserver as a generator in a shell pipe (see Graphviz case below)

showAnswer > *filename* : same as `showAnswer` but output is redirected to *filename*

who : show the list of users that have at any time been enrolled to this database; implemented by a query that displays the instances of the class `CB_User`

sub : show the list of visible submodules to which the user can branch via the `'cd'` command

show *name* : show the frame with the given name; shortcut for `ask getObject[name/objname]`

echo *string* : prints the string to standard output; use double quotes if the string has multiple words; a sequence `'\\n'` within the string is replaced by a newline character;

echo -n *string* : like the one-argument variant but no newline is printed after the string

nl : prints a newline character on standard output

exit : exits the shell (also stops a server which has been started in this shell)

Command arguments with white space characters have to be enclosed in double quotes (`""`). Command arguments may span multiple lines. Lines starting with `'#'` are comment lines². If an argument contains a string of the form `$PropName`, it will be replaced with the value of the corresponding Java property (which may be defined using the `-D` option of the Java Virtual Machine), if the property is defined. There are a couple of legacy commands that we support for backward compatibility:

startServer [*serveroptions*] : synonym for `'cbserver'`

enrollMe *host port* : synonym for `'connect'`

showUsers : synonym for `'who'`

showModules : synonym for `'sub'`

listModule *mod* : synonym for `'lm'`

listClass *class* : synonym for `'ls'`

getErrorMessages : synonym for `'why'`

getModulePath : synonym for `'pwd'`

setModule *mod* : synonym for `'cd'`

stopServer : synonym to `'stop'`

cancelMe : synonym for `'disconnect'`

newline : synonym for `'nl'`

quit : synonym for `'exit'`

² The last line of a CbShell script file should not be a comment line. Otherwise, CbShell fails to recognize the end of file correctly.

The CShell utility can be used in Unix pipes to extract textual output The CServer and pass it to subsequent programs as input. To do so, you should start the CServer with tracemode no and using the showAnswer command of CShell to specify the elements to be written to standard output. The CShell script below realizes the extraction of Graphviz [<http://graphviz.org>] specifications from ConceptBase.cc:

```
# File: myscript
connect
tellModel ERD-graphviz2
tellModel UniversityModel
ask ShowERD[UniversityModel/erd] OBJNAMES default Now
showAnswer
exit
```

The complete example is available from the CB-Forum at <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/2519759>. Another resource for CShell scripts is the list of test scripts at <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/2596438>.

7.3.1 Argument delimiters

If an argument to a CShell command spans multiple lines or contains several words separated by blanks, then you have to enclose it in double quotes, being the default argument delimiter. Double quotes are also used for String objects in Telos frames. These have to be escaped like in the following example:

```
tell "
Peter with
    comment about: \"This is a Telos string\"
end
"
```

If many such frames need to be told, the escaping of Telos strings is cumbersome. You can use the single quote as argument delimiter in such cases, making the escaping the double quotes for the Telos string obsolete:

```
tell '
Peter with
    comment about: "This is a Telos string"
end
'
```

7.4 Interactive use of CShell

The CShell can be used to run a script file or in can be used in interactive mode. In the interactive mode, the CShell shall directly display the response to a command, typically the answer of the ConceptBase server. Some of the shortcuts like 'why' and 'cd' are specifically defined to make the interactive mode more effective.

Another feature of the interactive mode is that queries that are represented as a single word can also be asked without the keyword 'ask' in front of it. CShell will then ask the query using 'default' as answer format, i.e. the ConceptBase server will decide in which answer format to use. Below is a sample session of CShell in interactive mode.

```
cbshell
This is CShell, the command line interface to ConceptBase.cc
[offline]>connect
```

```

[localhost:4001]>mkdir M1
yes
[localhost:4001]>cd M1
M1
[localhost:4001]>tell "Employee in Class"
no
[localhost:4001]>why
Syntax error 1 in line 1, parser message:
syntax error, unexpected ENDOFINPUT, expecting END or ENDMIT
Syntax error Unable to parse Employee in Class.
[localhost:4001]>tell "Employee in Class end"
yes
[localhost:4001]>tell "bill in Employee end"
yes
[localhost:4001]>ls Employee
bill
[localhost:4001]>show bill
bill in Employee
end
[localhost:4001]>1+2
3
[localhost:4001]>stop
[offline]>exit

```

The term `1+2` is an example of a query that is not preceded by the `'ask'` command. Note that such queries may not contain blanks since it would split it into several words. Use quotes in such cases. The prompt `[offline]` indicates that the CBShell is not yet connected to a CBserver. Once connected, it displays the hostname and port number of the connected CBserver as prompt. The CBserver is started with disabled tracing. The trace messages of the CBserver would otherwise be displayed in the output as well.

7.5 Positional parameters

A CBShell script can use variables `$0 ... $9` inside the script to refer to the positional parameters supplied via the call of `cbshell`. Assume the following content of the script file:

```

# File: pascript
cbserver -u nonpersistent -t low -port $1
tell "$2 in Class end"
ask find_instances[$2/class] OBJNAMES LABEL Now
ask find_instances[$3/class] OBJNAMES LABEL Now

```

and the command line call

```
cbshell pascript 4321 MyClass Integer
```

The CBShell interpreter will then replace `$1` with `4321`, `$2` with `MyClass`, and `$3` with `Integer`. If you supply less parameters than required by the script, it will issue an error message and quit. If the script had started a CBserver, then this CBserver is stopped before quitting. Likewise, if the script had enrolled to an existing CBserver process, it will unenroll before quitting.

The variable `$0` is bound to the name of the script file. CBShell uses the Java tokenizer to separate the positional parameters. The tokenizer uses white spaces (blanks, tabs) to separate tokens. Use double quotes if one argument consists of several words, e.g.

```
cbshell otherscript "MyClass isA Integer end"
```

An example of a script with positional parameters is provided in the CB-Forum at <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/d3364559/ticket307.cbs.txt>. You can also call CShell scripts within scripts/batch files of the host operating system, see section ??.

7.6 Executable CShell scripts

A CShell script can be made executable under Unix/Linux and can then be used pretty much like any other shell script. Assume that you have a CShell script `myscript` that you want to execute directly from the command line.

As a first step, create a link to the `cbshell` at a common directory for installed programs:

```
sudo ln -s $CB_HOME/cbshell /usr/bin/cbshell
```

where `$CB_HOME` is replaced by the installation directory of ConceptBase. As a second step include the following comment as the first line of `myscript`:

```
#!/usr/bin/cbshell
```

Then, make the script executable:

```
chmod u+x myscript
```

You can then directly call the script by simply typing its name:

```
myscript
```

The direct call is equivalent to the call

```
cbshell ./myscript
```

7.7 CShell scripts within regular shell scripts

CShell offers the basic commands to interact with a ConceptBase server. However, it lacks control structures such as loops and conditions. It also cannot invoke arbitrary programs. Regular shell scripts such as the Bourne shell of Unix/Linux do provide these capabilities, and thus it is a natural idea to integrate CShell scripts within regular scripts to accomplish more sophisticated automation tasks.

As a first step, you should make the CShell script executable and declare in its first line

```
#!/usr/bin/cbshell -q
```

This allows to pass parameters that include special characters to the CShell script. Assume that the CShell script `ascript` was declared that way. Then, it can be called in a Bourne shell like:

```
ascript 'Jet 400' MyClass
```

The option `-q` prepares the CShell to treat parameters with single quotes in a special way. Assume further that `ascript` has the following content:

```
#!/usr/bin/cbshell -q
# File: ascript
...
tell "$1 in $2 end"
...
```

CShell will internally expand the quoted parameter `'Jet 400'` to `"Jet 400"` and then execute

```
tell "\"Jet 400\" in MyClass end"
```

The resulting frame in ConceptBase shall be

```
"Jet 400" in MyClass end
```

Essentially, single quotes of CShell parameters are converted to double quotes within ConceptBase. An simple example is given in <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/3372687>. A more elaborate example is available in the CB-Forum at <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/3384265>. The main example is in the `fileSizeDemo`.

7.8 CShell in a pipe

Assume you have a program generator that analyzes some input data (e.g from files) and produces output in the form of CShell commands (`tell`, `ask`, etc.). Then, this output can be directly passed to CShell in a Unix pipe:

```
generator | cshell -p
```

The generator program must take care of generating all required commands, in particular making sure that CShell is connected to a CServer. Consider as example generator the script file `printfiles4cshell` from the CB-Forum at <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/3384578>:

```
#!/bin/sh
echo "connect localhost 4001"
for file in *
do
  if [ -f ${file} ]
  then
    fsize=$( stat -c %s ${file})
    frame="'${file}' in File with size s: ${fsize} end"
    echo tell "\"\\\"${file}\\\" in File with size s: ${fsize} end\"
  fi
done
```

It first generates a command to enroll to a CServer on localhost at port 4001. Then, it forms a frame for each filename in the current directory to tell the file's size to the CServer. You can run the script in a terminal to see the output generated by it.

Now, assume that you have started a CServer on localhost with port number 4001. You should tell at least the following frame to the CServer to define the class `File` to which the above script refers to:

```
File in Class with
  attribute
    size: Integer
end
```

Then execute in a terminal window:

```
printfiles4cshell | cshell -p
```

It will tell the file size information to the CServer. At the end, the end of file detection of the CShell will trigger the CShell to exit from the CServer.

It is also possible to continue the pipe to a post-processing program. In this case, the commands sent to CShell should include `ask` commands in combination with `showAnswer`. The following pipe shows the main idea:

```
generator | cbshell -p | postprocessor
```

An elaborate example of this usage is in the CB-Forum at <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/3421289>. The example shows how to extract module import statements from program source files, store them in ConceptBase and let ConceptBase produce a graph specification, layed out by GraphViz.

Chapter 8

The ConceptBase.cc Usage Environment

The ConceptBase.cc User Interface consists of two main applications:

CBIva is the ConceptBase.cc Interface in Java that supports the editing of Telos frames, displays instances of Telos objects, etc.

CBGraph is a graphical editor for Telos models (or modules). Telos models can be represented by different graphical types. Insertion and deletion of Telos objects is also supported.

The interface is based entirely on Java, so it should be usable on all platforms with a compatible Java runtime environment. The Java interface includes a graphical browser and editor. Both CBIva and CBGraph can be used as stand-alone Java application.

8.1 The workbench CBIva

CBIva is a textual interface to a CBserver, which emphasizes the use of the frame syntax of Telos. You can start the CBIva workbench by the command

```
cbiva
```

The command will start a script file with the same name. We assume that the installation directory of ConceptBase.cc¹ is added to the search path of executable programs. After a few seconds, the CBIva main window should pop up. CBIva will attempt to connect to a CBserver on localhost or to a public CBserver (if configured, see ??). Figure ?? shows the main window connected to a CBserver and gives a short description of the buttons in the tool bar.

The main window consists of a menu bar, a tool bar with a button panel, the area for subwindows and a status bar. The first sub window (Telos editor) contains the history window, which records all operations for later reuse. In the following, each component of the user interface is explained.

8.1.1 The tool bar

The tool bar is the button panel below the menu bar. All buttons have tool tips, i.e. small messages that show the meaning of these buttons. The tool tips appear, if you move your mouse pointer over the button and do not move your mouse for about one second. The buttons are shortcuts for some operations that are frequently used and are also available in the menu. The operations apply to the Telos Editor which has currently the focus.

¹In rare circumstances, you may have to edit the script files `cbiva`, `cbgraph` and `cbshell`, so that the correct Java Virtual Machine is used and the environment variable `CB_HOME` points to the installation directory. See Installation Guide in the CB-Forum for details.




Figure 8.1: Main Window of CBIva

- **Quick Connect:** if connected to a CBserver then disconnect; otherwise try to connect either to the public CBserver (if configured) or to a running local CBserver; if the connection to a running local CBserver fails, then attempt to start one in the background.
- **Clear:** Clear the text of the Telos Editor window
- **Cut, Copy, Paste:** operations on selected portions of the Telos Editor text
- **Tell, Untell:** The objects specified in the active Telos Editor will be added or removed from the object base.
- **Retell:** A new window will popup and ask you to enter the frames that should be untold and told. The contents of the current Telos Editor will be inserted as default into the two text areas.
- **Ask:** Evaluate the query specified in the Telos Editor.² If you specify the name of an ordinary class (i.e. not a query class), then ConceptBase.cc will interpret this as a query call to find all instances of that class. You can also enter an arithmetic expression like

$$100 * \text{COUNT}(\text{QueryClass}) / \text{COUNT}(\text{Class})$$
- **List module:** List the content of the current module as Telos frames
- **Exit:** Disconnect from the CBserver and exit CBIva.

8.1.2 The menu bar

The toolbar has button for the most frequent operations of CBIva. The menu bar offers the complete set of operations including options to open other windows such as CBGraph (see section ??).

²The query can either be a query call referring to an existing query or a frame representing a new or existing query definition.

File: Connect: Connect to a ConceptBase.cc server (CBserver) started in another shell/command window (see section ??).

Disconnect: Disconnect from a CBserver.

Load & Save Telos Editor: Load or Save the contents of the current Telos Editor.

Load Model: Load a source model file (*.sml, *.sml.txt) to the server. As client and server are not supposed to share the same file system, this method is now implemented as a normal Tell method. The client reads the contents of the file into a string and sends the string to the server.

Start CBserver: Opens a dialog and asks for the parameters to start a CBserver. If the information has been entered, the server will be started and its output will be captured in a separate window. The workbench will be automatically connected to that server³.

Stop CBserver: Stops the CBserver, only allowed for the user who started the server or who has been designated as administrator user (see option "-a" in section ??).

Close: Same as Exit if CBIva was not started from CBGraph. If it was started from CBGraph and CBGraph is still running, then the Close option only hides the CBIva window.

Exit: Exit CBIva, if a CBGraph windows was opened via CBIva, it shall be closed as well. Likewise, a CBserver started via CBIva will be closed.

Edit: Clear: Clears the text area of the currently activated Telos editor.

Cut,Copy,Paste: Cut, copy or paste text to/from clipboard.

Tell,Untell: (Un-)Tells the text in the currently activated Telos editor

Retell: A new window will popup and ask you to enter the frames that should be untold and told. The contents of the current Telos Editor will be inserted as default into the two text areas.

Ask Frame: Temporarily tells the content of the Telos editor, extracts the query names from the frames, asks them as query calls without parameters, and returns the result in the Telos editor window

Ask Query Call: The query calls⁴ (see section ??) listed in the Telos editor are asked and the result is returned in the Telos editor window.

Load Object: Load the Telos frame of an object into the Telos editor.

Browse: New Telos Editor: Opens a new Telos editor (see section ??).

Display Instances: Opens the display instances dialog (see section ??).

Frame Browser: Opens the frame browser window (see section ??).

Display Queries: Shows the 'visible' (user-defined) queries stored in the current database and provides a facility to call these queries (see section ??).

Display All Queries: Like above but includes the many built-in queries of ConceptBase.

Display Functions: Shows all functions stored in the current database (see section ??).

Query Editor: Opens the query editor (see section ??).

Graph Editor: Opens the CB Editor (graphical browser, see section ??). If the interface is connected to a server, CBGraph will also establish a connection to this server and ask for the graphical palette, the initial object to be shown, and the module context (see section ??). Otherwise, CBGraph will start with no connection.

³This is *not* the standard way to start the CBserver. Normally, the CBserver is started in a separate shell/command window as explained in section ?? . The standard way offers more options and control over the CBserver. See also the installation guide for a discussion on the various ways to start ConceptBase.

⁴Usually, one only asks a single query call like `Q[v1/p1]`. However, ConceptBase.cc also supports comma-separated lists of query calls like `Q1[v1/p1], Q2[c1:p2]`. Such lists of query calls are evaluated one after the other. The results is merged into a single answer. For technical reasons, calls to builtin query classes like `get_object[Class/objname]` may only occur as a singleton.

Options: Set Timeout: Set the number of milliseconds the user interface waits for a response of the server.

Select Module: Select the current module (see section ??).

Select Version: Select or create a new version. A version is a special object that represents the state of an object base at a specific time. By default, all queries are evaluated on the current state of the object base (version "Now"). By selecting another version, queries are evaluated wrt. to a previous state of the object base.

Pre-Parse Telos Frames: If enabled, the user interface parses the contents of a Telos editor before it is send to the server. Thus, syntax errors might be already detected at the client side.

Show Line Numbers: Enables the display of line numbers in the Telos Editor window. Use 'Save Options' to memorize this setting for the next session of CBIva.

Use Query Result Window: If enabled, the result of a query is shown in a separate window in a table view.

Look and Feel: You can switch the look and feel to an other style. Default and preferred is 'Metal'.

Save Options: Saves the current options for subsequent sessions of CBIva.

Edit Options Manually: Complete and editable list of options for CBIva. The contents is maintained in a file '.CBjavaInterface' of the user home directory.

Help: ConceptBase.cc Manual: Opens a window⁵ with the online-version of this ConceptBase.cc manual.

ConceptBase Tutorial I/II: Opens a window with the online-version of the respective tutorial.

CB-Forum: Opens a window to the public version of the ConceptBase Forum with lots of examples.

About: Shows a dialog with information about this program.

License: Displays the license of ConceptBase.cc in a new window.

CB-Team: Displays a page about the ConceptBase Team.

History: Load History: Load previously saved contents of the history window.

Save History: Save contents of the history window to a file.

Redo History: Redos certain operations which are currently in the history. The operations can be selected from a list.

Set History Options: Select the type of operations that should be displayed in the history window.

The preferred way to start a CBserver from CBIva is to use the 'Start CBserver' from the File menu. It provides an output window for the trace messages of the CBserver, but this window is only visible if the trace mode is set to 'minimal' or higher. Disconnecting from the CBserver will not stop it. It has to be stopped explicitly or CBIva will stop it if it is shut down itself. The 'Quick Connect' button of the toolbar provides a simplified way to start a local CBserver. It does not ask for CBserver parameters. Instead, the default values are used for all parameters except that the server mode is set to 'slave' and multi-user mode is disabled. This causes the CBserver to shutdown whenever the last connected client disconnects from it.

You can move the tool bar outside the main window or display it in vertical form, if you click on the leftmost area of the tool bar and drag it to another place.

⁵ Under Linux with Gnome, the default web browser is used to display the content. Otherwise, a Java window capable of displaying HTML is opened. Due to unknown reasons, the mechanism works for Linux only if at least two browser tools are installed. You may want to install Chromium as second browser besides Firefox to fulfill the constraint.

8.1.3 The status bar

The status bar contains some fields that display general information about the status of the application.

- Connection status, either connected or disconnected from serve, or a pending action
- Short message, usually about the result of the last action
- Current version, i.e. the rollback time (see section ??) specified for queries (default: Now)
- Current module, the database module to which TELL/ASK operations are applied

The field for the current version shall display the current time if the version is set to 'Now'. If set to a rollback time in the past, it displays both the version name and the time associated to the version. It will then also hight the background of the text field to make the user aware that queries are evaluated against a past state of the database.

8.1.4 Telos editor

The Telos Editor is an editable text area, where you can edit Telos frames. The operations can be executed from the menu bar or the buttons in the tool bar. Furthermore, the text area has a popup menu on the right mouse button with the following items.

Display Instances: Displays the instances of the currently marked object

Load Object: Loads the Telos frame of the currently marked object into the Telos editor

Display in Graph Editor: Shows the currently marked object in CBGraph.

Change Module: Changes the module context to the module path defined by the currently marked text

Clear,Cut,Copy,Paste: same as in the menu bar

Small,Large: sets the text font size to either small (12 point) or large (18 point)

If the Telos editor window is empty, then you can drag and drop Telos text files (file type *.sml or *.sml.txt) into it. CBIva will open the file and paste the content into the Telos editor window. This function requires that the option 'Show Line Numbers' is enabled (see section ??). You can also drag and drop URLs pointing to publicly accessible Telos text files.

Tell transactions

The text in the Telos editor window is typically a sequence of Telos frames. When you press the 'Tell' button, the text is sent to the CBserver and processed there by the 'TELL' method of the CBserver. This is a single transaction which may fail or succeed.

You can use the string "{---}" in the text window to indicate that the frames should be told in multiple transactions. Consider the example:

```
Shop in Class end
Guest in Class with
    attribute
        dept: Shop
end
{---}
GuestEmployee in Class isA Guest,Employee end
```

Activating the 'Tell' button instructs CBIva to split the text into two parts and tell them in two separate transactions to the CBserver. This feature is useful when some parts contain meta formulas (see section ??) that first need to be compiled before they are used in subsequent parts of the whole text. If you do not use meta formulas, then you can omit this feature.

8.1.5 History window

The history window is part of the main Telos editor. It stores all operations and their results, so that they can later be used again. The buttons scroll the history back or forward, copy the text into the Telos editor or redo the operation in the history window (see figure ??). If the current operation is an “ASK”, then a single click on the copy-button will copy the query to the Telos Editor, and a double click will copy the result of the query. The size of the history window can be reduced by using the slider bar between the Telos editor and the history window.




Figure 8.2: Buttons of the History Window

8.1.6 Display instances

This dialog displays the instances of a class. The class is entered in the text field. When you hit return or press the “OK” button, the instances of this class will displayed in the listbox.

If you double click on an item in the listbox, the instances of this item will be displayed. The frame of a selected item can be loaded into the Telos editor by clicking the “Telos Editor” button. A history of already displayed classes is stored in the upper right selection list box. The “Cancel” button closes the dialog.

8.1.7 Frame browser

The Frame Browser (see figure ??) shows all information relevant to one object in one window. The window contains several subwindows with list boxes that show super- and subclasses, the classes, the instances, attributes and objects refering this object. In the center of the window, a small window with the object itself is shown. To view the attributes of the object, you must first select the attribute category in the subwindow “Attribute Classes”.

The Frame Browser can be used with and without a connection to a CBserver. If it is not connected, it retrieves the information out of a local cache, which can be loaded from a file by using the “Load” button. The file has to be plain text file with Telos frames. All objects in the cache can be saved into a text file as Telos frames with the “Save” button. The contents of the cache can be viewed with the “Cache” button. The result of a query can be added to the cache by using the “Add query result” button.

The button “Telos Editor” inserts the Telos frame of the current object into the Telos Editor.

8.1.8 Display queries

This dialog displays all visible queries⁶ stored in the current object base (see figure ??). From the list box, you can select a query and “ask” it or load its definition into the Telos editor.

If you “ask” a generic query class with parameter, another dialog will ask you to specify the parameters. For each parameter, you can specify whether the value entered should be used as “substitute” for the parameter or as a “specialization” of the parameter class (see section ??). You can select a value for the parameter from the drop-down list if you have clicked on the “Show Values” button. Note that this list might be very long. Especially for the predefined queries it usually returns all objects in the database as any object can be used as a parameter for these queries.

⁶Visible queries are those queries that are not instantiated to the class `HiddenObject`. Functions and certain system queries are excluded from the display.




Figure 8.3: Frame Browser




Figure 8.4: Display queries dialog

8.1.9 Display functions

This dialog is similar to the previous dialog but displays instances of `Function`. Note that functions are formally special queries. Consequently, the dialogs for functions and queries are pretty much the same. The separation into two dialogs serves quicker handling.

8.1.10 Query editor

The Query Editor (see figure ??) allows the interactive definition of queries. The name of the query is entered in the upper left text field, the super class in the upper right field. After you have entered this information, the list box “Retrieved Attributes” will be filled with all available attributes.

Now, you can select the attributes you want to have in the result. For selection of more than one attribute, you must press the CTRL key and select the attribute with a mouse click at the same time. All attributes can be deselected by the popup menu.




Figure 8.5: Query Editor

In the right listbox you can add computed attributes. The right mouse button brings up a popup menu, which lets you add or delete an attribute.

In the text area below the two list boxes, you can add a constraint in the usual CBQL syntax. The constraint must be enclosed in \$ signs.

The text area below, shows the Telos definition of the query and is updated after every change you have made. If your query is finished, you can press the “Ask query” to test the query, i.e. it is told temporarily and the results are shown in separate window. If you are satisfied with the result you can press the Tell button to store the query in the object base.

8.1.11 Tree browser

Note: The tree browser has been removed from CBIva due to legacy dependencies on packages no longer supported by Java.

The Tree browser (see figure ??) displays the super classes, classes and attributes of an object in a tree. To start the tree browser, you must mark an object in the Telos editor and select the item “View Object as Tree” from the popup menu or the Edit menu.




Figure 8.6: Tree Browser

To expand an item, just double click on the icon. If you mark an object name in the tree, you can load into the Telos editor with the “Telos Editor” button or open a new tree browser with the button “View Object as Tree”.

8.2 The graph editor CBGraph

The CBGraph Editor is an advanced graphical modelling tool that supports the browsing and editing of Telos models. It supports user-definable graphical types, i.e. objects may be visualized by dedicated graphical layouts. In addition to predefined graphical types, the user can add his/her own graphical types by modifying and adding certain objects in the knowledge base. Furthermore, the standard components can be replaced by own classes implementing specific application-dependent behaviour.

In the following, we first give an overview of CBGraph application and then present the main components and functions of CBGraph. Details about the use of graphical types can be found in Appendix ?? . An example for the definition of graphical types of the Entity-Relationship model is given in Appendix ?? .

8.2.1 Overview

The CBGraph Editor is entirely written in Java. It can therefore be used on any platform with Java 1.4 or compatible successors of Java 1.4. CBGraph is integrated with CBIva, i.e. Telos frames of objects shown in CBGraph can be loaded directly into a Telos editor and vice versa. CBGraph allows to open several 'internal windows' in its main window. Each internal window has a separate connection to a CServer. Thus, within a CBGraph you can establish multiple connections to the same CServer or even to different servers.

The communication with the CServer is done using pre-defined ConceptBase.cc queries and a special XML-based answer format. CBGraph requests information about the objects (names, attributes, etc.) but also about their graphical type.




Figure 8.7: CBGraph with three internal windows

Figure ?? gives an overview of the CBGraph Editor. Three internal windows have been opened, the

two windows on the left have been connected with the same server running on “localhost”, port 4001. The small window in the upper right corner is not connected to a server, but a few objects have been created. The title of the internal windows displays the graphical palette and the current module (if connected to a ConceptBase server) plus the connection status (either ‘offline’ or the hostname and portnumber of the ConceptBase server). The content of an internal window is a graphical view on the database module of the ConceptBase server it is connected to. It is possible that different internal windows connect to different ConceptBase servers, though this is not a typical use of CBGraph. If you start CBGraph from CBIva, then you can only start a single instance of it. However, you can start any number of CBGraph instances via the ‘cbgraph’ command (see below).

The two left internal windows of figure ?? are connected to the same server and show the same model but in different representations. This is caused by the fact, that for the upper window, the default graphical palette has been chosen, and for the lower window, a customized graphical palette specifically designed for the ER model has been selected (see appendix ?? and ??).

Furthermore, you can notice that the object “QueryClass” is represented by two different components. In the upper view, the detailed *component view* has been activated by a double-click on the object. It shows the Telos frame of the object. By a double-click on the title bar of this component, one can switch back to the default view of this object. Thus, each object can be shown by a small component (the default view) and a large component which gives more detailed information. Components are in this context specific Java objects, namely instances of javax.swing.JComponent. Thus, different components can be provided to represent an object (e.g., tables, buttons, text fields). You can implement your own component and integrate into CBGraph by extending a specific Java class. Details about the customization of CBGraph using graphical types and other components can be found in appendix ??.

CBGraph can be used to edit Telos models (see section ??). It can also display implicit relationships between objects, which have been derived by rules or Telos axioms. For each type of relationship (instantiation, specializations, and attributes), one can choose to see only the explicit relationships or to see all relationships.

8.2.2 Starting CBGraph via CBIva

CBGraph can be invoked from CBIva via the menu item *Browse* → *Graph Editor*. If CBIva is connected with a CServer, CBGraph will be connected to the same server and you will be prompted to enter the object name you want to start with, the name of a graphical palette, and the database module. The graphical palette is a Telos object which represents a set of graphical types which will be used to visualize Telos objects (see Appendix ??). On startup, CBGraph retrieves all information about the graphical types from the CServer. If CBIva has no connection with a server, CBGraph will be started without a connection and no internal window will be opened within CBGraph.

The connection to a CServer can also be established via the File / Connect menu. The dialog box has two tabs. The first is for providing the host and port number number of the CServer. The second is for providing the start object (default “Class”) to be displayed, the graphical palette (pick from a list), and the database module (default “oHome”).

The most comfortable way to start CBGraph is to double-click the graph file. It is equivalent to starting it via the command

```
cbgraph graph.gel
```

To do so, you need to configure your desktop according to the instructions at <http://conceptbase.sourceforge.net/CB-Mime.html>.

8.2.3 The cbgraph command

You can also start CBGraph as a stand-alone utility. The command

```
cbgraph [options]
```

will start an unconnected graph editor. You can interactively connect it to a running CBserver and open new windows to display graphs. More interesting is the use with a stored graph file, or several graph files, resp.:

```
cbgraph [options] filename [filename ...]
```

The format of the graph file is called Graph Editor Layout (GEL). It stores not only the layout of nodes and links but all other data necessary to edit the graph objects. In particular, it contains connection details of the CBserver module from which the graph was created. You can open more than one GEL file. Each will be loaded in its own frame.

There are a few options for the 'cbgraph' command synchronizing the data stored in the graph file with the CBserver:

- +r** With this option CBGraph will be instructed to load the module sources on its current module path from the CBserver and save it to the graph file when the 'save' function is enacted. The 'System' module will not be saved since it is (typically) not changed. For example, if CBGraph displays objects of module "System-oHome-M1", then the save function will store the module sources of 'oHome' and 'M1' to the graph file. We say "CBGraph reads the sources from the CBserver and saves them to the graph file". This option disables writing the module sources from the graph file to the CBserver when CBGraph is started.
- +w** This option enables the reverse direction. It will instruct CBGraph to extract module sources from the given graph file (or any graph file that is loaded via CBGraph) and tell them to the CBserver that it is connected to. You have to start a CBserver as a separate process using the hostname and port number that is stored in the graph file. This option disables by default including the module sources to the graph file when saving the graph file. The CBserver can well have an empty database because all required user-defined objects are stored as sources in the graph file. If the database is not empty, the definitions from the sources are added to the current database. It may well be that the operation causes an integrity violation or fails because the current user has no write permission on a given module. In case of success we say "CBGraph loads the module sources from the graph file and writes them to the CBserver".
- +rw** Enable bidirectional synchronization (default).
- rw** No synchronization.
- +f** Enable bidirectional synchronization and also write the module sources to text files in the directory where you started CBGraph; useful for debugging.
- host hostport** Instruct CBGraph to use the hostname and portnumber from hostport rather than those defined in the graph file when loading the graph file.

If you supply these options when calling CBGraph, it will also store them in the graph file when you store it. If you later call CBGraph with such a graph file, it will apply the stored options unless you specify new options in the command line. Hence, the options stored in the GEL file serve as new defaults for synchronizing the graph file with the CBserver.

Consider the two calls of CBGraph below. The first call sets the synchronization option to "+r", i.e. stores the module sources from the CBserver in the graph file when it is saved. The second call has no such option. Hence, it shall adopt the "+r" that was stored earlier in the graph file `example.gel`.

```
cbgraph +r example.gel
cbgraph example.gel
```

This behavior is useful if you keep the models in a persistent CBserver and want to use the module sources in the graph file only as a backup storage. It will then not tell the module sources from the graph file to the CBserver. The TELL operation can be very costly and is redundant if the CBserver anyway has all definitions already stored.

If you provide a filename (or several), then it must have been previously created by another graph editor, e.g. a graph editor started via CBIva. The above command will then display the graph stored in the graph file in an internal window and attempt to connect to the same CBserver that was active when the graph file was created. Hence, the graph file is a materialized view on the CBserver database that is visualized with 'cbgraph'. You edit a graph file with 'cbgraph' like you are editing a drawing with a drawing tool. The only difference is that the graph is linked to the database.

If the CBserver module specified in the graph file is not accessible, the graph is still opened and you can edit it. You can then however not commit changes to the database or add new objects from the database. The CBGraph editor displays the connection status in the title of its internal window containing the graph. The hostname of a CBserver in a graph file is by default 'localhost'. Consequently, the graph editor will try to establish a connection to a CBserver running on localhost. If you want to connect to a remote CBserver, then you should specify in CBIva the full domain name of the CBserver, e.g. 'myhost.acme.com'. This long name will then be stored in the graph file that is created from a graph editor started via CBIva. Such graph files can then be copied to other computers and can be loaded with CBGraph to auto-connect to the remote CBserver specified in the graph file, provided that CBserver is running and accessible.

If no CBserver is accessible, CBGraph will attempt to start a local CBserver in the background provided that the graph file specifies "localhost" as hostname and the graph file contains module sources. In other cases, CBGraph switches to the offline mode. You can still change location and size of the graphical elements and store it back to the graph file. But you cannot delete objects and you cannot add objects to the database. The menus to show attributes, instances, and subclasses shall also not work.

An example on how to create and use graph files containing materialized database views is available in the CB-Forum at <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/3613919>.

8.2.4 Redirecting the CBserver location

You can use the command line argument '-host' to override the hostname and portnumber encoded in the graph file. Assume a graph file 'graph1.gel' was created from a a CBserver connection at 'localhost:4001'. Then, loading this graph file in a subsequent call will also connect to 'localhost:4001'. If you want to use another CBserver, e.g. running on 'myhost.acme.com:4002', then call CBGraph from the command line as follows:

```
cbgraph +rw -host myhost.acme.com:4002 graph1.gel
```

Note that the CBserver must be running at the remote location before you enter the above command. When you subsequently save the graph file, it will have 'myhost.acme.com:4002' encoded as its connection. The redirection also works in the reverse direction. So, assume that the graph file 'graph2.gel' was created for a connection at 'myhost.acme.com:4002'. Then, the following command will redirect the connection to localhost:

```
cbgraph +rw -host localhost graph2.gel
```

The default port number is 4001. If no CBserver is running at 'localhost:4001', then CBGraph shall start it in the background. Note that local CBservers can currently only be started on Linux hosts.

8.2.5 Moving objects

All explicit information is a proposition in ConceptBase (see section ??), including attributes, specializations and instantiations. You can move objects in the graph editor by pointing the cursor to the object's label and then dragging it while keeping the left mouse button pressed. If a graph has many nodes and edges, then it is recommended to first click once on the node or edge to be dragged and then to drag it. This will switch off the anti-aliasing while dragging and thus be faster.

You may also want to select a group of objects and then move it as a whole. To do so either span a selection box with the left mouse button around the object to be selected, or press the "shift" key and select multiple objects individually. After a move, CBGraph shall redraw dependent edges that might be misplaced.

Some edges like instantiation links have no label in CBGraph, depending on the graphical type associated to it (see appendix ??). In such cases, a small square dot is displayed on the edge. Click on this square dot to select the edge and to drag it. If the edge has a background color (see appendix ??) different from the edge color, then the square dot is drawn in the background color.

Moving nodes and edges can sometimes lead to quite messy edge curves where the edge middle point is distant from the middle of the straight line between the source and destination of the edge. You can clean up such graphs by pressing the "shift" and "control" keys together and then click on a node whose edges shall be straightened.

8.2.6 Menu bar

The menu bar provides access to the most important functions of CBGraph.

File menu

Connect to server: Connect to a new server. You can have multiple connections to one server or different servers at the same time. Each connection will be represented in one internal window with the graphical palette, database module, host name and port number in the title bar of the window. The connection dialog consists of two tabbed panes. In the first one (Address), you can enter the host address (name or IP number) and the port number. In the second pane (Initial Object), you can specify the initial object to start the browsing process, the graphical palette, and the database module.

Start CBIva Workbench: Start a CB workbench (aka CBIva). If you started CBGraph directly or if you have already closed the workbench window, you can (re-)start the workbench by this menu item. If connected to a CBserver, the new CBIva will list the content of the current module in its Telos editor window.

Save: Save the current graph into a file. The current nodes, their location and the links will be saved into a file which can be reloaded later. The graphical types will also be saved into the file. By default, the file will get the extension "gel" (Graph Editor Layout). A checkbox lets to select whether you want to save the module sources into the graph file. It overrides the '+/-rw' options of the cbgraph command.

Load: Load a graph that has been saved with the previous menu item. Existing nodes and links in the current window will be erased. If the graph file was previously created to include module sources, then the module sources are told to the CBserver before the graph is displayed⁷. A checkbox lets to select whether you want to load the module sources from the graph file and tell them to the CBserver.

Print: Print the current graph. If the graph is larger than the page size, it will be automatically reduced to fit into the page. Thus, all printouts will be on one page.

Save image of graph: Saves an image of the current graph as PNG or JPG file. The PNG file format should be preferred as it delivers better results. For technical reasons, the image saving has to set the background color of the node/edge labels. This can overlap with the node shapes when they are not rectangles. If so, then use the screenshot function of your operating system to get better results.

Close: Closes CBGraph. A CBIva window will be still available if it has not been closed before.

Exit: Exit CBGraph and CBIva. This operation will close both windows of the ConceptBase.cc User Interface (CBIva), if they were started in combination, and exit the program.

If the graph was loaded from a GEL file and it was edited in the session, CBGraph shall ask the user whether to save the edited graph to the file when the user terminates CBGraph.

⁷ If you load a graph file that was created with different graphical types than the ones defined in the CBserver module that CBGraph might be connected to, then the graph file is inconsistent with the current CBserver module. You may be able to repair the inconsistency via the menu option Current connection / change graphical palette. Likewise, some objects displayed in the graph could be undefined in the database. You can validate them by one of the validation tabs in the menu "Current connection".

Edit menu

The operations in this menu have an effect on the selected objects. You can select an object by clicking on it with the left mouse button. Multiple selection is possible dragging a rectangular area while holding down the left mouse button or by holding the Shift-key and clicking on objects.

Erase Selected: This option will remove the currently selected nodes and edges from the view. This operation has no effect on the database.

Selection: With this submenu, you can either select all objects, all nodes or all edges in the current frame. Furthermore, you can clear your current selection.

Options menu

The options will be stored in a configuration file (see section ??) when you exit CBGraph.

Language: The text for menu items and buttons is available in two languages (German and English). With this option you can switch between the languages. This option is currently without function. All menu labels are set to English.

Background Color: Here you can change the background color of the graph to your favorite color.

Component View: With this option, you can configure the view of an object if the component view is activated by a double click or by the popup menu. By default, a tree-like representation of the object with its classes, instances, super classes, subclasses, outgoing and incoming attributes is used. If you select "Frame", then the Telos frame of the object will be shown in a text area, (see figure ??).

Invalid Telos Objects: CBGraph can validate objects that are currently shown in the graph, i.e. it checks whether the object is still valid in the database or it has already been removed (see Current Connection menu). If you select "Mark" here, then the objects will be marked with a red cross as invalid. "Remove from display instantly" will remove the objects directly from the view.

Popup Menu: These options control the behaviour of the Popup Menu (see section ??). The delay is the time (in milliseconds) an item of the popup menu has to be selected before the submenu is shown. Note that the construction of a submenu might require a query to the CBserver. If the option "*Popup Menu blocks while waiting for server*" is activated, then the editor will block the UI while it waits for an answer of the CBserver. Otherwise, the query to the server will be executed in a separate thread, and interaction with the UI will be possible. If you have the problem that some submenus are still shown after you have used the popup menu, set the delay to 0 and activate the option "*Popup Menu blocks while waiting for server*".

Look & Feel: This option allows you to adapt the look and feel of CBGraph windows to your desktop environment. By default, the 'Metal' look and feel is enabled. CBGraph is not yet compatible with all Java look and feels. We thus recommend to stick to the default 'Metal' look and feel.

Enable Click Actions: Click actions automatically trigger an active rule (see section ??) when nodes with a certain graph type are pressed (see section ??). You can enable and disable this feature. If you want to move nodes inside a graph that has nodes with click actions, then you may want to disable click actions until you have finished the re-arrangement of nodes.

Enable Derived Links: By default, CBGraph can display derived links/relations to and from a given object. If the database is large and the rules are complex, this feature becomes almost unusable due to long delays. In such case, disable the display of derived links. The behavior of CBGraph is then different wrt. the pop menu displayed when right-clicking on a node or link. It would show all possible attribute categories rather than only the used ones. For technical reasons, a change in this flag becomes only effective after re-starting CBGraph.

If an object is displayed in component view, one can switch back to the node view by double-clicking its title section.




Figure 8.8: Component view of Telos objects: tree and frame

View menu

CBGraph has an experimental layout algorithm that may be useful to reorganize the layout of a complex graph. The heuristic of the layout algorithm is rather simple and does not minimize link crossings.

Enable automatic layout: Call a layout algorithm everytime the graph is changed, e.g. by expanding the attributes. Disabled by default. Enabling it is not recommended.

Undo last layout operation: Undo the last change of the graphical layout. This option is only available when automatic layout is enabled.

Layout graph: Call the layout algorithm.

Zoom: Set the zoom factor of the graph, e.g. 120 (20 percent enlarged).

200,100,75,50,25: : Set the zoom factor accordingly.

Current connection menu

These operations have effect on the current connection, i.e. the connection of currently activated internal window.

Query to server: This operation will open a dialog which prompts you to enter the name of a query (see figure ??). The query can also be parameterized. If you click on the “Submit Query” button, the query will be sent to the server and the result will be displayed in the listbox. You can select the objects that should be added to the graph. Selection of multiple objects is possible.

Validate and update shown objects: This operation will check for every object, if it is still valid (i.e. if it still exists in the database), update the graphical type of the object, and the internal cache of the object is deleted (see below, section ??). Depending on the option “Invalid Telos Objects” (see above), the objects will be either marked or removed from the current view.

Validate and update selected objects: Similar to the previous option but is only applied to objects in the graphical view that have been selected.

Change graphical palette: The current graphical palette (=assignment of graphical types to nodes and links) is replaced by another one.

Change module: Change the database module for the currently active internal window.




Figure 8.9: Query dialog

8.2.7 Tool bar

The tool bar (see figure ??) consists of a set of buttons that are mainly short cuts for some menu items. The right half of the tool bar provides buttons for the creation of Telos objects.




Figure 8.10: Tool Bar of CBGraph

Open a new frame (without connection): Opens a new internal window without a connection to a server. Within this internal window, you can create new Telos objects, load existing layouts and save new layouts. Information about the graphical types is loaded from an XML file included in the JAR file (\$CB_HOME/lib/classes/cb.jar).

Load graph: see *File Menu* → *Load*

Save graph: see *File Menu* → *Save*

Hide selected objects: Hides the selected objects from the current view. This operation has no effect on the current database, i.e. the objects will be not deleted from the database.

Open connection: Opens a new internal window with a new connection to a server. See *File menu* → *Connect to server*.

Show object: This operation adds an object defined in the database to the graph. You will be prompted to enter the object name of the object you want to add to the graph.

Show links between marked objects: This operation will search for relationships between the selected objects. Do not select too many objects for this operation, n^2 queries have to be evaluated for n objects.

Creation of objects: The following four buttons open the “Create Object” dialog to create new individual objects, attributes, instantiations, and specializations. See section ?? for more details.

Show added/removed objects: Shows the objects that have been added or removed since the last commit (or since the window has been opened). Here, you can also select objects to undo the change, i.e. remove added objects or re-insert removed objects.

Commit: Sends the changes to the server. The list of objects to be added or removed is transformed into a set of Telos frames and transferred to the server.

8.2.8 Popup menu

The popup menu is activated by a click on the right mouse while the cursor is located over an object.

- **Toggle component view:**
switches the view of this object. In the detailed component view, you can either see the frame of this object or tree-like representation of super- and subclasses, instances, classes, and attributes of this object (see figure ??).
- **Display in Telos Editor:**
This operation will load the frame of the object into the Telos editor. If the Telos editor contains already some text, then the frame is appended to the existing text.
- **Super classes, sub classes, classes, instances:**
for each menu item you can select whether you want to see only the explicitly defined super classes (or sub classes, etc.) or all super classes including all implicit relationships. The query to the CB-server to retrieve this information will be done when you select the menu item. So, the construction of the corresponding submenu might take a few seconds.
- **Incoming and outgoing attributes:**
CBGraph will ask the CBserver for the attribute classes that apply to this object. For each attribute class, it is possible to display only explicit attributes or all attributes as above. The attribute class “Attribute” applies for every object and all attributes are in this class. Therefore, all explicit attributes of an object will be visible in this category. However, there will be no attributes shown in the “All” submenu, as it would take too much time to compute the extension of all implicit attributes.
- **Add Instance, Class, SuperClass, SubClass, Attribute, Individual:**
These menu items will open the “Create Object” dialog where you can specify new objects that should be created in the database. Note, that these modifications are not performed directly on the database. The editor will collect all modifications and send them to the CBserver when you click on the “Commit” button. See section ?? for more details.
- **Delete object from database:**
This operation will delete the object from the database. As for the insertion of objects before, the modification will be send to the server when you click on the “Commit” button. Note that this operation has an effect on the database in contrast to the next operation.
- **Hide object from view:**
The object will be removed from the current view. This operation has no effect on the database, i.e. the object will not be deleted from the database.
- **Show in new Frame:**
A new internal window (within CBGraph) will be shown and the selected object will be shown in the new window.
- **Straighten attached edges:**
All edges starting from and ending in the selected node are made straight and the edge labels move to the middle of the edge. This is the same function as described in section ??.

- **Freeze / Unfreeze:**

The position of the selected object is frozen (if not yet frozen). A "frozen" object cannot be moved or double-clicked. If frozen, the selected object is "unfrozen". This function is useful when certain objects should not be moved, e.g. when they are serving as regions, in which other objects are positioned.

8.2.9 Editing of Telos objects

CBGraph supports also the creation and deletion of Telos objects. A Telos object in the context of CBGraph is a *proposition* as described in chapter ?? . As described there, there are four types of Telos objects:

- Individuals,
- Instantiations (InstanceOf),
- Specializations (IsA), and
- Attributes.




Figure 8.11: Create Object dialogs for Individuals, Attributes, Instantiations, and Specializations

For each object type, we provide a dialog to create this object type as shown in figure ?? . This dialog is opened by clicking on one of the "Create" buttons in the tool bar, or by selecting an "Add ..." item from the popup menu (e.g. "Add instance" or "Add subclass"). If there are some objects selected in the current internal window, then the object names of these objects will be inserted into the text fields of the dialog in the order they have been selected (i.e., the first text field will contain the name of the object which has been selected first). Furthermore, if you move the cursor into a text field in the "Create Object" dialog and select an object in the graph, then the name of this object will be inserted into the text field.

As changes might lead to a temporary inconsistent state of the database, we do not execute the changes directly on the database. They are stored in an internal buffer in CBGraph and executed when you hit the "Commit" button in the tool bar.

Creating Individuals: If you want to create a new individual object, you just have to specify the object name. You have to enter a valid Telos object name, for example it must not contain spaces. In addition, you can select a graphical type for the object. Note that the selection of the graphical type has no effect in the database, e.g. by selecting the graphical type of a class (ClassGT) the object will not declared as an instance of Class. If you have performed the commit operation, the object will get the "correct" graphical type from the server.

Creating Instantiations: In the dialog for instantiations, you have to enter the name of the instance and the name of the class in the two text fields. If the object entered does not yet exist, it will be created and represented in the default graphical type.

Creating Specializations: This dialog is similar to the one before except that you specify here the name of the subclass and the name of the superclass. Objects that do not exist yet, will be created and represented in the default graphical type.

Creating Attributes: This is the most complex dialog as you have to specify the source, the label, the value, and the category of the attribute. The source and the value (or destination) of the attribute are normal object names. The label may be any valid Telos label. The attribute category has to be a select expression specifying an attribute category (e.g., `Employee!salary`, see chapter ??). The attribute category can be selected from a listbox by clicking on the “Select” button next to the text field of the attribute category. All attribute categories that apply to the current source of the attribute will be shown. Note that the list will be empty if the source object does not yet exist in the database.

If you have specified the attribute category, you can also select the attribute value from a listbox by clicking on the “Select” button next to the text field for the attribute value. The listbox will show all instances of the destination of the attribute category (e.g. all instances of `Department` for the category `Employee!dept`).

If you select the radio button “Show Attribute Instantiation” then CBGraph will also show the instantiation link for the attribute. For example, if you create a new attribute for `John` with the label `JohnsDept` in the attribute category `Employee!dept`, then the instantiation link between `John!JohnsDept` and `Employee!dept` will also be shown. As the graph gets quite confusing with too many links, this radio button is not selected by default.

Deletion of objects is also possible. As this operation should not be mixed up with the removal of an object from the current view, this operation is just available from the popup menu (item “Delete Object from Database”).

As you might make mistakes while editing the model, there is the possibility of undoing changes. The button “Show added/removed objects” list all objects that have been added or removed (since the last succesful commit or since the connection has been established). A screenshot of the dialog is shown in figure ?? . The left list shows the objects that have been added, the right list shows the objects that have been removed. By clicking on the button “Re-Insert/Delete” object, the selected objects will be re-inserted in or deleted from the graph ⁸.




Figure 8.12: List of objects which have been added or removed

⁸You can unselect an object by holding down the Control key and clicking on the object.

If you are satisfied with the changes you have done, you can click on the “Commit” button. Then, CBGraph will transform the objects to be added or removed into a list of Telos frames and send them to the server using the TELL, UNTell, or RETell operation. If the operation was successful, all *explicit* objects will be checked if they are still valid and if their graphical type has changed (as in the “Validate and update” operation from the “Current Connection” menu). If there is an error, the error messages of the server will be displayed in a message box. The internal buffer with the objects to add or remove will be not changed in this case.

8.2.10 Caching of query results within CBGraph

To improve the performance of CBGraph, several caches are used. On the other hand, the use of a cache causes several problems which will be addressed in this section. In particular, the caches of CBGraph are not updated automatically if the corresponding data in the server is updated.

Graphical Palette and Graphical Types: When a connection to a server is established, CBGraph loads the graphical palette and all its graphical types including their properties and other information. If an object has to be shown, the server sends only the name of the graphical type, the information about the properties are taken from the cache. Thus, if you change the graphical palette or a graphical type after CBGraph has established the connection, this change will not be visible in CBGraph. There is currently no method implemented to update the cache manually.

Graphical Types of Objects: When an object is loaded from the server also the graphical type for this object is retrieved. The graphical type of the object is updated when you invoke the “Validate and Update” operation from the “Current Connection” menu.

Lists of super/sub classes, classes/instances, attributes: The lists in the popup menu or in the tree-like view of an object are produced by evaluating queries. To reduce the communication between client and server, each query will only be evaluated once (when the corresponding popup menu should be shown or when the part of the tree should be shown). The result will be stored in a cache for each object. This cache is emptied when you invoke the “Validate and Update” operation from the “Current Connection” menu.

8.2.11 Graph files

Graph files (extension ‘gel’) are binary files that store the current state of a graph displayed in CBGraph. Since they are constructed from a ConceptBase database, they are a (materialized) view on the database. The view consists of the nodes and links displayed in the window, their positions, their graphical types, the hostname, portnumber and module from which the graph was created, the size of the window, its background color and image, and the window’s zoom factor. You can thus save the current state of your graphical view in the *gel* file and load it in a subsequent session with CBGraph to continue editing it, much like with a drawing tool.

The graph file stores serialized Java objects in the following sequence

```
String title of the internal frame
Dimension size of the graph editor
Dimension size of the internal window
Dimension size of the drawing area of the internal window
Integer number of nodes (incl. edge objects)
{
  node1
    Rectangle bound of node1
  node2
    Rectangle bound of node2
  ...
}
Integer number of edges
```

```

    { String source node of edge 1
      String node object on the edge 1
      String destination of edge 1
      String source node of edge 2
      String node object on the edge 2
      String destination of edge 2
      ...
    }
  Color background color
  Float zoomfactor
  String hostname
  String port number
  String module context as absolute path
  String long title of the graphical palette
  String name of the graphical palette
  Integer saveflag (bit0: HAS_IMAGE, bit1: HAS_SOURCE, bit2: HAS_PARAMS)
  IF HAS_PARAMS
    { String[] params
    }
  IF HAS_SOURCE
    { String saved modules e.g. "mod1-mod2"
      { String source of module 1
        String source of module 2
        ...
      }
    }
  IF HAS_IMAGE
    { PNG image of the background image
    }

```

Since edges are also objects in Telos, the nodes stored in the graph file include the edge node that represents the edge itself. The graph file stores complex information about the nodes including the graphical types of the nodes, their dimension and location. By default, the graph files stores the Telos module sources needed to manipulate the graph. The graph file stores the sources of all modules that are on the path from the root module *System* to the current module (the module that is active when the graph file is saved). The module *System* is not saved since it is typically not updated. Note that *ConceptBase* database can include a tree of module (see section ??). Hence, a graph file does in general not store the Telos sources of the complete database. If you create graph files for all leave modules of a database, then the combination of the graph files is completely containing the database as sources models. The extraction of Telos sources uses the builtin query *listModule* (see section ??). This is in most cases a faithful listing. However, there are rare cases when the extracted cannot be told to a *CBserver*. For example, if a module contains deductive rules that are essential for satisfying integrity constraints for objects defined in the same module, then a single *TELL* operation could fail because *ConceptBase* requires the deductive rules to be compiled. See section ?? for more details.

If you specified command line parameters like "+r", "-r", "+rw", or "-rw" at the start of *CBGraph*, then these parameters are stored in the GEL file.

The background image is not stored as a serializable Java object but as a PNG image using the *ImageIO* class of Java. It is always stored as the last element since the input routines shall read it until the end of the file. Note that some strings can be just null.

8.3 An example session with *ConceptBase*

In this section we demonstrate the usage of the *ConceptBase.cc User Interface*, by involving an example model. It consists of a few classes including *Employee*, *Department*, *Manager*. The class

Employee has the attributes `name`, `salary`, `dept`, and `boss`. In order to create an instance of `Employee` one may specify the attributes `salary`, `name`, and `dept`. The attribute `boss` will be computed by the system using the `bossrule`. There is also a constraint which must be satisfied by all instances of the class `Employee` which specifies that no employee may earn more money than its boss. The Telos notation for this model is given in Appendix ??.

8.3.1 Starting ConceptBase

To start a *ConceptBase* session, we use two terminal windows, one for the *ConceptBase.cc* server and one for the usage environment. We start the *ConceptBase.cc* server by typing the command

```
cbserver -port 4001 -d test
```

in a terminal window of, let us say machine alpha⁹. The parameter `-port` sets the port number under which the CBserver communicates to clients and the parameter `-d` specifies the name of the directory into which the CBserver internally stores objects persistently. Then, we start the usage environment with the command `cbiva` in the other window.

It is also possible to start the CBserver from the user interface. To do so, choose “Start CBserver” from the “File Menu” of CBIva (see section ??) and specify the parameters in the dialog which will be shown (see figure ??). The option *Source Mode* controls whether the CBserver accessing the database via the `-d` parameter (database maintained in binary files), or via the `-db` parameter (database maintained both in binary files and in source files). See section ?? for more details. Once the information has been entered via the OK button, the server process will be started and its output will be captured in a window. This output window provides also a button stop the server. If you started the server this way then you can skip the next section, as the user interface will be connected to the server automatically.




Figure 8.13: Start CBserver dialog and CBserver output window

8.3.2 Connecting CBIva to another CBserver

By default, CBIva will automatically connect to a local or public CBserver when started. If you want to start a CBserver with dedicated parameters from CBNiva, then first select **File/Disconnect** and then establish a new connection between the *ConceptBase.cc* server and the user client CBIva. This is done by choosing the option **Connect** from the File menu of CBIva. An interaction window appears (see Figure ??) querying for the host name and the port number of the server (i.e. the number we have specified within the command `cbserver -port 4001 -d test`).

⁹A full list of all parameters is described in section ?? . Note that the script `CBserver` must be in the search path. It is available in the subdirectory `bin` of your *ConceptBase.cc* installation directory.




Figure 8.14: The connect-to-server dialog

8.3.3 Loading objects from external files

The objects manipulated by *ConceptBase.cc* are persistently stored in a collection of external files, which reside in a directory called **application** or **database**¹⁰. The actual directory name of the database is supplied as the `-d` parameter of the command `CBserver`.

The `-u` parameter of the `CBserver` specifies whether updates are made persistent or are just kept in system memory temporarily. Use `-u persistent` for a update persistence or `-u nonpersistent` for a non persistent update mode¹¹.

The database can be modified interactively using the editor commands `TELL/UNTELL`. Another way of extending *ConceptBase.cc* databases is to load Telos objects (expressed in frame syntax) stored in plain text files with the extension `*.sml`. Call the menu item **Load Model** from the **File Menu** to add these objects to the database. In our example the database (directory) `Employee` was built interactively and can be found together with files containing the frames constituting the example in the directory

`CB_HOME/examples/QUERIES`

where you have to replace `CB_HOME` with the *ConceptBase* installation directory. The following files contain the objects of the `Employee` example expressed in frame syntax: `Employee.Classes.sml`, `Employee.Instances.sml`, `Employee.Queries.sml`. An alternative to interactively building a database is to start the server with an empty database (`-d <newfile>`) and then add the objects in these files by using **Load Model**. Note, that the `*.sml` extension may be omitted. During the load operation of external models, *ConceptBase* checks for syntactical and semantical correctness and reports all errors to the history window as it is done when updating the object base interactively using the editor. This protocol field collects all operations and errors reported since the beginning of the session.

8.3.4 Displaying objects

To display all instances of an object, e.g. the class `Employee`, we invoke the Display Instance facility by selecting the item **Display Instances** from the menu bar. In the interaction window we specify `Employee` as object name. The instances of the class `Employee` are then displayed (see Figure ??).

After selecting a displayed instance we can load the frame representation of an instance to the Telos Editor or display further instances.

8.3.5 Browsing objects

The *Graph Editor* is the preferred tool for browsing the objects managed by a `CBserver`. `CBGraph` is started by using the menu item “Graph Editor” from the “Browse” menu of `CBIVA`. Select `Employee` as the initial object to be shown in `CBGraph`. After starting `CBGraph`, it will open an internal frame, connect it with the current server, and load the `Employee` object.

The *CBGraph Editor* (described in detail in section ??) allows you to display arbitrary objects from the current *onceptBase* server. Then, we select the `Employee` object and choose the **Sub classes** option from the context menu available via the right mouse-button. We choose to only display explicit subclasses from the submenu and select the `Manager` object. The displayed graph is now expanded (figure ??).

¹⁰Historically, we used the terms ‘application’ or ‘object base’ instead ‘database’. We now believe that ‘database’ is a much better term.

¹¹In nonpersistent update mode, the database is actually copied to a temporary directory. This copy will be removed when you shutdown the server.




Figure 8.15: Display of Employee instances




Figure 8.16: The resulting graph after expanding the node Employee with subclasses

Now we expand the node `Manager`, a subclass of `Employee`, by choosing the menu item **Instances** from the popup menu for `Manager`. We select the menu item “Show all” to display all instances of `Manager`. The resulting graph is shown in figure ??.

Note that different object types are represented by different graphical objects. The instances of `Manager` are shown only as grey rectangles, because they are normal individual objects. The nodes `Manager`, `Salesman`, `Employee` etc. are shown as ovals, since these nodes are instances of the system class `SimpleClass` (see for a full description of graphical object semantics: Appendix ??).

One can move nodes and links by selecting a node or a link and then holding down the left mouse button while moving the cursor to a different position. When the button is released the selected object will be located at the current position and the related links are redisplayed. Selection and movement of multiple nodes and links is also possible. Nodes and links are selected by clicking on its label, e.g. `Manager` in figure ?. Some links like the blue specialization link in the figure have no label. Then one can select the




Figure 8.17: The resulting graph after expanding with the instances of Manager

link by clicking on the small square dot in the middle of the link. This square dot is by default invisible. It becomes visible when you click on any other node or link label in the graph, e.g. on Manager.

We can further experiment with CBGraph by showing the classes and attributes of Employee. The classes of Employee are shown by selecting “Instance of” from the popup menu. Attributes of an object can be shown by selecting “Outgoing attributes” from the menu. The next submenu will show all attribute classes that apply for the current object. In our example, Employee is an instance of Class. Therefore, it has the attribute classes `constraint`, `rule`, and `mrule` (see figure ??). The attribute class `Attribute` applies to all objects as in Telos any kind of object can have an attribute. Furthermore, all attributes of an object are member of the attribute class `Attribute`. As we want to see all attributes, we select this attribute class and select “Show all” from the next submenu. All attributes and their values will be shown in CBGraph.

CBGraph can also show implicit relationships between objects, e.g. relationships deduced by rules or the Telos axioms. For example, if we select the object John and select “Instance of” from the popup menu, we can display the implicit classes of John by selecting “All” from the next submenu. As John is an instance of Manager and Manager is a subclass of Employee, John is also an instance of Employee. As there is no explicit object John→Employee, the instantiation link between John and Employee will be represented as an implicit link, i.e. a dashed line (see figure ??).

The same applies also to attribute links. For example, the employee Herbert has an implicit boss-attribute to Phil. This can be shown by selecting “Outgoing attributes” → “boss” → “All” → “Phil” from the popup menu. Note, that the submenu “All” for the attribute class `Attribute` will be always empty as only explicit attributes can be displayed in this category.

8.3.6 Editing Telos objects

Editing Telos objects with the Telos editor

Before we are able to edit a Telos object, we have to load its frame representation in to the Telos Editor field first. For loading a Telos object to the Editor field, we choose the **Telos Editor** Button from




Figure 8.18: The graph after expanding it with the classes and the attributes of the class Employee




Figure 8.19: The graph showing implicit instantiation and attribute links

either the Display Queries oder Display Instances Browsing facilities or the Load Frame button from the *ConceptBaseWorkbench* window (see Figure ??).




Figure 8.20: The TelosEditor Field with the object Employee

Now we add an additional attribute, e.g. *education*, to the class *Employee* (for the description of the Telos syntax see Appendix ??). We have added the line *education : String* as shown in figure ???. To demonstrate error reports from the *ConceptBase.cc* user interface and how to correct them, we have made mistakes in the syntax notation of the added attribute.

By clicking the left mouse button on the **Tell** icon, the content of the editor is told to the *ConceptBase* server. Syntactical and semantical correctness is checked and the detected errors are reported to the Protocol field. The report resulting from our mistakes by specifying the new attribute is also shown in figure ??. Note, that this syntax error would have been already detected at the client side without interaction with the server if we would have enabled the option “Pre-parse Telos Frames” in the options menu.

We correct the error by adding a semicolon to the previous line and choose the **Tell** symbol again. This time, since there are no further mistakes, the additional attribute is added to the class *Employee*.

Now we can choose again the item **Outgoing Attributes** from the popup of the *Graph Editor* window for the node *Employee*. If we select “Show all” attributes of the attribute class “Attribute” the new attribute will NOT be shown. CBGraph uses an internal cache which will only be updated on request. Therefore, we select the object *Employee* and select “Validate and update selected objects” from the menu “Current connection”. The cache for the object *Employee* will be emptied. Now, showing all attributes should add the new attribute *education* to the graph.

Editing Telos objects using CBGraph

Telos objects can also be edited graphically in CBGraph. In our example, we want to add another attribute named *address* to the class *Employee*. The attribute destination of this attribute should be a new class called *Address*.




Figure 8.21: Trying to add an attribute to the class `Employee` with the resulting error report

First, we select the object `Employee` and click on the “Create Attribute” button in the tool bar or select “Add Attribute” operation from the tool bar. As we have selected the `Employee` object, it should be already inserted as source of the attribute. We have to type the label (address) and the destination of the attribute in the text fields (see figure ??). As this attribute does not belong to a specific attribute category (it is just an attribute), we do not have to specify an attribute class.

By clicking on the Ok button, `CBGraph` will create a new object for `Address` represented by the default graphical type (a gray box). Then, it will create the attribute link from `Employee` to `Address` with the label address. The result is shown in figure ??.

Now, we want to declare `Address` as an instance of `Class`. Therefore, we select `Address`, hold down the Shift-key and select `Class`. Both objects should be selected now. We click on the “Create Instantiation” button and a dialog as shown in figure ?? should appear.

As we have selected the objects in the correct way, the dialog already specifies the object we want to create (`Address` in `Class`) and we can click directly on Ok. The new instantiation link will be added to the graph.

Now, we are satisfied with our changes and want to commit them in the server. So far, the changes have been stored in an internal buffer of `CBGraph` and have not been sent to the server. We click on the “Commit” button in the upper right corner. `CBGraph` generates now Telos frames for the added objects and sends them to the server. If we did not make an error, all changes should be consistent and accepted by the server. This is shown by the appearing message box “Changes committed”. If an error occurs, an error message will displayed instead. The graphical editing is an alternative to the textual editing via the Telos Editor. It is appropriate for incremental changes to a model. Larger changes should better be made via the Telos Editor or even to an external text file that is loaded via the *File / Load Model* facility of




Figure 8.22: Adding an attribute to Employee




Figure 8.23: Adding an instantiation link between Address and Class

CB1va. If the changes were successfully told to the CBserver, CBGraph reloads the information of every visible object. In particular, the graphical types of the objects will be updated. As the object Address is declared as instance of Class, it will get the graphical type of a class, i.e. a turquoise box. The result is shown in figure ???. The object Employee is shown in the detailed component view, in this case the frame representation of the object is shown. As you can see, the attribute address is now also visible in the frame representation.




Figure 8.24: The resulting graph after commit

8.3.7 Using the query facility

Lets assume that we need to ask the server for all Employees working for Angus. We open a new Telos Editor (see menu item Browse). Then, we define a new query class (AngusEmployees) as follows:

```
AngusEmployees in QueryClass isA Employee with
constraint
    c: $ (this boss Angus) $
end
```

We can tell this query, so that it is stored in ConceptBase.cc and we can reuse it later, or we can just *ask* the query, i.e. the query will told temporarily and evaluated. If we ask the query, the answer is displayed in the Telos Editor field as well as in the history window. Figure ?? shows the CBIva with the query class and the answer.

We can also execute this query from CBGraph. From the menu “Current connection” we select “Query to server”. A dialog we ask for the name of query class. If we have told the example query, we can type `AngusEmployees` in the text field and hit on the “Submit Query”. The query will be evaluated and the objects in the result will be shown in the list box. We can select the objects which should be added to the graph (multiple selection with the Shift-key is possible) and click on the “Show objects” button. The selected objects will be added to the graph, however with no connection to existing objects.

8.4 Configuration file

The configuration options are stored in a file “.CBjavaInterface” in the home directory of the user. The settings are stored automatically on exit. You can edit the file manually with a normal text editor. It contains name-value pairs in the format `variable=value`. All variables can also be modified CBIva via the “Options” menu.

- Variables related to CBIva

PathForLoadModel: Path used by the load model dialog (contains the most recent directory selected in a dialog).




Figure 8.25: Query class and its answer

RecentConnections: Comma-separated list of recent connections in the format host/port, applies also to CBGraph.

PreParseTelosFrames: Frames are be parsed on client-side before sent to the server (true/false).

UseQueryResultWindow: Use the query result window to display results of a query (true/false).

ConnectionTimeout: Number of milliseconds the interface waits for a response of the server

LPICall: Enable LPI-Call (internal use only).

ShowLineNumbers: If set to 'true', the Telos editor of CBIva shall display line numbers.

CBIvaSmallfont: The font size of the TelosEditor text and log area as regular (small) font, default "12f".

CBIvaLargefont: The font size of the TelosEditor text and log area as large font, default "18f".

PublicCBserver: Either 'none' (=disabled) or the domain name of a computer that hosts a publicly accessible ConceptBase server. A port number can be appended as well like in 'cb-server.acme.com/4002'. Default port number is 4001. The variable PublicCBserver also applies to CBGraph and CBShell. See section ?? for more details. If the value is different from 'none', then CBIva shall attempt to connect to the public CBserver at startup.

- Variables related to CBGraph

PathForLayout: Path used by the dialog to store and load graphs (contains the most recent directory selected in a dialog).

ComponentView: Default view for the detailed representation in CBGraph (might be "frame" or "tree").

SaveLayoutWithGraphType: Layouts of CBGraph are stored with all information about graphical types (true/false).

InvalidObjectsMethod: Specifies whether objects that have been identified as invalid should be marked or deleted (mark/delete).

DiagramDesktopBackgroundColor: Background color of the desktop of CBGraph (comma-separated representation of an RGB-value).

DebugLevel: Level for debug messages. Possible values are SEVERE, WARNING, INFO, CONFIG, FINE, FINER, FINEST (according to the java.util.logging package). Default is WARNING.

ModuleSeparator: Can be either '-' or '/'.

NodeLevelAware: Enables or disables the special behavior of nodes with negative level in CBGraph. See also section ???. Default is "true".

Appendix A

Syntax Specifications

A.1 Syntax specifications for Telos frames

```
<object>          --> <objectname> <objectname> <inspec> <isaspec>
                    <withspec> <endspec>
                    | <objectname> <inspec> <isaspec> <withspec> <endspec>

<objectname>      --> ( <objectname> )
                    | <label> <bindings>
                    | <objectname> SELECTOR1 <label>
                    | <objectname> SELECTOR2 <objectname>

<bindings>        --> <empty>
                    | [ <bindinglist> ]

<bindinglist>     --> <singlebinding>
                    | <bindinglist> , <singlebinding>

<singlebinding>   --> <objectname> / <label>
                    | <label> : <objectname>

<inspec>          --> <empty>
                    | in <classlist>

<isaspec>         --> <empty>
                    | isA <classlist>

<classlist>       --> <objectname>
                    | <objectname> , <classlist>

<withspec>        --> <empty>
                    | with <decllist>

<decllist>        --> <empty>
                    | <declaration>
                    | <decllist> <declaration>

<declaration>     --> <attrcatlist> <proplist>

<attrcatlist>     --> <label>
                    | <attrcatlist> , <label>

<proplist>        --> <property>
```

```

| <proplist> ; <property>

<property>  --> <label> : <objectname>
|           | <label> : <complexref>
|           | <label> : <enumeration>
|           | <label> : <pathexpression>

<complexref>  --> <objectname> <withspec> <endspec>

<enumeration>  --> [ <classlist> ]

<pathexpression>--> <objectname> SELECTORB <pathargument>

<pathargument> --> <label>
|                 | <label> SELECTORB <pathargument>
|                 | <restriction>
|                 | <restriction> SELECTORB <pathargument>

<restriction>  --> ( <label> : <enumeration> )
|                 | ( <label> : <pathexpression> )
|                 | ( <label> : <objectname> )

<endspec>      --> end

<label>        --> ALPHANUM
|                 | LABEL
|                 | NUMBER

```

Note: ConceptBase represents internally object identifiers as *id_NUMBER* where *NUMBER* is a sequence of digits. For this reason, labels matching this pattern are forbidden in the Telos frame syntax.

A.2 Syntax of the rule and constraint language

In the definitions below the term *literal* is a synonym for predicate.

```

<assertion>    --> <rule>
|               | <constraint>

<rule>         --> forall <variableBindList> ( <formula> ) ==> <literal>
|               | <formula> ==> <literal>
|               | <literal>

<constraint>    --> <formula>

<formula>       --> exists <variableBindList> <formula>
|               | forall <variableBindList> <formula>
|               | not <formula>
|               | <formula> <==> <formula>
|               | <formula> ==> <formula>
|               | <formula> and <formula>
|               | <formula> or <formula>
|               | ( <formula> )
|               | <literal>
|               | <literal2>

<variableBindList>--> <variableBind> <variableBindList>
|                   | <variableBind>

```

```

<variableBind>  --> <varList> / <objectname>
                  | <varList> / [ <objList> ]
                  | ALPHANUM / <selectExpB>

<varList>       --> ALPHANUM , <varList>
                  | ALPHANUM

<objectname>    --> <label>
                  | <selectExpA>
                  | <deriveExp>

<label>         --> ALPHANUM
                  | LABEL
                  | NUMBER

<literal>       --> FUNCTOR ( <literalArgList> )
                  | ( <literalArg> <infixSymbol> <literalArg> )
                  | ( <arExpr> COMPSYMBOL <arExpr> )
                  | ( <literalArg> <label>/<label> <literalArg> )
                  | BOOLEAN

<literal2>      --> ( <label> in <selectExpB> )
                  | ( <selectExpA> in <selectExpB> )
                  | ( <selectExpB> isA <selectExpB> )
                  | ( <selectExpB> = <selectExpB> )

<infixSymbol>   --> INFIXSYMBOL
                  | <label>

<literalArgList>--> <literalArg> , <literalArgList>
                  | <literalArg>

<literalArg>    --> <objectname>

<arExpr>        --> <arExpr> + <arTerm>
                  | <arExpr> - <arTerm>
                  | <arTerm>

<arTerm>        --> <arTerm> * <arFactor>
                  | <arTerm> / <arFactor>
                  | <arFactor>

<arFactor>      --> ( <arExpr> )
                  | <objectname>
                  | <funExpr>

<selectExpA>    --> <selectExpA> <selector> <selectExpA>
                  | ( <selectExpA> )
                  | <label>

<selector>      --> SELECTOR1
                  | SELECTOR2

<deriveExp>     --> <label> [ <deriveExpList> ]
                  | <label> [ <literalArgList> ]
                  | <funExpr>

```

```

<funExpr>      --> <label>()
                | <label>(<literalArgList>)

<deriveExpList> --> <singleExp> , <deriveExpList>
                | <singleExp>

<singleExp>    --> <literalArg> / <label>
                | <label> : <label>

<selectExpB>   --> <label> SELECTORB <label>
                | <label> SELECTORB <selectExpB2>
                | <label> SELECTORB <restriction>

<selectExpB2>  --> <selectExpB>
                | <restriction> SELECTORB <label>
                | <restriction> SELECTORB <selectExpB2>
                | <restriction> SELECTORB <restriction>

<restriction>  --> ( <label> : <label> )
                | ( <label> : <selectExpA> )
                | ( <label> : <selectExpB> )
                | ( <label> : [ <objList> ] )

<objList>      --> <objectname> , <objList>
                | <objectname>

```

A.3 Syntax of active rules

The event, condition and actions of an ECArule are specified as a special assertion. Therefore, the syntax is an extension of the *normal* assertion language, shown in the section before.

```

<ecarule>      --> <variableBindList>
                ON [TRANSACTIONAL] <ecaevent> [FOR ALPHANUM]
                <ifclause> <ecacondition>
                DO <actionlist>
                <optelseaction>

<ifclause>     --> IF
                | IFNEW

<ecaevent>    --> <eventop>(<literal>)
                | <eventop> <literal>
                | <askop>(<literalArg>)
                | <askop> <literalArg>

<eventop>     --> Tell | tell
                | Untell | untell

<askop>       --> Ask | ask

<ecacondition> --> <condformula>
                | true
                | false

<condformula> --> <literal>
                | not <condformula>
                | <condformula> and <condformula>
                | <condformula> or <condformula>

```

```

| ( <condformula> )

<actionlist> --> <action> , <actionlist>
| <action>

<action> --> <actionop>(<literal>)
| <actionop> <literal>
| noop
| reject

<actionop> --> Tell | tell
| Untell | untell
| Retell | retell
| Ask | ask
| Call | call | CALL
| Raise | raise

<optelseaction> --> ELSE <actionlist>
| <empty>

```

A.4 Terminal symbols

```

ALPHANUM --> [a-z|A-Z|0-9|ACCENTCHAR]+

ACCENTCHAR --> umlauts and accents that are included in the 8bit ASCII code

LABEL --> sequences of characters excluding .|' "$ ; ! ^ - > = , ( ) [ ] { } / and
| special characters like newlines, tabs, backspace, blanks
| any sequence of characters enclosed in double quotes ("),
| a double quote must be escaped by \ which must be escaped by \
| any sequence of characters enclosed in $ except $,
| which must be escaped by \

NUMBER --> REAL | INTEGER

REAL --> [-]?([0-9]+\.[0-9]*|[0-9]*\.[0-9]+)([Ee][-+]?[0-9]+)?

INTEGER --> [-]?[0-9]+

BOOLEAN --> TRUE | FALSE

FUNCTOR --> From | To | A | Ai | AL | In
| Isa | Label | P | LT | GT | LE | GE | EQ | NE | IDENTICAL
| UNIFIES | Known

COMPSYMBOL --> < | > | <= | >= | = | <> | == | \=

INFIXSYMBOL --> COMPSYMBOL | in | isA

SELECTOR1 --> "!" | "^"

SELECTOR2 --> "->" | "=>"

SELECTORB --> "." | "|"

```

A.5 Syntax specifications for SML fragments

This format is only internally used to represent Telos frames as Prolog terms. It is included only for historical reasons.

```
<SMLfragment> --> SMLfragment(<what> , <in_omega> , <in> , <isa> , <with> )

<what>          --> what(<object> )

<in_omega>      --> in_omega(nil)
                  |   in_omega([<classlist> ])

<in>            --> in(nil)
                  |   in([<classlist> ])

<isa>           --> isa(nil)
                  |   isa([<classlist> ])

<with>          --> with(nil)
                  |   with([<attrdecllist> ])

<classlist>     --> class(<object> )
                  |   <classlist> , class(<object> )

<attrdecllist> --> attrdecl(<attrcategorylist> , <propertylist> )
                  |   <attrdecllist> , attrdecl(<attrcategorylist> , <propertylist> )

<attrcategorylist>--> nil
                  |   [ <labellist> ]

<propertylist>  --> nil
                  |   [ <propertylist2> ]

<propertylist2>--> property(<label> , <propertyvalue> )
                  |   <propertylist2> , property(<label> , <propertyvalue> )

<propertyvalue>--> <object>
                  |   <selectExpB>
                  |   enumeration( [ <classlist> ] )
                  |   [ <SMLfragment> ]

<selectExpB>    --> selectExpB( <restriction> , <selectOperator> , <selectExpB> )
                  |   selectExpB( <restriction> , <selectOperator> , <object> )
                  |   selectExpB( <object> , <selectOperator> , <selectExpB> )
                  |   selectExpB( <object> , <selectOperator> , <object> )

<restriction>   --> restriction( <label> , <selectExpB> )
                  |   restriction( <label> , enumeration( [ <classlist> ] ) )
                  |   restriction( <label> , <object> )

<selectOperator> --> dot
                  |   bar

<labellist>     --> <label>
                  |   <labellist> , <label>

<label>         --> ALPHANUM
                  |   LABEL
                  |   NUMBER
```

```

<object>      --> derive([ <substlist> ])
                | <selectexp>

<substlist>   --> <singlesubst>
                | <substlist> , <singlesubst>

<singlesubst> --> substitute(<object> , <label> )
                | specialize(<label> , <label> )

<selectexp>   --> <label>
                | select(<selectexp> , SELECTOR1, <label> )
                | select(<selectexp> , SELECTOR2, <selectexp> )

```

Appendix B

O-Telos Axioms

O-Telos is the variant of Telos (originally defined by John Mylopoulos, Alex Borgida, Manolis Koubarakis and others) that is used by the ConceptBase system. This list is the complete set of pre-defined axioms of O-Telos and thus defines the semantics of a O-Telos database (without user-defined rules and constraints). The subsequent axioms are written in a first-order logic syntax but all can be converted to Datalog with negation (though there is some choice in the conversion wrt. mapping to rules or constraints).

- Axiom 1: Object identifiers are unique.
$$\forall o, x_1, n_1, y_1, x_2, n_2, y_2 \ P(o, x_1, n_1, y_1) \wedge P(o, x_2, n_2, y_2) \Rightarrow (x_1 = x_2) \wedge (n_1 = n_2) \wedge (y_1 = y_2)$$
- Axiom 2: The name of individual objects is unique.
$$\forall o_1, o_2, n \ P(o_1, o_1, n, o_1) \wedge P(o_2, o_2, n, o_2) \Rightarrow (o_1 = o_2)$$
- Axiom 3: Names of attributes are unique in conjunction with the source object.
$$\forall o_1, x, n, y_1, o_2, y_2 \ P(o_1, x, n, y_1) \wedge P(o_2, x, n, y_2) \Rightarrow (o_1 = o_2) \vee (n = in) \vee (n = isa)$$
- Axiom 4: The name of instantiation and specialization objects (*in*, *isa*) is unique in conjunction with source and destination objects.
$$\forall o_1, x, n, y, o_2 \ P(o_1, x, n, y) \wedge P(o_2, x, n, y) \wedge ((n = in) \vee (n = isa)) \Rightarrow (o_1 = o_2)$$
- Axioms 5,6,7,8: Solutions for the predicates *In*, *Isa*, and *A* are derived from the object base.
$$\begin{aligned} \forall o, x, c \ P(o, x, in, c) &\Rightarrow In(x, c) \\ \forall o, c, d \ P(o, c, isa, d) &\Rightarrow Isa(c, d) \\ \forall o, x, n, y, p, c, m, d \ P(o, x, n, y) \wedge P(p, c, m, d) \wedge In(o, p) &\Rightarrow AL(x, m, n, y) \\ \forall x, m, n, y \ AL(x, m, n, y) &\Rightarrow A(x, m, y) \end{aligned}$$
- Axiom 9: An object *x* may not neglect an attribute definition in one of its classes.
$$\begin{aligned} \forall x, y, p, c, m, d \ In(x, c) \wedge A(x, m, y) \wedge P(p, c, m, d) &\Rightarrow \\ \exists o, n \ P(o, x, n, y) \wedge In(o, p) \end{aligned}$$
- Axioms 10,11,12: The *isa* relation is a partial order on the object identifiers.
$$\begin{aligned} \forall c \ In(c, \#Obj) &\Rightarrow Isa(c, c) \\ \forall c, d, e \ Isa(c, d) \wedge Isa(d, e) &\Rightarrow Isa(c, e) \\ \forall c, d \ Isa(c, d) \wedge Isa(d, c) &\Rightarrow (c = d) \end{aligned}$$
- Axiom 13: Class membership of objects is inherited upwardly to the superclasses.
$$\forall p, x, c, d \ In(x, d) \wedge P(p, d, isa, c) \Rightarrow In(x, c)$$
- Axiom 14: Attributes are "typed" by their attribute classes.
$$\forall o, x, n, y, p \ P(o, x, n, y) \wedge In(o, p) \Rightarrow \exists c, m, d \ P(p, c, m, d) \wedge In(x, c) \wedge In(y, d)$$

- Axiom 15: Subclasses which define attributes with the same name as attributes of their superclasses must refine these attributes.

$$\forall c, d, a_1, a_2, m, e, f$$

$$Isa(d, c) \wedge P(a_1, c, m, e) \wedge P(a_2, d, m, f) \Rightarrow Isa(f, e) \wedge Isa(a_2, a_1)$$
- Axiom 16: If an attribute is a refinement (subclass) of another attribute then it must also refine the source and destination components.

$$\forall c, d, a_1, a_2, m_1, m_2, e, f$$

$$Isa(a_2, a_1) \wedge P(a_1, c, m_1, e) \wedge P(a_2, d, m_2, f) \Rightarrow Isa(d, c) \wedge Isa(f, e)$$
- Axiom 17: For any object there is always a unique "smallest" attribute class with a given label m .

$$\forall x, m, y, c, d, a_1, a_2, e, f (In(x, c) \wedge In(x, d) \wedge P(a_1, c, m, e) \wedge P(a_2, d, m, f)$$

$$\Rightarrow \exists g, a_3, h In(x, g) \wedge P(a_3, g, m, h) \wedge Isa(g, c) \wedge Isa(g, d))$$
- Axioms 18-22: Membership to the builtin classes is determined by the object's format.

$$\forall o, x, n, y (P(o, x, n, y) \Leftrightarrow In(o, \#Obj))$$

$$\forall o, n (P(o, o, n, o) \wedge (n \neq in) \wedge (n \neq isa) \Leftrightarrow In(o, \#Indiv))$$

$$\forall o, x, c (P(o, x, in, c) \wedge (o \neq x) \wedge (o \neq c) \Leftrightarrow In(o, \#Inst))$$

$$\forall o, c, d (P(o, c, isa, d) \wedge (o \neq c) \wedge (o \neq d) \Leftrightarrow In(o, \#Spec))$$

$$\forall o, x, n, y (P(o, x, n, y) \wedge (o \neq x) \wedge (o \neq y) \wedge (n \neq in) \wedge (n \neq isa) \Leftrightarrow In(o, \#Attr))$$
- Axiom 23: Any object falls into one of the four builtin classes.

$$\forall o In(o, \#Obj) \Rightarrow In(o, \#Indiv) \vee In(o, \#Inst) \vee In(o, \#Spec) \vee In(o, \#Attr)$$
- Axioms 24-28: There are five builtin classes.

$$P(\#Obj, \#Obj, Proposition, \#Obj)$$

$$P(\#Indiv, \#Indiv, Individual, \#Indiv)$$

$$P(\#Attr, \#Obj, attribute, \#Obj)$$

$$P(\#Inst, \#Obj, InstanceOf, \#Obj)$$

$$P(\#Spec, \#Obj, IsA, \#Obj)$$
- Axiom 29: Objects must be known before they are referenced. The operator $\preceq$ is a (predefined) total order on the set of identifiers.

$$\forall o, x, n, y P(o, x, n, y) \Rightarrow (x \preceq o) \wedge (y \preceq o)$$
- Axioms 30,31 (axiom schemas): For any object $P(p, c, m, d)$ in the extensional object base we have two formulas for "rewriting" the In and A predicates. The In is mapped to a unary predicate where the class name is forming part of the predicate name and the A predicates is mapped to a binary predicate that carries the identifier of the class of the attribute in its predicate name. Internally, user-defined deductive rules that derive In and A predicates will also be rewritten accordingly. This extends the choices for static stratification.

$$\forall o In(o, p) \Rightarrow In.p(o)$$

$$\forall o, x, n, y P(o, x, n, y) \wedge In(o, p) \Rightarrow A.p(x, y)$$

The following axioms are taken from papers on Telos (i.e. formulated by Mylopoulos, Borgida, Koubarakis, Stanley and Greenspan): axioms 2, 3, 4, 10, 12, 13, 14. Axiom 1 is probably also in an earlier Telos paper though we could not immediately find it there. The axioms 15 and 16 are similar to the structural ISA constraint of Taxis [?] for attributes. In O-Telos, we do however not inherit attributes downward to subclasses but rather constrain refined attributes at subclasses in the sense of co-variance. Moreover, attributes in O-Telos are objects as well, hence the notion of specialization is more complicated than for the Taxis case. Axiom 17 is needed to be able to uniquely match an attribution predicate to a most specific attribute. This is utilized in the compilation of logical expressions, in particular for generating triggers that only evaluate the affected logical expressions when an update occurs. The remaining axioms 18-28 are also specific to O-Telos. They define the five predefined objects in O-Telos. Axiom 29 takes care that objects cannot refer via its source/destination parts to objects that were defined later than the object itself. This virtually forbids to define a link between two objects when one of the objects is not yet defined.

While this sounds natural, we need to postulate it. Otherwise, we can't guarantee that we can refer to any object by a name. Axioms 30 and 31 are used to transfer instantiation and attribution facts from the extensional databases to the intensional database. They have more a technical purpose in the mapping of logical expressions to Datalog.

While O-Telos has just five predefined objects and 31 predefined axioms, the ConceptBase system has many more pre-defined objects to provide a better modeling experience and for representing concepts like query classes, active rules, functions etc. They are in a way also predefined but are less essential in understanding the foundations. So, O-Telos is the foundation of ConceptBase but ConceptBase has more pre-defined constructs than those mentioned in the axioms of O-Telos.

ConceptBase allows to add user-defined rules and constraints. The semantics of an O-Telos database including such rules and constraints is the perfect model of the deductive database with the $P(o, x, n, y)$ as the only extensional predicate and all axioms and user-defined rules/constraints as deductive rules. Note that integrity constraints can be rewritten to deductive rules deriving the predicate *inconsistent*.

This list of axioms is excerpted from M.A. Jeusfeld: Änderungskontrolle in deduktiven Objektbanken. Dissertation Universität Passau, Germany, 1992. Available as Volume DISKI-19 from INFIX-Verlag, St. Augustin, Germany or via <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/d340216/diskil9.pdf> (in German).

Axioms 19-21 have been corrected after Christoph Radig found an example that led to the undesired instantiation of an individual object to #Inst or #Spec, respectively.

Appendix C

Graphical Types

The concept of a *graphical type* enables the specification of an external graphical presentation for ConceptBase objects. The graphical type is declared using a special pre-defined attribute category. An application program then uses this information to determine the graphical presentation of an object.

The next subsection introduces the basic concepts behind graphical types, while section ?? presents the standard graphical type definitions for the ConceptBase Graph Editor. Section ?? describes the definition of application-specific types.

C.1 The graphical type model

A specific graphical type is defined as an instance of the object `GraphicalType`. `CBGraph` uses the subclass `JavaGraphicalType`. Instances of this class specify a graphical representation of an object by defining graphical attributes such as shape, color, line thickness, font etc. Since the actual attributes and their admissible value depend on the used visualization tool, the definition of `GraphicalType` looks very simple.

```
GraphicalType in Class
end
```

The declaration of a graphical type for a concrete object is done by using the attribute `graphtype` which is defined for `Proposition` and therefore available for all objects:

```
Proposition with
  attribute
    graphtype : GraphicalType
end
```

The attribute can be defined explicitly for an object or can be specified by using a deductive rule (see section ?? for an example). One can attach a priority value to each graphical type. If there multiple `graphtype` attributes defined for one object, the graphical type with the highest priority value will be used by `CBGraph`.

Many modeling applications require multiple notations to provide different perspectives on the same set of objects. Each perspective emphasizes a specific aspect of the world, such as the data-oriented, the process-oriented and the behavior-oriented viewpoint, and uses an aspect-specific notation. A graphical notation (as e.g. the Entity-Relationship diagram) typically consists of a set of different graphical symbols (as e.g. diamonds, rectangles, and lines). A *graphical palette* is used to combine the set of graphical types that together form a notation:

```
Individual GraphicalPalette in Class with
  attribute
    contains : GraphicalType;
    default : GraphicalType
end
```

In such a setting the same object may participate in different perspectives. ConceptBase offers the possibility to specify multiple graphical types for the same object. A tool can then provide different graphical views on the same object. To get the desired graphical type of an object under a specific palette, an application program specifies the name of the actual graphical palette as answer format when querying the ConceptBase server. Although this mechanism is available for arbitrary application programs we restrict our description to the CBGraph Editor.

The default specification serves as a catch all: an answer object, for which none of the graphical types of the current palette is specified, is presented using the default graphical type of that palette.

C.2 The standard graphical types

CBGraph is implemented using the Java Programming Language. It is entirely based on the Swing toolkit (package `javax.swing`). The graphical objects shown in CBGraph are all instances of the class `JComponent` in the `javax.swing` package. User-defined representations of objects can be provided by overwriting a specific class of CBGraph (details are given below).

C.2.1 The extended graphical type model

Based on our experience with a legacy graph browser for X11, we have extended the graphical type model for the CBGraph Editor. First, the class `GraphicalType` has been specialized by a class `JavaGraphicalType`:

```
Class JavaGraphicalType isA GraphicalType with
  attribute
    implementedBy : String;
    property : String;
    priority : Integer
  rule
    rPriority : $ forall jgt/JavaGraphicalType (not (exists i/Integer
      A_e(jgt,priority,i))) ==> A(jgt,priority,0) $
end

Individual DefaultIndividualGT in JavaGraphicalType with
  property
    bgcolor : "210,210,210";
    textcolor : "0,0,0";
    linecolor : "0,0,0";
    shape : "i5.cb.graph.shapes.Rect"
  implementedBy
    implBy : "i5.cb.graph.cbeditor.CBIndividual"
end
```

The object `DefaultIndividualGT` is an example for the instantiation of a graphical type. The attribute `implementedBy` specifies the full name of the Java class that provides the implementation for this graphical type. This class has to be a sub class of `"i5.cb.graph.cbeditor.CBUserObject"`. The `property` attribute specifies name-value pairs which will be used by the Java implementation to set certain properties, e.g. color, shape, font¹. The `priority` value is used to resolve ambiguity if multiple graphical types apply to one object. The graphical type with the highest priority will be used. The rule specifies a default value of 0 for the priority.

The graphical palette has also been extended. There are now defaults for different types of objects, and graphical types for implicit links can be defined. Thus, the default attribute defined in `GraphicalPalette` will not be used anymore. The `contains` attribute has still to be used, i.e. a graphical type will only be used if it is also contained in the current graphical palette. Although the attributes are not

¹Colors are given as RGB color value, e.g. 210,210,210 is light grey and 0,0,0 is black.

declared as single and necessary, each graphical palette should have exactly one value for the each of the default and implicit attributes.

```
Class JavaGraphicalPalette isA GraphicalPalette with
  attribute
    defaultIndividual : JavaGraphicalType;
    defaultLink : JavaGraphicalType;
    implicitIsA : JavaGraphicalType;
    implicitInstanceOf : JavaGraphicalType;
    implicitAttribute : JavaGraphicalType;
    palproperty : String
end
```

The attribute `palproperty` is used for declaring any number of properties of a palette. The properties are passed to the CBGraph Editor when it loads the palette at startup time. CBGraph supports the following properties for palettes:

bgcolor: sets the background color of the windows that displays the graph; format should be "`r, g, b`", e.g. "`255, 255, 255`" for white

bgimage: sets the background image for the graph windows; the image shall be specified by the URL to a PNG, GIF, or JPG image; it is scaled by the CBGraph editor to fit into the internal window showing a graph of this palette

longtitle: used for setting the title of the graph windows employing this palette; if no long title is specified, then CBGraph uses the name of the palette itself for forming the window title; if the longtitle is set to the empty string "", then it will cause CBGraph not to include it in the title of the graph windows

The purpose of the background image is to highlight regions of a graph, e.g. regions for instances, classes, and meta classes. Another typical use is to support canvasses like the business model canvas <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/d3595098/bmg-egadget.png> used in the Telos models described in the CB-Forum at <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/d3595098/>. Only "http" URLs are supported.

If a background image is specified for a palette shown in an internal window of CBGraph, then CB-Graph links it with the size and zoom factor of the internal window. Initially, the zoom factor is set to 100% and the internal window size is set to display the image in its original resolution, provided that it fits well to the screen size. You can then resize the internal window and the background image shall be resized accordingly. Analogously, the image is resized when the zoom factor changes. The background image is also stored in the GEL file, see section ??.

C.2.2 Default graphical types

For the standard objects, there are a number of predefined graphical types. There are contained in the graphical palette `DefaultJavaPalette` which is used by default by the CBGraph Editor.

type of object	graphical type	style
Individuals	DefaultIndividualGT	gray box
Links	DefaultLinkGT	thin black line with label
InstanceOf	DefaultInstanceOfGT	green line without label
IsA	DefaultIsAGT	blue line without label; white edge heads
Attribute	DefaultAttributeGT	black line with label
Class	ClassGT	turquoise box
SimpleClass	SimpleClassGT	pink oval
MetaClass	MetaClassGT	light blue oval
MetametaClass	MetametaGT	bright green oval
QueryClass	QueryClassGT	red oval
Implicit In	ImplicitInstanceOfGT	dashed green line
Implicit IsA	ImplicitIsAGT	dashed blue line; white edge heads
Implicit Attribute	ImplicitAttributeGT	dashed black line

The object `DefaultJavaPalette` has also some rules which define the default relationship between objects and graphical types, e.g. all instances of `Class` have the graphical type `ClassGT`.

If you want to customize the graphical types for your model, then you must define the new graphical types (see below) and then add them to a new graphical palette as instance of `JavaGraphicalPalette`. Take the default graphical palette as a starting point since you may want to reuse some of the existing graphical types. See file `03-ERD-GTs.sml` at <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/188651> for an example.

The safest way to define a correct graphical palette is to use the intermediate `XPalette` as superclass of the new palette. It can be found in the CB-Forum at <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/d4204468/StandardPalette.sml.txt>. With `XPalette` you can define a new palette without having to worry about values for required default graphical types.

```
BMG_Palette in JavaGraphicalPalette isA XPalette with
  palproperty
    bgimage: "http://conceptbase.sourceforge.net/CBICONS/bgimages/bmgcolor.png";
    longtitle: "Business Model"
  contains
    bmg1: Customer_GT;
    bmg2: Revenue_GT;
    bmg3: CustomerRelationship_GT;
    bmg4: Channel_GT;
    ...
end
```

An alternative to `DefaultJavaPalette` is the general-purpose `StandardPalette` available from the CB-Forum at <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/d4204468/StandardPalette.sml.txt>. It assigns white color to most nodes and uses broken green lines both for explicit and deduced instantiations. Just add the definitions to your Telos model to use this palette.

C.3 Customizing the graphical types

To support the user in defining his own graphical types we provide some examples and documentation of the properties.

There are two ways to customize the graphical types:

- Defining new graphical types with dedicated graphical properties using the provided implementations
i5.cb.graph.cbeditor.CBIndividual (for nodes) and *i5.cb.graph.cbeditor.CBLink* (for links)

- Defining new graphical types with a different implementation class which extends `i5.cb.graph.cbeditor.CBUserObject` (or `CBIndividual` or `CBLink`); this option requires changes to the Java source code of `CBGraph`

Both possibilities will be presented in the next two subsections.

C.3.1 Graphical properties of nodes and links

The easiest way to modify the representation of an object in the `CBGraph` Editor is to load an existing graphical type, modify its properties and store it as a new graphical type.

The properties available and their meaning are given in the following. Note that colors have to be given as RGB color value, e.g. "0,0,0" is black, "255,0,0" is red, "255,255,255" is white, etc. Furthermore, all attributes have to be strings, even if they are just numbers, e.g. use "1" instead of 1 as attribute value.

bgcolor: Background color of the shape (default: invisible).

textcolor: Foreground color of the shape (i.e. text color) (default: black "0,0,0").

linecolor: Color of the border of the shape (default: invisible).

linewidth: Width of the border of the shape (default: "1").

edgecolor: Color of the edge (default: black "0,0,0"); for `CBLink` only.

edgeheadcolor: Color of the edge head; the edge head is drawn in edge color if no edge head color is defined; for `CBLink` only.

edgeheadshape: Shape of the edge head at the destination side. If set to "none", then the edge head has no shape. Possible other values are listed in the table below; for `CBLink` only.

edgewidth: Width of the edge (default: "1"); for `CBLink` only.

edgestyle: possible values are: "continuous", "dashed", "dotted", "dashdotted", "ldashed" (dashed with longer intervals), "ldotted" (default: "continuous"); for `CBLink` only.

shape: The name of the class representing the shape and implementing the interface `i5.cb.graph.shape.IGraphShape` (default: no shape). The package `i5.cb.graph.shape` defines some useful default shapes, see below for details. The shape will be drawn in the background of the small component. In the default implementation, the small component is a transparent `JLabel`, thus the shape is completely visible. Note, that this might not be the case if you are going to change the implementation of a graphical type (see subsection ?? below).

image: The location (URL) of an image icon file that shall be used to display a node (`CBIndividual`) The image tag is a replacement for the shape attribute but can also be combined with a shape. The image icon can be in PNG, GIF, or JPG format. See subsection ?? for more details.

textposition: Relative position of the node's text label to the image icon. This property is only evaluated if a graphical type defines an image icon. Possible values are "center", "left", "right", "top", and "bottom" (default).

label: The label to be used for this object instead of the object name.

labellength: The maximum number of characters displayed as label of an object in `CBGraph` (default: "40"). A label that exceeds the length is truncated to the maximum length and the last four characters are replaced by "..." in the display in `CBGraph`.

align: Alignment of the label; possible values are "center", "left", "right", "top", "bottom", "topleft", "topright", "bottomleft", and "bottomright"; default is "center".

size: Initial size of the node in pixels, e.g. "20x20"; the non-numeric values "resizable" (node size can be resized) and "wrap" (node size can be resized and label will be wrapped) are allowed as well. The value "wrap_" works like "wrap" but shall also replace "_" in the node label by a blank. This allows to handle very long labels in combination with the `labellength` property. If the size property is set, the user can also resize the element via mouse actions (default: undefined, then the size is set by `CBGraph`).

location: Designate the initial location of a node or the label of an edge. The value shall be in the format "x,y". For example, the value "10,100" has the x coordinate 10 and the y coordinate 100. This property may be useful when certain nodes should have a given initial location, e.g. for canvas nodes that contain other nodes. (default: undefined, then the location is set by `CBGraph`).

freeze: The value "yes" indicates that the node (or the edge label) is fixed to its current location in the graph editor. The value "no" (default) indicates that the node can be freely moved. You can set this flag also via the "gproperty" attribute individually for each object (see section ??).

nodelevel: The level of the node relative to the standard node layer (=200) in the graph's diagram. Negative values put the node more in the background, positive values more in the foreground. Use this feature if you want to put certain nodes on top of each other (default: "0").

font: Name of the font to be used for the shape (e.g., "Arial", default: Default font of Java).

fontsize: Size of the font in pixels (default: default font size of Java).

fontstyle: The style of the font (e.g., "bold", "italic", "underlined", "bold,italic", ...).

clickaction: The name of a query class that shall be called directly when an object with this graphical type is clicked. See section ?? for details.

Edges with empty label (*anonymous* edge) are displayed with a square dot in the middle of the edge². The color of the square dot is by default the `edgecolor` and its size is set to 6 pixels. If the graphical type of an anonymous edge has `bgcolor` defined, then the square dot is adjusted to the `edgewidth` and displayed in `bgcolor`. If you set an explicit `bgcolor` for an edge, then the bounding box around the edge label shall be painted in that color. If you choose as `bgcolor` the same value as for the `bgcolor` of the palette, then edge labels appear more readable.

Do not forget to include the new graphical type into the graphical palette. It is not necessary to define a new graphical palette, you can extend the default palette. Furthermore, you have to define the `graphtype` attribute of some object in such a way that it refers to the new graphical type. Make sure, that the new graphical type has a higher priority than other graphical type which might apply (10 is the highest priority of the default graphical types).

The color strings in `bgcolor`, `textcolor`, `linecolor`, and `edgecolor` are encoded in the format "r,g,b", where r, b and g represent the red, green, and blue share of the color. All values must be from 0 to 255. The color string "0,0,0" results in black and "255,255,255" results in white. You can also add a so-called alpha value for the transparency of the color as fourth component of a color string. The value "255" stands for opaque (not transparent) colors. This is also the default. The smallest value "0" stands for maximal transparency, i.e. the color is not visible at all. Any value in between is a relative transparency. For example, "255,0,0,127" represents a red color that is about 50% transparent with respect to objects below such as the background.

An example of user-defined graphical types can be found in ??, see also the CB-Forum at <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/188651> for a complete specification of ER diagrams including graphical types.

Below are the supported edge head shape (property `edgeheadshape`). Theoretically, you can also use the node shapes like `Rect` but they are not configured specifically for edge heads and would be rendered in tiny sizes.

²An empty label is a label equal to "". Such labels only work for edges. If you want to display an node with an empty label, then set the label to " " (one blank character).

edge head shape	description
Arrow	triangular arrow head (default for thicker edges)
ArrowVee	vee-shaped arrow head (default for thin edges)
SmallArrow	small arrow head with straight base
RevArrow	reversed arrow (base pointing to the object))
HalfArrow	half arrow head
Karo	diamond-shaped arrow head
Square	square arrow head
Circular	circular arrow head (approximated)
Caret	caret shaped arrow head
Bar	small bar orthogonal to the edge line
Dot	small square arrow head
none	no arrow head

C.3.2 Node levels

CBGraph paints the nodes and edges of a graph in a so-called layered pane. This helps to separate nodes from edges and from interactive elements such as pop-up menus. The default absolute level for a node is 200 and the default absolute level for an edge is 100. That means that nodes are by default painted on top of edges, i.e. the node's shape is painted over an edge if they overlap. In some modeling languages, one may want to have certain elements always painted over some other elements. For example, the process elements of a BPMN process model should be painted on top of the pool, in which they are defined. Or consider a traffic light element that is composed of red, yellow and green lights. Then the symbol for the traffic light element should be painted below the symbol for the three part lights.

This ordering can be achieved by the *nodelevel* property for the graphical types. The *nodelevel* property is a relative increment to the default absolute node level 200. For example, by setting the node level to "-1", the resulting absolute node level shall be 199. By setting the relative node level to "-101", the resulting absolute node level would be 99, i.e. even below the level of edges.

As an example consider the traffic light scenario. The node level is set to "-1", hence it shall be painted behind the other node elements. The example is taken from the CB-Forum at <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/3762781>.

```
TrafficLight_GT in Class,JavaGraphicalType with
property
    textcolor : "255,255,255";
    linecolor : "0,0,0";
    linewidth: "3";
    bgcolor : "60,60,60";
    shape : "i5.cb.graph.shapes.RoundRectangle";
    size: "resizable";
    align : "bottom";
    nodelevel: "-1"
implementedBy
    implBy : "i5.cb.graph.cbeditor.CBIndividual"
priority
    pr : 22
rule
    gtrule: $ forall x/TrafficLight (x graphtype TrafficLight_GT) $
end
```

You can also use positive node levels to explicitly specify that the nodes are painted in the foreground of other nodes. The default relative node level is "0". If nodes have the same level, then they are painted in the order in which they are added to the diagram.

The node selection in CBGraph is adapted to take the node level into account. If a node with a negative level is selected by a left mouse click, then all nodes with a higher level whose center point is contained in the bounds of the first node also get selected. Nodes with negative levels are interpreted as a kind of a container. So, selecting and moving them is simplified by this behavior. You can also disable this behavior by the configuration variable "NodeLevelAware", see section ??.

C.3.3 Click actions

A click action is a property of a graphical type and contains the name of a query class as a string. A simple example is:

```
clickaction: "fireTransition";
```

If an object has a graph type with a click action, then the corresponding query class is called using the object name as single parameter. For example, if `t1` is the object name, then a click on the object in CBGraph will result in calling the query `fireTransition[t1]`. It is assumed that the ConceptBase server includes an active rule that is triggered by the query call. Hence, such calls can result in an update to the database. CBGraph shall refresh its graph after performing the query call to show the effect of the database update to the graph. You can also specify a click action with arity zero:

```
clickaction: "fire/0";
```

In this case the name of the clicked object is not included as a parameter of the query call. A click action like "fireTransition" is equivalent to "fireTransition/1".

Click actions let a graph directly interact with the ConceptBase server. Each click on a node whose graphical type has a click action will result in a corresponding query call that triggers active rules – assuming that there are active rules matching the query call. The active rule in the CBserver can change the database state, but it can also trigger calls to external programs.

You can also specify clickactions with two arguments like in

```
clickaction: "playMove/2";
```

In such cases, CBGraph will prepend the username before the object name of the node that has been clicked. The generated query call would look like `playMove[jonny,m1]`. Note that the query must have two arguments in this case, e.g.

```
GenericQueryClass playMove isA Position with
  parameter
    arg1: CB_User;
    arg2: Move
  ...
end
```

Note that the first argument for the username must have a label (`arg1`) that is lexicographically ordered before the label of the second argument (`arg2`). The username is the same that is used by the CBGraph tool to register to the CBserver. That user is then stored as instance of the predefined class `CB_User`.

Another option with click actions is to limit the scope of nodes and links in the current diagram that are refreshed after calling the click action. The click action can invoke an active rule which changes the database state. Consequently, certain objects in the diagram may get a new graphical type. By default, CBGraph shall refresh all nodes and links in the diagram after executing a click action. This can be rather slow when the displayed graph is large. The option "-n" allows to limit the refresh to the neighborhood of the selected object. The neighborhood is defined as the set consisting of the selected object, the direct neighbors object of the selected object, the direct neighbors of those neighbors, and all the links in between. Note that this only refers to the objects displayed in the graph!

You can enable the "neighbor" refresh by adding the string "-n" to a click action like in

```
clickaction: "fireTransition -n";
```

The "-n" option is not guaranteed to work correctly since some objects outside the neighborhood may be affected by the click action. Hence, only use this when you know that the effect is bound to the neighborhood and when the displayed diagram has all the required links displayed to compute the neighborhood.

You can enable and disable click actions by a checkbox in the options menu of CBGraph. The setting is also stored in the configuration file `.CBjavaInterface`. The entry is called "ClickActions" there.

See <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/3762781> for examples.

C.3.4 Providing a new implementation for a graphical type

If you got the sources of ConceptBase, then you can add your own implementation for graphical types. The sources are in files `CBPOOL-*.zip` downloadable via <http://conceptbase.sourceforge.net/CB-Download.html>. Its subdirectory `AdminPOOL` contains instructions on how to compile the ConceptBase system under Linux/Unix. Adapting the source code requires *considerable programming skills* in Java as well as an understanding of the way how CBGraph is implemented. Hence, this section is only meant for the programmers among you.

The `implementedBy` attribute of a graphical type specifies the class name of a Java class implementing this graphical type. This class has to be a subclass of "i5.cb.graph.cbeditor.CBUserObject". If you are going to implement your own class, it is most useful to extend `CBIndividual` or `CBLink` in the package `i5.cb.graph.cbeditor`, i.e. in directory `ProductPOOL/java/i5/cb/graph/cbeditor` of the `CBPOOL`.

The class `CBUserObject` provides several methods which can be overwritten to implement your own graphical types:

Component `getSmallComponent()` returns the small component (i.e. the component which is shown first). The return value should be an instance of `javax.swing.JComponent` although only `java.awt.Component` is required as return value. AWT Component will probably not work correctly in the Swing-based CBGraph Editor. If the method returns null, then the result of `getComponent()` will be used to visualize the object.

Component `getComponent()` returns the main component for this user object. This component is used when the small component is not shown. To be compatible with the CBGraph (which is implemented in JFC/Swing), the component should be a subclass of `JComponent`. This method may return null, but then `getSmallComponent` must return a value.

Shape `getShape()` returns the shape for this component. Shape is defined in the package `java.awt` and represents any type of shape, e.g. polygon, rectangle, ellipse, etc. As said above, the shape is shown in the background of the small component, so it may not be visible if the small component is not transparent. The default implementation in `CBIndividual` and `CBLink` provide a transparent `JLabel` as small component so that the shape is visible. The CBGraph Editor contains several predefined shapes (see below).

boolean `doCommit()` This method is called when the user has clicked on the "Commit" button. Changes that have been made within this component (e.g. within a form) can then be added to the list of objects to be removed/added from the database. If the method returns false, then the commit operation will be aborted.

Furthermore, the `CBUserObject` class provides several methods which might be useful for implementing new graphical types:

String `toString()` returns the full object name of the Telos object which is represented by this user object

TelosObject `getTelosObject()` returns the Telos object which is represented by this user object (see documentation of package `i5.cb.telos.object` for more details on `TelosObject`)

CBFrame getCBFrame() returns the frame in which this object is presented (a CBFrame is an extension of a JInternalFrame)

ObjectBaseInterface getObi() returns the ObjectBaseInterface (i.e. a connection to the server) of the current frame. This is useful if you want to execute your own queries to retrieve additional information from the server. (see documentation of package `i5.cb.telos.object` and `i5.cb.api` for more details)

String getProperty(String property) returns the value of a property

boolean hasProperty(String property) returns true if the property is defined for this graphical type

JPopupMenu getPopupMenu() returns the popup menu for this object. By using this method you can extend the popup menu with own operations. If you overwrite this method, the operations in the default popup menu will not be available.

Example: The following example defines a graphical type that uses a JButton to visualize the object. Only the method `getSmallComponent` is overwritten. The background color of the button will be changed if the property `bgcolor` has been defined. The example shows that the implementation of graphical types is quite simple. Place the source code of your Java implementation of a graphical type in the subdirectory `ProductPOOL/java/i5/cb/graph/cbeditor` of your CBPOOL.

```
import i5.cb.graph.cbeditor.*;
import javax.swing.*;

import java.awt.Component;

/**
 * An example for a user defined CBUseObject
 */

public class CButton extends CBUseObject {

    public Component getSmallComponent() {
        JButton jButton=new JButton(this.getTelosObject().toString());
        if (hasProperty("bgcolor"))
            jButton.setBackground(CBUtil.stringToColor(getProperty("bgcolor")));
        return jButton;
    }
}
```

The source code of the above example is included in the `cbeditor` directory in file `CButton.java`. So, you can indeed use `"i5.cb.graph.cbeditor.CButton"` instead `"i5.cb.graph.cbeditor.CBIndividual"` in graphical types to see the effect.

C.3.5 Shapes

The package `i5.cb.graph.shapes` contains several shapes which might be useful for the ConceptBase CB-Graph Editor. To use these shapes, you can either specify the full path, e.g. `"i5.cb.graph.shapes.Cloud"`, or just the last part like `"Cloud"` as value of the property `shape` of a graphical type.

class name	graphical representation
Arrow2, ArrowL, ArrowR, DoubleArrow, DownArrow	various arrows
Banner	a banner
Circle	a circle
Cloud	a cloud shape
Cross	a cross (like the red cross)
Diamond	a diamond/rhombus
DiRect, DiRectL, DiRectR	direction signs
DownPentagon	like Pentagon but rotated 180 degrees
Ellipse	an ellipse
FolderL, FolderR	folder shapes
House	a house shape
Pentagon, Hexagon, Sep-tagon, Octagon	as the name says
Rect	a rectangle
RoundRectangle	a rectangle with round corners
Page	a page shape
Star	a star
Triangle, TriangleL, Trian-gleR, DownTriangle	various triangles
Tube	a tube shape
UpHexagon	hexagon with pointed vertex on top/bottom
StadionCurve	variant of a round rectangle resembling a stadion curve
UpStadionCurve	variant of StadionCurve
XCross	a cross in the form of an X
PolygonShape	user-definable polygon

The user-defined polygon-curve shape allows you to specify any shape consisting of a set of points. The start point must be the same as the end point. Assume, you want to triangle pointing to the right, but the right extreme point being at the same height 0 as the upper left point. Then, the following graphical type would do the job:

```
MyTriangle_GT in JavaGraphicalType with
  property
    ...
    shape : "PolygonShape; 0,3,0,0; 0,0,4,0"
  implementedBy
    implBy : "i5.cb.graph.cbeditor.CBIndividual"
  priority
    pr : 22
end
```

In the shape string, the first part "PolygonShape" indicates that it is a user-defined polygon shape, the second part "0,3,0,0" are the x-coordinates of the polygon points, and the third part "0,0,4,0" are its y-coordinates. Note that the number of x-coordinates must be the same as the number of y-coordinates and that the polygon line ends in its starting point, here (0,0). The size of the bounding rectangle in the above example is 4x5 pixels. If your shape is more complicated, e.g. a curved shape, then you should embed it into a bigger rectangle. The polygon lines may not intersect each other.

Figure ?? visualizes the pre-defined graph shapes. Note that by default the dimensions of a shape are adjusted from the area that the object label occupies. This is fine for the shapes that are close to a rectangle. The other shapes should be used in combination with the size "resizable".

A variant of the "resizable" option is the "wrap"/"wrap_" option. It will additionally wrap the node label text according to the current node size. The "wrap"/"wrap_" option renders the node label with the




Figure C.1: Some of the standard shapes

HTML implementation of Java.

Examples for the use of resizable shapes graphical types can be found in the CB-Forum at <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/3596768>.

You can extend the shapes by using the parameterized graph type PologonShape as shown in the previous subsection. Use the "align" property to specify at which position the node's label should be displayed. Default is "center". The above link also contains examples of user-defined shapes.

C.3.6 Icons

You can specify an image icon that is displayed instead of a shape to be drawn for the small component of a node (CBIndividual). The syntax for specifying an image icon is

```
image: "<image file location>"
```

You can specify either the URL of the image file or the local path of the file in the URL syntax. For example

```
Class AgentGT in JavaGraphicalType with
  rule
    gtrule : $ forall a/Agent (a graphtype AgentGT) $
  property
    textcolor : "0,0,0";
    linecolor : "0,0,0";
    image: "http://myserver.comp.eu/images/AgentIcon.png"
  implementedBy
    implBy : "i5.cb.graph.cbeditor.CBIndividual"
  priority
    pr : 20
end
```

associates the graphical type AgentGT to the image icon AgentIcon.png. Note that only normal http addresses are supported, not https. You can also point to local files via the file protocol:

```

Class AgentGT in JavaGraphicalType with
  rule
    gtrule : $ forall a/Agent (a graphtype AgentGT) $
  property
    textcolor : "0,0,0";
    linecolor : "0,0,0";
    image: "file:///home/jonny/images/AgentIcon.png"
  implementedBy
    implBy : "i5.cb.graph.cbeditor.CBIndividual"
  priority
    pr : 20
end

```

Note that the image icon is looked up by CBGraph. Hence, the location must be in the file system of the computer on which CBGraph runs. If you place the image icon on a web server, then CBGraph will be able to fetch it from any computer provided that the access rights are set properly. Note that the URL must use the "http" protocol. CBGraph does not support "https" links for image files.

If you specify an image icon for a graphical type, then you can also set its textposition, for example:

```

...
image: "file:///home/jonny/images/AgentIcon.png";
textposition: "top";
...

```

By default, the node's text label is placed at the bottom of the image. In this case it shall be placed on top of it. Other possible values are "center", "left", and "right". Note that the property textposition is only evaluated in combination with an icon image. If a graphical type has no image icon, then any text position specified for it would be ignored. In most cases, the default value "bottom" is just fine.

You can also combine shapes with image icons. In such cases, the image icon plus the label are the "inner content" and the shape is drawn around it. In the example below, the label is placed left of the image icon. Both are aligned in the center of a circular shape with gray background and black line color. CBGraph shall compute the required size of the surrounding shape from the dimensions of the image icon and its label. An exception holds when the "size" property is set to a fixed dimension like "50x40".

```

...
image: "file:///home/jonny/images/AgentIcon.png";
textposition: "left";
shape : "i5.cb.graph.shapes.Circle";
align : "center";
bgcolor : "200,200,200";
linecolor : "0,0,0";
...

```

The location specified in the "image" and "bgimage" property can either be a URL to an image file (starting with "http://" or "file://", not "https://") or a relative file location such as "diaicons/icon1.png". In the latter case, CBGraph shall first check if a local directory "CBICONS" exists in the ConceptBase installation directory (environment variable CB_HOME). If that exists, it shall expand the relative path to an absolute path using the location of CB_HOME. If the local directory does not exist, CBGraph shall expand the relative path to a URL starting with "http://conceptbase.sourceforge.net/CBICONS/". You can add your own icons to the local directory CBICONS in your ConceptBase installation directory. Below is an example of a relative image location.

```

...
image: "images/AgentIcon.png";
...

```

Further examples on using image icons are provided in the CB-Forum at <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/3506150>.

C.4 Object-specific graphical properties

Nodes and links get their graphical properties from the graphical type assigned to them. The assignment is typically defined by deductive rules deriving a fact (`x graphtype gt`). In general, several such facts may be true for a given object `x`. Then, the priority of the graphical type is used to pick a unique solution. As a result, all objects with the same graphical types are also rendered in the same way, except of the name of the object.

In some situations, one may want to assign specific graphical properties to objects depending on the object state. The graphical type provides the general properties and specific graphical properties are derived from the object. For example, all employees could be displayed by a rectangular node with white background, but employees with a high salary are displayed in yellow color. Further, employees that are assigned to departments get a thicker line width.

In principal, the different cases can be realized by dedicated graphical types. In the example above, one would need at least four different graphical types (regular employees without department, high salary employees without department, regular employees with department, and high salary employees without department). The different cases thus lead to an explosion of graphical types.

ConceptBase thus provides a second mechanism to directly assign graphical properties to objects. The graphical properties are defined for any proposition:

```
Proposition with
  attribute
    gproperty: Proposition
end
```

Note that the target class is `Proposition` rather than `String`, as used for the `attribute` property of graphical types. The reason is to add more flexibility, e.g. to assign integers as values for certain `gproperty` attributes like line width.

Any object³ may have (derived or explicit) `gproperty` attributes. The values of these attributes overrule the corresponding values of the graphical type of the object:

```
bill in Employee with
  ...
  gproperty
    bgcolor: "240,240,0";
    linecolor: "0,0,220"
end
```

The properties may also be derived by rules, e.g.

```
Employee in Class with
  attribute
    salary: Integer
  rule
    rel: $ forall e/Employee s/Integer (e salary s) and (s > 1000)
        ==> (e gproperty/bgcolor "255,255,200") $
end
```

The labels of the `gproperty` attributes shall be taken from the list in section ???. The object-specific graphical properties are not assigned to any palette. They are global and overrule the properties from the graphical type. It could be that there are multiple rules that derive the same `gproperty` attribute, e.g.

```
Individual in Class with
  rule
    rx1: $ forall x/Individual (x gproperty/bgcolor "255,255,255") $
end
```

³Due to technical limitations, only node objects (=instances of `Individual`) get their `gproperty` feature scanned by CB-Graph. Hence, you can only use it for node objects, not for edges.

This rule may collide with rule `re1`. In such cases, both `"255,255,200"` and `"255,255,255"` as values of `bgcolor`. `CBGraph` will then pick any of them (actually the last one transmitted overrules any previous ones). Since the order is subject to the `CBserver` rule engine, one can hardly predict, which value prevails. Hence, write design the rules in such a way that such collisions are avoided.

The `gproperty` attribute `label` adds a new functionality to the system: you can overrule the node and link name displayed in `CBGraph`. For example, it may be useful to replace the name of a shelf with its current fill level. Another example are 'places' of petri nets. Instead of the place name, one can display the number of tokens of that place as the label of the place node.

Another interesting `gproperty` attribute is `'labellength'`. By default, `CBGraph` assumes a maximum label length of 40 characters. If the node label length exceeds this threshold, it will be truncated and the last four characters are set to `"..."`. If you need to have longer labels, then use the `'labellength'` property. An example shows how to use it:

```
Employee in Class with
  attribute
    name : String
  rule
    r1 : $ forall e/Employee n/String (e name n) ==> (e gproperty/label n)$;
    r2 : $ forall e/Employee (e gproperty/labellength 50)$
  end

bill in Employee with
  name
    n : "William the Conquerer from Abessinia della Cruz"
  end
```

You can easily check, which object-specific properties are currently assigned to objects by the following query:

```
ObjectProperty in QueryClass isA Proposition with
  retrieved_attribute
    gproperty : Proposition
  end
```

You can retrieve the objects with colliding `gproperty` attributes via the query

```
ObjectWithMultipleProperties in QueryClass isA Proposition with
  retrieved_attribute
    gproperty : Proposition
  constraint
    clash : $ exists L/Label p1,p2/Proposition (this gproperty/L p1) and
      (this gproperty/L p2) and (p1 <> p2) $
  end
```

We advise to use the `gproperty` feature in combination with graphical types. The graphical type of an object provides the graphical properties that apply to all objects that fall into the class covered by the graphical type. The object-specific properties then overrule certain properties of that graphical type or add properties that were not defined by the more general graphical type. Use the `gproperty` feature with great care. For example, assigning object-specific properties to all instances of the class `Individual` is not wise since `Individual` is a generic class: all node-like objects are instances of `Individual`.

C.5 Graphical types for derived links

Derived links (and attributes) are displayed by default with the graphical type `ImplicitAttributeGT`, i.e. a dashed line with the attribute label defined at the class level. Derived links have no object identity. Thus, one cannot attach a `graphtype` or `gproperty` attribute to them.

ConceptBase uses another method to allow user-definable graphical types for such links. Assume, there is a class definition as follows:

```

Person in Class with
  attribute
    knows : Person
  rule
    trrrule : $ forall x,y,z/Person
              (x knows y) and (y knows z)
              ==> (x knows z) $
end

```

Let the attribute `knows` be derived by some rules.

Then, one can define a graphical type `ImplicitGT_knows` that shall be applied to all derived using the class label `knows`, e.g.

```

ImplicitGT_knows in JavaGraphicalType with
  property
    textcolor : "20,20,220";
    edgecolor : "250,20,20";
    bgcolor : "255,255,255,100";
    edgestyle : "dashdotted";
    edgewidth : "3"
  implementedBy
    implBy : "i5.cb.graph.cbeditor.CBLink"
  priority
    p : 10
end

```

This graphical type then has to be added to the right graphical palette:

```

PersonPalette in Class,JavaGraphicalPalette with
  contains,defaultIndividual
    xx1 : DefaultIndividualGT
    ...
  contains
    xx14 : ImplicitGT_knows
end

```

Note that the label of the graphical type starts with the prefix `ImplicitGT_`, which is then followed by the label of the derived link. `CBGraph` shall assign this graphical type for derived links if the current graphical palette contains such a graphical type. Otherwise, the default (usually `ImplicitAttributeGT` or the graphical type listed as `implicitAttribute` in the graphical palette) is used.

The user-defined graphical types for derived links allows to create domain-specific visualizations of derived information. It is rather common to have multiple derived link types such as `knows`. It thus makes sense to distinguish them also in the graphical visualization.

Derived instantiations ("in") and derived specializations ("isA") are handled differently. Their dedicated graphical type is specified in the graphical palette as follows:

```

contains,implicitIsA
  c3 : ImplicitIsAGT
contains,implicitInstanceOf
  c4 : ImplicitInstanceOfGT

```

Their link name is fixed in the database is fixed. Derived attributes in contrast can have any link name.

C.6 Palette-specific methods to expand related objects

Nodes and links in the graph editor CBGraph can be expanded to show their instances/classes, subclasses/superclasses, and attributes/relations. For the latter, the default behavior of CBGraph is to determine which attribute/relation categories are actually used by the selected object and then create the suitable popup-menu for the object by only shows those categories that are actually used. The queries to compute these categories are:

```
find_used_attribute_categories
  in GenericQueryClass isA Proposition!attribute with
  parameter,required
  objname : Proposition
  constraint
    r : $ exists x/Proposition AD(this,~objname,x)  $
end

find_used_incoming_attribute_categories
  in GenericQueryClass isA Proposition!attribute with
  parameter,required
  objname : Proposition
  constraint
    r : $ exists x/Proposition AD(this,x,~objname)  $
end
```

This is convenient but can also be a very costly operation in case that the object is occurring in many derived facts (derived relations, derived attributes).

A way out of this dilemma are dedicated queries that computes the eligible categories of the derived outgoing and incoming attributes/relations. Consider the example below:

```
MyPalette in Class,JavaGraphicalPalette isA XPalette with
  contains
    gt1: THING_GT;
    ...
  palproperty
    outcatquery : "alt_used_attribute_categories";
    incatquery  : "alt_used_incoming_attribute_categories"
end
```

The two new palette properties outcatquery and incatquery specify the replacement queries for the default queries. Note that you need to define these queries as well, e.g.:

```
alt_used_attribute_categories
  in GenericQueryClass isA Proposition!attribute with
  parameter,required
  objname : Proposition
  constraint
    r : $ exists gt/JavaGraphicalType
        (~objname graphtype gt) and (gt forOutgoing ~this)  $
end
```

In this case, the applicable attribute categories are attached to the graphical types of objects:

```
THING_GT in Class,JavaGraphicalType with
  property
    ...
  rule
    gtrule : $ forall x/Thing (x graphtype THING_GT) $
  forOutgoing
    out1 : Thing!aproperty
end
```

The new feature allows to hide certain attribute/relation properties from the pop-up menu.

Appendix D

Examples

D.1 Example model: the employee model

The Employee model can be found in the directory `$CB_HOME/examples/QUERIES/`. It consists out of the following files:

Employee_Classes.sml: The class definition

Employee_Instances.sml: Some instances for this model

Employee_Queries.sml: Queries for this model

Note, that the files must be loaded in this order into the server.

D.2 A Telos modeling example - ER diagrams

D.2.1 The basic model

This example gives a first introduction into some features introduced in ConceptBase version 4.0. It demonstrates the use of *meta formulas* and *graphical types* while building a Telos model describing Entity-Relationship-Diagrams. The following model forms the basis:

```
{ ***** }
{ *                * }
{ * File: ERModelClasses * }
{ *                * }
{ ***** }

Class Domain
end

Class EntityType with attribute
    eAttr : Domain;
    keyeAttr : Domain
end

Class RelationshipType with attribute
    role : EntityType
end
```

```

Class MinMax
end

"(1,*)" in MinMax with
end

"(1,1)" in MinMax with
end

Attribute RelationshipType!role with attribute
    minmax: MinMax
end

```

The model defines the concepts of *EntityType*s and *RelationshipType*s. Each entity that participates in a relationship plays a particular *role*. This role is modelled as a Telos *attribute-link* of the object *RelationshipType*. The *attributes* describing the entities are modelled as Telos *attribute-links* to a class *Domain* containing the value-sets. Roles can be restricted by the “(min,max)”-constraints “(1,*)” or “(1,1)”. The next model contains a concrete ER-model.

```

{*****}
{*          *}
{* File: Emp_ERModel    *}
{*          *}
{*****}

Class Employee in EntityType with
    keyAttr, attribute
        ssn : Integer
    eAttr, attribute
        name : String
end

Class Project in EntityType with
    keyAttr, attribute
        pno : Integer
    eAttr, attribute
        budget : Real
end

Integer in Domain end

Real in Domain end

String in Domain end

Class WorksIn in RelationshipType with role, attribute
    emp : Employee;
    prj : Project
end

WorksIn!emp with minmax
    mProjForEmp: "(1,*)"
end

WorksIn!prj with minmax
    mEmpForProj: "(1,*)"
end

```

```

ConceptBase in Project with pno
    cb_pno : 4711
end

Martin in Employee with ssn
    martinSSN : 4712
end

M_CB in WorksIn with
    emp
        mIsEmp : Martin
    prj
        cbIsPrj : ConceptBase
end

Hans in Employee with ssn
    hans_ssn : 4714
end

```

The entity-types *Employee* and *Project* participate in a binary relationship *WorksIn*. The attributes *Employee!ssn* and *Project!pno* are key-attributes of the respective objects.

D.2.2 The use of meta formulas

The above model distinguishes *attributes* and *key attributes*. One important constraint on key attributes is monovalence. In the previous releases of ConceptBase it was possible to declare Telos-attributes as instance of the attribute-categories *single* or *necessary*, but the constraint ensuring this property could not be formulated in a general manner, because the use of variables as placeholders for Telos-classes e.g in an *In-Literal* was prohibited. To overcome this restriction, *meta formulas* have been integrated into the system. An assertion is a *meta formula* if it contains such a class-variable. The system tries to replace this *meta formula* by a set of semantic equivalent formulas which contain no class-variables. In previous releases properties as *single* or *necessary* had to be ensured “manually” by adding a constraint for each such attribute. This job is now performed automatically by the system.

Example: *necessary* and *single*

The following meta formula ensures the *necessary* property of attributes, which are instances of the category *Proposition!necessary*. The semantics of this property is, that for every instance of the source class of this attribute there must exist an instantiation of this attribute.

```

Class with constraint, attribute
    necConstraint:
        $ forall c,d/Proposition p/Proposition!necessary x,m/VAR
            P(p,c,m,d) and (x in c) ==>
                exists y/VAR (y in d) and (x m y) $
end

```

It reads as follows:

For each attribute with label *m* between the classes *c* and *d*, which instantiates the attribute *Proposition!necessary* and for each instance *x* of *c* there should exist an instance *y* of *d* which is destination of an attribute of *x* with category *m*.

One should notice that the predicates *In (x, c)* and *In (y, d)* cause this formula to be a meta formula. The instantiation of *x* and *m* to the class *VAR* is just a syntactical construct. Every variable in a constraint

has to be bound to a class. This restriction is somehow contrary to the concept of meta formulas and the VAR-construct is a kind of compromise. The resulting In-predicates are discarded during the processing of meta formulas. The VAR-construct is only allowed in meta formulas. It enables the user to leave the concrete classes of *x* and *m* open, without instantiating them to for example to *Proposition*.

The *single*-constraint can be defined in analogy. These constraints can be added to the system as if they were “normal” constraints. Their effect becomes visible, when declaring attributes as *necessary* or *single*. This is done in the following model.

```
{*****}
{*          *}
{ File: ERSingNec      *}
{*          *}
{*****}

{* necessary constraint (metaformula) *}
Class with constraint
  necConstraint:
    $ forall c,d/Proposition p/Proposition!necessary x,m/VAR
      P(p,c,m,d) and (x in c) ==>
        exists y/VAR (y in d) and (x m y) $
end

{* every Entity has a key *}
Class EntityType with
  necessary
    keyAttr : Domain
end

{* single constraint (metaformula) *}
Class with constraint
  singleConstraint :
    $ forall c,d/Proposition p/Proposition!single x,m/VAR
      P(p,c,m,d) and (x in c) ==>
        (
          forall a1,a2/VAR
            (a1 in p) and (a2 in p) and Ai(x,m,a1) and Ai(x,m,a2) ==>
              (a1=a2)
        ) $
end

{* every Entity key is monovalued ( = necessary and single) *}
Class EntityType with rule
  keys_are_necessary:
    $forall a/EntityType!keyAttr In(a,Proposition!necessary)$;
  keys_are_single:
    $forall a/EntityType!keyAttr In(a,Proposition!single)$
end
```

The effects of this transaction can be shown by displaying the instances of instances of the class *metaMSFOLconstraint*. The *single*- and *necessary* constraints are inserted into this class after adding them to the system. These constraints themselves can have specializations: constraints which are added automatically to the system when inserting objects into the attribute-category *single* resp. *necessary*.

For the ER-example, one of the created formulas reads as:

```
$ forall x/EntityType (exists y/Domain (x keyAttr y)) $
```

We observe the relationship to the `necConstraint`-formula: the formula has been generated by computing one extension of $\text{In}(p, \text{Proposition!necessary})$ and $P(p, c, m, d)$ and replacing the predicates $\text{In}(p, \text{Proposition!necessary})$ and $P(p, c, m, d)$ by this extension, which results in the following substitution for the remaining formula:

c	EntityType
d	Domain
m	keyAttr

Metaformulas defining sets of rules

Another use of *Metaformulas* is the formulation of deductive rules. Metaformulas defining deductive rules extend the possibilities of defining derived knowledge.

Assignment of graphical types: This example first demonstrates the use of meta formulas to assign graphical types to object-categories. The minimal graphical convention for ER-diagrams is to use rectangular boxes for entities and diamond-shaped boxes for relationships. In our modelling example these graphical types are assigned to objects which are instances of *EntityType* or *RelationshipType* and to instances of these objects.

```
{*****}
{*
{* File: ERModelGTs
{* Definition of the graphical palette for *}
{* ER-Diagrams for use on color displays *}
{*
{*****}

{* graphical type for inconsistent roles *}
Class InconsistentGtype in JavaGraphicalType with
  attribute,property
    textcolor : "0,0,0";
    edgecolor : "255,0,0";
    edgewidth : "2"
implementedBy
  implBy : "i5.cb.graph.cbeditor.CBLink"
priority
  p : 14
end

{* graphical type for entities *}
Class EntityTypeGtype in JavaGraphicalType with
property
  bgcolor : "10,0,250";
  textcolor : "0,0,0";
  linecolor : "0,55,144";
  shape : "i5.cb.graph.shapes.Rect"
implementedBy
  implBy : "i5.cb.graph.cbeditor.CBIndividual"
priority
  p : 12
end

{* graphical type for relationships *}
Class RelationshipGtype in JavaGraphicalType with
```

```

property
  bgcolor : "255,0,0";
  textcolor : "0,0,0";
  linecolor : "0,0,255";
  shape : "i5.cb.graph.shapes.Diamond"
implementedBy
  implBy : "i5.cb.graph.cbeditor.CBIndividual"
priority
  p : 13
end

(* graphical palette *)
Class ER_GraphBrowserPalette in JavaGraphicalPalette with
  contains,defaultIndividual
    c1 : DefaultIndividualGT
  contains,defaultLink
    c2 : DefaultLinkGT
  contains,implicitIsA
    c3 : ImplicitIsAGT
  contains,implicitInstanceOf
    c4 : ImplicitInstanceOfGT
  contains,implicitAttribute
    c5 : ImplicitAttributeGT
  contains
    c6 : DefaultIsAGT;
    c7 : DefaultInstanceOfGT;
    c8 : DefaultAttributeGT;
    c14 : EntityTypeGtype;
    c15 : RelationshipGtype;
    c16 : InconsistentGtype
end

EntityType with rule
  EntityGTRule:
    $ forall e/EntityType A(e,graphtype,EntityTypeGtype)$ ;
  EntityGTMetaRule:
    $ forall x/VAR (exists e/EntityType In(x,e)) ==>
      A(x,graphtype,EntityTypeGtype)$
end

RelationshipType with rule
  RelationshipGTRule:
    $ forall r/RelationshipType A(r,graphtype,RelationshipGtype)$ ;
  RelationshipGTMetaRule:
    $ forall x/VAR (exists r/RelationshipType In(x,r)) ==>
      A(x,graphtype,RelationshipGtype) $
end

```

To activate the *ER_GraphBrowserPalette* , select this graphical palette when you start the Graph Editor or make a new connection in the Graph Editor (see section ??).

Handling inconsistencies The *necessary* and *single* conditions on attributes from the previous section could also be expressed as deductive rule. The difference is, that if they are formulated as constraints, every transaction violating the constraint would be rejected. The definition of rules handling *necessary* and *single* enables the user to handle inconsistencies in his model. The following example demonstrates this

concept in the context of our ER-model. The example defines rules handling the restriction of roles by the “(min,max)”-constraint “(1,*)”. This restriction is not implemented using constraints. Instead a new Class *Inconsistent* is defined, containing all role-links which violate the “(1,*)” constraint. These inconsistent links get a different graphical type (e.g a red coloured attribute link) than consistent role links to visualize the inconsistency graphically.

```
{*****}
{*                                     *}
{* File: ERIncons                      *}
{* Definition of a Class "Inconsistent" *}
{* containing roles violating the "(1,*)" *}
{* constraint                          *}
{*                                     *}
{*****}

{* new attribute category revNec for attributes
   which are "reverse necessary" *}
Class with attribute
    revNec : Proposition
end

{* roles with "(1,*)" must fullfill revNec property *}
RelationshipType with rule, attribute
    revNecRule:
    $ forall ro/RelationshipType!role
        A(ro,minmax,"(1,*)") ==> In(ro,Class!revNec) $
end

{* definition of Class "Inconsistent" *}
{* forall instances "p" of Class!revNec:
   If there exists a destination class "d", there must
   be a source class "c" with an attribute instantiating "p",
   otherwise "p" is inconsistent
   *}
Class Inconsistent with rule, attribute
    revNecInc:
    $ forall p/Class!revNec
        (exists c,m,d/VAR y/VAR P(p,c,m,d) and In(y,d) and
         not(exists x/VAR In(x,c) and A(x,m,y))) ==>
         In(p,Inconsistent)$
end
```

To activate the different graphical representation of inconsistent roles, the definition of the graphical type for attributes has to be modified. Tell the following frame:

```
Inconsistent with rule,attribute
    incRule :
    $ forall e/Attribute In(e,Inconsistent) ==>
        (e graphtype InconsistentGtype)$
end
```

The effect of these transactions can be shown when starting the *Graph Editor* and displaying the attributes of the RelationshipType instance WorksIn. Be sure to switch the graph editor to the palette ER_GraphBrowserPalette before doing so (see section ??). The attribute WorksIn!emp is displayed as red link like in figure ?. If you had already started the graph editor before telling the last frame, you should synchronize it with the CBserver via the menu Current connection. By telling

```
H_CB in WorksIn with
```

```

emp
  hIsEmp : Hans
prj
  cbIsPrj : ConceptBase
end

```

the inconsistency is removed from the model. To see the update of the graphical type in the Graph Editor, you have to select the inconsistent link and select “Validate and update selected objects” from the “Current connection” menu.




Figure D.1: Graph Editor showing the example model with inconsistent link

D.2.3 Limitations and final remarks

Metaformulas extend the expressive power of defining rules and constraints for ConceptBase Objectbases. The implementation of this mechanism is not complete at the moment, but should enable the use of the most frequently requested concepts like *single* and *necessary*. Some limitations are listed below:

- **limited partial evaluation**

The partial evaluation procedure is limited to a conjunction of predicates, preceded by *forall*-quantors.

- **no source-to-source transformation**

Formulas are not converted automatically into the form mentioned above supported by the meta

formula mechanism, even if they could be transformed. If the class-variables can't be bound using partial evaluation of the input-formula, the formula is rejected, even if there exists an equivalent formula, in which partial evaluation could be used to bind the class-variables.

- **not all classes are supported**

Generated formulas where variables are quantified over instances of the following classes or attributes of them will be ignored. Those classes are:

Boolean, Integer, Real, String, TransactionTime,
MSFOLassertion, MSFOLrule, MSFOLconstraint,
metaMSFOLconstraint, metaMSFOLrule,
BDMConstraintCheck, BDMRuleCheck, LTruleEvaluator, ExternalReference,
QueryClass, BuiltinQueryClass, AnswerRepresentation,
GraphicalType, X11_Color, ATK_TextAlign,
ATK_Fonts, ATK_LineCap, ATK_ShapeStyle

The justification is twofold. Some of these generated formulas, e.g. a formula beginning with

```
$ forall x/1 ...$
```

are regarded as redundant, because the object *1* as instance of *Integer* should have no instances. Another justification is, that the use of meta formulas should be restricted to user-defined modeling tasks. Manipulation of the most system classes is disabled for reasons of efficiency and safety.

In the case of deductive rules, additional problems arise similar to the stratification problem. At the moment only monotonous transactions are allowed. This means that generated formulas can only be inserted or deleted during one transaction, both operations at the same time are not permitted.

The meta formula mechanism also influences the efficiency of the system: every transaction has to be supervised whether it affects the meta formulas or one of the generated formulas. If the preconditions of the generated formulas don't hold anymore, e.g. if the instances of *EntityType!eAttr* are no longer instances of *Proposition!necessary* in the previous model, the corresponding generated formula has to be deleted. If additional attributes are instantiated to *Proposition!necessary*, additional formulas have to be created. The process of supervising the transactions is quite expensive and if it slows down the overall performance too much, some of the meta formulas can be disabled temporarily (Untelling a meta formula removes all the code generated).

Many more examples for meta formulas, e.g. for defining transitivity of attributes, are available from the CB-Forum (<http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/1042523>).

Appendix E

Predefined Query Classes

This chapter gives an overview of the query classes which are predefined in a standard ConceptBase installation. The names of parameters of the queries are set in `typewriter` font. Most of the queries listed here are used by the ConceptBase user interface CBIva to interaction with the CBserver. A normal user typically formulates queries herself. In fact, most queries listed below are very simple and directly representation as query class. An exception are the functions for computation and counting. They cannot be expressed by simple query classes but extend the expressiveness of the system.

E.1 Query classes and generic query classes

These queries can also be used in the constraints of other queries.

E.1.1 Instances and classes

ISINSTANCE: Checks whether `obj` is instance of `class`. The result is either `TRUE` or `FALSE`.

IS_EXPLICIT_INSTANCE: Same as before, but returns `TRUE` only if `obj` is an *explicit* instance of `class`.

find_classes: Lists all objects of which `objname` is an instance.

find_instances: Lists all instances (implicit and explicit) of `class`.

find_explicit_instances: Same as before, but only explicit instances are returned.

E.1.2 Specializations and generalizations

ISSUBCLASS: Checks whether `sub` is subclass of `super`. The result is either `TRUE` or `FALSE`.

IS_EXPLICIT_SUBCLASS: Same as before, but returns `TRUE` only if `sub` is an *explicit* subclass of `super`.

find_specializations: Lists the subclasses of `class`. If `ded` is `TRUE`, then the result will also include implicit subclasses, if `ded` is `FALSE` only explicit information will be included.

find_specializations: Same as before, but for super classes.

E.1.3 Attributes

IS_ATTRIBUTE_OF: Returns TRUE if `src` has an attribute of the category `attrCat` which has the value `dst`.

IS_EXPLICIT_ATTRIBUTE_OF: Same as before, but only for explicit attributes.

find_all_explicit_attribute_values: Lists all attribute values of `objname` that are explicitly defined.

find_iattributes: Lists the attributes that *go into* `class`.

find_referring_objects: Lists the objects that have an explicit attribute link to `class`.

find_referring_objects2: Lists the objects that have an explicit attribute link to `objname` and for which the attribute link is an instance of `cat`.

find_all_referring_objects2: Same as before, but including implicit attributes.

find_attribute_categories: Lists all the attribute categories that may be used for `objname`. This is a lookup of all attributes of all classes of `objname`.

find_incoming_attribute_categories: In contrast to the previous query, this query returns all attribute categories that go into `objname` (i.e. attribute categories for which `objname` can be used as an attribute value).

find_attribute_values: Lists all objects that are attribute values of `objname` in the attribute category `cat`.

find_explicit_attribute_values: Same as before, but only for explicit attributes.

E.1.4 Links between objects

find_incoming_links: Lists the links that *go into* `objname` and are instance of `category`. Note that all types of links are returned, including attributes, instance-of-links and specialization links.

find_incoming_links_simple: Same as before, but without the parameter `category`.

find_outgoing_links: Lists the links that *come out of* `objname` and are instance of `category`. Note that all types of links are returned, including attributes, instance-of-links and specialization links.

find_outgoing_links_simple: Same as before, but without the parameter `category`.

get_links2: Return the links between `src` and `dst`.

get_links3: Return the links between `src` and `dst` that are instance of `cat`.

E.1.5 Other queries

find_object: This query just returns the object given as parameter `objname`, if it exists. Thus it can be used to check whether `objname` exists, but there is a builtin query `exists` which does the same. The query is mainly useful in combination with a user-defined answer format (e.g. the Graph Editor is using this query to retrieve the graphical representation of the object).

AvailableVersions: Lists the instances of `Version` with the time since when they are know. This query is used by the user interface to use a different rollback time (Options → Select Version).

listModule: Lists the the content of a module as Telos frames, see also section ??.

E.2 Functions

Functions may also be used within other queries. You may define your own functions (see section ??).

E.2.1 Computation and counting

COUNT: counts the instances of a class, this may be also a query class

SUM: computes the sum of the instances of a class (must be reals or integers)

AVG: computes the average of the instances of a class (must be reals or integers)

MAX: gives the maximum of the instances of a class (wrt. the order of < and >, see section ??)

MIN: gives the minimum of the instances of a class (wrt. the order of < and >, see section ??)

COUNT_Attribute: counts the attributes in the specified category of an object

SUM_Attribute: computes the sum of the attributes in the specified category of an object (must be reals or integers)

AVG_Attribute: computes the average of the attributes in the specified category of an object (must be reals or integers)

MAX_Attribute: gives the maximum of the attributes in the specified category of an object (wrt. the order of < and >, see section ??)

MIN_Attribute: gives the minimum of the attributes in the specified category of an object (wrt. the order of < and >, see section ??)

PLUS: computes the sum of two reals or integers

MINUS: computes the difference of two reals or integers

MULT: computes the product of two reals or integers

DIV: computes the quotient of two reals or integers

IPLUS: computes the sum of two integers; result is an integer number

IMINUS: computes the difference of two integers; result is an integer number

IMULT: computes the product of two integers; result is an integer number

IDIV: computes the quotient of two integers and then truncates the quotient to the largest integer number smaller than or equal to the quotient

ConceptBase realizes the arithmetic functions via its host Prolog system SWI-Prolog. Integer numbers on Linux are represented as 64-bit numbers, yielding a maximum range from $-2^{64} - 1$ to $2^{64} - 1$. SWI-Prolog supports by default under Windows and Linux64 arbitrarily long integers. Real numbers are implemented by SWI-Prolog as 64-bit double precision floating point numbers. ConceptBase uses 12 decimal digits.

E.2.2 String manipulation

ConcatenateStrings: concatenates two labels, typically strings; the two arguments can be expressions that are concatenated after their evaluation

ConcatenateStrings3: concatenates three labels; arguments may not be expressions

ConcatenateStrings4: concatenates four labels; arguments may not be expressions

concat: same as ConcatenateStrings

StringToLabel: removes the quotes of a string and returns it as a label (not an object), useful if labels should be passed as a parameter of a query

toLabel: evaluates the argument and that creates an individual object with the canonical representation of the argument result; the canonical representation is an alphanumeric where special characters are replaced by substrings like "C30_" for special characters

For example, the expression `toLabel(concat(" ", 1+2))` will return `C42_3`. The substring `"C42_"` is the canonical representation (ASCII number) of `" "`. The subexpression `1+2` is evaluated to 3.

E.3 Builtin query classes

These queries must not be used within other queries as they do not return a list of objects. They may only be used directly from client programs.

exists: Checks whether `objname` exists and returns `yes` or `no`.

get_object: Returns the frame representation of `objname`. This query may be either called with just one parameter (`objname`) or with four parameters (`objname`, `dedIn`, `dedIsa`, `dedWith`). The `ded*`-parameters are boolean flags that indicate whether implicit (deduced) information should also be included in the frame representation. Note that the order of the parameters has to be the same as listed above.

get_object_star: Returns the frame representation for all objects with a label that match the given wildcard expression. Only simple wildcards with a star (*) at the end are allowed.

rename: Renames an object from `oldname` to `newname`. This is a low-level operation directly on the symbol table that works directly on the symbol table. It only checks whether `newname` is not already used as label for a different object, no other consistency checks are performed. The parameters have to be given in the order `newname`, `oldname`.

Appendix F

CBserver Plug-Ins

An LPI plug-in (“logic plug-in”) is a small Prolog program that is attached to the CBserver (which is implemented in Prolog) at startup-time. It extends the functionality of the CBserver, for example for user-defined builtin queries. A plug-in can also interface to the services of the operating system¹.

You can create a file like `myplugin.swi.lpi` to provide the implementation for user-defined builtin queries or for call actions in active rules (see section ??). You can use the full range of functions provided by the underlying Prolog system (here: SWI-Prolog, <http://www.swi-prolog.org>) and the functions of the CBserver to realize your implementation. You can consult the the CB-Forum for some examples at <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/2768063>. You find there examples for sending emails from active rules, and for extending the set of builtin queries and functions.

F.1 Defining the plug-in

Once you have coded your file `myplugin.swi.lpi`, there are two methods to attach it to the CBserver. The first method is to copy the file into an existing database directory created by the CBserver:

```
cbserver -g exit -d MYDB
cp myplugin.swi.lpi MYDB
```

The first command creates the database directory MYDB if not already existing and initializes it with the pre-defined system classes and objects. The second command copies the LPI file to the database directory. This method makes the definitions only visible to a CBserver that loads the database MYDB.

The second method instructs ConceptBase to load your LPI code to *any* new database created by the CBserver. To do so, copy the LPI file into the system database directory that holds the definitions of predefined ConceptBase objects:

```
cp myplugin.swi.lpi <CB_HOME>/lib/SystemDB
```

where CB_HOME is the directory into which you installed ConceptBase. The number of LPI files is not limited. You may define zero, one or any number of plug-in files.

A couple of useful LPI plug-ins are published via the CB-Forum, see <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/2768063>. Note that these plug-ins are copyrighted and typically come with their own license conditions that may be different to the license conditions of ConceptBase. If you plan to use the plugins for commercial purposes, you may have to acquire appropriate licenses from the plug-in’s authors.

¹Since we currently supply the CBserver for Linux only, you need to run the CBserver on a local Linux computer in your network if you want to use the plug-ins. Note that we supply a ready-to-use virtual appliance that includes Linux and ConceptBase and that can be executed via a virtualization engine, see <http://conceptbase.sourceforge.net/import-cb-appliance.html> for details.

F.2 Calling the plug-in

There are two ways to trigger the call of a procedure implemented by an LPI plugin.

1. By explicitly calling a builtin query class (or function) whose code has been implemented by the LPI plugin. The LPI code would then look similar to the code in `SYSTEM.SWI.builtin` and you must have defined an instance of `BuiltinQueryClass` that matches the signature of the LPI code. The call to the builtin query class may be enacted from the user interface, or it may be included as an `ASKquery` call in an instance of `AnswerFormat`. Refer for more information to section ?? and to the directory `Examples/BuiltinQueries` in your ConceptBase installation directory.
2. By calling the implemented function as a `CALL` action of an active rule. See section ?? for an example. In that case, there does not need to be a definition of a builtin query class (or function) to declare the signature of the procedure.

If the code of an LPI plugin realizes a ConceptBase function, e.g. selecting the first instance of a class, then you can use that function wherever functions are allowed. As an example, consider the definition of the LPI plugin `selectfirst.swi.lpi` from <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/d2984654/selectfirst.swi.lpi.txt>:

```
compute_selectfirst(_res,_class,_cl) :-
    nonvar(_class),
    cbserver:ask([In(_res,_class)]),
    !.

tell: 'selectfirst in Function isA Proposition with
parameter class: Proposition end'.
```

The first clause is the Prolog code. The predicate must start with the prefix `compute_` followed by the name of the function. The first argument is for the result of the function. Subsequently, each input parameter is represented by two arguments, one for the input parameter itself and a second as a placeholder of the type of the input parameter. The second clause tells the new function as ConceptBase object so that it can be used like any other ConceptBase function. For technical reasons, the 'tell' clause may not span over more than 5 lines. Use long lines if the object to be defined is large. The function object (here: `selectfirst`) is stored in the `System` module of the database. This is the root module of the module hierarchy, thus functions defined in this way are visible and executable in all sub-modules. If you omit the 'tell' clause in the LPI file, then you need to tell it manually to the database. This can also be done in a module different to `System`. In this case, the function can only be called in those modules where the function object is visible.

If you just want to invoke the procedure defined in an LPI plugin via the `CALL` clause of an active rule, you do not need to include a 'tell' clause. Consider for example the `SENDMAIL` plugin from <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/2269675>.

You need to be very careful with testing your code. Only use this feature for functions that cannot be realized by a regular query class or by active rules. LPI code has the full expressiveness of a programming language. Program errors may lead to crashes of the CBserver or to infinite loops or to harmful behavior such as deletion of files. Query classes, deductive rules and integrity constraints can never loop infinitely and can (to the best of our knowledge) only produce answer, not changes to your file system or interact with the operating system. Active rules could loop infinitely but also shall not change your file system and shall not interact with the operating system unless you call such code explicitly in the active rule.

You can disable the loading of LPI plugins with the CBserver option `-g nolpi`. The CBserver will then not load LPI plugins upon startup. This option might be useful for debugging or for disabling loading LPI plugins that are configured in the `lib/SystemDB` sub-directory of your ConceptBase installation.

F.3 Programming interface for the plug-ins

The CBserver plug-ins need to interface to the functionalities of the CBserver, the Prolog runtime, and possibly the operating system. To simplify the programming of the CBserver plug-ins, we document here the interface of the module `cbserver`. We assume that the code for the plug-ins is written in SWI-Prolog and the user is familiar with the SWI-Prolog system.

cbserver:ask(Q,Params,A) asks the query `Q` with parameters `Params` to the CBserver. The answer is returned as a Prolog atom in `A`. The atom holds the answer represented in the answer format of `Q`, see also chapter ??.

Example: `cbserver:ask(find_classes,[bill/objname],A)`

cbserver:ask(Preds) evaluates the predicates in list `Preds`. The predicate can backtrack and will bind in case of success the free variables in `Preds`. We currently support only the following predicates: `In`, `A`, `AL`, and `Isa`.

Example: `cbserver:ask([In(X,Employee),A(X,salary,1000)],S)`

cbserver:askAll(X,Preds,Set) finds all objects `X` that satisfy the condition in `Preds` and puts the result into the list `Set`. Supported predicates in `Preds` are: `In`, `A`, `AL`, and `Isa`.

Example: `cbserver:askAll(X,[In(X,Employee),A(X,salary,1000)],S)`

cbserver:tellFrames(F) tells all frames contained in the atom `F`. The call will fail if there is any error in `F`.

Example: `cbserver:tellFrames('bill in Employee end')`

cbserver:makeName(Id,A) converts an object identifier to a readable object name.

cbserver:makeId(A,Id) converts an object name (Prolog atom) into an object identifier used by the CBserver to identify Telos objects. If `A` is already an object identifier, it is returned as well in `Id`.

cbserver:arg2val(E,V) transforms an argument (either an object identifier or a functional expression) to a value (a number or a string). The value can then be used in Prolog style computations such as arithmetic expressions.

cbserver:val2arg(V,I) transforms a Prolog value (number, string) to an object identifier, possibly by creating a new object for the value.

cbserver:concat(X,Y) concatenates the strings (Prolog atoms) contained in the list `X`. The result is returned in `Y`.

Note that ConceptBase internally manages concepts by their object identifier. The programming interface instead addresses concepts (and objects) by their name, i.e. the label of the object or the Prolog value corresponding to the label. You may have to use the procedure `makeName` and `makeId` to switch between the two representations. The two procedures `arg2val` and `val2arg` are useful for defining new builtin functions on the basis of Prolog's arithmetic functions. Assume for example, that the object identifier `id123` has been created to correspond to the real number `1.5`. Then, the following relations hold: `makeName(id123,'1.5')`, `arg2val(id123,1.5)`. Hence, `makeName` returns the label `'1.5'` whereas `arg2val` returns the number `1.5`.

The interface shall be extended in the future to provide more functionality. Be sure that you only use this feature if user-defined query classes cannot realize your requirements!

References

- [All83] Allen, J.F.: Maintaining knowledge about temporal intervals. *Comm. ACM*, Vol. 26, No. 11, Nov. 1983, doi <http://dx.doi.org/10.1145/182.358434>.
- [BBMR89] Borgida, A., Brachman, R.J., McGuinness, D.L., Resnick, L.A.: CLASSIC: a structural data model for objects. *Proc. ACM-SIGMOD Intl. Conf. Management of Data*, Portland, Or, 58-67, 1989, doi <http://dx.doi.org/10.1145/66926.66932>.
- [BCG*87] Banerjee, J., Chou, H.-T., Garza, J.F., Kim, W., Woelk, D., Ballou, N., Kim, H.-J.: Data model issues for object-oriented applications. *ACM Trans. Office Information Systems* 5, 1, 3-26, 1987, doi <http://dx.doi.org/10.1145/22890.22945>.
- [BMS84] Brodie M.L., Mylopoulos J., Schmidt J.W. (ed.): *On Conceptual Modeling*, Springer-Verlag, 1984, doi <http://dx.doi.org/10.1007/978-1-4612-5196-5>.
- [Chen76] Chen, P. P.-C.: The Entity-Relationship-Model – Towards a Unified View of Data. *ACM Transact. on Database Systems*. Vol. 1, No. 1, March 1976, pp. 9-36, 1976, doi <http://dx.doi.org/10.1145/320434.320440>.
- [CW96] Chen, W., Warren, D.S. : Tabled Evaluation with Delaying for General Logic Programs. Vol. 43, No. 1, January 1996, pp. 20-74, 1996, doi <http://dx.doi.org/10.1145/227595.227597>.
- [DPZ2002] Dahchour, M., Pirotte, A., Zimanyi, E.: Materialization and its metaclass implementation. *IEEE Trans. Knowledge and Data Engineering*, 14, 5, pp 1078 - 1094, Sept/Oct 2002, doi <http://dx.doi.org/10.1109/TKDE.2002.1033775>.
- [Eber97] Eberlein, A.: *Requirements Acquisition and Specification for Telecommunication Services*. PhD Thesis, University of Wales, Swansea, UK, 1997.
- [Gall90] Gallersdörfer, R.: *Realization of a Deductive Object Base by Abstract Data Types* (in German), Diploma thesis, Universität Passau, Germany, 1990.
- [HJEK90] Hahn, U., Jarke, M., Eherer, S., Kreplin, K.: CoAUTHOR: a hypermedia group authoring environment. In Benford, J.M., Bowers, S.D. (eds.): *Studies in Computer-Supported Cooperative Work*, North-Holland, 79-100, 1990.
- [Jark93] Jarke, M. (ed.): *Database Application Engineering with DAIDA*, Springer-Verlag, 1993, doi <http://dx.doi.org/10.1007/978-3-642-84875-9>.
- [JeJa91] Jeusfeld, M.A., Jarke, M.: From Relational to Object-Oriented Integrity Simplification. *Proc. 2nd Intl. Conf. Deductive and Object-Oriented Databases*, Munich, Dec. 1991, doi http://dx.doi.org/10.1007/3-540-55015-1_25; also as *Aachener Informatik-Berichte* 91-19, RWTH Aachen, Germany, 1991.
- [JeSt92] Jeusfeld, M., Staudt, M.: Query optimization in deductive object bases. In Freytag, Vossen, Maier (eds.): *Query Processing for Advanced Database Applications*, Morgan-Kaufmann, 1992; also as *Aachener Informatik-Berichte* 91-26, RWTH Aachen, Germany, URL <http://conceptbase.sourceforge.net/mjf/91-26-r.pdf>.

- [Jeus92] Jeusfeld, M.A.: Änderungskontrolle in deduktiven Objektbanken. Dissertation, Universität Passau, Germany, postprint <http://merkur.informatik.rwth-aachen.de/pub/bscw.cgi/d340216/disk119.pdf>.
- [Jeus09] Jeusfeld, M.A.: Metamodeling and Method Engineering with ConceptBase. In [?], preprint <http://conceptbase.sourceforge.net/mjf/Metamodeling-and-method-engineering-with-ConceptBase--preprint.pdf>.
- [JGJ*95] Jarke, M., Gellersdörfer, R., Jeusfeld, M.A., Staudt, M., Eherer, S.: ConceptBase: A deductive object base for meta data management. *Journal of Intelligent Information Systems*, Vol. 4, No. 2, pp. 167-192, March 1995, doi <http://dx.doi.org/10.1007/BF00961873>.
- [JJM*09] Jeusfeld, M.A., Jarke, M., Mylopoulos, J. (eds.): Metamodeling for Method Engineering. Cambridge, MA, 2009. The MIT Press, ISBN-10: 0-262-10108-4, <https://mitpress.mit.edu/books/metamodeling-method-engineering>.
- [JJN*98] Jeusfeld, M.A., Jarke, M., Nissen, H.W., Staudt, M.: ConceptBase - Managing Conceptual Models about Information Systems. In Bernus, P., Mertins, K., Schmidt, G., (eds.): Handbook on Architectures of Information Systems, Springer-Verlag, pp. 265-285, 1998, http://dx.doi.org/10.1007/3-540-26661-5_12.
- [JJQV99] Jarke, M., Jeusfeld, M.A., Quix, C., Vassiliadis, P.: Architecture and Quality for Data Warehouses: An Extended Repository Approach *Information Systems*, Vol. 24, No. 3, pp. 229-253, 1999, doi [http://dx.doi.org/10.1016/S0306-4379\(99\)00017-4](http://dx.doi.org/10.1016/S0306-4379(99)00017-4).
- [JJR88] Jarke, M., Jeusfeld, M., Rose, T.: A global KBMS for database software evolution: documentation of first ConceptBase prototype, Report MIP-8819, Universität Passau, Germany, 1988.
- [JJR90] Jarke, M., Jeusfeld, M., Rose, T.: A software process data model for knowledge engineering in information systems. *Information Systems* 15, 1, 85-116, 1990, doi [http://dx.doi.org/10.1016/0306-4379\(90\)90018-K](http://dx.doi.org/10.1016/0306-4379(90)90018-K).
- [JJS*99] Jeusfeld, M.A., Jarke, M., Staudt, M., Quix, C., List, T.: Application Experience with a Repository System for Information Systems Development. In R. Kaschek (ed.): *Proceedings Entwicklungsmethoden fuer Informationssysteme und deren Anwendung (EMISA'99)*. Reihe Wirtschaftsinformatik, Teubner Verlag, Stuttgart, Germany, pp. 147-174, 1999, postprint <http://conceptbase.sourceforge.net/mjf/itrs004.pdf>.
- [JN2016] Jeusfeld, M.A., Neumayr, B.: DeepTelos: Multi-level Modeling with Most General Instances. *Proc. 35th International Conference on Conceptual Modeling (ER 2016)*, Gifu, Japan, Springer, LNCS 9974, 2016, 198-211, postprint <http://conceptbase.sourceforge.net/mjf/deeptelos-er2016-postprint.pdf>.
- [JQJ98] Jeusfeld, M.A., Quix, C., Jarke, M.: Design and Analysis of Quality Information for Data Warehouses. *17th International Conference on the Entity Relationship Approach (ER'98)*, Singapore, 1998, postprint <http://conceptbase.sourceforge.net/mjf/er98.pdf>.
- [Lud2010] Ludwig, S.A.: Comparison of a Deductive Database with a Semantic Web reasoning engine. *Knowledge-Based Systems*, 23, 6, 2010, pp. 634-642, doi <http://dx.doi.org/10.1016/j.knosys.2010.04.005>.
- [MBJK90] Mylopoulos, J., Borgida, A., Jarke, M., Koubarakis, M.: Telos: a language for representing knowledge about information systems. *ACM Trans. Information Systems* Vol. 8, No. 4, 1990, doi <http://dx.doi.org/10.1145/102675.102676>.
- [MBW80] Mylopoulos, J., Bernstein, P.A., Wong, H.K.T.: A language facility for designing database-intensive applications. *ACM TODS*, Vol. 5, No. 2, June 1980, doi <http://dx.doi.org/10.1145/320141.320150>.

- [NJS*14] Neumayr, B., Jeusfeld, M.A., Schrefl, M., Schütz, C.: Dual Deep Instantiation and its ConceptBase Implementation. *Proc. Advanced Information Systems Engineering (CAiSE 2014)*, Thessaloniki, Greece, June 16-20, 2014, 503-517, postprint http://conceptbase.sourceforge.net/mjf/Neumayr-et-al-CAiSE2014_postprint.pdf.
- [NJJ*96] Nissen, H.W., Jeusfeld, M.A., Jarke, M., Zemanek, G.V., Huber, H.: Managing Multiple Requirements Perspectives with Metamodels. *IEEE Software*, pages 37-48, March 1996, doi <http://dx.doi.org/10.1109/52.506461>.
- [OMG11] Object Management Group: OMG Unified Modeling Language Infrastructure, Version 2.4.1 OMG Document Number: formal/2011-08-05, URL <http://www.omg.org/spec/UML/2.4.1/>.
- [Quix96] Quix, C.: Sichtenwartung und Änderungsnotifikation für Anwendungsprogramme in deduktiven Objektbanken. Diploma Thesis, RWTH Aachen, 1996.
- [QJ02] Quix, C., Schoop, M., Jeusfeld, M.A.: Business Data Management for Business-to-Business Electronic Commerce *SIGMOD Record*, Vol. 31, No. 1, pp. 49-54, March 2002, URL http://dl.acm.org/ft_gateway.cfm?ftid=76258&id=507348.
- [RaDh92] Ramesh, B., Dhar, V.: Supporting systems development by capturing deliberations during requirements engineering. *IEEE TSE*, Vol. 18, No. 6, 1992, doi <http://dx.doi.org/10.1109/32.142872>.
- [RJG*91] Rose, T., Jarke, M., Gocek, M., Maltzahn, C., Nissen, H.: A decision-based configuration process environment. *Software Engineering Journal*, Special Issue on Software Environments and Factories, July 1991.
- [RLGW14] Rossini, A., de Lara, J., Guerra, E., Wolter, U.: A formalisation of deep metamodelling. *Formal Aspects of Computing*, 26, 6, pp 1115-1152, November 2014, doi <http://dx.doi.org/10.1007/s00165-014-0307-x>.
- [SSW94] Sagonas, K., Swift, T., Warren, D.S.: XSB as an efficient deductive database engine. Proceedings of the ACM SIGMOD International Conference on the Management of Data, pages 442-453, Minneapolis, Minnesota, May 1994, doi <http://dx.doi.org/10.1145/191843.191927>.
- [SIGA91] Rich, C. (ed.): "Special Issue on Implemented Knowledge Representation and Reasoning Systems", *SIGART Bulletin* Vol. 2, No. 3, June 1991.
- [StJa96] Staudt, M., Jarke, M.: Incremental Maintenance of Externally Materialized Views. In *Proc. of the 22th VLDB Conference*, Bombay, India, 1996.
- [Soir96] Soiron, R.: Kostenbasierte Anfrageoptimierung in deduktiven Objektbanken. Diploma Thesis, RWTH Aachen, 1996.
- [Stan86] Stanley M.T.: *CML – a knowledge representation language with application to requirements modeling*, M.S.thesis, University of Toronto, Ontario, 1986.
- [Stau90] Staudt, M.: Anfragerepräsentation und -auswertung in deduktiven Objektbanken. Diploma thesis, Universität Passau, 1990.
- [TKDE90] Stonebraker M. (ed.): Special Issue on Database Prototype Systems, *IEEE Trans. on Knowledge and Data Engineering* Vol. 2, No. 1, March 1990.